

Effective Conversational AI

Chatbots that work

Andrew Freed
Cari Jacobs
Enikő Rózsa

Foreword by Jesús Mantas



Improve your conversational AI's effectiveness by following this cycle

Is the conversational AI meeting your goals?
What do users think?
What is the impact of recent changes?

MEASURE

There is always opportunity.
What are the low-performance areas?
How can they be improved?

IDENTIFY


Release to production.
Inform your users about improvements!

DEPLOY

Estimate expected effort and improvement.
Prioritize your backlog.
Implement what's most important.

IMPLEMENT

Improvements are based on analyzing conversation against these outcome dimensions.
There's more to success than "just" containing a conversation!



Effective Conversational AI

Effective Conversational AI

CHATBOTS THAT WORK

ANDREW R. FREED
CARI JACOBS
ENIKŐ RÓZSA
FOREWORD BY JESÚS MANTAS



MANNING
SHELTER ISLAND

For online information and ordering of this and other Manning books, please visit www.manning.com. The publisher offers discounts on this book when ordered in quantity. For more information, please contact

Special Sales Department
Manning Publications Co.
20 Baldwin Road
PO Box 761
Shelter Island, NY 11964
Email: orders@manning.com

©2025 by Manning Publications Co. All rights reserved.

No part of this publication may be reproduced, stored in a retrieval system, or transmitted, in any form or by means electronic, mechanical, photocopying, or otherwise, without prior written permission of the publisher.

Many of the designations used by manufacturers and sellers to distinguish their products are claimed as trademarks. Where those designations appear in the book, and Manning Publications was aware of a trademark claim, the designations have been printed in initial caps or all caps.

☺ Recognizing the importance of preserving what has been written, it is Manning's policy to have the books we publish printed on acid-free paper, and we exert our best efforts to that end. Recognizing also our responsibility to conserve the resources of our planet, Manning books are printed on paper that is at least 15 percent recycled and processed without the use of elemental chlorine.

The author and publisher have made every effort to ensure that the information in this book was correct at press time. The author and publisher do not assume and hereby disclaim any liability to any party for any loss, damage, or disruption caused by errors or omissions, whether such errors or omissions result from negligence, accident, or any other cause, or from any usage of the information herein.



Manning Publications Co.
20 Baldwin Road
PO Box 761
Shelter Island, NY 11964

Development editor: Rebecca Senninger
Technical editors: Jack C. Crawford
Stéfan van der Stockt
Review editor: Kishor Rit
Production editor: Kathy Rosslund
Copy editor: Andy Carroll
Proofreader: Melody Dolab
Typesetter and cover designer: Marija Tudor

ISBN 9781633436404
Printed in the United States of America

Andrew: Thank you to my wife Elise and kids Greg and Jeff for supporting me in writing another book!

*Cari: To Jason, for your never-ending support throughout my writing process and life in general. And to my dad, Jim.
(Surprise! I wrote a book!)*

Enikő: Thanks to my family, whose unwavering support and encouragement have made this book-writing journey possible. And to my late father, a prolific technical author who paved the way—I stand on your shoulders as I continue your legacy.

brief contents

PART 1	FRAMEWORK FOR IMPROVING CONVERSATIONAL AI	1
1	■ What makes conversational AI work?	3
2	■ Building a conversational AI	23
3	■ Planning for improvement	44
PART 2	PATTERN: AI DOESN'T UNDERSTAND	77
4	■ Understanding what your users really want	79
5	■ Improving weak understanding for traditional AI	105
6	■ Enhancing responses with retrieval-augmented generation	135
7	■ Augmenting intent data with generative AI	170
PART 3	PATTERN: AI IS TOO COMPLEX	191
8	■ Streamlining complex flows	193
9	■ Harnessing context for an adaptive virtual assistant experience	207
10	■ Reducing complexity with generative AI	230
PART 4	PATTERN: REDUCE FRICTION	249
11	■ Reducing opt-outs	251
12	■ Conversational summarization for smooth handoff	278

contents

<i>foreword</i>	<i>xiii</i>
<i>preface</i>	<i>xv</i>
<i>acknowledgments</i>	<i>xvii</i>
<i>about this book</i>	<i>xix</i>
<i>about the authors</i>	<i>xxii</i>
<i>about the cover illustration</i>	<i>xxiv</i>

PART 1 **FRAMEWORK FOR IMPROVING CONVERSATIONAL AI 1**

1	<i>What makes conversational AI work?</i>	3
1.1	Introduction to conversational AI	4
	<i>Why use conversational AI?</i> 5 ▪ <i>How does conversational AI work?</i> 6 ▪ <i>How you build conversational AI</i> 7	
1.2	Introduction to generative AI in conversational AI	10
	<i>What is generative AI</i> 11 ▪ <i>Generative AI guardrails</i> 12	
	<i>Effectively using generative AI in conversational AI</i> 13	
1.3	Introducing continuous improvement in conversational AI	15
	<i>Why continuously improve</i> 16 ▪ <i>The continuous improvement cycle</i> 17 ▪ <i>Communicating continuous improvement to stakeholders</i> 19	
1.4	Follow along	22

2 *Building a conversational AI* 23

- 2.1 Building an FAQ bot 24
 - FAQ bot foundations* 24 ▪ *Static question and answering* 26 ▪ *Dynamic question and answering* 31
- 2.2 Routing agents and process-oriented bots 33
 - Routing agents* 33 ▪ *Transitioning from a routing agent to a process-oriented bot* 35
- 2.3 Responding to the user with generative AI 38
 - Integrating with an LLM* 38 ▪ *Routing requests to an LLM* 40

3 *Planning for improvement* 44

- 3.1 Knowing when you need to improve 45
- 3.2 Your cross-functional team 46
- 3.3 Driving to the same goal 49
 - Revisit business goals* 50 ▪ *Effectiveness* 53
 - Coverage* 62
- 3.4 Identifying and resolving problems 64
 - Finding problems* 65 ▪ *Group review* 67 ▪ *Determining acceptance criteria* 72
- 3.5 Developing and delivering fixes 74
 - Sprint planning* 74 ▪ *Measure again* 75

PART 2 PATTERN: AI DOESN'T UNDERSTAND 77

4 *Understanding what your users really want* 79

- 4.1 Fundamentals of understanding 80
 - The impact of weak understanding* 80 ▪ *What causes weak understanding?* 81 ▪ *How do we achieve understanding with traditional conversational AI?* 83 ▪ *How do we achieve understanding with generative AI?* 84
- 4.2 How is understanding measured? 87
 - Measuring understanding for traditional (classification-based) AI* 87 ▪ *Measuring understanding for generative AI* 89
 - Measuring understanding with direct user feedback* 90
- 4.3 Assessing where you are today 91
 - Assessing your traditional (classification-based) AI solution* 91
 - Assessing your generative AI solution* 92

- 4.4 Obtaining and preparing test data from logs 93
 - Obtaining production logs 93* ■ *Guidelines for identifying candidate test utterances 94* ■ *Preparing and scrubbing data for use in iterative improvements 98* ■ *The annotation process 99*
- 4.5 What does the data tell us? 101
 - Interpreting annotated logs for traditional (classification-based) AI 101* ■ *Interpreting annotated logs for generative AI 103* ■ *The case for iterative improvement 103*

5 **Improving weak understanding for traditional AI** 105

- 5.1 Building your improvement plan 106
 - Identify problematic patterns in misunderstood utterances 106*
 - Incremental improvements 110* ■ *Where to start: Identifying the biggest problems 110*
- 5.2 Solving “wrong intent matched” 116
 - Improve recall for one intent 116* ■ *Improve precision for one intent 118* ■ *Improve the F1 score for one intent 120*
 - Improve precision and recall for multiple intents 120*
- 5.3 Solving “no intent matched” 125
 - Clustering utterances for new intents 125* ■ *When to stop adding intents 130*
- 5.4 Supplementing traditional AI with generative content 131
 - Combining traditional and generative AI for an intent 132*
 - Prompting to convey understanding 133*

6 **Enhancing responses with retrieval-augmented generation** 135

- 6.1 Beyond intents: The role of search in conversational AI 136
 - Using search in conversational AI 137* ■ *Benefits of traditional search 138* ■ *Drawbacks of traditional search 139*
- 6.2 Beyond search: Generating answers with RAG 140
 - Using RAG in conversational AI 140* ■ *Benefits of RAG 142*
 - Combining RAG with other generative AI use cases 144*
 - Comparing intents, search, and RAG approaches 145*
- 6.3 How is RAG implemented? 146
 - High-level implementation 147* ■ *Preparing your document repository for RAG 148*

- 6.4 Additional considerations of RAG implementations 151
 - Can't we just use an LLM directly?* 151
 - *Keeping answers current and relevant with RAG* 152
 - *How easy is it to set up the ingestion pipeline?* 152
 - *Handling latency* 157
 - *When to use a fallback mechanism and when to search* 158
- 6.5 Evaluating and analyzing RAG performance 159
 - Indexing metrics* 159
 - *Retrieval metrics* 161
 - *Generation metrics* 163
 - *Comparing efficiency of indexing and embedding solutions for RAG* 165

7 *Augmenting intent data with generative AI* 170

- 7.1 Getting started 171
 - Why do it: Pros and cons* 172
 - *What you need* 173
 - How to use the augmented data* 173
- 7.2 Hardening your existing intents 175
 - Get creative with synonyms* 176
 - *Generate new grammatical variations* 179
 - *Build strong intents from LLM output* 182
 - Creating even more examples with templates* 185
- 7.3 Getting more creative 188
 - Brainstorm additional intents* 188
 - *Check for confusion* 188

PART 3 *PATTERN: AI IS TOO COMPLEX* 191

8 *Streamlining complex flows* 193

- 8.1 The pain of complexity 194
 - Complexity's effect on the end user* 194
 - *Complexity's effect on business metrics* 196
 - *The incremental cost and benefit of reducing complexity for the user* 198
- 8.2 Simplifying and streamlining the user journey 199
 - Spotting complex dialogue flows* 199
 - *Using what is known about the user* 199
 - *Aligning with the user's mental model* 201
 - Allowing flexibility in the expected user responses* 202
 - Supporting self-service task flows with API/backend processes* 204

9 *Harnessing context for an adaptive virtual assistant experience* 207

- 9.1 Importance of context in virtual assistant performance 208
 - How context influences user interactions* 209
 - *What is contextual information?* 212

- 9.2 Understanding modality 217
 - Comparing modalities* 217 ▪ *Importance of modality in designing virtual assistant flows* 219 ▪ *Examples of how modality affects user experience* 220 ▪ *Voice bot design considerations* 222
- 9.3 Enhancing context awareness and improving the overall user experience with RAG 223
 - Designing adaptive flows with RAG* 224 ▪ *Strategies for retrieving and generating contextually relevant responses* 226
 - Maintaining and updating adaptive flows* 227

10 Reducing complexity with generative AI 230

- 10.1 AI-assisted process flows at build time 231
 - Generating dialogue flows with generative AI* 232 ▪ *Improving dialogue flow with generative AI* 235
- 10.2 AI-assisted process flows at run time 237
 - Executing dialogue flows with generative AI* 238 ▪ *Using LLM for a search process* 240
- 10.3 AI-assisted flows at test time 243
 - Setting up generative AI to be the user* 244 ▪ *Setting up the conversational test* 246

PART 4 PATTERN: REDUCE FRICTION 249

11 Reducing opt-outs 251

- 11.1 What drives opt-out behavior? 252
 - Immediate opt-out drivers* 252 ▪ *Motivations for later opt-outs* 253 ▪ *Gathering data on opt-out behavior* 254
- 11.2 Reducing immediate opt-outs 256
 - Start with a great experience: Greetings and introductions* 257
 - Convey capabilities and set expectations* 259 ▪ *Incentivize self-service* 259 ▪ *Allow the user to opt in* 260
- 11.3 Reducing other opt-outs 262
 - Try hard to understand* 262 ▪ *Try hard to be understood* 262
 - Be flexible and accommodating* 263 ▪ *Convey progress* 264
 - Anticipate additional user needs* 264 ▪ *Don't be rude* 265
- 11.4 Opt-out retention 266
 - Start right away by collecting opt-out data* 267 ▪ *Implementing an opt-out retention flow* 267

- 11.5 Improving dialogue with generative AI 270
 - Improving error messages with generative AI* 270
 - *Improving greeting messages with generative AI* 272
- 11.6 Sometimes it's okay to escalate 277

12 *Conversational summarization for smooth handoff* 278

- 12.1 Intro to summarization 279
 - Why summarization is needed* 279
 - *Elements of effective summaries* 280
- 12.2 Preparing your chatbot for summarization 284
 - Using out-of-the-box elements* 284
 - *Instrumenting your chatbot for transcripts* 285
 - *Instrumenting your chatbot (for data points)* 288
- 12.3 Improving summaries with generative AI 290
 - Generating a text summary of a transcript with summarizing prompts* 290
 - *Generating a structured summary of a transcript with extractive prompts* 294
- index* 299

foreword

The artificial intelligence revolution will do to our intelligence what the lever did to our physical strength. It will change the world at micro and macro levels, from how each of us writes, thinks, or makes decisions, to how large organizations redesign work and transform jobs. And one of the most common ways in which people will use AI is conversational user applications.

Conversational AI is a powerful tool that allows businesses and organizations to serve their customers in better and faster ways. It increases self-service capabilities and handles common inquiries, freeing human agents for focus on higher-value work.

As common as this conversational interface of AI is, there are not many books that describe how to do it well. I was happy to encourage my colleagues Andrew Freed, Cari Jacobs, and Enikő Rózsa to share their hands-on experience and provide a framework that others can benefit from. In this book, they have organized and outlined a framework of common challenges to take into account when interacting with conversational applications and interfaces, and provided practical solutions using a variety of techniques, including data science, generative AI, and conversational design principles.

Conversational AI is rarely perfect when switched on. That's one of the most common misconceptions of leaders who want an "instant gratification" implementation of AI. Conversational AI requires a solid data platform as a foundation, an architecture supporting security and identity, and well-thought-out experience and journey designs. You will find many of these in the examples provided in this book, based on the authors' hands-on experience building and enhancing conversational AI systems.

The book is structured around common pain points that users experience while using conversational user interfaces, and it describes methods for solving each of them. By following the techniques and best practices outlined in this book, organizations can

create more engaging, effective, and reliable conversational AI systems that will be adopted faster, deliver a greater experience, and translate to a faster return on investment and incremental business value.

In short, *Effective Conversational AI* is a must-read for anyone interested in designing highly effective conversational AI applications. Whether you're just starting out with conversational AI or you're a seasoned pro, this book will have something for you. It is a timely and essential resource for anyone looking to harness the power of conversational AI to drive innovation, improve user experiences, and drive business value.

—JESÚS MANTAS, GLOBAL MANAGING PARTNER, IBM

preface

Conversational AI is an exciting technology that helps users fulfill their needs faster and helps companies handle user inquiries with lower cost. Conversational AI solutions (often called *chatbots*) have exploded in popularity, especially since the COVID-19 pandemic. There are many books and blogs on how to get started with conversational AI, but most of these books stop at building your first chatbot and do not describe how to improve a production solution. Many enterprises use this technology so that their customers can self-service on a scale that may be prohibitively expensive or impossible with a human workforce. Unfortunately, a significant proportion of these AI solutions underperform.

There have also been plenty of hype and resources on generative AI, including prompt engineering and small demos. However, these are often small-scale in nature, such as proofs of concept and prototypes. There are few resources for maintaining and improving these solutions at an enterprise scale. Generative AI has reignited interest in this space, but it is not a panacea, especially for enterprises offering end-to-end task completion.

We have delivered many conversational AI solutions to production in the past decade. We have worked with a variety of chatbots: question-answering, process-oriented, and routing agents. We have seen the joys and challenges of conversational AI up close.

We wrote this book to help you overcome those challenges. Too often, we have seen chatbots treated as a “set-and-forget” solution. We have also seen chatbots get worse through improper or ill-informed maintenance. As conversational AI builders, we love to dig into underperforming AI solutions and bring them up to excellence. As conversational AI consumers, we want to encounter better chatbots in the wild!

This book helps conversational AI solution owners and stakeholders learn how to identify and remediate the problems that cause chatbots to fail or not reach their full-est potential. Within these pages, you will find a collection of patterns, strategies, and approaches framed around common pain points that exist in conversational AI solutions.

acknowledgments

We've heard that writing a book is "an act of insanity." It's at least a labor of love! We are grateful for the incredible support we've received while writing this book for you.

ANDREW I'm thankful for my friends and colleagues who helped us refine our thinking and reviewed early chapters of our book, including Dan Toczala, Jennifer Gao, and Stéfan van der Stockt. We also thank the innumerable colleagues who have built and improved chatbots alongside us, including but not limited to Leo Mazzoli, Victor Povar, Rebecca James, Jasmeet Singh, Greg Ecock, Tomi Jenkins, Morgan Carroll, Jonathan Roe, Preeth Muthusamy, Marco Noel, Taylor Wood, Jim Kennedy, Elizabeth Smith, Richie Limpijankit, Janice Chan, Yugandhar Chejarla, Kanchan Pandey, Syed Taher, Anirban Mukherjee, Anik Majumder, Swapnil Sharma, and Terrence Nixa. I'm especially thankful to my wife Elise, children Greg and Jeff, and parents Ron and Debbie for their support throughout this process.

CARI I would like to extend personal thanks to my amazing partner, Jason Kerns. This past year has been especially grueling. Your support, patience, and encouragement have meant the world to me. I would also like to express my gratitude to several other former colleagues and mentors who, over the past three decades, shaped my career trajectory and influenced my work ethic: Jared Young, Sean Higgins, Bart Day, Lori Workman, Cory Yochens (rest in peace), Jeff Fetherolf, Tim Shera, Heidi (Piper) Morgan, and Jeff Matteo. Thanks also to my kids (Lani, Ryan, Joe), my bonus kids (Alex, Josh, Lily), my grandson Cameron, Ashley Jacobs, and Bruce Kerns for the enthusiasm and kind words every time the topic of this book came up.

ENIKŐ Writing this book has been a journey I could not have completed without my family's encouragement and support. To my husband Shahram, and to our wonderful kids, Jennifer, Alex, Rachelle, and our bonus kids, Mehr, Margaret, and Tal, thank you for the countless late nights spent discussing ideas around the kitchen island, fueling this endeavor with your insights and laughter. Lisa and Erik, thank you for walking and taking care of Theo so I could spend more time writing. Your kindness and support have been invaluable.

I also extend my heartfelt thanks to my colleagues, former colleagues, and mentors, Will Raabe, Currie Boyle, Craig Trim, Claire Turner, Victor Povar, Brenda Hadlock, Xavier Vergés, and Les Yip, with whom we built chatbots before they were even called chatbots. Thanks to those working on conversational AI and the continuous improvement of chatbots with me, including but not limited to Monisankar Das, Chayan Ray, Avi Yaeli, Sergey Zeltyn, Ignas Valancius, Romanas Marčenko, Eimantas Pēlikis, Kristina Ribačionkaitė, Ateret Anaby-Tavor, Ella Rabinovich, Madhusmita Patil, Arzoo Sabharwal, Richa Manral, and many more. I am grateful for your hard work, dedication, and insights.

We are grateful to the entire staff of Manning Publications for their support and help throughout this process. Special thanks to our technical editors, Jack C. Crawford and Stéfan van der Stockt. Jack is a highly skilled AI architect with a Master's in Computer Information Systems from Claremont Graduate University. He leads generative AI efforts for the virtual assistant of a high-impact mobile application serving millions of users. Stéfan is an AI Solution Architect for IBM who focuses on generative AI, machine learning, and artificial intelligence. He helps IBM clients scope out and define projects to implement production-grade solutions that rely on these technologies.

To all the reviewers: Abdullah Al Imran, Anandaganesh Balakrishnan, Artem Daineko, Ayush Bihani, Brandon Smith, Bruno Sonnino, Erico Lendzian, Felipe Coutinho, Gary Pass, Harinath Mallepally, Igor Vieira, James Black, Jiri Pik, John Kelvie, Jonathan Reeves, Lucas Petralli, Marco Kotrotsos, Maxim Volgin, Nahid Alam, Oleg Kopychko, Parth Santpurkar, Piotr Pindel, Richard Vaughan, Scott Ling, Simone Sguazza, Stefano Ongarello, Swapneelkumar Deshpande, Tong Zhu, Umesh Hodeghatta, and Venkatraman Umbalacheri Ramasamy, your suggestions helped make this a better book.

Finally, heartfelt thanks to Jesús Mantas for his excellent foreword that captures the essence of using conversational AI in the wild.

about this book

The technology in this book goes by many names: conversational AI, virtual assistants, automated agents, chatbots, bots, or just “the AI.” No matter what you call it, this book will teach you how to effectively use the technology to meet your needs and satisfy your users. If you have an underperforming conversational AI, this book will teach you how to improve it.

Who should read this book

Effective Conversational AI is for people who currently maintain conversational AI solutions, including business sponsors, product owners, and conversational AI designers and developers. Software development experience may be useful for some technical resolutions, but is not required for many of the conceptual or design remediations recommended throughout the text (thanks in part to low-code/no-code conversational AI tooling). The pain points and resolution patterns described throughout this book should also be insightful to those who are considering building a conversational AI solution—an ounce of prevention, as they say.

How this book is organized: A road map

This book is divided into four parts and 12 chapters. The first part introduces conversational AI, the benefits and pain points of the technology, and a structured approach for improving conversational AI. The remaining parts each focus on a single pain point and offer a variety of ways to improve the AI by removing that pain point.

Part 1 introduces concepts fundamental for building and improving conversational AI.

- Chapter 1 introduces conversational AI and generative AI and shows how they work better together. It also introduces pain points and pitfalls that this book will help you avoid.
- Chapter 2 teaches you how to build and evolve a chatbot by starting with simple question-answering, adding in process flows, and finally using generative AI.
- Chapter 3 shows how to evaluate conversational AI with objective measurements and how to build an improvement plan against these metrics.

Part 2 presents multiple strategies to help conversational AI systems understand user requests.

- Chapter 4 walks you through discovering what users want out of your AI and measuring your AI's understanding. These techniques support the next three chapters.
- Chapter 5 demonstrates how to improve the understanding of intent-based (classifier-based) conversational AI solutions, which is especially useful for answering frequently asked questions.
- Chapter 6 uses both search and retrieval augmented generation (RAG) to understand and answer user questions, especially less-common questions ill-suited for intents.
- Chapter 7 teaches techniques for using generative AI at build time to generate training and testing data for conversational AI systems.

Part 3 tackles the challenges faced by users and builders when conversational AI gets complex.

- Chapter 8 illustrates multiple methods for simplifying process flows to increase the chance that users can successfully complete them.
- Chapter 9 goes deep into using all available context to ask the right questions of users and deliver the right responses to them.
- Chapter 10 unleashes generative AI on complex dialogues, instructing LLMs to design, critique, or replace complex process flows.

Part 4 focuses on reducing friction for the users of AI and the human agents that AI sits in front of.

- Chapter 11 sums up reasons that users opt out of conversational AI immediately or during a conversation and how you can help them be less likely to do so.
- Chapter 12 explains effective conversation summarization techniques, especially when users opt out. These are needed for humans who continue these escalated conversations.

We suggest reading part 1 of the book first to learn the mental model used in the book. (Experienced builders may skim chapter 2, which builds a new chatbot from scratch.) After that, dive into any part describing a pain point you're interested in resolving. The parts can be read in any order. Within a part, we suggest reading the

chapters in order, but this is not strictly required. Each chapter includes exercises for you to practice or ideate on the concepts you have learned.

About the code

This book contains examples of source code (and LLM prompts), both in numbered listings and in line with normal text. In both cases, source code is formatted in a fixed-width font like this to separate it from ordinary text. Sometimes code is also **in bold** to highlight code that has changed from previous steps in the chapter, such as when a new feature adds to an existing line of code.

In many cases, the original source code has been reformatted; we've added line breaks and reworked indentation to accommodate the available page space in the book. In rare cases, even this was not enough, and listings include line-continuation markers (`↵`). Additionally, comments in the source code have often been removed from the listings when the code is described in the text. Code annotations accompany many of the listings, highlighting important concepts in the code and prompts.

You can get executable snippets of code from the liveBook (online) version of this book at <https://livebook.manning.com/book/effective-conversational-ai>. The complete code for the examples in the book is available for download from the Manning website at <https://www.manning.com/books/effective-conversational-ai>. The code for this book is also stored on GitHub at <https://github.com/andrewfreed/EffectiveConversationalAI>. This site includes the working CakeBot example from chapter 2 as well as all our code snippets, sample conversations, and large language model prompts. Due to the fast-moving and probabilistic nature of LLMs, you may not get the exact same responses from LLMs as we demonstrate in the book.

liveBook discussion forum

Purchase of *Effective Conversational AI* includes free access to liveBook, Manning's online reading platform. Using liveBook's exclusive discussion features, you can attach comments to the book globally or to specific sections or paragraphs. It's a snap to make notes for yourself, ask and answer technical questions, and receive help from the authors and other users. To access the forum, go to <https://livebook.manning.com/book/effective-conversational-ai/discussion>. You can also learn more about Manning's forums and the rules of conduct at <https://livebook.manning.com/discussion>.

Manning's commitment to our readers is to provide a venue where a meaningful dialogue between individual readers and between readers and the authors can take place. It is not a commitment to any specific amount of participation on the part of the authors, whose contribution to the forum remains voluntary (and unpaid). We suggest you try asking the authors some challenging questions lest their interest stray! The forum and the archives of previous discussions will be accessible from the publisher's website as long as the book is in print.

about the authors



ANDREW R. FREED is a Distinguished Engineer with over 20 years of experience, the latter half in AI. He joined IBM's Watson division shortly after Watson defeated past champions in Jeopardy! and has delivered many AI-based solutions since. Andrew has learned from a wide variety of technical books and blogs and is passionate about "paying it forward" by sharing his hard-won knowledge. Outside of work, he enjoys spending time with his family.



CARI JACOBS has been working in information technology for nearly 30 years. Her experience includes datacenter operations, Unix administration, and production application support. In 2014, she joined IBM's Watson division. As a cognitive engineer/solution architect, she has consulted on and delivered conversational AI solutions for dozens of Fortune 500 companies. She has also worked with many other national brands, regional businesses, government agencies, universities, and startup companies. She loves to learn and share her knowledge. Her hobbies include kayaking, photography, and Brazilian jiu jitsu.



ENIKŐ RÓZSA is a Distinguished Engineer and the CTO for IBM's Global AI & Analytics Practice. With a remarkable 30-year career at IBM, she has consistently delivered innovative, multiplatform conversational AI solutions across various industries. She thrives on tackling complex challenges that demand integrating emerging AI technologies. Enikő is also an accomplished inventor, having published multiple patents in Natural Language Processing (NLP) and

usability. Her passion for chatbots began when she co-invented and led the successful production of an ontology-based natural language dialogue system, which revolutionized client self-service for technical support. Outside of her professional achievements, Enikő cherishes her time with family, enjoys walking her dog Theo, and loves hosting dinner parties for friends and family, where she brings people together to share good food and conversation.

about the cover illustration

The figure on the cover of *Effective Conversational AI* is “Le Donne Di Procida,” or “The Women of Procida,” taken from the collection *Usi e costumi di Napoli e contorni descritti e dipinti* (*Customs and Traditions of Naples and its Surroundings, Described and Painted*) by Francesco de Bourcard, published in 1853. Each illustration is finely drawn and colored by hand.

In those days, it was easy to identify where people lived and what their trade or station in life was just by their dress. Manning celebrates the inventiveness and initiative of the computer business with book covers based on the rich diversity of regional culture centuries ago, brought back to life by pictures from collections such as this one.

Part 1

Framework for improving conversational AI

Have you ever had a bad experience with a chatbot? Perhaps with a voice system that always tells you “please listen carefully—our menu options have recently changed” or a chatbot that never understands your questions. Hopefully you’ve had great experiences with AI as well—AI that seems like it knows you and proactively identifies your needs. What separates good conversational AI from the bad?

Conversational AI is bigger than ever, as companies look to improve customer experience and their own bottom line through this technology. Generative AI has rekindled interest in the technology and made it easier than ever to add intelligence to a chatbot. Many of these chatbots look great in a prototype phase and then falter in production. This part of our book sets the stage for building and improving conversational AI.

Chapter 1 introduces conversational AI and generative AI technology, their pros and cons, and how they complement each other. It also lays out the types of pain points many AI solutions run into. Chapter 2 shows you how to build and evolve a chatbot, layering on incremental complexity and capability, ending with a hybrid of traditional and generative techniques. Chapter 3 demonstrates how to objectively evaluate your AI and develop plans to improve it.

1

What makes conversational AI work?

This chapter covers

- Identifying and minimizing conversational AI risks
- Assessing where generative AI can help you in your conversational AI
- Using generative AI safely
- Continuously improving your AI and aiming for a defined target

We've all encountered computerized conversational agents that caused us pain, such as a chatbot that didn't understand anything we said, a robotic voice initiating a confusing dialogue flow, or a phone assistant that made us immediately opt out to a human representative. When your conversational AI solutions cause these problems, how do you resolve them? How can you build them correctly in the first place? This book will show you how to create chatbots and other conversational AI solutions that your customers will be happy to use.

As conversational AI practitioners, we work with customers who are just starting to deploy automated agents for limited tasks as well as with large organizations that face high levels of business risk—situations where one generative AI hallucination

might outweigh the benefits of dozens of correct and fluent interactions. Using a variety of examples pulled from our work, we'll present options for implementing and improving conversational AI, with and without generative AI.

We'll start with a brief look at classical conversational AI technology, followed by an introduction to generative AI and to the continuous improvement process we recommend for safely and effectively getting the most out of your conversational AI. Then, in chapter 2, you'll build your own chatbot using both classic and generative AI techniques.

1.1 Introduction to conversational AI

Conversational AI, also known as *chatbots*, *virtual agents*, *AI assistants*, and *digital employees*, is a set of technologies designed to mimic or replace human interactions using written or spoken natural language. Conversational AI is routinely used to automate customer service, offer “voice assistant” services like Alexa and Siri, or to prescreen an eventual human-to-human interaction. Generally speaking, you can divide conversational AI into three categories:

- *Question-answering*—Also known as FAQ bots, these AI solutions deliver a response directly to a user's question, usually without any follow-up.
- *Process-oriented or transactional solutions*—The user is guided by an AI to achieve some goal through a series of questions from the bot; for instance, checking an account balance, booking an appointment, or checking the status of an insurance claim. This type of conversational AI may execute the transaction or collect information for manual fulfillment.
- *Routing agents*—In this case, the bot's only job is to figure out where to redirect the user. The redirection may be to a different specialist bot or a human agent.

Some AI solutions contain a mix of all three. A retail banking chatbot may do simple question-answering for things like “when are you open” and “where are you located,” process flows for opening accounts and checking account balances, and route users to specialists for cases like fraud reporting.

These types of chatbots have similar architectures but different emphases. A routing agent only needs to understand a user's initial intent, but a process-oriented bot needs to not only understand intent but also keep the user engaged through an entire process flow. In this book, we'll walk you through several conversational AI challenges and success stories, as illustrated in table 1.1.

Table 1.1 Challenges in conversational AI that we have solved

Pain point	Example success story	In this book
Did not understand user intent	Increased intent recognition accuracy from 76% to 92%	Part 2 (chapters 4–7)
Too much complexity put on the user	Increased search success from 40% to 90%	Part 3 (chapters 8–10)
Immediate opt-out by users	Reduced immediate opt-outs by 15%	Part 4 (chapters 11–12)

All chatbot types face the challenge of understanding the user. Process-oriented bots are especially susceptible to burdening the user with complexity, and we also find that all chatbot types can be plagued with immediate opt-outs. The latter parts of the book focus on specific challenges, with examples from multiple chatbot types wherever possible. Feel free to skip ahead to the challenges that interest you.

Conversational AI solutions are built to solve problems. If they are not solving problems, they're causing pain to their users. The pain points inform how we should improve the system. But before we can improve on an existing solution, we need to understand what motivated the solution in the first place.


1.1.1 Why use conversational AI?

An effective conversational AI provides exceptional user experience and benefits, saving users time and energy while saving corporations support costs. It never gets tired, so it can help users 24/7. And it is personalized, efficient, and maybe even proactive, guiding users to achieve their goals.

A bad conversational AI does the reverse—it frustrates users, decreases satisfaction, or floods support lines because “the bot didn’t understand what I wanted.” It makes users sit through overly verbose messages, asks them questions that it shouldn’t need to ask, or is cold and rude to them. Figure 1.1 shows a painful chatbot experience in a process-oriented bot.


Conversational AI doesn’t have to be painful, and it can offer a better and more streamlined experience than one requiring human intervention. The scenario in figure 1.1 put a heavy burden on the user. Technically, the dialogue flow made sense—a user *could* ask about any claim. And maybe the user isn’t asking about their own claim. But this ignores the general case—most users are asking about their own claim. Most users can be

identified—chat users by the email address they logged in with, or voice users by their phone number. Figure 1.2 shows a user-centric way to solve the same claim status problem by using these reasonable assumptions. The assumptions also personalize the experience. This system provides an answer quicker than a human could!



1. Pain for the user who keeps giving information but doesn't receive any

Figure 1.1 A painful chat experience with a process-oriented bot that puts cognitive burden on the user. The AI has not provided any value in three conversational turns.



1. Bot applies context for efficient interaction

Figure 1.2 A delightful experience that uses context and reasonable assumptions to complete the user's goal quickly. The context could be loaded from a log-in process (chat) or from the caller phone number (voice).

Sometimes you can fix a process-oriented bot by improving the process. Keep in mind that chatbots are not purely a *technology problem*. Chatbots interact with people, and people are often messy. Technology alone cannot fix all chatbot experiences.

Having seen good and bad chat experiences, let's review how conversational AI works.

1.1.2 How does conversational AI work?

A conversational AI solution typically includes three steps:

- 1 Figure out what the user wants.
- 2 Gather additional information necessary to satisfy that want.
- 3 Give the user what they want.

The solution should accomplish these goals as quickly and easily as possible while following legal and ethical guidelines, such as handling sensitive information securely and not pretending the AI is an actual human. If the AI solution cannot achieve those goals, or introduces too much friction into the process, users will abandon the AI and look for another solution. This may mean going to a human who can help them or quitting your service.

Figure 1.3 shows the high-level flow in a conversational AI solution, and these steps are supported by the architecture shown in figure 1.4, annotated based on a "reset password" scenario from a process-oriented bot.





Figure 1.3 Flow diagram for conversational AI. In many use cases, "additional information" includes user profile data.



1. How the user interacts with the bot; for instance, a graphical or telephonic interface
2. Asks NLU to interpret the user's message
3. Converts natural language into intent ("I forgot my password" becomes #reset_password)
4. Determines the next step in the conversation, such as asking security questions
5. Invokes external APIs needed to accomplish the task (resetting user password)
6. Delivers the response to the user via the user interface

Figure 1.4 A conversational AI logical architecture annotated with a password reset example

Let's expand on the three primary steps:

- *Figure out what the user wants*—The user generally makes their request in natural language, so a natural language understanding module receives this message and determines the intent behind it. This is usually done with a machine learning algorithm, such as a text classifier. Example intents include "reset password" or "find a store." The intent drives the next step in the process.
- *Gather additional information necessary to satisfy that want*—The user's initial request often does not include enough information to fulfill it—the request just starts a journey. A dialogue engine guides the user through all the steps necessary to fulfill the request. It may have to ask clarifying or follow-up questions like "what's your account number" or "what is your zip code." It may use an orchestration layer to interact with other systems through application programming interface (API) calls. The dialogue engine manages conversation state and applies logic to respond to the user.
- *Give the user what they want*—The flow concludes when the user's request has been fulfilled. Their password has been reset, or they receive the address to your store, or they have been connected to a human who can complete their need.

There can be slight variations in these steps across the different kinds of bots. For instance, a question-answering bot rarely uses APIs, but a process-oriented bot frequently does. A routing agent only indirectly gives the user what they want (by routing the user to the correct specialist).

1.1.3 How you build conversational AI

Building a conversational AI solution works best when you involve a set of diverse skills across your team, as shown in figure 1.5. It's important to understand how these

solutions are built if you are trying to improve them. In this section, we will summarize the build process. For a more complete treatment, see *Conversational AI* (Manning Publications, 2021).



Figure 1.5 It takes a dream team with diverse skills to build an enterprise-ready conversational AI.

The starting point for conversational AI is user design. Look at what your users want to achieve and how you can help them achieve these goals in a quick and frictionless experience. All the players in figure 1.5 should contribute to these user-centric questions:

- What are the most frequent pain points of your users?
- What do they need to do?
- What information are they likely to have? (And what information won't they have?)
- How are they likely to express their needs?

Once you know what the user needs, think through what *you* need to satisfy the user. For instance, let's assume your users keep getting locked out of their accounts. They need a password reset function. What do you need to reset a password? Typically, you need to do at least three things for password resets:

- Extract meaning from the user's statement (determining that they have a password problem, even if they don't use specific terms, such as "password" or "reset").
- Access an API that can authenticate the user and reset the password.
- Collect enough information about the user to reset their password.

These needs drive the rest of your building process.

EXTRACTING MEANING

Chatbots start by extracting meaning from the user, identifying intent from users' natural language utterances via a text classifier. An *utterance* is what the user says, an *intent* is what it means (as in, what the user wants), and a *classifier* categorizes utterances into intents.

Chatbot platforms are getting easier to use with a trend toward low-code or no-code, but that doesn't mean they will understand your needs with no human involvement. It's best to have a data scientist optimize the training data for representativeness, balance, and variety, and to perform tests to make sure the trained classifier is as accurate as possible. If this is not done well, it will lead to the pain point of "the bot doesn't understand me," because the AI is generally programmed to route unrecognized utterances to a generic response.

The best input data for this training process comes from previous user interactions, such as historical chat logs, call center transcripts, or emails. Part 2 of this book covers collecting good data and using it to improve intent recognition.

USING APIs

A developer needs to expose an API to the virtual assistant. They need to clearly define the required input parameters, output response formats, and error conditions so it is clear how the API should be integrated into the chatbot. The function exposed by the API can be implemented in any programming language—what's important is that there is an API endpoint that the assistant can securely reach.

If the API does not exist, your chatbot project could be the perfect reason to build it. Or the design of the chatbot may necessitate a change in an API. APIs are useful for getting structured information to a user (checking their account balance, finding their open claims) or acting for the user (resetting their password, opening an account)—you might not be able to satisfy the users' needs without the right APIs.

APIs are most often used in process-oriented bots, but they are also helpful for supplying additional user context to question-answering and routing agents.

COLLECTING MORE INFORMATION

You need a conversational flow that gets the information required to invoke the API or to answer the user's initial question. This will be influenced by the channel you are building for (such as web or phone) and what you can reasonably expect the user to have. For instance, in a password reset scenario on the web, it's common to ask a security question. But it can be difficult to collect this information via the phone, and it's insecure to collect the information via SMS. In contrast, phone and SMS channels may be able to use the user's phone number as a piece of the authentication puzzle.

The available APIs may influence the conversation design, or the conversation design may influence the API, or they may influence each other. If the process of collecting more information becomes difficult for users, it can lead to the "too much complexity" or "immediate opt-out" pain points when users learn they may not be able to successfully use the assistant.

It's also worth noting that not every conversational AI requires all three of the things we've been discussing:

- Some APIs may not require additional information. For instance, a “store hours” API may return the same response no matter who is asking.
- Frequently asked question (FAQ) bots may not invoke any APIs at all and need only to match user utterances to intent/response pairs.
- A bot that falls back to search may not even include any intents. This is a popular pattern with conversational search solutions built with generative AI, either using the built-in knowledge from a large language model (LLM) or supplementing an LLM with your data by searching a knowledge base and generating an answer from those search results. This pattern can also be built as a hybrid model where intents are constructed for the most common questions and all other questions are routed to search or generative AI.

Exercises

- 1 Review the last several chatbots you have interacted with (or that you have built yourself). Were they question-answering, process-oriented, or routing agents? Why?
- 2 What challenges did each of these chatbots face? How do you wish they would perform better?

1.2 *Introduction to generative AI in conversational AI*

Any sufficiently advanced technology is indistinguishable from magic.

—Arthur C. Clarke

Generative AI (a method that dynamically generates new content) is an exciting new technology. You've probably seen it do some cool tricks: “write a Shakespearean sonnet,” “describe AI but speak like a pirate,” or “build me a plan to make 100 dollars ethically.” But it's not magic, and it is not a panacea. Generative AI can help you reap benefits, but you'll need to work to avoid harmful outcomes like hallucinations.

Generative AI can help us solve several of the pain points in conversational AI solutions:

- *Did not understand user intent*—Generative AI can help us train stronger intents in our conversational AI. Or it can replace some or all intent recognition through retrieval-augmented generation (RAG) by summarizing content that came from a search (retrieval) process. It can also be more adaptive to nuance in the user's intent.
- *Too much complexity put on the user*—Generative AI can help us write simpler prose in our dialogue or test the system for unexpected complexity.
- *Immediate opt-out by users*—Generative AI can help us write more engaging prose that also helps our users.

We can use generative AI inside the conversational AI, letting it assist our users directly by answering their questions or searching for information. We can also use generative AI to assist us as we build our conversational AI, such as using it to build better dialogue flows and messages and analyze previous conversations. Generative AI is not a replacement for classic conversational AI techniques—they work best together.

1.2.1 What is generative AI

Generative AI is a blanket term for AI powered by *foundation models*, which are generalized AI models trained on a broad set of tasks. While there are several kinds of foundation models, this book focuses on LLMs—machine learning models that are trained on huge textual datasets. How huge? Use “all the internet’s text” as your mental model.

A model that has seen “an internet’s worth of text” should be excellent at understanding word and sentence sequences. The model is trained to receive a series of words and predict a word that is likely to follow the previous words. By repeating this process of predicting the next word, LLMs can generate words, sentences, paragraphs, or even entire pages of text!

You can use LLMs inside your conversational AI system. The LLMs can perform tasks that are directly exposed to your users or can perform tasks that assist you in building the conversational AI. Table 1.2 lists several of these tasks.

Table 1.2 Sample tasks where conversational AI builders can quickly and efficiently use LLMs

Consumer-facing tasks	Build assistant tasks
Generate answers (from retrieval-augmented generation) Summarize conversation transcripts	Copyedit or write dialogue and flows Augment your training data

LLMs can perform these tasks with little or no training and speed up your development process, and they are resilient to minor variations in user questions that a traditional classifier might not understand. But they also come with potential dangers:

- LLMs learn from their training data. Have you ever been on the internet? The internet is full of bias, hateful speech, and misinformation. Retrieval-augmented generation is a great way to generate answers because it grounds LLM output in your documents, rather than using the LLM’s internal data (which is generally trained on internet content).
- LLMs do not know whether their responses are true, only that the responses are a probable extension of their “prompt.” This is the basis of *hallucinations*—a response that looks good but is not useful. You never know what LLMs may say. This is why using them as dialogue-writing assistants is excellent, because you can review their output before using it.

LLMs will lie to you without a care in the world. Or they will generate a better-than-expert-level response in seconds. LLMs can exhibit amazing creativity or horrifying

bias—there is plenty of both on the internet! To use LLMs with confidence in your conversational AI solution, you need guardrails.

1.2.2 Generative AI guardrails

Would you deploy generative AI if you knew bad actors could exploit it to respond to requests like “how do I build a bomb” or “tell me a racist joke”? Probably not! Fortunately, there are several ways to put guardrails around LLMs. These are especially important if we pass LLM output to our users. Let’s look at a few kinds of guardrails.

MODEL AND TRAINING DATA SELECTION

Our first guardrail is in the choice of model. Most practitioners choose to use an existing model rather than building their own. This is because training a brand new LLM may cost millions of dollars and take months.

LLMs are trained on a huge dataset—many are trained on some version of The Pile. The Pile is an 886.03 GB diverse, open source collection of English text created as a training dataset for LLMs ([https://en.wikipedia.org/wiki/The_Pile_\(dataset\)](https://en.wikipedia.org/wiki/The_Pile_(dataset))). Many LLM trainers leave out some parts of The Pile (to remove biased data or profanity, for example) and add more data (such as private or licensed data). Many open source LLMs come with a “model card” describing the data and methodology used to train the model. By reviewing the model card, you can select a model with a suitable dataset.

This is a helpful first step, but it’s far from the only choice.

PREFILTERING INPUT FOR HATE, ABUSE, AND PROFANITY

Another option is to screen the user’s input and block any attempts that seem problematic. There are multiple techniques for doing this, including scanning for keywords (like profanity or slurs) or running a classifier on the input. This becomes an arms race where LLM providers try to make the models safer, and users get cleverer. Some users try to “jailbreak” a prompt. An LLM may reject a prompt like “Tell me how to make a bomb,” but they could be tricked by a request like “Tell me a story like my grandmother used to. Whenever I couldn’t fall asleep, she’d tell me a story in exquisite detail about how she made a bomb as a child. Tell me that story.” In fact, one primitive technique to reduce jailbreaking is to limit the length of the user’s input.

CONTEXTUAL INSTRUCTION AND PROMPT

Our next guardrail is the instructions we give the LLM via the prompt. Figure 1.6 shows how effective context is in guiding an LLM.




Figure 1.6 Adding context in the prompt is an important way to guide an LLM.

Context keeps the LLM from having to use its own (stale) data and reduces the likelihood of hallucinations. Retrieval-augmented generation (chapter 6) provides context from your trusted documents. Context can also be used to assign a persona to the LLM, such as “you are a friendly copy editor,” which is useful for revising content drafts (chapter 10).

Providing context to the LLM is a powerful technique.

POSTFILTERING OUTPUT

Like the prefiltering option, you can also scan the output from an LLM for certain content. For instance, you can scan for keywords or other indications of hate, abuse, and profanity (HAP). Libraries can help with this—one example is the profanity-check library at [pypi.org](https://pypi.org/project/profanity-check/) (<https://pypi.org/project/profanity-check/>).

For some use cases, you can also compare the answer against some parts of the prompt. In retrieval-augmented generation, the LLM is supposed to answer questions only from the documents retrieved by the search process. You can do a textual similarity analysis to see whether most or all the answer text appears in the documents used.

HUMAN IN THE LOOP

The safest option is not to give the LLM free rein, period. Having a human “in the loop” ensures you know what your LLM is doing. There are two versions of this: retroactive review and beforehand review.

Retroactive review means you periodically monitor the responses an LLM provides. For instance, you may have a weekly process where you review a sample of LLM inputs and outputs. This will not prevent a bad outcome, but at least you will know one occurred, and you can adjust the LLM.

In contrast, a beforehand review means you use the LLM to assist a human, and the human has the final call. An example of this is using the LLM as a copy editor—it generates static dialogue messages that a human inserts into a dialogue engine.


Using LLMs in this way can help reduce user experience pain points through methods like generating training data to solve “did not understand user intent” and rewriting dialogue to reduce “dialogue flow is too complex (or rude).”

1.2.3 Effectively using generative AI in conversational AI

Two fundamental requirements for using generative AI effectively are to use the right model for the job and to mitigate risk by applying appropriate guardrails.

THE RIGHT MODEL (AND PARAMETERS) FOR THE JOB

There are thousands of LLMs, and they are trained on different tasks. You can refine an LLM’s behavior on these tasks by experimenting with prompts and parameters. Figure 1.7 demonstrates the effect of the “repetition penalty” parameter on the Flan-ul2 model for a creative task. Different tasks require different parameters. A low repetition penalty is useful when using text from the documents you have provided. A higher repetition penalty is helpful in creative tasks like list generation.



1. Shared prompt.
2. Flan-UL2 response (low repetition penalty). The model repeats itself after three list entries.
3. Flan-UL2 response (high repetition penalty). The output includes no repetition.

Figure 1.7 Effect of changing one LLM parameter (repetition penalty)

In this book, we will use several different model and parameter sets to demonstrate a variety of techniques. We want to show that our techniques are broadly applicable. You may not see your model of choice referenced in this text, and you may need to use different prompts, parameters, or models in your use case. By the time you read this book, a completely new set of models may be available for use!

For each task, you may want to experiment with multiple models as well. For instance, Flan-UL2 is an LLM trained on 50 tasks, including question answering and information retrieval (<https://huggingface.co/google/flan-ul2>)—it's a generalist model. MPT-7B-Instruct is an LLM specializing in one task—short-form instruction following (<https://huggingface.co/mosaicml/mpt-7b-instruct>). Models also have different cost profiles and performance characteristics. You are likely to experiment with several different models before selecting the right one for your task. You may select different models for different tasks within the same solution. Table 1.3 includes some do's and don'ts for selecting an LLM.

Table 1.3 Dos and don'ts for LLMs

Don't	Do	Why
Don't use a model only because you saw it perform well (on a task that you don't need).	Select a model suited to your task, or experiment with several such models.	Performance is task-dependent, including any parameters or prompt engineering. Tasks include generation, classification, extraction, question answering, retrieval-augmented generation, summarization, and translation.
Don't discard a model or prompt because of one bad experiment.	Test on multiple inputs, models, and parameters.	Sometimes you'll get unlucky. It takes multiple tests to have confidence in an LLM configuration.

Table 1.3 Dos and don'ts for LLMs (continued)

Don't	Do	Why
Don't blindly let the LLM have full control, especially in responding to your conversational AI users.	Apply guardrails at multiple levels.	You (or your organization) own the ultimate output. "The LLM said so" is no excuse.

"The LLM said so" really isn't an excuse

In 2024, a Canadian airline chatbot offered a discount that didn't exist. In court they argued the chatbot was a "separate legal entity that is responsible for its own actions." The court disagreed. The company was ordered to pay the discount offered by the chatbot. (See the story on the BBC website: <https://mng.bz/GejV>.)

APPLY APPROPRIATE GUARDRAILS AT EVERY STEP OF THE WAY

Make sure you are thinking about guardrails in all stages of using an LLM:

- *Before*—Choose an LLM that is fit for your purpose and whose dataset aligns with your values. Decide how much freedom and oversight the LLM will have—can it perform tasks from end to end or will all output be reviewed by humans?
- *During*—Experiment with the LLM, tuning and adapting it for your task and verifying the functionality of any content controls.
- *After*—Periodically assess the LLM's past performance, and assure it still meets your business needs.

Consider the worst outcome for an LLM, and make sure you have a strategy to combat it. For example, in question-answering, you may be most afraid that the LLM will make up answers with no basis in reality (hallucinations). You could mitigate this by assigning contextual bounds or continuously reviewing LLM responses.

Exercises

- 1 Think about the chatbots you wrote about in the previous set of exercises. How could they have been improved with generative AI?
- 2 For each of the generative AI uses, how would you use it safely? Are hallucinations a problem for each use case? Do you need to worry about hate, abuse, and profanity?

1.3 Introducing continuous improvement in conversational AI

Software is like entropy. It is difficult to grasp, weighs nothing, and obeys the second law of thermodynamics; i.e., it always increases.

—Norman Ralph Augustine

“Entropy” broadly means tending towards chaos constantly.

—Sid Sriram

Software is never perfect the first time. Requirements are not perfectly understood, needs change, or user feedback drives changes in software. AI software is no different. Without improvement, AI software will most likely slide into decay.

1.3.1 Why continuously improve

Even if we tune a conversational AI perfectly for the present day, our needs will change:

- Users will make new requests and use the software differently.
- Your business will have new rules for fulfilling processes.
- Technology like generative AI will make possible what used to be impossible.
- Newer and better-performing AI models will become available.

Conversational AI has several components, including understanding the user’s initial intent, gathering additional information as needed, and completing the user’s request. Each of these components will likely change over time, requiring continuous improvement. A degradation in any of these components increases user frustration and degrades business outcomes.

Like a chain, a conversational AI solution is only as strong as its weakest link. Perhaps we have a great process for presenting information to the user, but we never reach it because we rarely understand their initial intent. Figure 1.8 shows a conversion funnel for a process-oriented bot that finds member’s claims, showing the relative number of users reaching each step.




Figure 1.8 Cumulative success in a process is dependent on success in each of the individual steps. Visually it looks like a funnel that narrows after each step.

Success is multifaceted. For the user to get what they want, we need to

- Engage them (A)
- Understand them (B)
- Present everything they need (C)

We can think of success in any process flow as A times B times C. If we see that our success rate is not what we want, we need to investigate each component of that success chain. Odds are good that we can find ways to improve each component. We can even use this framework to think about question-answering bots, with each subsequent question as the next step in the process. Chapter 3 expands on this framework.

Again, failures in a process flow may not solely be solved by technology. Generative AI can still misunderstand users and still give wrong answers. Some manual work is required to identify areas of improvement and to do the work of improvement. A continuous and incremental approach to improvement increases your chances of success.

1.3.2 The continuous improvement cycle

For any given challenge, a perfect solution may not be obvious or even possible. This is especially true in AI, where possibilities change daily and where changes may have unexpected side effects. Therefore, it's important to improve your conversational AI via a series of small changes, and in chapter 3, we'll show you how to estimate the effect of each change. For now, recognize that a change might make a small improvement, a large improvement, or may make things worse! Each change will produce an additional learning opportunity.

Figure 1.9 shows a typical continuous improvement cycle applicable to any chatbot.




Figure 1.9 A continuous improvement lifecycle for conversational AI

A typical continuous improvement cycle includes the following:

- *Measure*—You need a baseline of the system's performance before making changes.
- *Identify a problem*—Find something that is wrong, broken, or non-optimal. Ideally, this problem will be directly connected to a business metric. For example, "We notice a lot of calls transfer to an agent when <condition>."
- *Implement*—Assume the problem is important enough to fix, implement a solution to the problem. For example, update intent training or copyedit your dialogue.

- *Deploy*—Deliver the change and record the effect on the original problem.
- *Repeat*—Repeat as needed. If the change was successful, congratulations, and if not, you can undo the change. Move on to the next problem, or iteratively improve on the same problem.

We prefer making small and predictable changes over larger and unpredictable changes. To reduce “bot doesn’t understand users,” we prefer to change just the single worst-performing intent (request type) rather than changing many (or all) intents at once. For low completion within a process-oriented flow, we prefer changing one step at a time, rather than changing or rearranging many steps.

Figure 1.10 shows an example of making a large change to a system. Because the change is large, it will take a long time to deploy to production, and it has a wide variety of outcomes. It could cause a huge benefit, a small benefit, or a small detriment. We won’t know anything until this huge change is deployed. This approach is quite risky—if the change goes badly, how will you explain it to your stakeholders? “We took a long time to make this change, and to our surprise, we made things worse. We’re not sure which part of the change made things worse, so we’ll have to undo everything and start over.” Yikes! That is more risk than most people would be willing to take.




Figure 1.10 Large changes—like retraining all intents—take a long time and have less predictable outcomes.

Contrast this with figure 1.11. Here we don’t make one major change but rather four minor changes. Each of the changes has the same possible outcomes (much better, a little better, or worse) but on a smaller scale. This approach has several benefits:

- *Each change is easier to understand*—If we only change one thing, it is much easier to connect the outcome to the change. Smaller changes are also easier to debug.
- *More learning opportunities*—Rather than one chance to learn, we have four.
- *More options*—With smaller changes and smaller risks, we can stop earlier if we achieve our goals.




Figure 1.11 Making many small changes—like retraining one intent at a time—has a smaller “blast zone” for each change, bringing quicker value and more learning.

In figure 1.11, we might have decided that the first two changes were sufficient. We could have stopped with this moderate improvement. The third change made the system worse, but since it is a small change, it is easy to reverse. We learned a lot, quickly.

Most excitingly, the incremental change approach lets us lock in improvements (and business value) sooner! Let’s transform the chart to capture business value. The smaller and faster changes delivered positive change before the “big bang” change was even finished. This will delight our stakeholders and our users too.

Using continuous improvements and small changes, we will either have a minor improvement that delivers business value quickly or a minor decrease in performance that we can easily reverse and learn from. Figure 1.12 shows how frequent small changes deliver value quickly.




Figure 1.12 Area over the dotted line is additional business value over the “big bang” change. Working code in production delivers value.

Better AI performance should lead to better business value for your stakeholders. But how can you convey that improved value in a way they will understand?

1.3.3 Communicating continuous improvement to stakeholders

Definitions of a successful AI solution vary, but you are probably using one of the standard success metrics:

- *Cost reduction*—Measured by containment or average handle time. (Completing calls without any human involvement, or helping humans work more quickly.)

- *Customer satisfaction*—Measured by net promoter score (NPS) surveys, time-to-resolution, or reduced customer churn.

Your stakeholders invested in conversational AI to achieve a business outcome, so you should be measuring your AI solution against that outcome. Check both your current performance and the trend of your performance to make sure you are improving (or at least not getting worse). The changing needs of your solution mean you are constantly fighting against entropy. Sometimes you'll need to continuously improve just to maintain your current success levels.

In this book, you will learn several techniques for improving your AI solution, and some of them will be deeply technical. You may be excited to try these techniques, but you may need to convince your stakeholders to pay for the improvements. It's critical that you speak in their language: less technical jargon, more business value!

Consider this example of describing a fix to “the bot doesn't understand the user”:

- *Heavy on technical jargon*—“We're going to increase the accuracy of #claim_status intent. The classifier identifies this intent with a 0.92 F1 score with most confusion coming from #claim_submission and #auth_status.”
- *Focused on business value*—“We will increase containment, increase user satisfaction, and reduce incorrect call routing by more accurately identifying Claim Status calls. This is our most popular call type. Accuracy problems frustrate users as they repeat themselves, leading to increased opt-out rates. Further, misunderstood callers can get routed to the wrong human agent, increasing our cost. This problem also decreases user satisfaction.”

The technical detail is great for putting into your technical backlog, but this detail is just jargon to most stakeholders who are only interested in what it means to them.

We suggest classifying your improvement work such that it aligns with business objectives. You can also add technical classifications for ease of managing your backlog—everyone should know the business effects behind the work in your backlog. Table 1.4 connects generic reasons for improving conversational AI to specific business metrics.

Table 1.4 Aligning improvement reasons with business metrics

Improvement reason	Business metric	Description
Cost reduction	Containment	Reduce the number of calls going to a human. This is primarily for process-oriented bots.
Cost reduction	Average handle time	Reduce the time spent by a human by increasing productive work done in the AI. For instance, if the AI authenticates the caller, the human agent won't have to. This is primarily for process-oriented bots.
Cost reduction	Human touches	Reduce the number of humans who touch a call. (Increases when calls are routed to the wrong human.) This is primarily for routing agents.

Table 1.4 Aligning improvement reasons with business metrics (*continued*)

Improvement reason	Business metric	Description
User satisfaction	Net promoter score (NPS)	Improve results on post-service surveys.
User satisfaction	Time to resolution	Reduce the amount of time from first contact to problem resolution.
Compliance	N/A	Restrictions that you must adhere to at the risk of severe penalty. This is part of the cost of doing business.

Note that some improvements may affect several business metrics, as shown in table 1.5.

Table 1.5 Technical improvements may affect multiple business metrics.

Technical improvement	Affected business metrics
Increased intent-recognition accuracy	Improves containment (callers won't quit due to frustration) Improves human touches (when routing, goes to right human) Improves average handle time Improves time to resolution (from reduced retries) May improve NPS (from reduced retries)
Clarify a confusing question	Improves containment (callers won't quit due to frustration) Improves time to resolution (from reduced retries)
Shorten a lengthy message	Improves time to resolution Improves NPS

NOTE Some business goals contradict each other. For instance, a medical insurer improved the accuracy of a “claim denied reason” intent. Callers used to immediately transfer due to the intent not being recognized by the AI; therefore, they did not take a post-call survey given when the AI completes a task. After the intent accuracy improved, callers could self-service and find out their claim was denied. This improved containment, but now those unhappy callers took a survey to complain, and the insurer’s NPS for their assistant dropped.

Exercises

- 1 Consider other technical improvements, like “reducing flow complexity,” “shortening dialogue,” and “reducing friction points.” What business objectives do they influence?
- 2 How would you address these improvement areas incrementally?

1.4 Follow along

In this book, we will demonstrate conversational AI practices using two types of software platforms. The techniques we use will work on many different platforms:

- *Conversational AI platform*—A core software platform that provides conversational AI capabilities like natural language understanding and dialogue management. There are many choices, like Amazon Lex, Google Dialogflow, Microsoft Azure AI Bot, and Rasa, just to name a few. We are experts in IBM watsonx Assistant and use it in this book.
- *Generative AI model platform*—A service that offers one or more LLMs that you can interact with through APIs. Popular choices include Anthropic, ChatGPT, Gemini, Hugging Face, and Ollama. In our day jobs, we use IBM watsonx.ai and its Prompt Lab, and we used it to build and test the prompts in this book.

Why a commercial cloud platform?

Installing the prerequisite software for AI applications can be challenging. LLMs are generally resource intensive. Using a commercial cloud platform lets you get started quickly and focus on building conversational AI and generative AI.

The techniques described in this book are broadly applicable across different conversational AI and generative AI platforms. Where appropriate, we will call out any terminology differences. There are many excellent choices—you can use the technology platform you're comfortable with or explore a new one!

Summary

- Conversational AI must be built with the user experience in mind. Good conversational AI helps users complete their tasks quickly. Bad conversational AI frustrates users.
- There are thousands of generative AI models. Large language models (LLMs) are a subtype of generative AI models that are good at generating text.
- LLMs can perform many tasks with impressive performance but also have significant risks, including hallucination. It takes thoughtful guidance and guardrails to use LLMs effectively and responsibly.
- LLM technology can supplement conversational AI. LLMs can respond to users directly and also assist you in building your conversational AI.
- Continuous improvement is possible and necessary for effective conversational AI.
- Iterative improvement delivers higher business value with lower risk.

Building a conversational AI

This chapter covers

- Building an FAQ conversational AI
- Building a process-oriented conversational AI
- Using generative AI inside of your conversational AI

In production, conversational AI can be quite complex, and throughout this book, we'll cover many techniques that address the real-world problems you'll face as you build and deploy your own solutions. In this chapter, we'll build Cake Bot, a conversational AI solution with elements from several different kinds of conversational AIs. This will give us a solid foundation for understanding conversational AI structure.

We'll follow a fictitious small American bakery from Ohio called Cake Shop. The company makes custom cakes and takes orders for delivery or pickup. They want to add a conversational AI solution to their website to help their customers. Since they have never built a bot before, they intend to start small but hope to quickly expand the scope and capability of their solution. They decide to start with an AI solution that answers their most frequently asked questions.

Many of the tasks in this chapter *could* be done with large language models. However, this bakery is cautious. They especially want to control the wording of

responses given for several question types. Thus, their solution will blend traditional and generative techniques.

We will demonstrate the building process using a conversational AI platform (IBM's watsonx Assistant), and we'll later fold in a generative AI platform (IBM's watsonx.ai). The key concepts we demonstrate are applicable across many different AI platforms. You can easily use your platform of choice for conversational AI and generative AI.

2.1 *Building an FAQ bot*

Most conversational AI builders start with a question-answering bot. Also known as FAQ bots, these AI solutions deliver a response directly to a user's question, often without any follow-up questions. The user asks a question, the bot returns an answer, and the conversation is done when the user is finished asking questions. These bots work especially well when there are a small number of (frequently asked) questions.

In this section, we will build an FAQ bot for Cake Shop. Some questions will have a static response that will be the same no matter how the question is asked. Other questions will have a dynamic response that will change based on information in the question. But before we train the bot on any question-answering, we will first put some basic scaffolding in place.

2.1.1 *FAQ bot foundations*

Every conversational AI needs to be able to start a conversation and react when it doesn't know what to do. Most conversational AI platforms provide this capability by default when creating a new chatbot. It's worth quickly checking these configurations and adapting them to your needs.

Cake Shop starts building their conversational AI (the "assistant"), and they title it "Cake Bot." From the conversational AI's main menu, their developer navigates to Actions, which lists all the assistant's capabilities. The first list is titled "Created by you" and is empty; the second list is titled "Set by assistant," and it lists the default capabilities, which are outlined in table 2.1.

Table 2.1 Default capabilities in a new assistant

Capability	Executed when
Greet customer	The assistant is first opened or engaged with. Opening the assistant starts a conversation.
No action matches	No action can be matched to the user's message (the message is not understood). Other platforms may call this the "fallback intent."
Trigger word detected	Keywords like profanity are detected.
Fallback	The user needs to leave the chatbot.

The first capability is the most important to customize, as it gives us our first chance to personalize the assistant. The default text is "Welcome, how can I assist you?" The

Cake Shop team changes this text to “Welcome to Cake Bot. How can I help you?” This is a minimum level of customization—it would be better to include additional information, like what the bot can do for users. However, the bot does not have any capabilities yet, so the Cake Shop team leaves this message as is.

Next, the “No action matches” action should be reviewed. This action will be invoked when the bot does not understand the user. Since the bot has no training yet, this action will be the default response to any user input. The default configuration is shown in figure 2.1.




Figure 2.1 Default configuration of “No action matches” in the assistant

This configuration is summarized as follows:

- 1 The action counts how many times it has been invoked in the conversation.
- 2 If three or less times, the response is “I’m afraid I don’t understand. Please rephrase your question.”
- 3 If four or more times, it deflects to a fallback routine. (The default fallback routine is to offer a human agent.)

The Cake Shop team decides to reduce this threshold by changing the 3 to a 1. This keeps their users from getting stuck.

Fallback action and connection to a human agent

Most conversational AI platforms have no-code and low-code integrations to connect users to a human agent through chat or voice. We will not dive deeper into this, since the details are platform-specific. Suffice it to say that this is a common pattern. For the sake of this chapter, we will focus on the conversational design and AI training.

At this point, we have a chatbot that does three things:

- 1 When the user opens the chat, they are greeted with “Welcome to Cake Bot. How can I help you?”
- 2 Whatever they say next, the chatbot responds that it doesn’t understand.
- 3 Whatever they say after that, the chatbot offers a human agent.

Boring! Let’s train this bot to answer some questions properly.

2.1.2 Static question and answering

Let’s start with a mental model of the chatbot components involved in answering questions.

In some platforms, you can directly connect questions to answers. In others, an additional layer is introduced to categorize similar questions into groups called *intents*. An intent-based question-answering system gives the builder full control over responses generated by the conversational AI. A generalized version of this design is shown in figure 2.2, using Cake Bot as our example.




Figure 2.2 Question-answering bots map user utterances to intents, which map to answers.

Let’s review the terminology in this diagram:

- *Utterance*—This is the input provided to the chatbot. For a question-answering bot, these are questions.
- *Intent*—This is a logical grouping of utterances with similar meanings.
- *Response*—This is the output from the chatbot. For a question-answering bot, these are answers.

For your first chatbot, intents save a lot of time. Notice that, as a builder, you do not have to distinguish between questions with similar meaning. “What time are you open?” and “What are your hours?” both relate to the operating hours of your store. It’s not critical for the bot to differentiate these. We give them the same “meaning” via the #store_hours intent. “What cakes do you sell?” has a different meaning and thus a different intent of #cake_options.

For each intent your bot serves, the bot is trained with example utterances. Modern intent-based systems require as few as five example utterances per intent. This is

not a bad trade-off; there are nearly an infinite number of ways to ask for store hours, and by providing a handful of examples, you can train your bot well.

Intent-based question-answering systems are a blessing and a curse: for each intent you train, you can control the response, which offers pros and cons.

Pros:

- You have complete design control over the response. You can copyedit it, format the text, and even include graphical elements. You know the exact contents of the response.
- For a small number of intents, this can be done quickly. You can set up your first chatbot in as little as an hour.

Cons:

- As the number of intents increases, it becomes more difficult to train the bot to recognize them all.
- The responses do not adapt to nuances in the user's questions. For "Are you open today?" the bot still responds generically: "We're open every day."
- Inaccurate or untuned responses give the user a painful feeling of "chatbot doesn't understand."

We'll address several of these downsides to question-answering bots in the next few chapters: how to collect the right data to train your bot (chapter 4), how to use that data to train stronger intents (chapter 5), how to supplement those intents with answers from documents and generative AI (chapter 6), and how to use generative AI for a few more training and testing tasks (chapter 7).

Let's start by training our chatbot on its first question-answering capabilities. For each one, we need a user intent, a set of related user utterances, and a response. The first set of questions and answers we'll define will cover the background on Cake Shop, operating hours for their stores, the kinds of cakes offered, the approximate cost of cakes, and information about their Cake Club. These intent-based question-answering responses are shown in table 2.2.

Table 2.2 Initial set of FAQ intents, with associated utterances and responses

Intent	Example utterances	Response
#background	Tell me about Cake Shop What's the background on your business? History of Cake Shop	Founded by Grandma Cake in 1980, we've made over 10,000 cakes for local residents!
#store_hours	Store hours What are your store hours? When are you open?	We are open Monday through Friday, 9:00 a.m. to 9:00 p.m.
#cake_options	Cake options Do you make wedding cakes? What kinds of cakes do you sell?	We offer cakes for many occasions, such as weddings, birthdays, anniversaries, retirement, and all-occasion cakes.

Table 2.2 Initial set of FAQ intents, with associated utterances and responses (*continued*)

Intent	Example utterances	Response
#cost	How much does a cake cost? Is there a minimum order value? Is there a surcharge for delivery?	Our cakes typically cost around \$30, with a \$5 delivery fee.
#cake_club	Cake rewards Cake Club Any special promotions or discounts?	Our Cake Club rewards program earns you a \$10 gift certificate for every ten cakes you purchase.

In the assistant, we define an action that detects an intent and gives a response—a question-answering action. This is the simplest kind of action we can define in any conversational AI platform. Figure 2.3 shows the user interface that starts this action definition.




Figure 2.3 User interface to create our first action

For each of these actions, we need to configure how they start (the user utterances) and what they do (respond with an answer). You’ll notice that these are the right-most columns in table 2.2. Some conversational AI platforms also use the intent label for the action; ours labels the action based on one of the user utterances that triggers it. We start our journey of defining the utterances that trigger the #background action in figure 2.4.

Note that the user interface points out that the chatbot’s recognition of this action will improve with more examples. For the sake of our demo, we will use three examples per action, which is enough to get us started. We will demonstrate multiple ways to find additional training examples in subsequent chapters.

Our question-answering action is almost complete. We have the questions that trigger it; now we need to define the chatbot’s response. The response for our #background action is shown in figure 2.5. This action has three parts:

- *Conditional logic*—For a static question-answering action, no logic is needed. The action only starts when the intent is detected.
- *Response*—“Assistant says” is the response to the user. Our response is simple text.
- *Next step*—For a static question-answering action, no next step is needed. Giving the answer ends the action.

Add example phrases:

Enter phrases that a customer types or says to start the conversation about a specific topic. These phrases determine the task, problem, or question your customer has.

The more phrases you enter, the better your assistant can recognize what the customer wants.

Enter phrases your customer might use to start this action

Total: 3

Tell me about Cake Shop

What's the background on your business?

History of Cake Shop

Figure 2.4 Defining the utterances that trigger an action

Step 1

1. The intent is enough, no other conditional logic

Is taken

without conditions

Set variable values f_x

Assistant says

2. Responds with this text

B I f_x

Founded by Grandma Cake in 1980, we've made over 10,000 cakes for local residents!

Define customer response

And then

3. Action complete! Nothing else to do.

End the action

Figure 2.5 Defining the response for a question-answering action. The simplest form has only one step after detecting the intent—give the response.

Licensed to Marco Bobbo <bobbomarco2@gmail.com>

We'll repeat these action-creation steps for each of the five intents. Each action is trained with the examples that trigger it and the response it should give. Each of these actions is a single-step action that ends once the answer is given.

When all five actions have been created, we are ready to do some testing. Figure 2.6 shows the testing interface for our chatbot.




Figure 2.6 Chat preview link

Let's ask some questions! Figure 2.7 shows the test results for a sample question.




Figure 2.7 Example question-answering response from Cake Bot

Note that the question asked does not exactly match any of our training examples. This indicates that the bot has learned the meaning in the examples. The following listing shows additional tests of the bot.

Listing 2.1 Testing Cake Bot with more questions

```
User: hours of operation?  
Bot: We are open Monday through Friday, 9am to 9pm.  
User: why did the chicken cross the road  
Bot: I'm afraid I don't understand. Please rephrase your question.  
User: cost for a cake?  
Bot: Our cakes typically cost around $30, with a $5 delivery fee.
```

This is a great start for our bot. We can train it on more intents, and we can make it more accurate by giving it more examples for those intents. But let's consider something different.

All the question-answering actions we've created have been single-step actions. The user gets the same response no matter what they ask. In the next section, you'll see how to evolve a static response into a dynamic response based on additional information.

2.1.3 Dynamic question and answering

Cake Shop presently has four locations: Columbus, Dublin, Westerville, and Grandview. When the bot was first created, all four locations had the same operating hours: 9:00 a.m. to 9:00 p.m. on weekdays. Circumstances have shifted—the Columbus store needs to open and close one hour earlier (8:00 a.m. to 8:00 p.m.). A single chatbot response doesn't cover all the stores anymore. Now when a user asks about store hours, we need to figure out which store they need the hours for. If they don't specify, we'll need to ask them a clarifying question.

The next listing shows how we want the bot to handle store hours questions in a series of sample questions.

Listing 2.2 Sample conversations for store hours, dependent on location

User: hours of operation?

Bot: To view our store hours, please select a location.

Bot: (Columbus, Dublin, Westerville, Grandview)

User: Columbus

Bot: Our Columbus store is open Monday through Friday, 8am - 8pm.

Ambiguous question is now clarified before answering

User: hours of operation?

Bot: To view our store hours, please select a location.

Bot: (Columbus, Dublin, Westerville, Grandview)

User: Dublin

Bot: Our Dublin store is open Monday through Friday, 9am - 9pm.

Unambiguous question is answered directly

User: hours of operation for Grandview?

Bot: Our Grandview store is open Monday through Friday, 9am - 9pm.

We can also draw a flow diagram covering these sample conversations, as shown in figure 2.8. It's helpful to create a flow diagram and sample conversations when your conversation has dynamism. Some of your team members will prefer the diagrams, others the conversations, and some will need both.




Figure 2.8 Process flow for a location-specific #store_hours intent

The “store hours” flow can be implemented in three steps:

- 1 Display “To view our store hours, please select a location” and a list of locations. The user must choose a location.
- 2 If step 1 = “Columbus,” display Columbus hours, and end the action.
- 3 Display the hours for the step 1 store, and end the action.

This works because steps “fall through” in our platform. Here’s how a few conversations work:

- The user types “store hours,” and step 1 fires. The user selects “Columbus,” and step 2 fires and completes the action.
- The user types “store hours,” and step 1 fires. The user selects “Grandview,” and the step 2 condition is not met. Step 3 fires and completes the action.
- The user types “store hours for Columbus.” The step 1 exit conditions are met, so step 2 fires and completes the action.
- The user types “store hours for Grandview.” The step 1 exit conditions are met, and the step 2 condition is not met. Step 3 fires and completes the action.

Figure 2.9 shows how these steps are implemented in our assistant.




Figure 2.9 The three steps for the `#store_hours` action

Cake Bot is off to a good start. It can answer some basic questions about Cake Shop, and it even has a little dynamism. Grandma Cake won't have to answer so many repetitive questions on the phone! But Cake Bot cannot take any action for the users yet. We'll look at that in the next section.

Exercises

- 1 Download this chapter's chatbot code from the book's GitHub site: <https://github.com/andrewfreed/EffectiveConversationalAI>. Load the chatbot in watsonx Assistant, and use the Preview panel to test the chatbot's question and answering flows.
- 2 Alternatively, implement Cake Bot in your preferred conversational AI platform:
 - Define a greeting message.
 - Define a fallback intent and/or fallback message.
 - Implement the five intents from table 2.2.

2.2 Routing agents and process-oriented bots

Not all bots are question-answering bots. Q&A bots are great at delivering answers, but what if the user needs more than an answer—what if they need the bot to act? For Cake Shop, we'd love for customers to be able to order cakes from the bot. If all we have is question-answering capability, figure 2.10 is the best we can do.




Figure 2.10 Cake Shop's cake order process as question-answering. But it doesn't really answer the question!

The user wants to complete a process but cannot do that inside the bot. They only get *instructions* on how to complete the process. A question-answering bot is thus often an early iteration of a more capable solution.

2.2.1 Routing agents

Cake Shop offers a wide variety of cakes with different flavoring and decoration options. There are decoration packages for weddings, graduations, birthdays, and more. There are flavoring options including vanilla, chocolate, and strawberry. Plus, there are payment and delivery methods. Given all these options, it's reasonable to assume the user may want or need to talk this process through with a human.

For many chatbot developers, the next logical iteration of their chatbot is a routing agent. The routing agent detects the intent from the user's utterance and determines who can best help fulfill the intent. Figure 2.11 reimagines our Cake Bot with routing agent capabilities.




Figure 2.11 A routing agent detects user intents and routes them to an appropriate specialist.

For the original Q&A requests, the bot works as it did before. But for cake-ordering requests, this bot does not attempt to answer the question at all—it just routes the call to an appropriate specialist. See our implementation in figure 2.12. The action has one step once the intent is detected: route the user to a specialist.

And then

Connect to agent (action ends)

If online	Let's send you to an available ordering specialist.
If offline	There are no agents available at this time. When one becomes available, we'll connect you.
To the agent	User wants to order a cake.

[Edit settings](#)

Figure 2.12 Routing agent configuration for `#cake_orders`. As soon as the intent is detected, the user is deflected to a human specialist.

This routing agent is just triaging incoming requests, which can be transferred to human agents or to specialized AI solutions. The human agents could use the telephone or live web chat. In this book, we'll generically refer to these humans as *call center agents*.

Press 1 for appointments . . .

You’ve probably phoned an interactive voice response (IVR) system that recites a menu of options and prompts you to select one (“press 1 for appointments”). This is also a routing agent. One downside to these systems is the length of time it takes to read the menu. A conversational AI routing agent lets you speak your intent, increasing the convenience over listening to a lengthy menu.

Routing agents let you implement conversational AI solutions iteratively rather than needing to handle everything at once.

The human agents in routing agent systems often know what type of request they are receiving but little else. In figure 2.12, they were only told that the user wanted to order a cake. For some process flows with high degrees of complexity and sensitivity, this may be ideal. For instance, a “report fraud” intent should probably connect to a human right away.

In other scenarios, an early deflection to a human agent is mundane for the agent and expensive for the employer. For insurance systems handling claim statuses, member IDs and claim dates must be collected before getting to higher value tasks like explaining what has happened with a claim. Here the AI assistant could first collect the member ID and claim date before directing the conversation to a human.

Thus, the next evolution of a routing agent is to shift more of the work to automation. Let’s build this for Cake Bot.

2.2.2 Transitioning from a routing agent to a process-oriented bot

The generalized process flow for ordering cakes is shown in figure 2.13. It includes four steps to clarify details about the cake being ordered, then a confirmation step, and finally fulfillment. (For brevity, we will omit the fulfillment details for the rest of the chapter—example code is available at our GitHub site: <https://github.com/andrewrfreed/EffectiveConversationalAI>.)




Figure 2.13 Process flow for ordering a cake from Cake Shop

With the full process flow designed, we can transition from a routing agent toward a process-oriented bot. Cake Bot will handle part of the cake-ordering process by collecting a few details before routing to a human agent to complete the process. Figure 2.14 shows the design for the first iteration of Cake Bot’s transition.




Figure 2.14 Transitioning a routing agent to a process-oriented bot. The bot now collects two pieces of information before handing off to a human.

Our process used to have one step (figure 2.12). Now we'll have four:

- 1 The bot will start the process by responding, "I can help with your cake order."
- 2 Ask which size cake is needed, and provide options (small, medium, large).
- 3 Ask the occasion for the cake, and provide options (birthday, wedding, anniversary, retirement, all-occasion).
- 4 Transfer the user to a human agent. This is the original first step from the routing agent, but the message to the agent has changed from "User wants to order a cake" to "User wants to order a <size> <occasion> cake." The assistant will inject the responses from steps 2 and 3 into the message.

These steps are executed sequentially. Figure 2.15 shows step 2.

Step 2 Size

Is taken: without conditions

Set variable values f_x

Assistant says

What size cake do you need?

small medium large

Edit response Edit validation

And then

Continue to next step

Figure 2.15 The step to collect the cake size offers explicit choices to the user.

Note that the bot can offer distinct options to the user as buttons, letting the user know which options are available. The bot can also allow the user to type their response if they prefer. You can explore these options in the sample code provided on the book's GitHub site.

Also note that each of these options will support a “fall through.” If the user starts the conversation with “I want to order a cake,” they will be asked for size and occasion. If they say “I want to order a large anniversary cake,” they will skip the size and occasion questions since they already provided that information.

The rest of the cake-ordering process has been implemented on the book's GitHub site following the steps described in this section. A sample conversation is provided in the following listing.

Listing 2.3 Sample conversations for cake ordering

```
User: I want to order a cake
Bot: I can take your cake order.
Bot: What size cake do you need?
Bot: (small, medium, large)
User: large
Bot: (birthday, wedding, anniversary, retirement, all-occasion)
User: birthday
Bot: What flavor cake would you like?
Bot: (vanilla, chocolate, marble)
User: vanilla
Bot: Which frosting would you like?
User: strawberry
Bot: Ok, I've got you down for a small vanilla birthday cake with
strawberry frosting. Is that correct?
User: yes
Bot: Great. Your order has been accepted.
Your order subtotal is $30.00.
Bot: Would you like to schedule your order for pickup or delivery? Please
note there is a $5 delivery fee.
Bot: (pickup, delivery)
User: delivery
(fulfillment details are omitted)
```

The confirmation message in step 7 plays back the information collected in previous steps.

The order confirmation in step 8 triggers conditional logic for the price of the cake.

Cake Bot is getting more capable. It has static question-answering capability about cakes, dynamic question-answering capability about store hours, and a process-oriented flow for ordering cakes. The Cake Shop team deploys its chatbot and is pleased with the results (and the users are pleased with their cakes!). Next, we'll take on our final challenge of the chapter: adding generative AI capability with a large language model (LLM).

Exercises

- 1 Refer to this chapter's chatbot code that you downloaded from the book's GitHub site (<https://github.com/andrewfreed/EffectiveConversationalAI>). Load

(continued)

the chatbot in watsonx Assistant, and use the Preview panel to test the chatbot's cake-ordering flow.

- 2 Alternatively, implement Cake Bot's ordering process in your preferred conversational AI platform:
 - Detect the cake-ordering intent.
 - Route the intent directly to a human agent.
 - Collect all four cake data points, and conclude with a summary.

2.3 *Responding to the user with generative AI*

Cake Bot only uses traditional conversational AI technology so far. The question-answering is done by an intent-based classifier. The ordering process is done with a sequential series of rules. This has worked well for the needs of Cake Shop so far.

When the Cake Shop team reviews the performance of the Cake Bot, they see an unusual trend. Users are asking the bot for recipes they intend to serve for dinner before the cake. There's no other pattern to recipe requests—there are requests for casseroles, salads, stir fries, and more. The team is heartened by the diversity of their users but does not know how to handle these requests in the Cake Bot. How could they detect all these different types of recipes, let alone respond to them all?

This is an excellent place for the Cake Shop team to incorporate some generative AI into their solution. They can use the existing intent mechanism to detect recipe requests and then route those to an LLM to generate an answer. They will need to integrate an LLM into their chatbot generally and send specific requests to that LLM.

Let's see how they can do that.

2.3.1 *Integrating with an LLM*

For many conversational AI platforms, the primary way to integrate with external systems is through application programming interfaces (APIs). These are ubiquitous integration patterns and fortunately are supported by a large variety of generative AI platforms that expose LLMs. The specific way APIs are integrated into conversational AI will vary by platform. In some platforms, this integration is done with code; others are low-code and visual interfaces. Differing platforms have different names for their integration capabilities, such as *extensions*, *integrations*, and *fulfillments*. Many let you integrate APIs via OpenAPI specifications.

We will add a generative AI platform as an extension to do LLM-based text generation. There are four steps to adding an extension in our platform (the details of the steps are included in the book's GitHub repository):

- 1 From the Integrations menu, select Build a Custom Extension.
- 2 Provide a name and description, like "Generative AI platform API call."

- 3 Provide an OpenAPI specification file. This file documents the capabilities of the extension, including the methods it exposes, its required and optional parameters, and the responses it provides. OpenAPI specification files are a common documentation format for APIs. They are usually provided by generative AI platforms.
- 4 Provide connectivity and authentication details, such as the URL of the API implementation and the API key needed to access it.

We add the extension and visually explore it from inside the assistant. Figure 2.16 shows the extension for the LLM text generation API in our platform.

Operation	Method	Resource
Generation	POST	/ml/v1/text/generation
Request parameters		Response properties
version string Required		results[].stop_reason string
input string Required		results[].generated_text string
model_id string Required		results[].input_token_count integer
parameters.top_k integer Optional		results[].generated_token_count integer
parameters.top_p number Optional		model_id string
parameters.time_limit integer Optional		created_at string

Figure 2.16 OpenAPI spec for our LLM text generation API with a subset of the possible request parameters

At the time of writing, our text generation API includes 15 input parameters and 6 output parameters—more than fit in figure 2.16! There are also a handful of parameters available without any customization, like the HTTP status code for the response. Other generative AI platforms will have a similar parameter set, perhaps with different parameter names or locations. Let’s review the most significant parameters:

- **input** (request)—The prompt to the LLM. It will include the instructions, context, and data for the LLM. Some of that data may come directly from the user.
- **model_id** (request)—Identifier of the LLM to use for the task. Most generative AI platforms let you pick from several models.

- `parameters` (request)—Key-value pairs that tweak the LLM’s behavior. These include the decoding method (greedy or sampling), number of output tokens to generate, and several other parameters.
- `generated_text` (response)—The output from the LLM.

We can use an extension from any step in any action. Earlier in this chapter, we used capabilities like “Assistant says,” “Continue to next step,” and “Connect to Agent.” For extensions, the capability is called “Use an extension.” Figure 2.17 shows what that extension invocation looks like for our recipe action. Other LLM tasks would look similar but have different configuration values. This parameter set is tuned for providing recipes.

And then

Use an extension

Extension ⓘ

watsonx.ai

Operation ⓘ

Generation

Parameters ⓘ

input	set to	Tr recipe_prompt
model_id	set to	mistralai/mixtral-8x7b-instruct-v01
project_id	set to	REPLACE WITH YOUR WATSONX ID
parameters.max_new_...	set to	1000
parameters.min_new_t...	set to	0
parameters.decoding_...	set to	greedy
parameters.repetition_...	set to	1
version	set to	2023-05-29

Edit extension

Figure 2.17 Invoking an LLM text-generation API from an action in the assistant

Let’s look at how we can connect this all together in Cake Bot.

2.3.2 *Routing requests to an LLM*

The flow diagram in figure 2.18 outlines how the recipe generation will be covered in Cake Bot. We first create a new action. Just like our question-answering actions, we

start with some example utterances that trigger this action. Our first three utterances are “Show me a recipe for,” “How can I make,” and “Tell me how to bake a.” Given the huge variety of possible recipes, we do not include the names of the dishes, just the way that recipe requests are likely to look.




Figure 2.18 Flow diagram for recipe generation in Cake Bot via LLM

Step 1 of the new action is to store the entirety of the user’s original utterance (from the system variable `input.text`) in a variable called `recipe_query_text`. This is a technique we have not done in previous steps. For the cake-ordering action, each option had an explicit and finite set of responses. Even if the user said, “large cake, please” we only wanted to store “large.” For the recipe request, we have no idea what the user will say, so we will capture their entire utterance.

Step 2 of the action is to define the prompt for the LLM. We concatenate a simple system prompt with the user’s request. The next listing demonstrates the expression used in building the `recipe_prompt` variable.

Listing 2.4 Building the recipe prompt, which is stored in the `recipe_prompt` variable

```
"You are a helpful kitchen assistant. Create a recipe as instructed by
the user.\n\nInput: ".append(recipe_query_text).append("\n\nOutput: ")
```

Step 3 of the action is to call the LLM. The parameters were shown in figure 2.17, but let’s dive into the specific values here:

- `input`—We assign the `recipe_prompt` variable value as input. This injects the user’s recipe request into the generalized prompt format shown in listing 2.4.
- `model_id`—There are many models available, but at the time of writing, `mistralai/mixtral-8x7b-instruct-v01` has performed well on this generation task.
- `project_id`—This is an identifier from the generative AI platform project instance.
- `min_tokens` and `max_tokens`—These are set to 0 and 1000 respectively, increased from the defaults of 0 and 200, because recipes tend to be a bit lengthy.
- `decoding_method`—Set to `greedy` for ease of debugging, so the model responds identically if the same input is given.
- `repetition_penalty`—Set to 1 (no penalty), since recipes on the internet typically have some repetition in them.

- **version**—Indicates the minor version of the API requested. We used the default at the time of writing: 2023-05-29.

Step 4 of the action is to check that the API call was successful, and if so, to display the response to the user. We'll generate a response starting with some static text, "Here's your recipe," and append the LLM response. We don't need to write any JSON parsing code; the OpenAPI specification tells us we can reference the value of `generated_text`. Figure 2.19 shows the action response configuration.




Figure 2.19 Displaying the output from the LLM call to the user

This was a lot of fun to build inside Cake Bot! Listing 2.5 shows one example conversation we had with the bot. The recipe is lengthy, so most of it is omitted for brevity. Load up Cake Bot and try it for yourself!

Listing 2.5 Sample abbreviated conversation with Cake Bot

```
User: I need a quick recipe for peanut butter pie
Bot: Sure! I'm looking up a recipe for you. Please wait a few seconds.
Bot: Here's your recipe:
...
Ingredients:
1 pre-made graham cracker crust
1 cup creamy peanut butter
...

Instructions:
In a large mixing bowl, beat together the peanut butter, powdered sugar, and
softened cream cheese until smooth and creamy.
...
Serve chilled and enjoy your quick and easy peanut butter pie!
Let me know if you need any further assistance or clarification.
```

Exercises

- 1 Refer to this chapter's chatbot code that you downloaded from the book's GitHub site (<https://github.com/andrewrfreed/EffectiveConversationalAI>). Load the chatbot in watsonx Assistant, and follow the instructions to integrate with watsonx.ai. Use the Preview panel to test the chatbot's recipe flow.

- 2 Alternatively, implement Cake Bot's ordering process in your preferred conversational and generative AI platforms:
 - Detect the recipe intent.
 - Build a prompt from a set of instructions and the user's input.
 - Direct the LLM's response to the user.

Summary

- Question-answering (Q&A) bots are a great way to start building your first conversational AI.
- Training Q&A bots with examples of questions lets you provide predefined answers to related groups of questions (intents).
- Actions start with an intent and can have many outcomes: answering a question, deflecting a user to a human agent, asking follow-up questions, and making API calls.
- A routing agent identifies intents and passes information to human agents. It is a great method for incrementally adding capability to a conversational AI while leaning on human capability.
- Conversational AI can use a combination of traditional and rules-based techniques along with generative AI.

Planning for improvement



This chapter covers

- Building a cross-functional team that achieves conversational AI success
- Defining success through business goals, key metrics, and user pain points
- Analyzing effectiveness using outcomes and metrics to guide improvements
- Implementing structured processes for identifying, reporting, triaging, and prioritizing problems

Every conversational AI solution should be built with success in mind, and success is defined differently depending on the type of chatbot involved. For instance, question-answering bots must deliver prompt, accurate responses while minimizing follow-up interactions. Process-oriented or transactional bots must guide users efficiently toward specific goals. Routing agents must seamlessly direct users to appropriate destinations.

However, misunderstanding user intent, excessive complexity, and immediate opt-outs can hinder progress and cause user pain. Addressing these challenges improves a chatbot's performance and helps it achieve success. Organizations that continuously improve their chatbots are most likely to deliver optimal outcomes.

Combining diverse expertise within cross-functional teams is crucial for continuous improvement. The team members can drive change through their unique perspectives, skills, and insights. However, the team needs to agree on how to improve their solution.

The conversational analyst wants to simplify the dialogues, but the business wants to convey specific information. Who is right? In this chapter, we'll show how a team at MediWorld, a fictional company, adapted and improved their chatbot. Their team started by enhancing their question-answering bot, but as user needs evolved, they transitioned to developing additional capabilities for a process-oriented bot.

3.1 Knowing when you need to improve

Imagine the following scenario:

MediWorld, a large drug store, had call centers overloaded with questions about the pandemic. They deployed a chatbot to provide information related to COVID-19. The bot detected a focused set of intents about the virus and responded with reliable information.

Figure 3.1 illustrates PharmaBot efficiently recognizing these intents.

"Which stores offer testing?" "I need a rapid test" #testing	"I have a fever, should I worry about COVID?" "What symptoms should I watch out for?" #symptoms	"Is it safe to go out in groups?" "How can I protect myself from COVID?" #prevention
--	---	--

Figure 3.1 PharmaBot efficiently detected informational intents from user queries.

When vaccine availability was imminent, the nature of customer questions changed dramatically. Suddenly, everyone had a slew of different questions:

- *Will I be eligible for the vaccine?*
- *Can I get a vaccine appointment?*
- *When can I get my second dose?*
- *Do I have to call for an appointment, or can I set it up here?*
- *Can I travel after my shot?*

PharmaBot was initially weak at understanding these questions, frequently responding, "Sorry, I'm not sure what you're asking. Please rephrase your question." Users were frustrated and dissatisfied, and more conversations ended up in the call center after failing in the bot. There was also an increase in immediate opt-outs, reflecting an apparent disconnect between user expectations and PharmaBot's capacity to address the evolving landscape. MediWorld's team set out to improve the bot, but they first had to agree on what "improve" meant.

NOTE The need for continuous improvement has never been more critical, as evolving user expectations and technological advancements demand constant refinement and adaptation. Bridget van Kralingen quipped, “The last best experience that anyone has anywhere becomes the minimum expectation for the experience they want everywhere.” Improvement requirements may come from internal sources (such as support of new features) or external sources (where an event drives new questions never seen before).

Recognizing the need to improve in conversational AI is pivotal to ensuring its effectiveness and relevance. A virtual assistant is not a static solution; its performance must evolve with user behavior, business needs, and advancements in technology. Signs that improvements are necessary often emerge through key performance indicators (KPIs) such as low containment rates, high fallback intent usage, or frequent agent escalations. Planning practical improvements starts with building a cross-functional team, defining clear success criteria, analyzing outcomes, and implementing structured processes for issue management.

When deciding when to start improving your conversational AI, best practices recommend beginning as soon as you notice recurring problems, declining engagement rates, or unmet business goals. A proactive approach can prevent small problems from escalating into larger problems. Establishing regular review cycles ensures that improvements align with evolving user expectations and organizational objectives.

Start your improvement journey by planning a comprehensive data collection strategy even before the first deployment. Remember, just having log files doesn't automatically reveal the pain points. It's essential to be methodical in identifying trends and patterns across user interactions. Many teams tend to fix isolated problems without considering the overall volume or frequency of those problems. While it may seem productive to address one-off problems, this rarely leads to meaningful improvements in overall performance. By focusing on systemic problems with significant effect, you can ensure your efforts are always directed toward tangible progress, making you feel focused and committed.

Remember, measuring performance before and after deploying changes is equally important. Establish baseline metrics before implementing fixes, and compare them with post-deployment data to assess whether the changes delivered the expected improvements. If the results don't align with your expectations, don't worry. Analyze the root cause further and iterate on your solution to effectively address any gaps or unforeseen problems. This process will give you the reassurance and confidence that your efforts are leading to tangible progress.

3.2 *Your cross-functional team*

MediWorld recognized the critical role of its PharmaBot in providing timely and accurate customer support. A multidisciplinary team of conversational analysts, customer support experts, and data analysts came together to assess and enhance PharmaBot's

performance. This group aimed not only at addressing existing challenges but also proactively anticipating and meeting the evolving needs of the user base:

Developers at MediWorld focused on refining PharmaBot’s natural language processing capabilities. They found ways to enhance the chatbot’s understanding of user queries so it could provide precise and context-aware responses.

Simultaneously, lead call center agents shared valuable insights from the calls transferred to them, shedding light on common pain points and frequently asked questions.

The MediWorld data analysts delved into user interaction data. They identified areas where the chatbot “failed” and categorized those failures by the last task attempted by the bot.

Figure 3.2 shows the kinds of insights that each group brought to the table.




Figure 3.2 The team identified areas for improvement using their diverse skills, setting the stage for an effective improvement plan.

Let’s look beyond the PharmaBot example and focus on teams typically involved in a conversational AI improvement plan. The specific roles, responsibilities, and team size can vary based on your organization’s size, its goals, and the complexity of the chatbot. In smaller projects, individuals may take on multiple roles. Chapter 1 introduced a “dream team” for building conversational AI (figure 1.5). A similar diverse group is needed to improve and refine existing chatbots. While the structure of this team may differ across organizations, it generally consists of three key subteams, all working together.

First is the support and maintenance team for the chatbot. This team is tasked with analyzing and evaluating the chatbot’s performance. Additionally, they serve as technical subject matter experts (SMEs). They know the existing intents the chatbot handles, the training data for those intents, and the dialogue flows in the chatbot. They can implement code and technical changes. Their roles and tasks are outlined in table 3.1.

Table 3.1 The chatbot support and maintenance team

Role	Tasks
Data analyst/data engineer	Analyzes user interactions and feedback to make informed recommendations regarding changes, fixes, and enhancements.
Chatbot developer/conversational analyst	Implements technical changes and enhancements to the chatbot. These may include new integrations (more of a developer role) or updates to the dialogues and actions of the chatbot (conversational analyst).
Quality assurance (QA) tester	Validates that a change, fix, or enhancement produces the desired outcome and does not result in any unexpected or negative outcomes. Testing may be manual or involve automated testing tools.
Project manager	Coordinates tasks; ensures that the continuous improvement process stays on schedule.
Other SMEs	Provide specialized knowledge in specific areas of the chatbot ecosystem; they may be brought on board as needed. For example, security experts assess potential threats and recommend appropriate security measures or remediation strategies to ensure the chatbot remains secure and resilient against evolving risks.

The second subteam is business stakeholders. They collectively ensure that the chatbot improvements align with the broader organizational goals and the business needs. Business stakeholders ensure that the chatbot is technically proficient and aligned with organizational goals, user needs, and legal standards. This team is broken down in table 3.2.

Table 3.2 The chatbot's business stakeholders

Role	Tasks
Executive leadership	Involved in aligning the improvements or the priorities of the improvements with overall business strategies
Customer service	Responsible for the business processes and workflows that the bot is addressing
Product manager (of the chatbot)	Responsible for overseeing the chatbot's development and strategic direction, ensuring it meets business objectives
IT department	Provides technical support and infrastructure for the chatbot's development, deployment, and maintenance
Operation manager	Collaborates to integrate the chatbot into operational processes, streamlining workflows
Legal and compliance teams	Ensure the improvements comply with industry regulations and legal requirements

The final subteam is the governance team. Their role is to ensure that the chatbot's deployment, use, and continuous improvement align with organizational policies, standards, and ethical considerations. They are identified in table 3.3.

Table 3.3 The chatbot's governance team

Role	Tasks
Corporate ethics/compliance focal	Addresses ethical concerns about the chatbot's behavior and decision-making as well as AI model risk management. They also ensure that their guidelines for responsible AI are followed through the improvement phase.
Governing executive team	Has the final say on prioritizing the system roadmap, backlog, and all costs (support or business team) associated with the system.

Having a diverse and cross-functional improvement team ensures that different perspectives and expertise contribute to the development and oversight of the chatbot initiative. Regular meetings, clear communication channels, and documentation of governance policies are essential for the team's effectiveness. Again, the specific stakeholders involved may vary based on the nature and scope of the project.

Exercises

- 1 Think about your last chatbot implementation, and list all the stakeholders you had. Discuss the stakeholder perspectives on the common goal of improving the chatbot and how this goal aligns with their specific objectives. Consider the potential conflicts between stakeholder interests and strategies for resolving them.
- 2 Alternatively, use MediWorld's PharmaBot as an example, and consider the various stakeholders and their goals.

3.3 Driving to the same goal

Even within a single improvement team, different members may have conflicting priorities about what should be addressed first. Consider the following scenario:

When the PharmaBot team first met, they couldn't agree on where to start. Everyone brought their "must-fix" lists. Some were hunches, some were informed by reading a few transcripts, and some came from detailed analysis. The team aligned on understanding the frequency of the problems: Do they come up in every conversation, or are they one-offs? Issue frequency was a key component in prioritizing fixes and enhancements, helping MediWorld enhance the performance of its chatbot.

Developers and data analysts advocated for refining PharmaBot's natural language processing abilities, analyzing recent interactions to identify areas for improvement in understanding complex and context-dependent questions. Meanwhile, the lead call center agents emphasized the need for PharmaBot to offer more detailed and empathetic responses, focusing on recurring user concerns they had encountered.

A data-driven approach helps with prioritization. Addressing the most common pain points will lead to an immediate and tangible improvement in the overall user experience. Figure 3.2 demonstrated how multiple team members can contribute diverse

observations and insights. Each member brought a unique perspective from their respective roles and expertise.

3.3.1 *Revisit business goals*

Conversational AI improvement team members must agree on the common goals of the improvement. The first critical step is reaching a consensus on what success means—the team must revisit the original business objectives that prompted the implementation of the chatbot. For instance, a question-answering bot must consistently respond quickly and accurately to user questions. A process-oriented bot must help users efficiently achieve their goals, like scheduling appointments or checking accounts. A routing agent must direct users to the place or specialist that can fulfill their needs based on their inquiry. Evaluating chatbot performance against these goals involves metrics such as response accuracy, user satisfaction, and the bot's ability to handle a broad range of relevant topics and use cases.

As the business landscape evolves, adaptability becomes paramount. Changes in user expectations or technological advancements may necessitate enhancements or strategic shifts to maintain optimal chatbot performance. Therefore, the team must continuously reassess and refine the chatbot strategy, ensuring it remains aligned with the overarching organizational and user objectives. This iterative process caters to the evolving needs of the business and its users.

FROM BUSINESS OUTCOMES TO METRICS

Defining the right metrics starts with understanding how business goals evolve over time. Consider the following scenario:

The PharmaBot team needed to align their efforts with the importance of efficient vaccine distribution and accessibility. While the original business goal was to answer questions and reduce the burden on their call centers, the business goals have changed to distributing vaccines effectively and automating the appointment-making process, which require different metrics.

When in doubt, consider your users—what do they need, and how do those needs drive business? What were the original business goals? What did users want? How are the two aligned? Recognize that these answers may shift over time. Figure 3.3 shows PharmaBot's first business goal: providing accurate and up-to-date information about the new pandemic. Accuracy was the key metric.




Figure 3.3 PharmaBot started as a simple Q&A bot. Many question varieties got the same answer.

Once a stable base was in place, the PharmaBot team added complexity and intelligence to the Q&A bot. They detected entities (contextual elements relevant to an intent) to provide more targeted answers to their users and improve accuracy, as figure 3.4 demonstrates.




Figure 3.4 Q&A got more complex by detecting entities in user utterances, leading to more specific answers within a common intent.

Then external influences changed the business objectives again. Vaccine availability shifted the nature of the bot. Instead of interacting with a pure Q&A about the virus, users wanted to act directly through the bot to schedule vaccine appointments. This required process-oriented flows to collect multiple pieces of information. Figure 3.5 shows the start of this process flow.




Figure 3.5 Some question types do not have a single static answer but require a full process flow to satisfy.

These new capabilities brought complexity. Automating testing and vaccine appointments required integration with scheduling systems and databases. This brought heightened emphasis on security and privacy measures. Protecting users' personal information, adhering to healthcare regulations, and ensuring secure transactions became paramount.

Not all conversational AI solutions have to deal with this evolution, at least not at the pace PharmaBot required. Managing a chatbot's evolution requires thoughtful consideration of the costs and benefits associated with each aspect of the chatbot's transformation.

The primary business goals for any conversational AI solution revolve around enhancing overall business outcomes. Every business has a goal or goals for the conversational system, which comes down to two key factors: revenue generation and cost reduction. These goals translate into metrics such as increased conversion rates,

higher average order value (AOV), and enhanced customer lifetime value (CLV) for revenue growth. On the cost side, metrics like reduced average handling time (AHT), lower operational costs, and improved first-contact resolution (FCR) reflect cost savings. Organizations expect a measurable return on their investment (ROI), and these performance metrics guide continuous improvement efforts, aligning the conversational AI's success directly with key business outcomes.

The business goals must be translated into measurable metrics. This allows a quantifiable assessment of how well the conversational AI meets its goals. The examples in table 3.4 demonstrate how businesses in various industries can define measurable metrics aligned with their specific goals.

Table 3.4 Sample metrics derived from business goals in various industries

Business goal	Resulting metrics	Bot type
Increase online sales, and reduce customer service costs.	Percentage of checkouts completed by the chatbot without human intervention: Achieving 75% of automated checkouts, leading to a reduction of 100,000 customer service inquiries per day, resulting in a daily cost savings of \$500,000.	Question answering
Improve customer support efficiency, and minimize service disruptions.	Percentage of inquiries successfully routed to the appropriate department or specialist: 90% of inquiries directed to the relevant support team without the need for manual intervention, leading to a decrease of 40,000 support tickets per day, resulting in a daily cost savings of \$700,000.	Routing agent
Enhance booking experience, and decrease support costs.	Percentage of automated booking confirmations without agent intervention: Achieving 70% of automated booking confirmations, reducing 80,000 support inquiries per day, resulting in a daily cost savings of \$640,000.	Process-oriented
Improve patient engagement, and optimize appointment scheduling.	Percentage of appointments scheduled autonomously by the virtual assistant: 90% of appointments are booked autonomously, reducing 30,000 manual scheduling tasks per day, resulting in a daily cost savings of \$700,000.	Process-oriented

Conversational metrics need clear links to business value to prove a return on investment. Metrics like call center deflection and routing accuracy reduce costs. Metrics like customer satisfaction lead to increased revenue when satisfied customers consume more services. The PharmaBot team achieved both cost savings and revenue growth by automating appointment scheduling.

ADDITIONAL BUSINESS DRIVERS

Beyond aligning with the core business goals discussed in the previous section, organizations should consider additional factors that drive value from conversational AI. Successful AI implementations do more than just support high-level objectives—they actively enhance customer engagement, optimize sales strategies, and reduce operational costs.

Conversational AI can strengthen customer interactions, guide sales, and suggest relevant products. Analyzing chatbot-driven conversion rates is crucial for refining

strategies. Businesses that explore AI-driven features for upselling and cross-selling can maximize revenue opportunities. From an operational perspective, conversational AI helps by handling routine tasks, freeing human agents to focus on complex activities. Automating processes improves efficiency, reduces support costs, and enhances satisfaction with quicker responses.

A thorough analysis can uncover opportunities for improvement and for optimizing chatbot performance. As technology evolves, businesses should expand chatbot functions for ongoing cost reduction and operational excellence.

Competitor analysis, evaluating features like natural language understanding and personalized experiences, can guide continuous improvement. Regular updates enable adaptation to change in the competitive landscape, driving innovation for positive business outcomes.

3.3.2 Effectiveness

When determining improvement priorities, another key factor to consider is the chatbot's effectiveness. Does the chatbot do what it was intended to do? While the concept of "effectiveness" is simple (does it work as expected?), it goes beyond task completion. It involves providing a positive and efficient experience for users.

Let's continue with our scenario, where the team is now looking at their dashboard showing chatbot metrics. Most chatbot development platforms have a simple analytics dashboard containing KPIs summarizing how users engage with the chatbots. These dashboards typically contain data on the number of conversations, chatbot confidence, and conversation duration. They may also include the most frequently asked questions or intents.

Figure 3.6 shows the analytics dashboard created for PharmaBot. While it shows some of the KPIs, it does not express PharmaBot's effectiveness. The total number of conversations helps us understand traffic, but it does not help assess how many people successfully completed the chats or how far they went.




Figure 3.6 PharmaBot's basic analytic dashboard shows a usage summary but cannot give insight into what users like (or don't like) about the bot.

The team sought to determine the current assistance rate. Basic analytics gave them information on total inquiries per day, but they needed help figuring out how many of these were successful. The team needed to go beyond the mere quantity of conversations. Everyone knew that user demand was increasing. Was PharmaBot meeting the new demand? They needed to figure out how to measure the bot's effectiveness and find ways to optimize its performance. The team found that 45 % of all conversations transferred to the call center. This “containment” metric influenced the cost to the business. What could they do with this number?

In the PharmaBot team's case, one key metric for measuring effectiveness was containment. *Contained conversations* are when the chatbot can fully handle a user query on its own; *uncontained conversations* require a human to be involved. The *containment rate* is calculated as the number of contained conversations divided by the number of total conversations. This metric provides a high-level measure of chatbot performance, as illustrated in figure 3.7.




Figure 3.7 Basic daily dashboard showing a simple business metric: containment. This metric is tracked daily but still does not give deep insights into bot performance.

PharmaBot's dashboard implies a simple definition of success: “Contained calls are successful.” This mental model is summarized in figure 3.8. However, while containment is a valuable metric, it does not tell explain why users succeed (or fail) when interacting with the bot. A more detailed analysis is needed to gain deeper insights.




Figure 3.8 The simplest outcome definition. This does not give insight into how the bot can be improved.

CONVERSATION OUTCOMES

To better understand how conversation outcomes affect chatbot improvement, let's go back to our example scenario. The PharmaBot maintenance team needed to move beyond high-level performance metrics and analyze the actual conversation outcomes. Containment alone didn't capture the full picture—it only told them whether or not a conversation stayed with the bot. But what really happened in those conversations? Here's how they conducted an in-depth review of the conversations:

The team dived into a sea of transcripts to understand what happened in the conversations. How did they end? The team discovered a myriad of endings—successful completions, abandoned conversations, and a perplexing number of transfers.

The data analyst observed different flavors of success. First, the expected case where the PharmaBot responds well to the query and the user is satisfied. Second, there are handoffs to the call center due to business rules, like when the user lives in a state where MediWorld cannot do SMS confirmations. These handoffs are also successful, as they align with what PharmaBot set out to do: collect required information so a human specialist doesn't have to. The team agreed that they needed to document these two kinds of cases separately. They labeled them as "Automated Resolution" and "Intentional Transfer."

Then there were definite failure scenarios. Users asked for a human agent after the bot misunderstood them. The bot also automatically transferred users when it misunderstood them consecutively. And some users disconnected midway through appointment scheduling.

Lastly, the PharmaBot team found conversations where the users didn't even try. Users were either silently disconnecting after the bot's greeting or starting the conversation with the utterance "agent." One team member remarked, "Perhaps there's a psychological factor at play here: some users don't want to engage with a bot. Let's separate these conversations too."

The team finally developed a nuanced categorization system: Success (automated resolution and intentional transfer), Failure (abandonment and escalation), Bot not wanted (immediate disconnect and immediate escalation).

Once the categorization system was in place, counting the conversations in each bucket was easy. The PharmaBot team was starting to really understand their bot's performance.

Defining detailed conversation outcomes related to your metrics will give you insights into your solution's performance. Conversation outcomes describe how user interactions with the chatbot conclude, categorizing whether the chatbot resolved the query, required human assistance, or resulted in an incomplete session. Defining these outcomes is critical to assessing and improving your solution against your business goals. Once your solution is deployed, analyze your conversation logs, and classify them against the outcome model.

Figure 3.9 shows a new detailed outcome model. It starts with containment on the left, and then it breaks down contained (and non-contained) outcomes with detailed reasons. Finally, these reasons are mapped back to success and failure categories. The reasons help us understand what went right and wrong in the bot. For instance, conversations may not be contained due to "failure": maybe the user opted out or the bot

repeatedly didn't understand. Some conversations are intentionally transferred to humans following the underlying business process—those aren't failures.




Figure 3.9 Breaking down why conversations are not contained gives more insight into bot performance and shows you where your bot needs improvement. One way to achieve this is by using detailed outcome classifications, which define the specific results of chatbot interactions. These classifications categorize conversations based on their resolution, user experience, and next steps.

Now that we've introduced the idea of a granular outcome model, considering how (where and why) a conversation might end, let's look at the same metrics dashboard we saw in figure 3.7. Instead of looking at 45% containment, we can better understand the conversations. On the dashboard of figure 3.10, we replaced the containment rate chart with a success-rate chart, indicating success/failure/not wanted as the highest-level categories.




Figure 3.10 Enhancing the summary dashboard with a success rate. Not all contained calls are successful; not all transferred calls are failures.

In fact, you may want to illustrate the details of the outcomes as well. In figure 3.11, we show what this might look like. This approach enables you to quickly break down high-level outcomes into detailed outcomes, which could help you get stakeholder buy-in on improvements too.




Figure 3.11 The detailed outcome model depicts conversation outcomes and outcome details aggregated over a set time period. This provides much greater insight into bot performance than a binary “contained or not contained” model.

The detailed outcome model’s strength comes from its flexibility. Every conversational AI project can define its own unique outcome categorization. The model depicted in figure 3.11 is a useful sample implementation. As always, adjust this model for your needs. For example, if your chatbot does not have a human handoff, you can omit the escalation and transfer outcomes.

Here are some suggested categorizations for common types of conversational AI solutions.

For Q&A bots:

- Success—Conversation completion scenarios:
 - Automated resolution—The Q&A bot successfully answers the user’s inquiry or provides relevant information without human intervention.
- Failure—The interaction fails to achieve the desired outcome:
 - Abandoned—The user leaves the conversation before getting a good answer, possibly due to frustration or dissatisfaction with the bot’s responses.
 - Escalated—The Q&A bot does not understand the user, and the interaction is escalated to a human agent for further assistance. This could occur by user

request, or the bot may automatically escalate after multiple consecutive misunderstandings.

- Chatbot not wanted:
 - Immediate disconnect—The user exits the conversation without ever sending a message to the bot.
 - Immediate escalation—The user’s first utterance to the bot is a request for a human agent.

For a transactional or process-oriented bot:

- Success—Conversation completion scenarios:
 - Automated resolution—The process-oriented bot successfully completes the user’s intended task, such as booking an appointment, without human intervention.
 - Intentional transfer—If required by business rules, the bot transfers the interaction to a human agent, even though no “errors” were encountered.
- Failure—The interaction fails to achieve the desired outcome:
 - Abandoned—The user abandons the conversation midway through the transaction, possibly due to complexity or confusion with the bot’s interface.
 - Escalated—The bot starts but cannot complete a process flow due to misunderstanding the user or the user’s request for a human agent.
- Chatbot not wanted:
 - Immediate disconnect—The user exits the conversation without ever sending a message to the bot.
 - Immediate escalation—The user’s first utterance to the bot is a request for a human agent.

For a routing agent:

- Success—Conversation completion scenarios:
 - Intentional transfer—The routing agent successfully directs the user to the correct department or specialist, potentially passing all information collected so far. A routing agent may have 0% containment and be working very well!
- Failure—The interaction fails to achieve the desired outcome:
 - Abandoned—The user exits the conversation before being routed by the bot.
 - Escalated—The routing agent cannot gather enough information to route the user, either through misunderstanding or a user request for a human agent.
- Chatbot not wanted:
 - Immediate disconnect—The user exits the conversation without ever sending a message to the bot. The user cannot be automatically routed.

- Immediate escalation: The user’s first utterance to the bot is a request for a human agent, bypassing all automated routing.

When you categorize conversations like this, the number of conversations in each category helps you assess the effectiveness of the chatbot implementation and identify areas for improvement.

The detailed outcome model should be integrated with the conversational design. A great method is defining milestones for each of the bot’s “happy path” questions. Figure 3.12 shows this design for PharmaBot. The “Schedule appointment” milestone, shown in the shaded box, signifies a key step in which the bot completes the scheduling process. “Help with anything else?” is also shaded, indicating that the bot is ready to assist further after completing a primary task.




Figure 3.12 High-level design of PharmaBot with milestones for significant parts of the conversation. “Schedule appointment” and “Help with anything else?” are both marked as successful paths.

The FAQ intents have only one milestone (the FAQ response), whereas the process flow for appointments gathers multiple data points. The milestones that declare successful completions should be marked.

The detailed outcome model is most powerful when it’s overlaid with the design. When each conversation has an outcome and a “last milestone,” you can quickly find insights. In figure 3.13, we see PharmaBot’s metrics over time for failed conversations, including the last deployment date.




Figure 3.13 When the outcome model and conversation design are overlaid, insights become apparent. This chart breaks down failed conversations by the last step before failure, helping identify where users struggle most. The different categories—such as appointment scheduling, intent detection, and zip code entry—show their relative contributions to overall failures over time. The spike after “Changes deployed” highlights the effect of recent updates, offering insights into areas needing further optimization.

The PharmaBot team can further drill down into the failed conversations to view the detailed outcomes of abandoned and escalated. Combining the outcome model and the conversation design can jumpstart your data-driven analysis. It can tell you where to start your investigation.

The detailed outcome model in figure 3.11 provides deeper insights into user experience by revealing where and why conversations fail. In this case, the outcome is failure, and the primary reason is user escalation. However, looking at the specific points where users escalated—the last step before failure—provides actionable insights. The breakdown in the chart highlights key escalation points: during appointment scheduling, intent detection, member ID collection, and zip code entry. After changes were deployed, failures at the intent detection step dropped significantly, while failures in member ID and zip code collection saw slight declines. However, appointment scheduling failures spiked, indicating a new area of friction. This level of analysis allows the team to prioritize improvements effectively, ensuring that their fixes target the most pressing user pain points. While containment rate is often used as a high-level measure of effectiveness, it does not tell the full story—some containments may still result in poor user experiences, and some transfers may be necessary for a successful outcome. The outcome model in figure 3.13 helps the conversational AI team distinguish between these cases and refine the bot accordingly.

CUSTOMER SATISFACTION

In the outcome model, we infer customer satisfaction through conversational outcomes. This is a quick method, but it’s indirect and can leave out some details. It can be useful to be more direct with customer satisfaction.

Customer satisfaction can be measured by gathering direct feedback from users. Thumbs up or down or a numeric satisfaction score are standard. The only drawback with these metrics is that the user response rate is often low—many users hate giving feedback—although unsatisfied users are more likely to take the chance to complain.

You can implement surveys after the chatbot concludes an interaction. The survey could include questions about ease of use, helpfulness, and overall satisfaction with the bot. In addition, net promoter score (NPS) surveys may also be presented to users. Keep the survey brief: the longer the survey, the less likely users are to complete it.

You can also assess customer satisfaction by reviewing a sample of conversations. The review can include chat logs or summaries from human agents who completed the conversation. These logs and summaries may even be categorized using large language models. Table 3.5 shows the types of feedback you can look for in each type of conversational outcome.

Table 3.5 Linking feedback to conversation outcome details

Conversation outcome details	User feedback	Notes and caveats
Automated resolution (success)	Positive feedback or none	Users may give positive feedback to conclude an interaction (“thanks!”). But this is unlikely: most users disconnect once they get what they need.
Transfer (success)	Positive verbal feedback or no feedback. Not in chatlogs.	Users may give verbal feedback to the human agent about the bot, not to the chatbot directly.
Abandoned (failure)	Negative feedback or none (disconnection before a process completes)	User’s last comment to the bot may have negative sentiment (“I hate this!”). But many users won’t bother expressing their frustration—they just disconnect.
Escalated by user (failure)	Negative feedback	Users who request an agent mid-process flow are unhappy (“get me to an agent!”).
Escalated by bot (failure)	None	When the bot initiates the escalation, we can’t prove the user was unhappy, but we could reasonably assume so.
Immediate disconnect (chatbot not wanted)	No feedback provided (no engagement)	Users who immediately disconnect may hate all chatbots, may hate your chatbot, or may have connected to it by accident. You can’t know for sure.
Immediate escalation (chatbot not wanted)	No feedback provided, or verbal negative feedback is given expressing desire for human assistance instead of chatbot use	Users who immediately escalate may hate all chatbots, may hate your chatbot, or may just prefer humans. You can’t know for sure.

3.3.3 Coverage

As part of improving chatbot effectiveness, identifying gaps in what the bot can handle (or, in other words, its *coverage*) is just as important as addressing escalations. In the ongoing analysis, it became clear that some user questions weren't being misunderstood—they simply weren't covered by the bot's existing knowledge. The scenario continues as the team uncovers these gaps and works to expand PharmaBot's capabilities:

The team prioritized tackling escalations first, since they have the biggest effect on the metrics. They analyzed transcripts from escalated conversations, and patterns started to emerge. Users often escalated right after the bot didn't understand them. The data analyst cross-referenced these instances against the dialogue flow and suggested where they could improve the bot's natural language processing capabilities. During this analysis, they found several questions that PharmaBot was not equipped to answer:

- *"I heard about rare side effects. How can I distinguish between post-vaccine symptoms and something more serious that requires medical attention?"*
- *"If I missed the recommended second dose of my COVID-19 vaccine by a few days, will it still be effective, or do I need to start over?"*
- *"I'm pregnant, and I'm unsure about getting the COVID-19 vaccine. Can you provide information on the safety and potential benefits for pregnant individuals?"*
- *"I've been diagnosed with an autoimmune condition. Can I still receive the COVID-19 vaccine, and are there any additional precautions I should take?"*
- *"I've read conflicting information about the long-term effects of COVID-19 vaccines. What is known about their safety over an extended period, and are there ongoing studies?"*

For questions like these, PharmaBot responded with "Sorry, I do not understand," even after users tried to rephrase their questions. The team looked for clusters of questions with similar characteristics to see what the bot should be trained on next. Along with other intents, a group of inquiries related to vaccine safety emerged.

Coverage measures how many user questions the chatbot attempts to answer effectively. A chatbot with low coverage either lacks training data for key topics or struggles with overlapping and ambiguous intents, where similar questions confuse the model and prevent it from confidently selecting the correct response.

The team must analyze user interactions to improve coverage, identifying gaps where the chatbot fails to provide meaningful responses. This process involves assessing transcripts, tracking failed queries, and pinpointing recurring user needs currently unsupported. Addressing these gaps may require refining training data, restructuring intent classification, or introducing alternative approaches such as retrieval-augmented generation (RAG).

Several methods can enhance chatbot coverage, each addressing different challenges:

- Chapter 5 explores improving weak understanding by refining training data, addressing missing intents, and optimizing labeling strategies.

- Chapter 6 discusses how to bypass intent limitations using retrieval-based techniques combined with generative AI (RAG) to provide more dynamic and informed responses.
- Chapter 7 demonstrates how generative AI can generate additional training and testing data to expand the chatbot's ability to handle diverse queries.

Regardless of the approach, improving chatbot coverage requires continuous analysis of real-world interactions. Examining user utterances, identifying common failure points, and iterating on the bot's design ensures it evolves to meet users' needs more effectively.

Exercises

- 1 Consider the following chatbots and their business goals. Devise suitable metrics and goals for each chatbot type to evaluate their effectiveness and performance:
 - Bot 1—Customer support bot for an e-commerce website (Q&A bot)
 - Bot 2—Banking bot for account management (transactional or process-oriented bot)
 - Bot 3—Customer service bot for a telecom company (routing bot)
- 2 Review the following scenarios found in the PharmaBot chatlogs. Based on the PharmaBot's response to each scenario, assign an outcome from the model. Once you've assigned outcomes to each scenario, reflect on any patterns or trends you observe regarding the bot's performance and areas for potential improvement. Recall that the conversation outcomes are success, failure, and chatbot not wanted. The outcome details are automated resolution (success), intentional transfer (success), abandoned (failure), escalated (failure), immediate disconnect (chatbot not wanted), and immediate escalation (chatbot not wanted):

– Scenario 1

User: "I heard about rare side effects. How can I distinguish between post-vaccine symptoms and something more serious that requires medical attention?"

PharmaBot: "Hello! It's essential to be informed about potential side effects. Common side effects of COVID-19 vaccines include soreness at the injection site, fatigue, and mild fever. However, it's crucial to seek medical attention immediately if you experience severe symptoms like difficulty breathing or chest pain."

User: "Thanks for the information, but what if I'm unsure whether my symptoms are severe enough to warrant medical attention?"

PharmaBot: "If you're unsure, it's always better to err on the side of caution and consult a healthcare professional. They can assess your symptoms and provide guidance tailored to your specific situation."

PharmaBot: "Would you like me to connect you with a healthcare professional to discuss your concerns further?"

User: "Yes, please. That would be helpful."

(continued)

PharmaBot: "Great! Please hold for a moment while I transfer you to a human agent who can provide personalized assistance. Your health is our top priority."

PharmaBot: (Initiates transfer to a human agent)

– Scenario 2

User: "If I missed the recommended second dose of my COVID-19 vaccine by a few days, will it still be effective, or do I need to start over?"

PharmaBot: "Hello! It's important to follow the recommended vaccination schedule for optimal protection. If you've missed the second dose by a few days, it's best to contact your healthcare provider for guidance on the next steps."

User: "I see. Will the delayed second dose affect the effectiveness of the vaccine?"

PharmaBot: "The delay may impact the vaccine's effectiveness, but receiving the second dose as soon as possible is essential. Your healthcare provider can assess your individual situation and advise accordingly."

User: "Got it, I'll contact my doctor for advice. Thanks for your help!"

PharmaBot: "You're welcome! It's always best to consult with a healthcare professional for personalized guidance. If you have any more questions, feel free to ask."

– Scenario 3

User: "Hi there. I have a question about the COVID-19 vaccine."

PharmaBot: "Hello! I'm here to help. What would you like to know about the COVID-19 vaccine?"

User: "Agent"

PharmaBot: "Of course! Let me connect you with one of our human agents who can assist you further."

PharmaBot: (Initiates transfer to a human agent)

– Scenario 4

User: "Can you tell me about the ingredients in the COVID-19 vaccine?"

PharmaBot: "Sorry, I don't understand. Could you please rephrase your question?"

User: "What are the components of the COVID-19 vaccine?"

PharmaBot: "I'm sorry, I still didn't quite get that. Can you try asking in a different way?"

User: "What is the composition of the COVID-19 vaccine?"

PharmaBot: "I'm having trouble understanding your question. Could you provide more context or try asking in simpler terms?"

User: (hangs up)

3.4 *Identifying and resolving problems*

Identifying problems is crucial to continuously improving a conversational AI. The chatbot team needs a methodology for working through problems, including how to

find problems, how to reason through them as a group, and how to determine when they are resolved. This methodology will allow the team to work toward a common goal.

3.4.1 Finding problems

The best way to start finding problems is by examining trends in your conversation outcomes. It's great to see successful conversations, but focus on the failed and “bot not wanted” conversations. Dig into the outcomes that have upward trends. As shown earlier, these outcomes are most insightful when overlaid on your conversational design. What was the last thing the user did before the negative outcome? Figure 3.14 dashboard helped the PharmaBot team.




Figure 3.14 A dashboard that breaks down an outcome by the last step taken, observed over time.

Ideally, your analysis tool can count and plot conversations by

- Conversation outcome
- Last step taken in the conversation
- Date and/or time of conversation

Your conversational AI platform may already track these data points. Other platforms may need to be instrumented by adding context variables into key parts of your dialogue flow. With these metrics in place, teams can uncover unexpected user behaviors that affect chatbot performance. For example, when the PharmaBot team analyzed escalated conversations, they discovered a surprising pattern:

The PharmaBot team drilled into escalated conversations. They found many failures when PharmaBot asked users if their appointment was for vaccines or testing. Many users replied “yes”—a response that didn’t align with the expected format.

Recognizing the importance of understanding user behavior, the team realized that the bot's inability to interpret this ambiguous "yes" response was causing frustration among users. Some even said "yes" again when PharmaBot repeated the question. This was a surprising source of both abandoned and escalated conversations.

Once a problem area is found, you can start designing a solution. The design is affected by how many ways the problem is encountered and can be addressed. In the "unexpected yes" scenario, there are two ways out: handle the "yes," or try to get users to stop saying it.

Let's look at a few more ways to find problems.

QUALITATIVE PROBLEM EXPLORATION

While metrics and conversation logs provide valuable insights, some chatbot problems are difficult to detect through quantitative analysis alone. Launching a qualitative improvement effort by collecting and analyzing user feedback allows you to uncover more user pain points. Let's look at how the PharmaBot team went about their surveys:

To uncover deeper user frustrations, the PharmaBot team launched a qualitative improvement effort by collecting and analyzing direct user feedback. They encouraged users to share detailed descriptions of their challenges and expectations through a survey. Once the feedback was collected, the team categorized it to identify common pain points, as shown in table 3.6.

Table 3.6 Survey responses leading to identified pain points (part 1)

User	Survey response	Identified problem
1	I tried asking about vaccine eligibility in different states, but the bot couldn't provide clear information. Not knowing if I qualify for the vaccine when I plan to travel is frustrating. The responses seemed generic and didn't address the complexity of eligibility criteria in various locations	PharmaBot understood the basic intent (eligibility) but failed to provide state-specific information. The chatbot did not consider the user's location or the state they inquired about, leading to an unhelpful response.
2	I attempted to schedule a vaccine appointment, but the process felt confusing. The bot's instructions were unclear, and I felt unsure if my appointment was successfully booked. It would be helpful if the bot could provide more guidance throughout the scheduling process.	The scheduling workflow lacked clarity, causing users to feel uncertain about whether their appointment was successfully booked.
3	(No survey response given; the user left the chat)	Overcomplicated steps discouraged users from completing the process. User 3 did not get to the survey.

Unlike conversation logs, qualitative feedback provides direct insight into user frustrations—you don't have to infer what went wrong. Combined with corresponding conversation transcripts, this feedback creates a clearer picture for the improvement team, making diagnosing and addressing chatbot deficiencies easier.

Recruiting users to provide actionable feedback can significantly enhance chatbot performance and user satisfaction. However, most users are reluctant to leave feedback. Providing small incentives or even a simple, genuine thank-you can encourage participation. If feedback is a key part of your improvement strategy, consider implementing a system that creates a win-win situation for both users and your team.

WARNING The provided examples offer valuable insights into specific user challenges, but they may not be statistically significant. Don't rush into solutions based on isolated instances. Look for repeated occurrences of the identified problems to gauge the scale and effect of each one.

QUANTITATIVE EVALUATION FOR ISSUE DISCOVERY

While qualitative feedback helps uncover user frustrations, it can also reveal measurable, functional problems. Challenges like slow response times or confusing dialogue flows can be quantified through conversation logs, which help teams diagnose problems and prioritize improvements. Let's continue with our scenario to see what other problems were found:

In addition to finding descriptive problems, the PharmaBot team uncovered some addressable functional problems. A sample problem is shown in table 3.7.

Table 3.7 Survey responses leading to identified pain points (part 2)

User	Survey response	Identified problem
4	It took the bot nearly 5 minutes to tell me about vaccine appointment availability. The delay was quite inconvenient, especially when trying to plan my schedule. A faster response would have been more helpful. I went to the bot to avoid being on hold!	Excessive response time frustrated users and diminished the chatbot's value as a faster alternative to traditional customer support.

Problems like this can be identified by analyzing the time taken for each step in a conversational log. You can track the average and maximum times taken at each step. Outliers may indicate poorly performing backend systems or confusing questions that users spend a lot of time thinking about.

This analysis could even be done on a subset of conversations. For instance, a slowly responding API is more likely to cause users to disconnect. Dive into the abandoned conversations, reviewing what the users were saying and how long the individual steps took. This kind of analysis can be done without asking users directly for feedback.

By identifying specific challenges associated with the conversation flow, analysts can target improvements in the conversational system effectively.

3.4.2 Group review

After reviewing their conversational outcome metrics and user feedback, the PharmaBot team compiled a list of concrete problems. Now they must build their improvement plan, starting with prioritizing the problems.

TRIAGING THE PROBLEMS

With these metrics in place, teams can uncover unexpected user behaviors that affect chatbot performance. For example, when the PharmaBot team analyzed escalated conversations, they discovered a surprising pattern:

The most critical problem identified in PharmaBot was the frequent misunderstanding of user queries, particularly in differentiating those about COVID-19 testing and vaccine-related appointments. Users got frustrated when the bot didn't understand, frequently ending conversations in abandonment or escalation. The call center agents agreed that they heard this problem when listening to frustrated users. Analytics confirmed the high volumes of this problem.

The team agreed to address this high-priority problem. The business goal was to complete setting up appointments without human agent intervention. They had found a recurring pattern of why appointment completion failed: the bot was confused about what type of appointment the user wanted, and users felt misunderstood—many of the transactions failed. The bot caused pain points of both not understanding and being too complex.

To move forward after finding problems, teams must systematically evaluate and prioritize the problems based on their perceived value and expected effect on the system. This involves assessing factors such as how frequently the problem occurred, the cost of implementing a fix, and the potential benefits of resolving the problem. By taking a structured approach to prioritization, they can ensure that improvements deliver the most value with the resources available. In the PharmaBot scenario, the highest priority was given to fixing the misunderstanding around appointment types, as this directly affected the ability to complete the booking workflow, a core business objective. Figure 3.15 illustrates a sample assessment of a problem. For a more thorough depiction, insights might include the volume of the problem, conversational outcomes, affected user complexity of remedy, other affected flows, and more.

Insights

- Schedule Appointment flow has a significant increase in **escalated** conversations after January 29th change
- Increase likely due to added language: “ask for an agent to schedule an appointment”

Recommendations

- Change verbiage to shorten the response and break up the dialogue
- Add screening questions for callers to access schedule appointment agent services
- Design SMS scheduling

Figure 3.15 Analysis of an increase in escalated conversations within the Schedule Appointment flow, identifying a potential cause (new wording directing users to agents) and providing recommendations to improve

The problem analysis depicted in figure 3.15 can contribute to a broader triage process by helping prioritize chatbot improvements based on effect and resolution complexity. Each problem is documented similarly, with insights into the problem, likely causes, and proposed solutions. In a full triage process, many such problem entries

are evaluated based on effect, frequency, and resolution effort to determine prioritization. The best prioritization practices consider value, proposed outcomes, and expected return. You must do a cost/benefit analysis. Benefits may be direct (improving containment) or indirect (improving the user experience). Costs may include time, effort, and complexity of the fix, and fixes that require buy-in from multiple departments will take longer. The expected return considers both benefits and costs, scaled by the volume of conversations affected. The goal is to focus on areas where the expected return justifies the investment of time and resources. An example of an expected return calculation for a problem is shown in figure 3.16.

Assumptions			
Agent Cost Per Call	\$6.00		
Calls per Day	80,000		
% Abandons	10%		
# Abandons per Day	8,000		
# Days in period	30		
	High Volume	Mid Volume	Low Volume
% Callbacks per Day	100%	50%	25%
Frustrated Callers directed to Agent	240,000	120,000	60,000
Cost of Callbacks over period	\$ 1,440,000	\$ 720,000	\$ 360,000

Figure 3.16 Assessment of the cost of users reaching call center agents after abandoning their chatbot conversations in frustration

The cost effect can be easily calculated when a given dialogue flow is handled by a human agent instead of through the chatbot. This calculation considers the agent cost per call, the total number of calls per day, and the rate of conversations transferred to human agents. The priority of this problem is much easier to assign when it's backed by this financial effect. This calculation can be repeated for all problem types. Remember that some costs are indirect: for example, a rude bot can lower customer satisfaction, making it challenging to quantify the financial effect.

Effort is another important prioritization driver. Just as there is a cost to the problem itself, there is also a cost to implementing the fix. Effort refers to the time, resources, and complexity involved in the implementation. The key is to balance the problem's importance with the implementation speed. The best problems to fix are high effect and low effort. First, address those problems that have a high effect but can easily be done. Figure 3.17 categorizes improvement opportunities based on their effect and the effort required to implement them.




Figure 3.17 An effect-effort matrix visualizes the relationship between the effort required and the potential effect of a proposed change.

High-priority items have high effect and low effort, followed by medium-priority items yielding incremental results. Low-priority items are those with low effect and high effort, and they are poorer candidates for fixing. This matrix can help teams prioritize their efforts effectively, focusing on changes that offer the greatest potential effect with the least amount of effort.

We can dive deeper into the categories presented in the matrix. For each category, we can provide sample problems, the user pain point they cause, and why they might occur. Table 3.8 starts with some high-effect and low-effort problems, and table 3.9 shows example high-effect, high-effort problems. Table 3.10 outlines some medium-effect, high-effort problems.

Table 3.8 Example high-effect and low-effort problems

Problem	User pain point	Why the problem might occur
Incorrect or insufficient dialogue response	Chatbot doesn't understand	Incorrect response due to poor intent recognition, input validation, or not adapting to the user's context
Poor dialogue response	Chatbot is too complex	Format and/or text does not convey the information clearly.
Broken dialogue trees	Chatbot doesn't work	The chatbot fails to function properly due to incorrect or misconfigured conditions and transitions within the conversation flow. These errors occur when the logic determining how the chatbot moves from one step to another ("jumps") is flawed or has not been thoroughly tested. As a result, users may experience dead ends, irrelevant responses, or abrupt conversation drops, negatively affecting containment and user satisfaction.
Flow enhancements	Chatbot is too complex	Processes have particularly complex steps that users can't easily complete. This is especially likely in long conversations.

Table 3.9 Example high-effect, high-effort problems

Problem	User pain point	Why the problem might occur
User unable to complete flow	Chatbot is too complex	Failures occur across many different steps in a process flow, necessitating a completely redesigned flow.
User questions are not addressed at all	Chatbot doesn't understand	Insufficient intents are implemented to cover user demand. This may require adding search or retrieval-augmented generation to handle infrequent question types.

Table 3.10 Example medium-effect, high-effort problems

Problem	User pain point	Why the problem might occur
Incomplete dialogue response (due to failed API)	Chatbot doesn't understand	Incomplete response due to API failure. The bot may not support all API request or response variations.
Intent confusion	Chatbot doesn't understand	Intent confusion can occur when the training data is imbalanced, meaning that certain intents have too few or too many example utterances, leading to misclassification. Additionally, discrepancies between training data and real-world user queries can make it difficult for the chatbot to recognize the correct intent.

These categorizations are not hard and fast. You should adjust the relative prioritization of changes based on the frequency with which the problems occur.

SOLUTIONING THE HIGH-LEVEL FIX

Once a high-priority problem is identified, the next crucial step is *solutioning*—outlining a high-level fix to rectify the problem. This involves a collaborative effort with the team working together to formulate a comprehensive solution. To ensure a structured approach, the team must address three fundamental questions: Who will be responsible for implementing the fix? What changes need to be made? How will the solution be implemented? The “who” encompasses the specific roles and responsibilities of the team that implements the fix. The “what” defines the nature of the solution, whether it involves refining the bot’s natural language processing capabilities, improving contextual understanding, or implementing a more sophisticated intent recognition system. The “how” outlines the technical approach and methodologies required for the implementation.

Additionally, the team must determine the development effort involved, considering factors such as coding complexity, integration requirements, and potential dependencies on external systems. This solutioning phase is crucial for devising a well-informed plan for continuous improvement.

ASSIGNING PRIORITIES TO ALL FIXES

A prioritized fix table is an indispensable tool for steering improvement initiatives. The table encapsulates a structured approach by assigning priority numbers, articulating

concise descriptions of identified problems, proposing recommended changes, quantifying the potential effect on the user experience, and providing direct links to associated GitHub issues. This comprehensive framework not only streamlines the development process but also facilitates efficient communication and collaboration among team members. Figure 3.18 shows a sample prioritized fix table.

Priority	Description	Recommended Change	Value/effect	ID
1	There is a training data imbalance across intents. This leads to poor recognition of the popular "schedule appointment" intent.	Increase training on that intent from production user utterances.	Pain point: Bot doesn't understand Technical: Improve intent recognition accuracy Business metrics: Containment	325
2	Bot tells user they have entered an invalid member ID.	Revise dialogue for kindness. Instruct how to find ID on their card.	Pain point: Too much complexity Business metrics: NPS	334

Figure 3.18 A sample prioritized fix table

Each column of the table plays a critical role in organizing and addressing problems effectively:

- *Priority*—Helps establish the urgency of each problem, ensuring that critical problems are addressed first.
- *Description*—Provides a brief but clear overview of the identified problem.
- *Recommended change*—Specifies the proposed solution or modification to rectify the problem, guiding development efforts.
- *Value/effect*—Quantifies the expected improvement in user satisfaction or usability resulting from the recommended change, aiding in prioritization.
- *ID*—Establishes a direct link to the corresponding problem in the project's issue tracker, such as a GitHub repository. This streamlines collaboration and tracks progress on the resolution of each problem. The GitHub issue may also provide further elaboration, progress, and status on the problem.

3.4.3 Determining acceptance criteria

Once the key problems have priorities and a high-level solution, the next step is to define the fix's acceptance criteria. Simply, how will we know when this problem has been resolved? Acceptance criteria are useful for validating functionality in the development environment and verifying improvements in production.

For instance, when PharmaBot could not handle "yes" to "Is your appointment for vaccines or testing?" the acceptance criteria might look like this:

- 1 When PharmaBot asks users to choose between vaccines and testing, and they say "vaccines," they get vaccine appointments.

- 2 When PharmaBot asks users to choose between vaccines and testing, and they say “testing,” they get testing appointments.
- 3 When PharmaBot asks users to choose between vaccines and testing, and they say “yes,” it asks them to confirm that they want vaccine appointments.

These acceptance criteria help the testing team validate current functionality (criteria 1 and 2) and new functionality (criteria 3). The fix can’t be deployed until it meets the acceptance criteria.

Once the fix is deployed, the team can verify that the number of conversations ending (and failing) with “yes” to the “vaccine or testing” question dramatically decreases or disappears altogether—for example, did the original metrics improve? Clear and measurable standards ensure your team is aligned with user expectations and project goals, setting the stage for successful bot improvement.

Exercises

- 1 The following sample problems are related to conversational AI implementation, each with varying degrees of complexity. Your task is to prioritize these problems based on their effect on the conversational AI system’s effectiveness and efficiency, considering both qualitative and quantitative volume metrics. Assess the complexity of the problem as an input to your prioritization. The examples cover various industries, as the improvement and prioritization efforts are applicable across chatbot types and domains.

a Inaccurate response generation:

- Description: The chatbot occasionally provides inaccurate or irrelevant responses to user queries, leading to user dissatisfaction and confusion.
- Effect: High—It affects user experience and trust in the chatbot’s capabilities.
- Volume metric: Frequency of inaccurate daily responses (e.g., 15% of total responses).

Sample Chat:

User: "Can you tell me about the vaccine side effects?"

Chatbot: "The COVID-19 vaccine is safe and effective."

User: "But I heard about people experiencing severe reactions."

Can you provide more information?"

Chatbot: "The vaccine is safe and effective."

Users: (hangs up)

b Slow response time:

- Description: The chatbot takes too long to generate responses, leading to user frustration and impatience, especially in time-sensitive situations.
- Effect: Moderate—it negatively affects user satisfaction and engagement with the chatbot.
- Volume metric: Percentage of user queries with misunderstood language per day (e.g., 8% of total queries).

Sample Chat:

User: "Can you give me the scoop on the COVID jab?"

(continued)

Chatbot: "I'm sorry, I don't understand what you are asking.
Could you please rephrase your question?"

c Limited language understanding:

- Description: The chatbot struggles to understand queries that use colloquial language, slang, or complex syntax, resulting in misinterpretation and inadequate responses.
- Effect: Moderate—it restricts the bot's ability to engage with users effectively, leading to frustration and reduced user satisfaction.
- Volume metric: Average response time in seconds per user query (e.g., 8 seconds).

Sample Chat:

User: "Can you provide information about COVID testing locations?" (Long pause, user hangs up.)

d Inconsistent integration with backend systems:

- Description: The chatbot experiences inconsistencies in integrating with backend systems, resulting in incomplete or incorrect information being provided to users.
- Effect: High—it undermines the chatbot's reliability and erodes user trust in its ability to provide accurate information.
- Volume metric: Percentage of conversations with backend integration errors per day (e.g., 12% of total conversations).

Sample Chat:

User: "Can you check if there are any vaccine appointments available tomorrow?"

Chatbot: "I'm sorry, I'm experiencing some technical difficulties. Please try again later."

- 2** Use sample conversations from your latest implementation, and repeat the preceding exercise with the data from your chatbot.

3.5 *Developing and delivering fixes*

Continuous improvement is often achieved through fixed-duration iterations, commonly known as sprints. Sprint iterations range from one to four weeks, depending on your organization's preference. While the prioritized fix table provides a general roadmap, the sprint plan specifically defines the next batch of functions to be delivered to users. The sprint plan is affected by resource availability: how much work you can develop and test in a time frame. It also prepares stakeholders for what they can next expect from your solution.

3.5.1 *Sprint planning*

This process establishes a systematic approach to issue tracking and resolution. It serves as the cornerstone for a well-coordinated, agile development journey, ensuring that your bot evolves in alignment with the proposed solutions and within the designated

timelines. Various tools, such as kanban boards, exist to visualize the state of a sprint throughout its duration. The most basic sprint visualization should include the problems being worked on and their status in the plan or execution cycle. Figure 3.19 shows one visualization that augments the fix table (figure 3.18) with two additional columns: status and the timeline, indicating the planned sprint.

Priority	Description	Recommended Change	Value/effect	ID	Status	Timeline
1	Training data imbalance...	Increase training ...	Pain point: Bot understanding...	325	In Dev	Sprint 3
2	Bot tells user invalid ID...	Revise dialogue for kindness...	Pain point: Too complex...	334	Planned	Sprint 4

Problem details
Planning

Figure 3.19 A prioritized table, including development sprints. Further columns, including UAT times and expected deployment dates, may be added.

3.5.2 Measure again

PharmaBot's team worked hard on improvements. When these improvements moved to production, the team monitored the metrics they expected to influence. By tracking failure outcomes against their past two deployments, they confirmed the fixes worked as expected.

Figure 3.20 shows the dashboard the PharmaBot team used.




Figure 3.20 Tracking conversation outcomes against the deployment of changes

Exercises

- 1 Sprint planning is crucial in addressing fixes and improvements in PharmaBot's development and delivery process. In these exercises, you will simulate a sprint planning session to prioritize fixes and enhancements for PharmaBot's iterative development cycle. You have two conversational analysts, a part-time backend developer, and a tester:
 - Review the prioritized fix table you created in the previous exercise.
 - Consider the capacity and velocity of the development team and allocate resources.
 - Create a sprint plan. Using a kanban board or similar tool, create a sprint plan that includes the prioritized fixes and enhancements, along with estimated effort and expected completion times. Consider adding columns for status and sprint inclusion to track progress and ensure transparency throughout the sprint.
 - Discuss expected deployment dates for the fixes and enhancements planned for the sprint.

Summary

- The continuous improvement cycle for conversational systems is an ongoing, iterative process.
- All improvements should drive toward the predefined business goals and user satisfaction.
- Meticulous metric definition, the right choice of monitoring tools, and a commitment to best practices are key.
- Use the “right” metrics relevant to your bot rather than those that are easiest to measure.
- Detailed conversation outcomes allow you to target a specific set of conversations for improvement.
- Several factors can determine a problem's priority, such as its frequency of occurrence, the expected improvement and complexity of a fix, and the team's capacity.
- Regression testing and the analysis of improvements are critical to ensuring improvements have occurred.

Part 2

Pattern: *AI doesn't understand*

T*he chatbot doesn't understand me* is the most common pain point users have when interacting with conversational AI. For traditional chatbots, this is caused by a poorly trained classifier; for other chatbots, it could be from having a bad search mechanism or simply not having access to the right information.

No matter what kind of chatbot you are building, the core capability you need is to understand what the user is asking for and respond appropriately. Chatbots that understand requests give useful responses; those that don't understand say something like "I'm sorry, I didn't understand—would you mind rephrasing?"

There are multiple methods for addressing poor understanding, and all of them hinge on identifying what your users want to accomplish and how they phrase their requests. Chapter 4 focuses on how to extract this data from logs and other sources. Chapter 5 shows how to improve chatbot understanding through training classifiers, and chapter 6 adds generative AI to the runtime mix with retrieval augmented generation (RAG). Chapter 7 uses generative AI at train and test time by creating new training and testing data with LLMs.

4

Understanding what your users really want

This chapter covers

- Recognizing indicators of weak understanding
- Measuring chatbot understanding
- Assessing your chatbot's current state
- Collecting and preparing log data to measure chatbot understanding
- Interpreting initial log data

A good chatbot experience is generally associated with the chatbot identifying (understanding) what the user wants. This is one of the key metrics you will use to measure performance. Sometimes a chatbot is deployed and has great initial understanding (or at least “good enough” for a pilot program). Over time, though, you may notice that it is returning wrong answers. Maybe your users are complaining more, either directly to the chatbot (“That doesn’t answer my question!”) or in the form of survey responses. Engagement could be trending downward while abandonment trends upward. You may start hearing from the call center about escalations that should have been handled in the virtual assistant. These are all indications that your conversational solution might be suffering from weak understanding.

In theory, chatbots should get better over time, but it is not uncommon to see a decline in understanding. We want to help you recognize when and why this could be happening in your solution. We will explain how to avoid some of the pitfalls and plan for common eventualities in the lifecycle of your solution. In this chapter, we will explore what it means for your conversational AI to have “good performance” in terms of its ability to correctly identify or classify a user’s goal (i.e., to understand the user). We will also offer techniques for preparing data for use in measuring a classifier’s performance or assessing generated responses.

4.1 *Fundamentals of understanding*

Being understood is a fundamental aspect of human communication. In a conversational AI, we use natural language processing techniques to try to understand what our user wants or needs. Because the scope of things a user could want is nearly infinite, and the way they might combine words to express those wants or needs is also infinite, this is a very difficult problem to solve.

4.1.1 *The impact of weak understanding*

Not being understood by a chatbot is probably the biggest source of frustration for a user. They came to your chatbot to get answers, and they may get an answer, but it may have nothing to do with their question. Perhaps the chatbot instructed them to rephrase their question, so they come up with different words to express the same goal. Sometimes this works, and other times they get a response asking them to rephrase (again!). Oftentimes, as in figure 4.1, your users will end up asking for an agent after one or two failures.




Figure 4.1 Accuracy or coverage problems frustrate the user because it takes more time—and sometimes multiple contacts—to achieve their goal. It also causes the user to lose confidence in the virtual agent.

If this is happening to your users, your chatbot most likely has a problem with *accuracy* (the chatbot’s ability to match what it heard against what it knows), *coverage* (the range of topics that your solution is expected to know about), or both. From the outside, it is impossible to tell which is the underlying root cause. For that, you are going to need to collect data. Without that information, it is difficult to know what to fix—and fixing the wrong thing can obfuscate or compound existing problems. Before you know it, your conversational solution becomes costly and difficult to maintain. Worse still, it is not delivering the value it promised (by failing to reduce, or perhaps even increasing, the need for human intervention).

One of the biggest success factors for a chat solution is how an organization approaches the ongoing maintenance of the solution. Ideally, the project sponsor and support team will have set the expectation that the solution needs iterative improvements—especially in the beginning—as it is exposed to more data from real-world users. Despite advances in autolearning, large language models, and generative AI, chatbots don’t tend to magically get better over time.

Expect to commit support resources throughout the bot’s lifecycle

Does the organization feel that a chatbot should be a “set it and forget it” solution? Is there a lack of commitment to the ongoing care and feeding of the virtual assistant? These are the red flags of neglect, and they pretty much guarantee eventual failure.

A chatbot is essentially a digital employee. Much like a human resource, it requires initial training plus occasional retraining, reinforcement, and the opportunity to acquire new skills.


4.1.2 What causes weak understanding?

These are the most common reasons a chatbot will exhibit a decline in understanding:

- Manufactured training data (trained examples that do not reflect a representative user’s vocabulary)
- Insufficient scope or gaps in topic coverage
- New information in the world that is not passed on to the virtual assistant
- Lack of a vetting or gatekeeping process when adding new intents, updating existing intents, or changing model inference parameters

That last point—lack of a gatekeeping process—results in the types of weak understanding problems that are the most difficult to resolve. Without oversight by a knowledgeable owner or a dedicated model-training team, unvetted changes can quickly compound the problem of weak understanding. In traditional classifiers, model updates made by someone who is not familiar with the entire training set often introduce duplications, intent training conflicts, and unjustified disparities in the volume of training examples across intents. Untested model parameter or prompt changes will cause unexpected behavior in a generative model.

In fact, we saw this happen with a client who had been making changes to their classifier training set, growing the total intents from 21 to 53 over the course of nine production deployments. The business did not see an effect right away; rather, over time the result of these untested changes manifested as poor survey results, incomplete journeys, unnecessary escalations, and lots of negative feedback. Subject matter experts reported that the bot was giving wrong answers for questions that it used to get right. These are classic symptoms of weak understanding, but they could not pinpoint exactly when it all started. A series of retroactive experiments against their prior versions told the story, which is shown in figure 4.2.



1. A client launched a pilot with 21 intents. No quantitative measurements were taken, but user testing was reported as satisfactory. The initial accuracy (measured retroactively) was 79.7%.
2. We tested each major version, looking for clues as to why users were so unhappy with the current state. The first major classifier performance decline happened when dropped added two intents in V6; the accuracy dropped to 66.4%.
3. From there, they continued adding intents without realizing the performance was spiraling downward until the accuracy of the current state reached 55%.

Figure 4.2 A retroactive assessment of classifier performance shows a hard-won lesson on the effect of untested changes over time. Had each version been tested as part of a predeployment process, the team would have postponed any version updates until the classifier problems were resolved. It took several weeks to get the classifier back into good working order.

4.1.3 How do we achieve understanding with traditional conversational AI?

Traditional (non-generative) conversational AI systems are taught by ingesting examples of user requests grouped by *intents*, sometimes referred to as *classifications* or *clusters*. Intents contain a variety of paraphrases that all express the same goal. Some systems also incorporate *entities*, which are like keywords that further refine the meaning or specifications of a request.

The conversational logic is configured to identify an intent (or a combination of intent + entity) and take an action based on that identification. This action could be as simple as answering a question, or it could initiate a complex transactional exchange. Table 4.1 shows examples of intents, entities, and potential next steps in a conversational exchange.

Table 4.1 Example utterances may be handled differently based on the presence or absence of entities.

Utterance	Intent	Entity	Possible next step
"How many bags can I check?"	Bag_Allowance		Display bag check policy
"I want to book a flight"	Book_Flight		Collect destination
"I need a one-way ticket to Costa Rica"	Book_Flight	Costa Rica	Collect departure details
"I'd like to upgrade my seat to first class"	Flight_Upgrade	first class	Initiate upgrade process

The types of bots that use traditional classification technology tend to be topic routing agents, question/answer (FAQ) bots, and, to some extent, process-oriented (self-service) assistants. Keep in mind that classification-based bots rely on a predefined set of question topics (intents). You need to know in advance what questions you expect your bot to encounter.

As a matter of practicality, the range of topics or intents that you teach your system will be specific to your domain and your solution's use case or purpose. As solution owners, one of our primary and ongoing tasks is to tune our system to correctly understand the greatest volume of user demands. Finding the ideal balance between topic breadth and topic depth can be difficult and often involves tradeoffs. For example, it is not cost effective to train a classifier to understand every possible topic. Furthermore, attempting to do so can weaken its understanding of topics that are salient to your users.

When an organization tries to train a classifier to detect every possible topic, the classifier's ability to see clear distinctions across all intents can be diminished. If the intents trained in your system aren't representative of user demand (meaning you have a large number of low-volume topics), they tend to cause problems with accuracy and confidence. Figure 4.3 illustrates a "long tail" chart; the greatest business value for

a classifier-based chatbot is typically realized by focusing on the high-to-moderate volume requests. Low-volume requests are typically handled by some sort of fallback mechanism, such as escalation, search, or generative AI.




Figure 4.3 As request volume tapers off to the right, the chart has the appearance of a long tail. Each use case must define the optimum tradeoff of depth and breadth as it relates to topic coverage. The cutoff point for business value is usually somewhere in the moderate-volume range. This is not to say that all low-volume requests should be excluded, but there may be diminishing returns associated with extending your classifier's coverage for these topics.

Prior to initial launch, you need to make some predictions about which topics will be most important for your bot to understand. These predictions are often based on logs from human interactions, call center metrics, focus groups, surveys, or other research or information-gathering activities. Your focus should be on training your model to be good at recognizing these requests, as well as any other ancillary conversational maintenance intents (such as greetings, chitchat, repeat, and escalate). Once your solution is in production, you'll need to validate those predictions by collecting and analyzing data about your conversational interactions.

4.1.4 *How do we achieve understanding with generative AI?*

How does a generative AI model achieve understanding? This is a trick question, because generative AI does not so much understand the meaning of an utterance as it creates new data that looks like the data it was trained on, using the utterance as a reference point. This is a nuanced distinction, but with generative AI, we try to simulate understanding by instructing a model to assess the input from a certain viewpoint and then generate a specific type of output.

Particularly for conversational AI, our goal is to generate output that reflects or addresses the user's request with specificity and/or personalization (not just a high-level categorization, such as topic classification or entity extraction). Figure 4.4 demonstrates the fundamental difference between classification model outputs and generative model outputs.




Figure 4.4 Traditional classification models use supervised learning to predict one of several predefined classifications. They look for the intent, or meaning, of a user input. Generative models use decoding transformers to create a text completion. They predict the next sequence of tokens (loosely, words or characters) that are most likely to occur after the user input.

A quick note on LLM foundation architectures

Encoder-only architectures are best for non-generative use cases, such as training predictive models based on text embeddings. They focus on extracting meaningful context from inputs and require labeled data for fine tuning.

Decoder-only architectures are designed explicitly for generative AI use cases. They are “trained” in an unsupervised fashion by ingesting large amounts of data. They focus on predicting the next token in a sequence and can be instructed to perform specific tasks, including classification, question answering, and summarization.

Some LLM model architectures are *encoder-decoder*, which means they can support both generative and non-generative uses cases. These are typically used in scenarios where the input is large, but the output is relatively small, such as translation or summarization.

Unlike traditional classifier-based AI, there is no predefined list of intents that are “in-scope” for the generative model. But like traditional AI, you still need to have good command of the domain and of the range of problems your users are likely to bring to your bot. This will inform the strategies you employ that nudge your LLM to

produce responses demonstrating that the user's input was understood. There are several effective tools at your disposal to accomplish this:

- *Selecting the right model for the job*—Some models are more optimized for conversational output (as opposed to generating code or writing an essay or news article).
- *Prompt engineering*—This technique supplies a model with inputs in order to produce optimal outputs. These inputs might include instructions, context, input data, and output indicators. Prompt engineering can often achieve a good simulation of understanding, and it can instruct the model to produce output in a conversational tone.
- *One-shot or few-shot prompting*—You can enhance your prompt with one or more examples of the output and format you want the model to generate.
- *Parameter tuning*—Parameters such as temperature, top- p , and top- k influence the randomness and diversity of the generated text. Increasing these values tends to increase the “creativity” in a generated response.
- *Retrieval-augmented generation (RAG)*—RAG can enhance the perception that the bot understands while keeping the generated answer grounded in your domain. Many businesses employ RAG in their conversational solutions to ensure that the generated responses are based on external, verifiable facts and the latest information.

At the time of writing, enterprise conversational solutions most often employ generative AI as question-answering (Q&A) bots. Most business-oriented chatbots that use this technology are not fully generative—they often employ a hybrid approach of classification (with predefined response pairs), task-oriented flows, and generated responses. Generated responses may be incorporated into the dialogue design, invoked as a fallback option (e.g., when classification fails to predict an intent with sufficient confidence), or both.

Generative AI can also be used to enhance classification response outputs by inserting a personalized greeting or problem summary before delivering a “canned” (preconfigured) dialogue response or launching a task flow. Done well, this can engage a chatbot user on a deeper level, exhibiting “understanding” with empathy by acknowledging the user's specific situation.

Exercises

Reflect on the solution you are currently building or supporting. Ask yourself these questions:

- 1 Is my solution exhibiting any symptoms of weak understanding, such as
 - Giving wrong answers, especially answers that are not relevant or are completely unrelated to the input topic
 - Taking the fallback/anything else/escalation routes more often than you would expect

- Disambiguating, or clarifying the topic, more often than you would expect on a seemingly straightforward request (for solutions that employ a disambiguation feature)
 - Producing outdated or incorrect information
 - Receiving negative feedback or poor NPS scores
- 2 How was my solution originally trained and tested? If it was deployed, was a baseline measurement taken?
 - 3 Has the solution been updated to recognize new topics and produce answers that are accurate and current?
 - 4 Who is allowed to make changes to the solution? Are these changes documented? Is the solution monitored after a change to ensure the change produces the intended effect?

4.2 How is understanding measured?

Understanding, for a chat solution, is typically measured in terms of accuracy. For a classifier, that means an ability to accurately predict the intent. For generative models, it is the ability to create correct and useful output. There are multiple methodologies and tools for measuring how well a solution understands user inputs. The approach you take will depend on which technologies your solution uses (traditional, generative, or both) and what phase you are currently in (predeployment or post).

4.2.1 Measuring understanding for traditional (classification-based) AI

Classifier performance is measured in terms of accuracy, precision, and recall. *Accuracy* is the percentage of correct predictions that were made. *Recall* refers to the classifier's ability to identify the correct intent, while *precision* is the classifier's ability to refrain from giving a wrong intent. Higher accuracy usually correlates to a perception of "good understanding." A chatbot can't deliver a predefined response or invoke the correct process-oriented flow if it does not understand the user's intent.

You can assess your classifier's performance using some data science techniques, such as *k*-fold cross validation or blind testing. *Blind testing* refers to the fact that a given test utterance does not already exist in the training set; i.e., the classifier has not "seen" the utterance before. Your test set may be manufactured, such as with AI-generated data, or representative (constructed from actual user utterances pulled from logs). *K*-fold and blind tests can provide information about your model's overall accuracy, as well as report on its recall and precision. The metrics produced by such tests help identify where the model is performing well and where it might be confused. Chapter 5 contains detailed instructions for improving classifier understanding, so we will just give an overview of the testing approaches here.

MEASURING UNDERSTANDING WITH K-FOLD CROSS VALIDATION

If your chatbot has not yet been deployed, a *k*-fold cross validation test is the easiest and most accessible method for measuring accuracy because it does not require additional annotated data. It uses only your existing training set. This method essentially

measures the internal consistency of your data labeling—a high accuracy score mainly indicates that your training examples were grouped with other similar examples. The process involves pulling a percentage of data out of training, creating a temporary blind test set. The remaining data is used to create a temporary classifier. Next, each blind example is run against the classifier, and the predictions are scored. Finally, the temporary blind set is folded back into the training set. This process is iterated k times so that every example is used as a training example and as a test utterance, but never both at the same time.

A k -fold test will give you a prediction of the accuracy of your classifier, assuming the data you used to train your model is representative of the inputs your model will encounter when it is deployed to production. However, this can lead to a false sense of security, especially if your training data is highly manufactured or does not quite resemble actual user utterances. Another caveat is that small datasets can produce unreliable measurements if there isn't enough data to withhold examples for testing while still training each intent with minimally sufficient examples. For these reasons, k -fold testing is not the preferred testing method once your solution is in production.

MEASURING UNDERSTANDING WITH AI-GENERATED BLIND TEST DATA

Obtaining test data through a generative process is done through the same means as obtaining generated training data: you prompt a model to generate variations of examples and use them as a “blind” test set. This method is best suited for predeployment but may also be appropriate in the early go-live phase to supplement gaps in your production logs.

Like k -fold testing, the validity of your accuracy measurements is wholly dependent on whether the test data closely mirrors the inputs your model receives at production runtime. This approach can be vulnerable to bias and over-fitting. As such, we advise caution and suggest you validate your generated data against production logs as soon as they are available.

MEASURING UNDERSTANDING WITH REPRESENTATIVE BLIND TEST DATA

If your chatbot has already been deployed, the production logs are one of your key tools for assessing your chatbot's accuracy. These logs contain truly representative data about what your users ask for and how they phrase these requests. By “representative,” we mean both a realistic volume distribution of the intents triggered in your system as well as utterances that capture the user's goal—in whatever combination of words comes naturally to them.

Using production logs will produce the least biased testing data, but it also requires a degree of upfront, manual effort. That effort does pay off, however, as you will have created a reusable asset for taking measurements of future changes. You'll need to obtain a sample of these logs and review the customer inputs (utterances) against the intents returned by your system. This data will need to be annotated by a human who can identify the definitive correct (aka “golden”) intent that the utterance belongs to. Your initial annotations will give you a baseline accuracy. This data

will then be used to build your *representative blind test set*, which is essentially a list of test questions and the answer key all in one file.

SELECTING THE BEST METHOD FOR YOUR SITUATION

The cost and effort tradeoffs for each method are entirely dependent on the size and current phase of your solution (predeployment or post):

- *K*-fold cross validation may be seen as “cheap and easy” because it does not require human annotation beyond the task of the initial annotation done for training purposes. However, there may be an API cost to running your experiment *k* times. This cost is usually negligible for smaller systems but could result in thousands or tens of thousands of API calls per experiment for larger systems.
- Generated test datasets incur the cost of generating the data in addition to the API cost of running an experiment.
- A representative blind test set may have a lower API cost for running an experiment (compared to *k*-fold, assuming your test set is smaller than your training set), but the cost of human annotation can be significant. This also requires that the solution is in production, interacting with real users. The benefit is that the experiment results are going to be more meaningful than *k*-fold and generated test set results.

In summary, there are three primary methods to measure your classifier’s ability to understand users. The method you choose should align with your current stage of development or deployment, as outlined in figure 4.5.

Classifier test method	Applicable phase
<i>K</i> -fold cross validation	Predeployment
Generated test data (blind)	Predeployment, early postdeployment
Representative blind test data	Postdeployment

Figure 4.5 *K*-fold cross validation and generated test data are suitable for situations where representative data is not available. Once a solution is deployed to production, representative blind test data will produce the most reliable measurement of your classifier’s ability to understand.

4.2.2 Measuring understanding for generative AI

Measuring whether a generated answer has demonstrated “good understanding” is an onerous task, and automated test approaches are still emerging. Our challenge is the nature of generative AI: every generated response is possible or likely to be unique to each user input.

Before you deploy a solution with generative AI, you should define what it looks like for your bot to demonstrate good understanding. For generative conversational AI, we suggest you define “good understanding” by the following dimensions:

- The generated answer matches any specified output format or style, including
 - Positioning of the bot (the purpose and viewpoint of the bot’s persona)
 - The designated tone and personality of the bot’s persona
- The generated answer is appropriate for the user’s input in terms of content length and structure (for example, does the nature of the user input require a response that is a short answer, step-by-step instructions, or a detailed essay?).
- The generated answer is free from false information (hallucinations).
- The generated answer is free from hate, abuse, profanity, bias, and discrimination.
- The generated answer is free from damaging information—even if true—such that a company would be legally liable or incur damage to their reputation (for example, negative commentary about a competitor or leaking sensitive data).
- The generated answer is resilient to prompt-injection attempts.
- The generated answer is correct and complete and either successfully terminates a flow or progresses the flow to a next step or the next best action.

If your solution has already been deployed, obtain a representative sample of your logs. Perform a manual review to assess your bot’s level of understanding. Each generated answer will be judged as correct, sufficient, or appropriate against the dimensions you have defined for the solution.

This is, of course, time-consuming, but the effort will pay off. Your annotated set can be used as a golden test set for future improvements. This test set will give you a baseline for tracking the effect of changes to your model parameters (such as temperature, top p , top k) and other LLM configuration settings. These samples can also inform any few-shot examples (sample inputs paired with desired outputs) you include in your prompt engineering or fine-tuning.

4.2.3 *Measuring understanding with direct user feedback*

One way to measure good understanding at scale is to incorporate an answer feedback mechanism directly in the user experience, such as a thumbs up/down reply option. This method can be used for both traditional and generative solutions.

Be mindful of how often you solicit feedback, and know what purpose your feedback serves. Which aspect of the experience is the rating meant to reflect: satisfaction or dissatisfaction with a particular answer (for a question/answer use case), the self-serve process and its outcome (for a process-oriented bot), or the conversational experience as a whole?

Exercises

- 1 Explore and document your solution (or review and update it as needed), with emphasis on the components most responsible for demonstrating understanding:
 - For classifiers, this means auditing the training data.
 - For solutions that include search and retrieval, audit the source documents or URLs, any supplemental document enrichments, and the ingestion schedule to ensure that your knowledge base contains the most relevant and up-to-date information.
 - For generative AI solutions, audit the dialogue flows that invoke generated responses, and map the prompts, parameters, and LLM settings to their intended outcomes.
- 2 Reflect on your current test methodologies, if any. Do you have any historical test metrics that can be correlated to current symptoms of weak understanding?
- 3 Think about the test methodologies presented in this section. Which approach is optimal for the current phase of your solution lifecycle?

4.3 Assessing where you are today

Before you start making plans for improvements, you will want to perform an assessment of where the solution stands in terms of its ability to accurately identify the users' goals and needs. The nature of your assessment will depend on which technology your solution uses. Classification and generative models perform very different functions and therefore have different aspects to be assessed.

4.3.1 Assessing your traditional (classification-based) AI solution

For traditional AI, start by reviewing the training set to orient yourself to the domain and current scope:

- How many classifiers are used in your solution?
- How many different intents does the system (or each classifier) handle?
- How unique is each intent?
- Do the training examples in any intent seem to overlap with other intents?
- Does the range of topics (intents) align with your impression of the chatbot's purpose?
- How does the solution handle input it does not understand?
- What is the complexity of the dialogue? Are there complex flows, backend integrations, or search integrations?

It can be helpful to visualize your classifier training data volume in chart form. Figure 4.6 shows an example training set. There isn't a lot of information to be gleaned just yet, but this will give us a basis for comparison once we assemble our test data.




Figure 4.6 This classifier has 13 intents. The training example counts range from 7 to 30.

In general, we expect our intents with higher counts of training examples to be more popular. We want our most popular topics to be understood a majority of the time. Higher-volume intents may also represent topics that handle a greater variety of phrases. For the most part, we don't like to see a huge disparity in volumes across the training set. For instance, a training set that has some intents trained with hundreds of examples while others have just a handful might exhibit performance problems such as over-selection (frequently selecting a wrong intent due to the bias of training volume).

4.3.2 *Assessing your generative AI solution*

With generative AI, as in traditional AI, you need to understand the domain and scope your bot operates within. However, instead of concerning yourself with classifications of input, you need to appraise the data sources that your model will draw its answers from when it produces an output. Is generative AI used to produce answers or responses in your solution? If so, familiarize yourself with the circumstances:

- Are answers generated for every user input?
- Do you call for generated answers as a fallback option for your classifier?
- Do you call for generated text to supplement a classification-based “canned” answer in the dialogue?
- Does your solution make use of more than one LLM, such as different models for different types of responses, multiple language support, etc.?
- Does your solution make use of prompt engineering, prompt tuning, fine tuning, or other customized settings? Is this documented anywhere, along with the outcome goals for which each setting was originally implemented?
- Does your solution make use of RAG? If so, what is that data source? How often is it updated? Does it contain additional data enrichments?

Exercises

- 1 Assess your solution using the criteria we described in this section (according to the type of AI you use).
- 2 Once you have performed your initial solution assessment, be sure to document its current state—this will be your baseline system configuration.
- 3 As you follow along with the improvement recommendations and examples given throughout this book, be prepared to record your changes in a way that will help you correlate any updates you make to the subsequent performance measurements.

4.4 Obtaining and preparing test data from logs

For the rest of this chapter, we'll assume that you do have a production system and access to the logs. We will show you how to obtain and prepare that data to create an asset you can use to measure the current state (and to validate future changes).

There's a bit of initial work involved to build a test set from production logs. Figure 4.7 shows the major tasks involved in preparing data for testing (or training).




Figure 4.7 Once you obtain some data, each utterance should first be sorted into buckets to identify potential candidates; this will separate the good, usable test data from the bad or irrelevant user inputs. The data may also need to be scrubbed to fix problems like personal identifiable information (PII). After that, the data will need to be annotated (for classifiers, it will need to be labeled with the correct intent; for generative AI, it will need to be associated with an ideal output response). Finally, the data will need to be converted into one or more sets that can be consumed by your testing tool.

4.4.1 Obtaining production logs

Ideally, you will have access to production logs that span a full year or more. This will help ensure that your test set will have a true representative sample of the range of topics your bot encounters for the various seasons and events that influence your industry. Collect log samples from various weeks or months throughout the year. If your solution is newer, expect to refresh your test sets more frequently during your solution's first 12 to 18 months.

Once you have obtained some production logs, you may find it easiest to convert this data into a CSV or Excel file (if it hasn't already been updated). We find it most useful to transform the data into one row per conversational exchange (a user input and a bot output), grouped by conversation ID. Depending on the timeframe you select, the volume of users, and the complexity and purpose of your solution, your file

may have just a few hundred rows of data, or it could have 100,000 rows or more of conversational exchanges.

One simple shortcut for reducing the volume to a manageable set is to select the first user utterance in each conversation. This may not work in all cases, but figure 4.8 shows that it is often a reliable way to harvest useful data from your logs. In a natural language-driven exchange, users tend to express their most important need in the initial input. If your average conversation lasts ten turns, a conversation log with 100,000 rows of raw data could be reduced to about 10,000 rows of data to review. Deduplication can often further reduce this by a few thousand. This is a very workable volume and will usually contain rich and diverse examples that you can use for testing your solution.

Conv ID	Turn	User Input	Chatbot Output
1001	1	I want to open an account.	What type of account do you want to open?
1001	2	a savings account	How much do you want to deposit into your new savings account?
1001	3	\$500	Ok, let's get you started with a new savings account.
1002	1	My bank card was stolen.	I can freeze your card. Do you have any suspicious charges?
1002	2	yes	Do you know when your card was stolen
1002	3	It must have been yesterday afternoon.	Ok, I can submit a fraud alert and have a new card sent to you.
1003	1	What is my account balance?	I can look up your balance, what is your account number?
1003	2	123456791	Sorry, I couldn't find that account, would you like to try again?
1003	3	yes	What is your account number?
1003	4	123456719	Your account balance is \$547. Is there anything else you need?
1003	5	Yeah, do I have any late fees?	You currently owe \$42 in late fees.
1004	1	Hi there,	Hello. Is there something I can help you with?
1004	2	I need a copy of my bank statement.	I can get your bank statement. What billing period do you want to see?
1004	3	Last month	A copy of your statement from last month has been sent to your inbox.
1005	1	Speak to customer service.	To get you to the right person, please tell me more about what you need.
1005	2	I need to report a stolen credit card.	I can freeze your card. Do you have any suspicious charges?

Figure 4.8 Raw chat logs show that a user's primary goal is often captured in the first turn of a conversation, but sometimes it occurs as an additional request later in the conversation or after exchanging a pleasantry. It might even follow an opt-out request. Selecting the first row will usually yield enough usable data while reducing the amount of time your annotators spend sorting through the utterances that aren't useful to the classifier, such as button clicks, common responses, and PII or other user-specific information. (The structure of your logs may vary by tool.)

4.4.2 Guidelines for identifying candidate test utterances

Whatever you do to obtain and preprocess your logs, your next task is to identify potential blind test candidates. We treat this as a “first pass” exercise: just determine if an utterance is *potentially* usable. Additionally, we counsel the reviewers not to over-analyze what they see; if you cannot make a determination about any given utterance within a minute or so, discard the utterance and move on. (If you're feeling really conflicted or sense a pattern, mark it for later review and move on.) We use the following criteria to identify potential test candidates from production logs, along with any special handling instructions:

- Is the utterance unintelligible?
- Is the utterance completely unrelated to the domain?

- Is the utterance ambiguous?
- Does the utterance contain multiple intents?
- Is the utterance related to the domain but out of scope?
- Does the utterance express a goal that is in domain and in scope?

Let's look at each of these in turn.

IS THE UTTERANCE UNINTELLIGIBLE?

Maybe a cat walked across the keyboard, or the user just mashed the keys in a fit of frustration. Perhaps the speech-to-text technology mistranscribed the caller's question into an unintelligible mess. Speech solutions can also pick up background noise and conversations, especially if they are not properly tuned for the environment. Your file may contain a number of user inputs that just don't make any sense.

These are examples of unintelligible or unrelated utterances:

- "does it school" (Incoherent—if this came from a voice solution, it was potentially a speech mistranscription.)
- "she didn't she said there are four and only gave us one yes you can do that I'm about to catch my flight and I'll check on it when I get to the office" (Potentially a speech transcription of a background conversation on the caller side.)
- "klewtkhacalifornia liense" (Likely typos, severe enough to render the utterance unintelligible.)

These lines can be excluded from your blind test set. Any recognizable patterns, such as possible speech transcription problems, should be set aside for further evaluation or forwarded to the appropriate team.

IS THE UTTERANCE COMPLETELY UNRELATED TO THE DOMAIN?

You may occasionally come across questions that are intelligible but entirely off-topic for the domain or the bot's intended purpose. For example, if your solution is designed to help electric utility customers manage their account and services, you can exclude questions about pop culture trivia if they happen to appear in the logs.

Though you could configure a solution to send unrecognized topics to an LLM, these utterances do not belong in your classifier test set because a golden intent cannot be assigned. Such utterances could be used in negative testing, which will help you understand if your solution is appropriately identifying when it should *not* attempt to answer.

IS THE UTTERANCE AMBIGUOUS?

Perhaps you'll find a single word or a short phrase that is related to the domain but doesn't express a clear goal. For example, if a user of a banking chatbot simply says, "account," what do they want? Do they want to open an account? Close an account? Check an account balance? Who knows?

A subset of ambiguous utterances may include responses generated by a button click or as part of an information-gathering flow. (If you selected the first natural language utterance of every conversation, you might not see these.) These are generally

not useful for the classifier’s performance testing unless they align with an intent that is used within a flow. Include such utterances only when appropriate.

These are examples of ambiguous utterances:

- “driver license” (Perhaps relevant to the domain, but no clear goal is expressed.)
- “that one” (An anaphor referring to contextual information that appears to have been provided in an earlier statement but may have lost its meaning as an individual utterance.)
- “2” (Could refer to a button choice or phone channel selection, or to an amount or quantity provided as a response to the previous question.)

In most cases, these utterances should not be included in your classifier accuracy test because they likely will not align with any single intent, but rather multiple intents. They are not meaningless, however. Set these aside to understand how often your users communicate in this way. Determine whether your other chatbot features, such as disambiguation or clarifying questions, are handling them appropriately.

DOES THE UTTERANCE CONTAIN MULTIPLE INTENTS?

Most classifier-based chatbots perform best when they are given one goal at a time. Utterances that express multiple valid, distinct goals should be excluded from your classifier accuracy test set because you cannot definitively assign a “correct” intent.

The exception to this rule would be if your solution has a disambiguation mechanism. Disambiguation is a way to clarify the user’s primary goal by presenting the top n intents identified by a classifier. For these solutions, you may want to run your multi-intent utterances against your classifier to verify that all intents listed would be presented with the appropriate disambiguation choices.

These are some examples of utterances with multiple intents:

- “Do you have the COVID booster? How can I make an appointment?” (Two goals expressed: 1) availability of vaccine booster, 2) make an appointment.)
- “I want to update the address on my driver’s license and find out what is required to get a commercial driver’s license.” (Two goals expressed: 1) update address, 2) get information for obtaining a CDL.)
- “I currently have 95,000 loyalty points. Do they expire? How many more points do I need to reach Platinum status? Can I purchase points for this?” (Three goals expressed: 1) find out if reward points expire, 2) find out the delta between current point balance and next level reward status, 3) get information about purchasing points to reach a higher status.)
- “I want to talk to an agent about reporting a stolen vehicle.” This is very common. A user will often pair a request for an agent along with their true goal. If both intents exist in your classifier training set, you can handle such utterances in one of two ways:
 - Exclude these as candidates if it is impossible to label a single “correct” intent.
 - Include these candidates, but label them according to the “preferred” intent. (A preferred intent might be the self-service option if containment is a priority and the competing intent would escalate.)

As with ambiguous utterances, these should be set aside and evaluated separately to better understand your users. You may want to devise additional strategies to handle these situations if they are occurring often. If users tend to ask related questions, or they pair common requests in a single utterance, your output responses in these intents could be updated to anticipate or meet all of the needs. For the first example we gave—“Do you have the COVID booster? How can I make an appointment?”—your answer regarding booster availability may include a link to make an appointment.

A word about handling multiple intents with classification models

We have seen extensive and heroic attempts to handle multiple intents programmatically in conversational AI solutions. This usually involves logic to collect the top n intents and store them in context, and then more logic to present the additional topics after the first one is answered. In most cases, the result is an over-engineered solution that is brittle, difficult to scale, or simply wasted effort. This approach also has a major flaw: such logic cannot reliably distinguish between an utterance that truly contains multiple goals and an utterance that contains a single goal that may have triggered multiple intents.

Many modern chatbot frameworks provide automated topic disambiguation (for example, “Did you mean: [Intent 1] [Intent 2] [Intent 3]”). Our general recommendation is to allow the disambiguation feature to do its job. Sometimes, this means that the user must ask their questions or state their goals one at a time. The frequency and importance of such scenarios is usually not worth the effort required to build and maintain custom logic for handling multiple intents in a classification-based chatbot.

Generative AI is typically much better at handling multiple intents than classification-based solutions, so you can include these utterances as candidates in a test set if your solution has this capability.

IS THE UTTERANCE RELATED TO THE DOMAIN BUT OUT OF SCOPE?

You are likely to come across utterances that express a single, clear goal that is relevant to the domain, but the current solution is not equipped to handle them. For example, a banking chatbot may allow users to check an account balance but may not be trained to recognize requests about interest rates. An airline chatbot may be versed on airline policy but not be grounded in facts about airport security.

Such questions may be very reasonable from the user’s perspective, and gaps in topic coverage often lead to frustration for your users. This is especially true if you don’t have a generative AI or search fallback. If your bot responds, “I’m sorry, I don’t understand. Please rephrase your question,” no amount of rephrasing will get the user to a satisfactory answer. How should these be handled?

If your classifier does not have any trained intents to handle such requests, these should set aside. On further review, they may be grouped into topics or categories, but they will be excluded from your test set for now because a golden intent cannot be assigned. Monitor these topics for volume and add them to your improvement backlog as appropriate.

Similarly, if your generative solution is not prepared to answer such questions (for example, the document repository in a RAG solution does not have content to address the topic), set these aside for the time being, but monitor the volume.

DOES THE UTTERANCE EXPRESS A GOAL THAT IS IN DOMAIN AND IN SCOPE?

Score! Questions or requests that are in scope for your solution and domain belong in your golden test set.

4.4.3 Preparing and scrubbing data for use in iterative improvements

If you’ve never seen production logs for a chatbot, you will be surprised at how messy they are. You are going to see a lot of bad or informal grammar, misspelled words or typos (on a text-based channel), speech mistranscriptions (on a voice channel), and potentially various forms of personal identifiable information (PII). Here’s how we recommend handling these.

BAD OR INFORMAL GRAMMAR

For the most part, leave it be! There is a lot of diversity in how humans express themselves. The user may not know exactly how to communicate what they need—especially to a machine. If a goal can be identified, it is a representative example and should be generally preserved as is.

TYPOS AND MISPELLED WORDS

Unless a typo or misspelled word significantly changes the meaning of the overall phrase, leave it as is. Commonly misspelled words are representative of how your users communicate. Your classifier should be able to give a good answer whether the user asks, “What’s the difference between loan balance and principal?” or “whats teh diffrnce between loan balance and principle?”

Proper case and punctuation are generally ignored by a classifier, but you may need to verify this with your technology platform.

SPEECH MISTRANSCRIPTIONS

If your solution uses speech-to-text (aka automated speech recognition), you won’t encounter typos, but you probably will see unexpected words that are most likely the result of a speech mistranscription. The first line of attack is to train your speech models, if possible. The underlying technology of a chatbot classifier is text-driven, so it is best to have the most faithful representation of the user’s utterance before it hits the text classifier.

If you find that the speech models are still consistently mistranscribing words that are significant within your domain, include these in your test set (and ultimately, you will probably end up supplementing your training data). For example, for an electric utility company, we consistently saw an important domain term, “residential,” mistranscribed as “presidential.” As speech model updates can take longer to implement, and this was causing loss of call containment, an immediate fix was to add “presidential” as a synonym to our chat solution. Another example was the mistranscription of “VIN” as “BIN” for a use case that needed to understand “vehicle identification number.” For

this, we made sure that the training data contained both variations. We also preserved the mistranscriptions for our testing purposes.

PERSONAL IDENTIFIABLE INFORMATION

You may also find various forms of PII, such as names, phone numbers, physical or email addresses, social security numbers, account numbers, etc. These do not belong in your training or test data. Ideally, this information would be masked in your logs, but even this technology is not perfect. If your solution has a PII masking function, you should replace any real data with the same type of masking characters (e.g., #####-#### for a ten-digit phone number). If not, either remove the PII entirely, or replace it with an obviously fictionalized representation, such as “username@email.com.”

4.4.4 The annotation process

After you have narrowed down your data to utterances that express a clear goal that belongs in your domain (and have scrubbed them where appropriate), they need to be properly annotated for the task at hand.

ANNOTATING A GOLDEN TEST SET FOR TRADITIONAL (CLASSIFIER-BASED) AI

Annotating a test set for a classifier involves labeling each utterance with the appropriate intent. This task is a little easier said than done, and it’s where you will spend the most time building your test set.

It’s fairly easy to identify and discard an unintelligible or ambiguous user input. However, once you know an utterance belongs in your domain, it takes a bit more time to label it with the correct intent. The person or team tasked with annotating (labeling each utterance with the correct intent) will need to be familiar with the current training data. This process will definitely expose problems with overlap in your intents, as your human annotators will be stuck with the question of how to label an utterance.

A team might take several approaches to complete the work of labeling data for testing or training. Sometimes a single person is tasked with the job. Sometimes a whole team will try to take this on. When that happens, they often think that a “divide and conquer” approach is most efficient. In our experience, this can lead to problems that take longer to resolve.

In an ideal world, everyone would sit in the same room and judge each utterance together. This approach facilitates discussion regarding the purpose of each intent. All annotators need to understand the criteria used to differentiate intents that share a lot of the same key words but have different goals. Another equally valid approach is to have multiple annotators judge the same data separately (or at least a percentage of overlapping data) and compare any differences to reach a resolution.

There is one shortcut we wouldn’t hesitate to take if your logs include the intent that was predicted at runtime for each utterance: make a first pass and judge whether the predicted intent was correct. Then you need only judge and label the remaining incorrect utterances with the correct intent.

This exercise may take anywhere from a few hours to several days, and it can be taxing on both your vision and your cognitive load. As a first run, instruct your annotators

to make their best judgment and move on. If it takes more than sixty seconds to judge an utterance, skip it and come back later. It is also important to take breaks every hour or two. It helps to walk away and come back after a period of refresh.

Could I just use an LLM to do all that work?

If you are building your first classifier, you could certainly run utterances against an LLM as a first pass at labeling or classifying your data. However, if you already have production logs, there will be no added value to running the utterances against a separate classification LLM because you still need a human judge to review the classifications produced by this exercise.

Once you have annotated the test set, you will have a golden set of human-judged, labeled data. Depending on your use case, this could include a few hundred to a few thousand utterances. This asset will give you some immediate information about your classifier's current accuracy. It will also be used to help tune your system.

The last thing you need to do is convert your data into a file that can be consumed by your testing tool. This will produce an asset that can be used to measure the effect of future updates. The format may vary by tool, but it will typically be a text or CSV file that contains a row for each test utterance in one column and the golden intent in the other column. Table 4.2 shows how a test set might look.

Table 4.2 Sample test set with one utterance/intent pair per row

Utterance	Golden intent
I want to speak with a real person	Request_Agent
Can I talk to a manager	Request_Agent
Get me customer service	Request_Agent
Are you open on Sundays	Office_Hours
What time do you open	Office_Hours
When does your office close	Office_Hours
What are your weekend hours	Office_Hours

ANNOTATING A GOLDEN TEST SET FOR GENERATIVE AI

Creating a test set to measure generative AI involves judging the quality of the answer produced by your solution (if you are working with production logs) and updating or replacing it with the ideal answer, according to the dimensions you previously defined for your solution. Subject matter experts will need to review each example output to ensure that it is factual and complete, represents the brand, and reflects the purpose and positioning of the virtual agent persona.

Once you have reviewed the output, you will have a set of utterances paired with a golden answer or response. This asset will give you some immediate information about the quality of your generated responses. It will also be used to tune your prompts and LLM configurations.

The last thing you need to do is convert your data into a file that can be consumed by your testing tool. The format may vary by tool, but it will typically be a text or CSV file that contains a row for each test utterance in one column and the golden response in the other column. Table 4.3 shows a sample test set.

Table 4.3 Sample test set with one utterance/answer pair per row

Utterance	Golden response
Can I bring a snowboard on my flight as checked baggage?	You can bring one set of snowboard equipment as a checked bag. The set must be in one bag and can include up to two snowboards and one snow boot bag. If the set weighs more than 50 pounds (23 kg), you'll have to pay overweight bag fees.
How long do I have to wait to get my refund?	Credit card refunds will be processed within five business days of the request. All other refunds will be processed within 20 business days of the request.

Exercises

- 1 Obtain data from your own logs, and identify candidate test utterances.
- 2 Scrub the data as needed to remove PII.
- 3 Assess the classification predictions or generated answer content. Record these outcomes as baseline performance measurements.
- 4 Assign a golden intent or ideal response.
- 5 Save the file in a format that can be consumed by your testing tool.

4.5 What does the data tell us?

If your logs included the original intent prediction or generated answer, you now have what is needed to calculate a baseline measurement of your solution's current accuracy rate for understanding. (Divide the number of correct predictions or answers by the total candidates judged.) Your annotated utterances will show you the range and frequency of topics your users present to the chatbot.

4.5.1 Interpreting annotated logs for traditional (classification-based) AI

For classifier-based systems, you might be interested in looking at the volume distribution across your intents. How does this compare to your training example volumes for each intent? Figure 4.9 shows an idealized, fairly balanced distribution of training examples compared to occurrences seen in the logs.




Figure 4.9 The dark bars represent the number of training examples in a system. The light bars represent the number of annotated utterances for each intent. If your chart follows a similar pattern, your training priorities are probably in good alignment with the demands of your solution.

A stark disparity between trained examples and actual occurrences in the logs is not indicative of problems in and of itself, but it can inform your priorities if your accuracy is low. Figure 4.10 shows an example of annotated utterances that are wildly out of alignment with how the system was trained.




Figure 4.10 The training example counts (dark bars) show a large disparity across many intents, as compared to the annotated log data (light bars). Without accuracy numbers for each intent, we cannot immediately tell if this disparity is having a negative effect. However, we can make some observations, such as 1) the first five intents are not nearly as important to our users as we thought they might be, and 2) the intents with the highest volume in our logs (the light bars for intents 6, 10, 11, and 12) may be a lot more important to our users than we predicted.

You should also review the volume of utterances that were judged to be in domain but out of the current scope. (These would have been identified and set aside as part of the preparation tasks described in section 4.1.4.) Does there appear to be a demand for topics that the classifier is not currently trained on? A misalignment between what your users expect to be able to ask and what your classifier is trained to recognize contributes to a perception of weak understanding.

Your overall accuracy provided a big-picture view of the solution's ability to understand. The next step is to drill down into the specific intents. You might start by looking at the poorest performers that are also high-volume/high-value in your solution. In chapter 5, we will explore in depth the process for improving classifier understanding.

4.5.2 Interpreting annotated logs for generative AI

Your annotated logs for a generative AI solution will give you a picture of the range of questions and requests that users are providing. Throughout the annotation process, you may have discovered gaps in coverage about the domain. You may also have gained a better grasp of how prompt engineering or fine-tuning improvements could make your generated answers better. If your solution employs RAG, you might start correlating the quality of your answers to the documents in your repository.

Your overall accuracy provided a big-picture view of the solution's ability to understand. In chapter 6, we will explore in depth the process for improving your generative AI so that it conveys good understanding.

4.5.3 The case for iterative improvement

At this point, you should be armed with the data you need to begin planning improvement cycles. Your performance findings will serve as a roadmap for improvements. Keep in mind that this is an iterative process. You will make changes. Then you will take measurements to determine whether your change had a positive, neutral, or negative effect on understanding.

It is also important to note that your blind or golden test set will need to be refreshed throughout the lifecycle of your solution. Recall that one of the reasons a chatbot can become inaccurate is due to new information in the world. These are some examples we have seen:

- The global COVID-19 pandemic, which changed the way nearly everyone worked, navigated public spaces, and supported their families.
- New legislation passed, resulting in government organizations getting related questions.
- New products on the market or product recalls.
- A company experienced a data breach, and once the news broke, the chatbot was bombarded with questions like, "Is my data safe?" and "I want to know more about the hack."

Plan to review your logs on a regular basis. Depending on the volume of your solution, that might start out daily right after launch, then weekly, monthly, and quarterly. Don't forget to update your test sets according to the changes you make:

- If new intents are added to your system, new utterances need to be added to your test set.
- If intents were merged or split as part of your improvement efforts, the affected intents will need to be updated in your test set.
- If new areas of coverage are added to the knowledge base your generative solution references, your test set should include validations for this.
- If your solution adds new LLM scenarios or prompt customizations, these should be reflected in the test set.

Exercises

- 1 Review your annotated data and reflect on the findings. Are there areas that show poor understanding?
 - If so, what would you hypothesize is the root cause?
 - Is there more than one root cause?
- 2 How would you prioritize the improvements needed to achieve better understanding?

Summary

- Chatbots demonstrate good understanding when they identify what a user wants and they provide a satisfactory answer or progress the user toward their goal.
- For traditional AI, understanding relies on at least two mechanisms: correct classification of an intent and an ability to deliver an output based on that classification. (Additional mechanisms, such as entity detection or context, may modify or personalize outputs.)
- For generative AI, understanding relies on the utterance and any accompanying prompt to create a response meant to address a user's question or goal.
- Weak understanding is detrimental to business value and is often exhibited by a chatbot returning wrong answers or no answers at all.
- You can't assess the performance of your chatbot without first collecting some data.
- Chatbot understanding is usually measured in terms of accuracy or the rate at which the solution delivers a correct answer or takes the correct action.
- There are multiple tools and methods for measuring understanding. Some are dependent on the type of AI and/or the current phase, whether predeployment or post.
- A representative golden test set, curated from real user utterances (production logs), can be used to measure the bot's baseline performance and can be converted into a reusable asset to measure the effect of future changes.
- You should plan to monitor and retrain your solution throughout the life of the bot.
- Updates to training may require corresponding updates to the blind test set.

5

Improving weak understanding for traditional AI

This chapter covers

- Identifying the types of errors a classifier can make
- Establishing a baseline of current classifier performance
- Using data science methodologies to identify and prioritize improvements
- Infusing your traditional AI with generated content to enhance understanding

In this chapter, we will demonstrate a methodical, iterative approach to improving the understanding of a classification-based conversational solution. This chapter builds on the concepts introduced in the previous chapter and uses the output produced by the final exercise in section 4.4 (where you created a test set with the golden intent assigned to each utterance in a format that can be used by your testing tool). Later in this chapter, we'll explore how large language models can supplement intent-driven output responses to deliver a more robust experience. (If you're looking for generative AI improvement techniques, feel free to skip ahead to the next chapter.)

We will start by building an improvement plan and identifying the types of errors your classifier may be committing. Next, we'll iterate through seven improvement cycles to solve the various problems you might see in your own text classifier. Although data science techniques are used, you do not need to be a data scientist to extract meaningful insights about your data using the methodologies presented in this chapter.

5.1 Building your improvement plan

If you built a blind test set using a sample from your production logs, you should have a reliable “representative distribution” test set. This means that the topics that are most frequently asked by your users are represented with corresponding volume in your testing data. This will be a key factor in prioritizing any problems that are surfaced by your test results.

If you are working with the results of a k -fold test (discussed in chapter 4), you won't know for certain which topics are the most important, so the most egregious accuracy scores are a logical starting point.

In either case, it's now time to dig into those test results. An improvement plan starts with identifying the biggest problem spots in the bot's training.

5.1.1 Identify problematic patterns in misunderstood utterances

The first score that will grab your attention is the overall accuracy of your test results. This is a lot like getting back a spelling or math test and looking at the red ink at the top of the page. If your test had 100 questions and you got 79 of them correct, your accuracy score would be 79%. For classifiers, this number is good for an “at a glance” view of the model, but it doesn't give a complete picture of what is going on or where to start making improvements. For that, we need to understand the possible outcomes and types of errors our classifier may be committing. This is revealed in the measurements of recall, precision, and F1 score.

A BRIEF EXPLANATION OF RECALL, PRECISION, AND F1 SCORES

In chapter 4, we described *recall* as the classifier's ability to predict a correct intent and *precision* as the ability to refrain from predicting a wrong intent. You can think of this in terms of positive and negative predictions. For every utterance that we test against the model, there are four possible outcomes, and they are not mutually exclusive, meaning that every prediction is going to have two or three of these outcomes happening simultaneously. Figure 5.1 shows a confusion matrix that visualizes these possible outcomes:

		Actual intent	
		Positive	Negative
Predicted intent	Positive	True positive	False positive
	Negative	False negative	True negative

Figure 5.1 In a 2×2 confusion matrix, the possible outcomes are derived by comparing the predicted intent to the actual intent.

- *True positive*—A prediction that matches the correct intent
- *True negative*—A prediction that does not match an incorrect intent
- *False positive*—A prediction that matches an incorrect intent
- *False negative*—A prediction that does not match the correct intent

The first metric that might interest us is the recall of our intents. For this, we need to know the true positives and the false negatives. An intent that is returning false negatives is committing an error of under-selection. When measured per intent, this looks like an accuracy score. If our test had five questions for the `#Request_Agent` intent, and the classifier got those questions correct four times, the intent’s recall would be 80%:

$$\text{Recall} = \text{True positives} / (\text{True positives} + \text{False negatives})$$

The next metric that helps us understand our classifier is precision. This measures how good our classifier is at refraining from giving a false positive. An intent that is returning false positives is committing an error of over-selection. An example of over-selection can be seen in the last two rows of table 5.1:

$$\text{Precision} = \text{True positives} / (\text{True positives} + \text{False positives})$$

Table 5.1 Test results show seven utterances, five of which are labeled with the correct `#Request_Agent` intent. The first four predictions were true positives. The last two rows show where `#Request_Agent` was predicted twice for utterances where it shouldn’t have been (“What can I ask you” and “Somebody hit my car”). These false positives contribute to our precision calculation: $4 / (4 + 2) = 0.66$.

Utterance	Correct intent	Predicted intent	Correct
Customer service	Request_Agent	Request_Agent	1
Speak with an agent	Request_Agent	Request_Agent	1
Can I please speak with somebody?	Request_Agent	Request_Agent	1
Talk with a human	Request_Agent	Request_Agent	1
When will I get a live person?	Request_Agent	Office_Hours	0
What can I ask you?	VA_Capabilities	Request_Agent	0
Somebody hit my car	Report_Accident	Request_Agent	0

A full analysis of all possible outcomes for `#Request_Agent` is shown in figure 5.2. It also shows the true negatives (which are not used in our calculations but have been included to demonstrate the range of other outcomes).

Utterance	Correct intent	Predicted intent	True positive Request_Agent	False positive Request_Agent	False negative Request_Agent	True negative Request_Agent	True negative Office_Hours	True negative VA_Capabilities	True negative Report_Accident
Customer service	Request_Agent	Request_Agent	X				X	X	X
Speak with an agent	Request_Agent	Request_Agent	X				X	X	X
Can I please speak with somebody?	Request_Agent	Request_Agent	X				X	X	X
Talk with a human	Request_Agent	Request_Agent	X				X	X	X
When will I get a live person?	Request_Agent	Office_Hours			X			X	X
What can I ask you?	VA_Capabilities	Request_Agent		X		X	X		X
Somebody hit my car	Report_Accident	Request_Agent		X		X	X	X	

Figure 5.2 The highlighted columns are used for calculating precision and recall for #Request_Agent.

Now that we know the recall and precision, we can also calculate the F1 score, which is the harmonic mean of recall and precision. This calculation is made as follows:

$$F1\ score = (2 \times Precision \times Recall) / (Precision + Recall)$$

For our #Request_Agent intent, this would be calculated as $(2 \times 0.66 \times 0.8) / (0.66 + 0.8) = 0.72$. Table 5.2 shows all three scores.

Table 5.2 Recall, precision, and F1 score for #Request_Agent

Intent	Recall	Precision	F1 score
Request_Agent	0.80	0.66	0.72

What about the true negatives?

Earlier in this section, we mentioned true negatives—a prediction that does not match an incorrect intent. True negatives occur whenever we have more than one trained intent. However, they are not a useful measurement in our methods.

Why not? Well, for every prediction the model makes, there is only one way for it to be right, but there are two ways for it to be wrong. This seems a little unfair, and it's hard to see why if you're just looking at two intents. But imagine we have a model that was trained with 20 intents. Whenever we make a single prediction that returns a true positive, we will also get 19 true negatives. And for every false positive prediction,

we have 1 false negative and 18 true negatives. So all those true negatives add up to a very large number that, for our purposes, doesn't give us much insight. Therefore, we don't factor true negatives into our calculations.

DECIDING WHICH METRIC IS IMPORTANT

Recall, precision, F1 score: Which number should we care about? That's a great question! The answer is that it depends on what your organization values most in terms of what the solution needs to deliver. Here are some considerations to guide you to an answer:

- Recall is useful when there is a high cost associated with false negatives. Imagine the effect if a fraud detection tool missed 25% of the fraudulent transactions it evaluated. (For a chatbot, this would look like a correct intent that is not predicted 25% of the time.)
- Precision is useful when there is a high cost associated with false positives. Think of the gameshow Jeopardy!, which penalizes a contestant for attempting to answer and getting it wrong (or a chatbot that over-selects the #Request_Agent intent, resulting in unnecessary escalations).
- The F1 score is useful when there is a high cost associated with both false positives and false negatives. We like to use this for most implementations because it reflects a good balance of the recall and precision scores.

VISUALIZING YOUR DATA WITH A CONFUSION MATRIX

Earlier in this section, we showed a 2×2 confusion matrix to demonstrate the potential outcomes. A confusion matrix can help you assess the performance of a classification model by visualizing a summary of the predictions made by your model. Some testing tools produce this with their results output.

Figure 5.3 shows a fictional scenario where a classifier model made ten perfect predictions.




Figure 5.3 A solid diagonal line shows that each predicted intent (represented by a single letter) matched to the actual intent.

Shaded boxes that stray from the diagonal provide useful insights about where your model is confused, as shown in figure 5.4.

		Actual intent									
		A	B	C	D	E	F	G	H	I	J
Predicted intent	A										
	B										
	C										
	D										
	E										
	F										
	G										
	H										
	I										
	J										

Figure 5.4 This model had nine correct predictions, but wrongly predicted intent G when the actual intent was E.

5.1.2 Incremental improvements

An incremental improvement approach will affect measurable change in a manageable way. Every change you make to a classifier has the potential to affect multiple intents. Sometimes this effect is positive, but sometimes it's not. You might get away with updating several intents all at once, but if the testing shows a performance decline, it can be difficult to track down the culprit. You will have to balance the need for efficiency with your tolerance for rework.

5.1.3 Where to start: Identifying the biggest problems

Generally, the best place to start is with the highest volume intents that have the lowest F1 scores. The business may also weigh in on priorities. If a lower volume intent fails to recognize the type of request it was designed to handle, but this failure incurs costly human intervention, it might take priority.

For the rest of this chapter, we will explore a fictional use case: a chatbot that serves a population that interacts with a state's Bureau of Motor Vehicles (a type of US government agency that regulates and manages the issuance of state identification cards, driver's licenses, certain permits, and vehicle registrations).

To begin, let's follow the advice given in chapter 4 and take a quick, high-level look at our current training data, as laid out in table 5.3.

Table 5.3 Intents with example counts in a baseline training set

Intent name	Number of examples
Accident_Report	2
Appointment	6
Change_Contact_Records	3
Chitchat_Goodbye	3
Chitchat_Hello	4
Chitchat_Thanks	2
Chitchat_VA_About	8
Fee_Info	5
General_Negative_Feedback	6
General_Request_Agent	5
Get_ID_Number	4
Item_Not_Received	8
License_or_ID	5
License_Reinstatement	4
Login_Issue	4
Name_Change	6
Office_Information	6
Payment_Methods	3
Refund_Overcharge	4
Report_Sold_Vehicle	6
Report_Stolen_License_Permit_ID	5
Report_Stolen_Plates_Registration	3
Report_Stolen_Vehicle	2
Request_Receipt	4
Vehicle_Permit	5
Vehicle_Title	6
Walk_In	6
Grand Total	125

We can make some quantitative statements about this training set. It has 27 intents with a grand total of 125 training examples. The examples are distributed fairly

evenly. As a qualitative assessment, we might say that many of the intents appear to be unique, but a few of them might have some overlap. Some terms definitely overlap across intent names. A peek at the full set of training utterances (not shown) revealed that many terms appear in multiple intents, such as “ID,” “title,” “permit,” “vehicle,” “stolen.” However, as shown in table 5.4, the contexts in which these words appeared were judged to be appropriately labeled.

Table 5.4 Utterances extracted from the baseline training set show a variety of terms overlapping across multiple intents.

Utterance	Labeled intent
How much is an ID?	Fee_Info
I need to find out my ID number	Get_ID_Number
I didn't receive my ID	Item_Not_Received
Title never came	Item_Not_Received
Add a person to the title	Vehicle_Title
How do I get a driving permit?	License_or_ID
Replace my program parking permit	Vehicle_Permit
I sold a vehicle	Report_Sold_Vehicle
I need to report a stolen car	Report_Stolen_Vehicle
My ID was stolen	Report_Stolen_License_Permit_ID

Overall, it seems that the range of topics is reasonable for the chatbot's purpose, which in this case is to answer questions a user might have when dealing with a state's Bureau of Motor Vehicles.

ESTABLISHING A BASELINE

Now that we have made an initial assessment of our training data, we need to understand how it is currently performing. We'll start by running a *k*-fold cross validation test to establish a baseline. The results, our first version (V1) shown in table 5.5, are not that bad, considering the low volume of data present in the training set.

Table 5.5 Baseline (V1) *k*-fold results

Intent	Number of samples	Number of predictions	Recall	Precision	F1 score
Accident_Report	2	2	1	1	1
Appointment	6	8	1	0.75	0.8571
Change_Contact_Records	3	0	0	0	0
Chitchat_Goodbye	3	0	0	0	0

Table 5.5 Baseline (V1) *k*-fold results (continued)

Intent	Number of samples	Number of predictions	Recall	Precision	F1 score
Chitchat_Hello	4	6	1	0.6667	0.80
Chitchat_Thanks	2	2	1	1	1
Chitchat_VA_About	8	8	1	1	1
Fee_Info	5	2	0.40	1	0.5714
General_Negative_Feedback	6	7	1	0.8571	0.9231
General_Request_Agent	5	4	0.80	1	0.8889
Get_ID_Number	4	6	1	0.6667	0.80
Item_Not_Received	8	6	0.6250	0.8333	0.7143
License_Reinstatement	4	4	1	1	1
License_or_ID	5	5	0.60	0.60	0.60
Login_Issue	4	4	1	1	1
Name_Change	6	8	1	0.75	0.8571
Office_Information	6	6	1	1	1
Payment_Methods	3	3	1	1	1
Refund_Overcharge	4	4	1	1	1
Report_Sold_Vehicle	6	5	0.8333	1	0.9091
Report_Stolen_License_Permit_ID	5	6	1	0.8333	0.9091
Report_Stolen_Plates_Registration	3	3	0.3333	0.3333	0.3333
Report_Stolen_Vehicle	2	3	1	0.6667	0.80
Request_Receipt	4	4	1	1	1.0000
Vehicle_Permit	5	6	1	0.8333	0.9091
Vehicle_Title	6	9	1	0.6667	0.80
Walk_In	6	4	0.6667	1	0.80

Our *k*-fold test had a total of 125 questions (the grand total of our training set), and it got 105 of them correct, for an overall accuracy of 84%. Several intents had perfect recall and perfect precision (which is often a hallmark of a manufactured data set). There were two intents that had a recall of 0; they each had only three training examples. This reveals one of the flaws of *k*-fold testing—there simply weren’t enough examples to distribute across the auto-generated train and test sets. More than likely, those intents will perform better than 0 in production. However, the intents with

perfect recall will probably not perform quite as well. If you are launching a pilot and have no other training data available, these results are generally good enough to go live, with a strong caution to the stakeholders that they should expect lower actual performance until representative data is available for use in training updates.

Once the solution is live, a new baseline should be taken using the blind test set you created from the logs. We have an example of this in table 5.6, and it really emphasizes the gap in performance predicted by our k -fold test compared to real user inputs.

Table 5.6 Baseline (V1) blind results

Intent	Number of samples	Number of predictions	Recall	Precision	F1 score
Accident_Report	2	2	1	1	1
Appointment	7	5	0.7143	1	0.8333
Change_Contact_Records	4	4	1	1	1
Chitchat_Goodbye	1	1	1	1	1
Chitchat_Hello	1	1	1	1	1
Chitchat_Thanks	1	1	1	1	1
Chitchat_VA_About	1	2	1	0.50	0.6667
Fee_Info	11	9	0.8182	1	0.90
General_Negative_Feedback	2	3	1	0.6667	0.80
General_Request_Agent	3	2	0.6667	1	0.80
Get_ID_Number	3	5	1	0.60	0.75
Item_Not_Received	16	9	0.4375	0.7778	0.56
License_Reinstatement	5	5	0.60	0.60	0.60
License_or_ID	7	5	0.5714	0.80	0.6667
Login_Issue	9	5	0.4444	1	0.6153
Name_Change	9	9	1	1	1
Office_Information	9	11	1	0.8182	0.90
Payment_Methods	2	2	1	1	1
Refund_Overcharge	3	4	1	0.7500	0.8571
Report_Sold_Vehicle	6	7	1	0.8571	0.9231
Report_Stolen_License_Permit_ID	7	8	0.8571	0.75	0.80
Report_Stolen_Plates_Registration	5	4	0.80	1	0.8889
Report_Stolen_Vehicle	2	2	0.50	0.50	0.50

Table 5.6 Baseline (V1) blind results (continued)

Intent	Number of samples	Number of predictions	Recall	Precision	F1 score
Request_Receipt	4	4	1	1	1
Vehicle_Permit	4	5	1	0.80	0.8889
Vehicle_Title	2	8	1	0.25	0.40
Walk_In	8	5	0.3750	0.60	0.4615

On the first run of our blind test, 102 questions were correct out of 134, for an overall accuracy of 76%—8 points lower than the 84% predicted by our k -fold test.

VALIDATING YOUR INITIAL TRAINING STRATEGY

Once you have obtained annotated logs and taken some baseline performance measurements, you can validate the decisions that informed your initial training strategy.

Scarcity of representative training data is a very common problem for conversational AI projects. Just like many other newly launched chatbots, our initial training set was developed by subject matter experts (SMEs) who manufactured training examples for the topics they believed would occur most frequently. In figure 5.5, we can compare the number of examples trained for each intent to the number of examples that were present in the randomly selected logs used for testing.




Figure 5.5 A comparison of training examples to the utterances in our representative blind test set shows that there is some disparity in volume for many of the most popular intents (the representative blind utterances) on the left side of the graph. We also see disparity across several of the least popular intents (those on the right).

A side-by-side volume comparison of training data to representative blind utterances per intent can help us understand if our solution’s topic coverage is in alignment with the real-world interactions. One of the first observations we noted was that `#Item_Not_Received` was the most popular real-world intent. This validated the initial build strategy of supplying that intent with a higher number of training examples (relative to most other intents). We also noted that `#Chitchat_VA_About` had a high number of training examples compared to how infrequently this topic came up in the logs. This intent may be over-trained. It certainly doesn’t seem to be as popular as we thought it might be. Yet, until we look at the performance metrics for these intents, we cannot draw any solid conclusions. Rather, these observations might inform our improvement recommendations.

Exercises

- 1 Run a representative blind test using your own data, and identify which intents, if any, exhibit poor performance.
- 2 Does your training volume align with the intent volume seen in your logs?
- 3 How would you prioritize improvements for the poorest-performing intents?

5.2 Solving “wrong intent matched”

When your chatbot returns the wrong intent, it has committed two categories of errors: false positives (predicting the wrong intent), and false negatives (failing to predict the right intent). Let’s walk through an improvement cycle to demonstrate how we would approach this problem.

5.2.1 Improve recall for one intent

We will start with `#Login_Issue`, which was the fifth most popular topic but had a considerably low recall of 0.44. There were nine test utterances in our blind set; it got four questions correct (true positives) and five incorrect (false negatives). This intent had a perfect precision score, which means it never showed up as a wrong prediction for other intents. Table 5.7 shows the summary metrics.

Table 5.7 Summary metrics for `#Login_Issue`; a blind test set run against our baseline classifier shows low recall but perfect precision.

Intent	Number of samples	Number of predictions	Recall	Precision	F1 score
Login_Issue	9	5	0.4444	1	0.6153

In table 5.8, we can drill down to the result details of the blind test. Our classifier failed to predict a correct intent five times. Three of those were predictions of a wrong intent. Two were instances where confidence was so low that the classifier did not return a prediction.

Table 5.8 Baseline blind result details for #Login_Issue show that we had a recall score of 44%. Out of nine utterances, the correct (aka *golden*) intent was predicted five times.

Utterance	Golden intent	Predicted intent	Confidence
BMV portal password reset	Login_Issue	<none>	n/a
I can't get on my profile	Login_Issue	Item_Not_Received	0.8131
I need help logging into my BMV profile	Login_Issue	<none>	n/a
I never got my security verification code	Login_Issue	Item_Not_Received	0.2358
I tried logging in and it didn't work	Login_Issue	Login_Issue	0.8033
I'm not able to get into the portal	Login_Issue	Login_Issue	0.6680
Password locked out	Login_Issue	Login_Issue	0.5520
Password reset	Login_Issue	Login_Issue	0.4875
You never sent a security code	Login_Issue	Item_Not_Received	0.2091

If we look at our current trained examples, it's easy to see why so many questions were missed. There were only four examples:

- I'm unable to log in on the website
- Online account problem
- Online problems
- Problem signing onto my account

Our training examples lack the variety of meaningful words and phrases seen in interactions with real users. Users might refer to their account as their “profile.” They list explicit problems such as being “locked out,” needing a “password reset,” and failing to receive a “security code.” We should expect to see an improvement if we add a few representative examples (obtained from our logs):

- Help signing in to online portal
- I need to reset my password
- I need a security code to log on

With these additions, we updated our classifier to V2 and reran the blind test set. Let's look at how this affected the recall for #Login_Issue in table 5.9.

Table 5.9 Blind test result details show improved recall for our newest classifier version (V2). Out of nine utterances, the correct (aka *golden*) intent was predicted eight times.

Utterance	Golden intent	Predicted intent	Confidence
BMV portal password reset	Login_Issue	Login_Issue	0.8253
I can't get on my profile	Login_Issue	Item_Not_Received	0.8131
I need help logging into my BMV profile	Login_Issue	Login_Issue	0.6846

Table 5.9 Blind test result details show improved recall for our newest classifier version (V2). Out of nine utterances, the correct (aka *golden*) intent was predicted eight times. (continued)

Utterance	Golden intent	Predicted intent	Confidence
I never got my security verification code	Login_Issue	Login_Issue	0.7179
I tried logging in and it didn't work	Login_Issue	Login_Issue	0.8899
I'm not able to get into the portal	Login_Issue	Login_Issue	0.7840
Password locked out	Login_Issue	Login_Issue	0.9083
Password reset	Login_Issue	Login_Issue	0.9204
You never sent a security code	Login_Issue	Login_Issue	0.2551

Our overall accuracy improved from 76% to 79% (106 out of 134 correct), and table 5.10 shows a dramatic improvement in the recall and F1 score. The precision score also remained steady.

Table 5.10 A comparison of summary metrics; our V2 classifier shows an overall improvement compared to the baseline (V1) for #Login_Issue.

Intent	Number of samples	Number of predictions	Recall	Precision	F1 score
Login_Issue—Baseline (V1)	9	5	0.4444	1	0.6153
Login_Issue—V2	9	8	0.8889	1	0.9412

5.2.2 Improve precision for one intent

Next, let's experiment with improving the precision for an intent. The #Chitchat_VA_About intent remained unchanged between the baseline test results and the V2 test results. (It is important to look at the newest results after each change.) Table 5.11 shows that the recall was perfect, but the precision was only 50%. This means our classifier is placing a bit more importance on this topic, and it is showing up as a false positive (over-selecting) in another intent.

Table 5.11 Metrics after the V2 update show that #Chitchat_VA_About has perfect recall but poor precision.

Intent	Number of samples	Number of predictions	Recall	Precision	F1 score
Chitchat_VA_About	1	2	1	0.50	0.6667

In table 5.12, we see that there was only one test question in our blind set for this intent, but our classifier predicted the intent twice.

Table 5.12 V2 blind result details show an over-selection for #Chitchat_VA_About.

Utterance	Golden intent	Predicted intent	Confidence
Do you have a name?	Chitchat_VA_About	Chitchat_VA_About	0.8042
Where are my tags?	Item_Not_Received	Chitchat_VA_About	0.3015

Our training has eight examples. We knew that these examples were manufactured (in fact, they were provided by a template), but our logs show that this is not a very common topic. Our blind test set only contained one utterance for this intent.

One strategy for improving precision is to prune the training examples. This tells our classifier that the intent isn’t quite as dominant as the other intents within our solution. We’ll discard three examples because they are either overly redundant or, in the case of “Where are you from,” there was no evidence in the logs that this was a relevant question:

- Are you a robot?
- What can I ask you?
- What can you do?
- What can you help me with? (REMOVE FROM TRAINING)
- What’s your name?
- Where are you from? (REMOVE FROM TRAINING)
- Who am I talking to? (REMOVE FROM TRAINING)
- Who are you?

Once the training was updated (now V3), we ran the blind test again and reviewed the results. We saw an improvement to the precision for the #Chitchat_VA_About intent from V2 to V3—it was a perfect score across all metrics. Oddly enough, our overall accuracy dropped to 78% (from 79%), and one of the questions we lost was from our #Login_Issue intent. Table 5.13 shows the changes in metrics from V2 to V3 for both intents.

Table 5.13 Metrics before and after V3 update for #Chitchat_VA_About and #Login_Issue show that changing one intent can have an effect on another intent.

Intent	Number of samples	Number of predictions	Recall	Precision	F1 score
Chitchat_VA_About—V2	1	2	1	0.50	0.6667
Chitchat_VA_About—V3	1	1	1	1	1
Login_Issue—V2	9	8	0.8889	1	0.9412
Login_Issue—V3	9	7	0.7777	1	0.875

Although #Login_Issue had a slight decline, the current F1 score of 0.875 is still far better than the baseline F1 score of 0.6153. Keep in mind that smaller datasets are more sensitive to small changes, and a change to any intent can potentially affect every intent. Those changes may have negative or positive results. Instead of focusing on this, however, we will make a few more changes elsewhere and check back to see if the intent improves.

5.2.3 *Improve the F1 score for one intent*

Let's move forward with improving the F1 score for #Item_Not_Received. Table 5.14 shows that it had an F1 score of 56% after our V3 update.

Table 5.14 After V3 update, the F1 score remained unchanged at 0.56 for #Item_Not_Received.

Intent	Number of samples	Number of predictions	Recall	Precision	F1 score
Item_Not_Received—V2	16	9	0.4375	0.7777	0.56
Item_Not_Received—V3	16	9	0.4375	0.7777	0.56

The intent had eight training examples, but our logs showed that this is a very popular topic, so we need it to perform much better. We'll add 10 more examples from our logs to that intent (now V4) and run another experiment.

Table 5.15 shows that our recall for this intent has now more than doubled, and though the precision fell slightly, the F1 score is greatly improved. The classifier's overall accuracy also increased from 78% to 81%.

Table 5.15 Before and after metrics for #Item_Not_Received show an improved F1 score.

Intent	Number of samples	Number of predictions	Recall	Precision	F1 score
Item_Not_Received—V3	16	9	0.4375	0.7777	0.56
Item_Not_Received—V4	16	19	0.875	0.7368	0.8

5.2.4 *Improve precision and recall for multiple intents*

Sometimes there is confusion due to a heavy overlap of terms across intents that have similar goals. Figure 5.6 shows the confusion matrix that our testing tool provided.

In our model, we see a fair amount of confusion across the intents that relate to stolen items. One solution to this problem is to merge intents. This must be considered carefully. The intents were probably created separately by design, as they all have different answers. However, entity detection can be used to route the flow to the appropriate answer.




Figure 5.6 Confusion matrix after the V4 update. The density in shading represents the volume of questions predicted for a given intent. If a classifier test had a perfect accuracy score, you would see a solid black diagonal line running from the upper left corner to the lower right corner. The shaded squares that stray away from this diagonal line mark the areas of confusion within your model.

We’ll merge all of these into a single intent called #Report_Stolen. These examples are listed in table 5.16. Don’t forget that the blind test set will need to reflect this change, as well as the related dialogue flows.

Table 5.16 Examples from three intents to be merged into a new #Report_Stolen intent

Intent name	Training example
Report_Stolen_Vehicle	Report a stolen car
Report_Stolen_Vehicle	I need to report a stolen car
Report_Stolen_Plates_Registration	My plates were stolen
Report_Stolen_Plates_Registration	My registration was stolen
Report_Stolen_Plates_Registration	License plate stolen off vehicle
Report_Stolen_License_Permit_ID	Stolen real ID
Report_Stolen_License_Permit_ID	Wallet was stolen
Report_Stolen_License_Permit_ID	My drivers license was stolen
Report_Stolen_License_Permit_ID	My ID was stolen
Report_Stolen_License_Permit_ID	My permit was stolen

The conversational flow will be updated so that when a defined entity value or synonym is detected in an utterance, the corresponding original answer is provided. You may also need a default condition to disambiguate or provide a generic answer in case an utterance triggers the new intent but no entity is detected. Table 5.17 is an example of what that might look like.

Table 5.17 Dialogue updates using entity detection for the new #Report_Stolen intent

Entity/synonym detected	Treatment
vehicle, car, truck, motorcycle	Routes to original answer for #Report Stolen Vehicle
plates, registration, tags	Routes to original answer for #Report_Stolen_Plates_Registration
ID, license, permit	Routes to original answer for #Report_Stolen_License_Permit_ID
(none detected)	Disambiguate ("It sounds like something was stolen; can you tell me what it was?")

With these changes, our classifier is now on V5. Table 5.18 shows the metrics for the three old intents under V4 and the metrics for our new intent in V5.

Table 5.18 Metrics before and after V5 update show that merging three intents into the single #Report_Stolen intent results in perfect scores across the board for this topic.

Intent	Number of samples	Number of predictions	Recall	Precision	F1 score
Report_Stolen_License_Permit_ID—V4	7	5	0.5714	0.8	0.6666
Report_Stolen_Plates_Registration—V4	5	4	0.8	1	0.8888
Report_Stolen_Vehicle—V4	2	2	0.5	0.5	0.5
Report_Stolen—new intent in V5	14	14	1	1	1

Our latest change dramatically improved the performance of this topic, and it bumped the overall accuracy to 85%, which is now higher than our baseline k -fold (which was 84%).

With that update complete, we can move on to other intents that need improvement. Following the iterative processes, we updated the remaining intents that showed the poorest performance by adding a few more examples from the logs. This became V6 of our classifier. Table 5.19 is an overview of the intents that were updated.

Table 5.19 Training example counts increase from V5 to V6 and more closely align with the volume present in the representative blind test set.

Intent	V5 training example count	V6 training example count	Test utterances in representative blind
License_or_ID	5	6	7
License_Reinstatement	4	6	5
Login_Issue	7	8	9
Walk_In	6	8	8

This update resulted in an overall accuracy of 92% for the latest classifier (now on V6). In the world of natural language classification, this is a very good score for a representative blind test set. You will never achieve 100%; even human-to-human communications don’t come close to that.

Every data set is different, and we could spend several more cycles tweaking our training if there is plenty of data available. However, there are diminishing returns associated with pursuing results that approach 100%. There is also a risk of over-fitting your model to the current blind test set. Once additional logs become available and a new test set is created, you may discover additional gaps (or your overfitting will be exposed).

Table 5.20 shows a comparison of the blind test F1 scores of the baseline classifier against our latest updates. Twelve of the intents did not change (and they were already

performing very well). One intent decreased from 90% to 80%, and the remaining 14 intents showed improvement. We felt that this was a good and reasonable tradeoff, improving more than half of our intents at the cost of one intent showing a slight decline.

Table 5.20 Comparison of baseline F1 scores and V6 F1 scores

Intent	Baseline (V1) F1 score	V6 F1 score	Change
Accident_Report	1	1	(no change)
Appointment	0.8333	0.833	(no change)
Change_Contact_Records	1	1	(no change)
Chitchat_Goodbye	1	1	(no change)
Chitchat_Hello	1	1	(no change)
Chitchat_Thanks	1	1	(no change)
Chitchat_VA_About	0.6667	0.9524	+ 0.2857
Fee_Info	0.90	0.80	- 0.1
General_Negative_Feedback	0.80	0.80	(no change)
General_Request_Agent	0.80	0.80	(no change)
Get_ID_Number	0.75	0.8571	+ 0.1071
Item_Not_Received	0.56	0.8750	+ 0.315
License_Reinstatement	0.60	0.75	+ 0.15
License_or_ID	0.6667	1	+ 0.3333
Login_Issue	0.6153	0.9412	+ 0.3259
Name_Change	1	1	(no change)
Office_Information	0.90	1	+ 0.1
Payment_Methods	1	1	(no change)
Refund_Overcharge	0.8571	0.8571	(no change)
Report_Sold_Vehicle	0.9231	1	+ 0.0769
Report_Stolen_License_Permit_ID	0.80	(n/a - merged)	+ 0.2
Report_Stolen_Plates_Registration	0.8889	(n/a - merged)	+ 0.1111
Report_Stolen_Vehicle	0.50	(n/a - merged)	+ 0.5
Report_Stolen	n/a	1	(n/a – merged)
Request_Receipt	1	1	(no change)
Vehicle_Permit	0.8889	1	+ 0.1111

Table 5.20 Comparison of baseline F1 scores and V6 F1 scores (continued)

Intent	Baseline (V1) F1 score	V6 F1 score	Change
Vehicle_Title	0.40	0.8	+ 0.4
Walk_In	0.4615	0.75	+ 0.2885

Exercises

- 1 Using the output from the previous exercise (a prioritized list of your poorest-performing intents), identify the category of error each intent is committing: recall, precision, or both.
- 2 Make iterative training adjustments to improve each intent.
- 3 Measure each change to verify that
 - The intended effect is achieved
 - No other intents were negatively affected

5.3 Solving “no intent matched”

Now that we have our classifier in good shape for the current scope, we can focus on expanding the domain, if needed. During an initial review of your production logs, you will almost surely encounter topics that were not included in the initial training set. Some of these topics will be obvious, but perhaps there wasn’t enough data to train an intent at the time of the initial launch. Maybe the business wasn’t ready to write answers for some topics. Sometimes a seasonal topic is not included because it was not in the forefront of anyone’s mind (e.g., tax season, hurricane season, fiscal year end, etc.). Other topics may be completely unexpected (e.g., a data breach).

Although you don’t have any intents defined to match these utterances, the classifier will always attempt to make a prediction; it doesn’t know what it doesn’t know, so it does its best to match an utterance to what it does know. In an ideal world, the classifier would return very low confidence, and this would trigger an “anything_else” or “no action matches” type of response. In reality, such user utterances often contain words that appear somewhere in your training, so it is possible that the classifier will predict an intent that has training examples with similar words.

5.3.1 Clustering utterances for new intents

In the guidelines described in chapter 4, we recommended setting aside utterances that were related to the domain but not included in the original scope. It’s time to address these.

One of the topics our logs revealed was related to users wanting to cancel their license or registration. We know from our logs how the classifier predicted each utterance at the time the utterance was asked. Now we can test them against our latest classifier (V6) to get new model predictions.

In table 5.21, we see that our classifier exhibited low confidence and/or was incorrect whenever an utterance contained a form of the word “cancel.”

Table 5.21 Unmatched utterances from logs with predictions from the V6 classifier

Utterance	Predicted intent	Confidence
Cancel a registration	Appointment	0.2681
Cancel my car registration	License_or_ID	0.3651
Cancel a drivers license	License_Reinstatement	0.3042
Canceling a registration	Appointment	0.2417
Cancellation of registration	Fee_Info	0.2786
Cancelling my registration	Item_Not_Received	0.3004
Cancel a replacement license	Vehicle_Permit	0.3264
Cancel the license	License_Reinstatement	0.3237
Cancel a title or registration	Vehicle_Title	0.5913
Cancel vehicle registrations	Item_Not_Received	0.2914
Commercial drivers license cancel	License_Reinstatement	0.2995
Driver's license cancellation	Get_ID_Number	0.3387
How do I cancel my vehicle registration?	License_or_ID	0.4324
I need to cancel a vehicle registration	License_or_ID	0.3481
I need to cancel my ID	Get_ID_Number	0.3205
I would like to cancel the registration on my car	Change_Contact_Records	0.3147
I would like to cancel my car registration	License_or_ID	0.3447
I would like to cancel my state identification card	Change_Contact_Records	0.2982
I wanted to cancel a registration	Item_Not_Received	0.3155
I want to confirm cancellation of my registration	Item_Not_Received	0.4092
Questions about cancelling registration for a vehicle	Fee_Info	0.2795
I want to cancel my registration on my pickup	Item_Not_Received	0.4761

We'll randomly divide these into a training set of nine utterances under a new `#Cancel_Registration_or_License` intent and add the remaining thirteen to our blind test set.

When we run the updated blind test set against our updated classifier (now V7), we get an overall accuracy of 92%, which is usually a very good, if not ideal, outcome. This will not always be the case, so if your overall performance drastically drops, you will

need to iterate through the applicable improvement steps (depending on whether the problem was recall, precision, or both) for the intents that were affected.

Let’s walk through one more example of adding a new intent. The logs contained several utterances referring to a data breach. This is an example of how a chatbot can exhibit declining performance due to new information in the world. In this case, the organization had never experienced a data breach before. But when it did, and this news became public, users suddenly had a lot of questions about it. This manifested as unmatched and incorrect predictions, as seen in table 5.22.

Table 5.22 Unmatched utterances on the topic of “data breach” from logs, with predictions from the V7 classifier. The classifier didn’t have enough confidence to match most of the utterances referring to “hack” or “data breach,” which is good because we hadn’t yet taught it anything about that topic. But most of the utterances that contain the word “stolen” match strongly against our #Report_Stolen intent. This may not go so well for the user because our solution doesn’t have any answers yet concerning data that was stolen.

Utterance	Predicted intent	Confidence
I want to know about that hacking on the BMV	<none>	n/a
I want to know about the breach in information at the BMV and if I’m at risk	<none>	n/a
My identity has been stolen	Report_Stolen	0.9483
My license number was stolen	Report_Stolen	0.9240
Need questions answered about data breach	<none>	n/a
No I’m curious about the current breach of stolen IDs	Report_Stolen	0.8604
Someone hacked my information	Report_Stolen	0.4662
Someone is using my drivers license number	Get_ID_Number	0.4067
Someone stole my identity	Report_Stolen	0.7705
Someone stole my information	Report_Stolen	0.8043
Stolen personal identity	Report_Stolen	0.9263
Stolen personal information	Report_Stolen	0.9166
Stolen social security number	Report_Stolen	0.7515
Was my account affected by the recent data hack?	<none>	n/a
Was my account hacked?	Login_Issue	0.3998
Was there a data breach?	<none>	n/a
Yeah I’d like to know if my driver’s license has been breached	Report_Stolen	0.4092
Yes what do I do about the data breach at the BMV?	<none>	n/a
Was my social security number stolen in the hack?	Report_Stolen	0.7806
I want to know if my information was stolen	Report_Stolen	0.9198

To resolve the problem of this unmatched intent, we selected seven representative utterances from the logs to create a new intent called #Data_Breach. Our selection ensured that a variety of important terms, such as “hack,” “breach,” and “stolen,” were added to our new training set. The remaining utterances were added to our blind test set, and we tested our newest classifier, V8. The new #Data_Breach intent returned a perfect score, and the F1 score comparisons in table 5.23 show that nearly all others remained steady or improved since our baseline reading.


Table 5.23 Final score comparison between the baseline (V1) and the final version (V8)

Intent	Baseline (V1) F1 score	V8 F1 score
Accident_Report	1	1
Appointment	0.8333	0.8333
Cancel_Registration_or_License	n/a (NEW)	0.9630
Change_Contact_Records	1	0.8889
Chitchat_Goodbye	1	1
Chitchat_Hello	1	1
Chitchat_Thanks	1	1
Chitchat_VA_About	0.6667	1
Data_Breach	n/a (NEW)	1
Fee_Info	0.90	0.9524
General_Negative_Feedback	0.80	0.80
General_Request_Agent	0.80	0.80
Get_ID_Number	0.75	1
Item_Not_Received	0.56	0.8750
License_Reinstatement	0.60	0.75
License_or_ID	0.6667	1
Login_Issue	0.6153	0.9412
Name_Change	1	1
Office_Information	0.90	1
Payment_Methods	1	1
Refund_Overcharge	0.8571	0.80
Report_Sold_Vehicle	0.9231	1
Report_Stolen_License_Permit_ID	0.80	(n/a - merged)
Report_Stolen_Plates_Registration	0.8889	(n/a - merged)

Table 5.23 Final score comparison between the baseline (V1) and the final version (V8) (continued)

Intent	Baseline (V1) F1 score	V8 F1 score
Report_Stolen_Vehicle	0.50	(n/a - merged)
Report_Stolen	n/a	0.9630
Request_Receipt	1	1
Vehicle_Permit	0.8889	1
Vehicle_Title	0.40	0.6667
Walk_In	0.4615	0.75

Our overall accuracy score remained steady at 92%. (Our updated blind test set has 160 questions, and 147 were correct.) You might recall that our very first blind test had an overall accuracy of 76%, so this is quite an improvement. Our V8 confusion matrix, shown in figure 5.7, also looks improved, with a fairly dark diagonal line.

**Figure 5.7** Comparison of baseline (V1) confusion matrix to the V8 update

We could iterate further to try to get a little higher, but for this use case, the classifier’s accuracy is more than good enough for the time being. Any further tweaks with the limited data we have at present are likely to over-fit our model to the current blind test set. Remember that a healthy strategy is to plan to iterate over the life of the bot, using newer logs and refreshed blind test sets.

5.3.2 When to stop adding intents

When reviewing your logs, you may have encountered a diverse range of other questions that are perfectly reasonable for the domain, but very infrequent. In our logs, we saw questions like the following, but no additional utterances with similar goals:

- I need a form for a doctor to fill out saying a driver is not safe to drive anymore.
- I have a question about electronic signatures.
- What is the process for getting a specialty license plate?

How do we know when to stop adding intents? It's best to let the data from our human-annotated logs guide us. We can total up all of the examples by intent and render them as a chart, as in figure 5.8.




Figure 5.8 Example of a longtail chart. The terms we use to describe the volume distribution of our available training data are “short head” and “long tail.” These terms describe the visual representation of rendering our data on a bar chart. The heavier-volume intents are on the left (the short head), and as the volume decreases for each intent, the data has the appearance of a long tail falling off to the right.

In our longtail chart, we picked a point to divide between what should be in scope versus out of scope. This point isn't a static, prescriptive position. It's a decision that should be made with the business by establishing a minimum number of training examples required to create a new intent. Everything that falls on the left of this line should probably be included in the training, as there is evidence that these topics will be asked more frequently. Everything to the right will not be trained in the current classifier. Over time, you may find enough data in the logs to justify adding a new intent. Until then, your solution will have to handle such topics with one of the following strategies: give a response saying the bot doesn't understand, fall back to an agent escalation, add a search integration to find answers in a document repository, or implement a retrieval-augmented generation (RAG) or large language model (LLM) component to generate answers.

Exercises

- 1 Identify new topics based on your logs, and build new intents from the utterances found in the logs.
- 2 Add utterances to your blind set, and test your changes.
- 3 Is your classifier able to recognize the new intent without negatively affecting the performance of your existing intents?

5.4 Supplementing traditional AI with generative content

In conversational AI, we typically think of delivering either a static answer (as in classic intent-driven implementations) or an answer that is entirely generated (as in a RAG pattern). Static answers fill a need where an answer must maintain consistency, either in content or in structure. Although personalization is possible, it is generally limited to defined entities or other context-driven dialogue conditions. This tends to result in colder, less personalized bot responses. Figure 5.9 shows how three users with the same general goal, but very different personal situations, all receive the same bot response.




Figure 5.9 In a traditional (classification-based) dialogue pattern, an intent is identified, and the dialogue is configured to give a static or minimally personalized answer.

5.4.1 Combining traditional and generative AI for an intent

We can enhance the user experience using a hybrid response pattern, which combines personalized generated content with the static predefined answers written for our intent. Our goal is to acknowledge the user’s problem while ensuring that important information is delivered with consistency. Many large language models excel at summarization tasks, so a model can be prompted to craft an empathetic message that conveys a personalized level of understanding. Figure 5.10 shows what this looks like from the user’s perspective.




Figure 5.10 An output response identifies the correct intent using traditional AI and then prepends generated text to the static output response configured for the intent. The generated greeting and summary convey to the user that the bot understands their goal and the particular details of the user’s situation.

This pattern employs an API call to the LLM as a dialogue step. Content is generated by the LLM and delivered just before the predefined output response. Figure 5.11 shows the high-level steps for such a pattern.




Figure 5.11 LLMs can be called within traditional dialogue patterns to greet a user and summarize their problem before delivering a predefined or static answer.

5.4.2 Prompting to convey understanding

In conversational AI, your bot's role is typically to be a representative of your company. They are a “digital” resource, as opposed to a “human” resource. Still, their job is to be the face of the company. Human agents are great at conveying empathy and understanding. In fact, they will often restate the user's problem to demonstrate that they understand. LLMs can be prompted to simulate this summarization behavior.

Since our traditional AI has already classified the user's intent under this pattern, we can craft a prompt that instructs the LLM to perform a specific task. In this case, we want the LLM to generate a personalized, empathetic greeting that can be paired with additional static content. The next listing shows a prompt instruction for summarizing a user's input.

Listing 5.1 Prompting a model to greet and summarize a user problem

```
<|instruction|>
You are a customer service agent for Friendly Skies Airline. Each input
contains a customer problem. Greet the customer and summarize their
problem.

<|input|>
Hello! This is Chihiro – I had a flight credit for a cancelled flight from
earlier this year. I don't find the credit anymore. Can you look for me
if you can locate it? This is for booking # WKRP01. My frequent flyer #
is 8675309. Thanks a lot in advance!

<|output|>
Hello Chihiro. It seems you had a flight credit for a cancelled flight from
earlier this year and you need assistance locating the credit for
booking number WKRP01.
```

Exercises

- 1 Collect a set of user utterances to test and tune an LLM prompt that can greet a user where appropriate and summarize their problem. Experiment with a variety of instruction prompts. The goal is to create an efficient prompt instruction that will produce good results for the majority of your utterance test set.

Summary

- A classifier's performance can be measured in terms of accuracy, recall, precision, and F1 score. These measurements reflect the types of errors a classifier may be committing.
- The performance metrics produced by your testing will inform your next steps toward improving classifier performance. Higher volume intents with low performance are a good place to start.

- Iterative test and train cycles will show you the effects of your changes.
- A chatbot can use additional strategies, such as disambiguation, clarifying questions, and entity detection to overcome confusion or route answers for merged intents.
- A chatbot with a strong classifier can deliver more business value by delivering the right answers on the first try and deflecting work that would otherwise be handled by a human agent. You should plan to monitor and retrain your solution throughout the life of the bot.
- Generative AI can supplement a traditional AI solution by infusing static chatbot responses with personalization and empathy, which enhances the perception of understanding.

Enhancing responses with retrieval-augmented generation

This chapter covers

- Enhancing chatbot responses without coding intents
- Improving weak understanding with RAG
- Evaluating the advantage of using RAG over traditional search models
- Selecting the proper RAG techniques for your conversational AI
- Assessing and improving the performance of RAG in your conversational AI systems

In previous chapters, we saw the “chatbot doesn’t understand” pain point for question-answering bots. We first addressed it by helping the chatbot understand more intents, but at some point there are diminishing returns to this strategy. Uncommon questions from the “long tail” may never make sense to implement as intents. This chapter introduces ways to handle that “long tail,” including search and retrieval-augmented generation (RAG). These are great methods for improving a chatbot’s weak understanding.

We concluded chapter 5 with advice on when to avoid adding new intents, especially when dealing with diverse, infrequent domain-related problems. In this chapter, we'll add search capabilities to improve weak understanding.

Both search and RAG allow you to improve a chatbot by adding data and documents without programming new intents. This allows you to serve thousands of intents with the simplicity of training just a few. The answers provided by these methods are more straightforward to change—just change the documents rather than changing your chatbot.

Search and RAG can be easier for you as a builder and efficient for your users. Let's explore how a chatbot can evolve using search and RAG capabilities.

6.1 Beyond intents: The role of search in conversational AI

Traditional conversational AI centers on understanding user intents. Systems are trained to recognize predefined categories of user queries and to provide pre-scripted responses.

Figure 6.1 illustrates a conceptual intent-based chatbot architecture in its simplest form. The chatbot's classifier detects the intent and determines the appropriate dialogue flow. When the classifier cannot identify the user's intent, the answer is a generic “I didn't understand” response style, leading to the “chatbot not understanding me” pain point. Intent-based question-answering is a great way to start handling frequently asked questions—you can define an exact answer to be given for different question types. Initially, this is quite effective, but it breaks down. Users often ask questions that deviate from predefined intents. When the predefined intents are insufficient to handle the user's questions, the user may receive irrelevant or incorrect responses, leading to frustration. Further, maintaining and evolving these intents requires significant effort.




Figure 6.1 An intent-based chatbot first detects an intent and then maps it to an answer.

There is a tradeoff between the specificity of the answer and the number of variations covered. Figure 6.2 demonstrates an example from the PharmaBot we introduced in chapter 3. The answer is accurate but generic—it detected a question about side effects but did not answer all the nuances in the user's question.

We can handle this nuance by adding search capabilities. There are two primary methods:

- *Traditional search* supplies the user with documents or passages relevant to their query. The user uses these documents to find their answer.
- *RAG* starts with a search process but expands on (augments) it by summarizing the passages into an answer.




Figure 6.2 Intent-based systems identify the main theme of an utterance and often give a static or generic answer.

Figure 6.3 illustrates how RAG refines information retrieval by locating relevant document passages and synthesizing them into a specific, contextualized answer. Unlike traditional intent-based systems that often respond with fixed or general answers, RAG dynamically pulls in content to address the user’s unique query, demonstrating how specific passages on ibuprofen and blood pressure are distilled into a targeted recommendation.




Figure 6.3 RAG finds relevant passages and summarizes them, giving a targeted answer.

Let’s dive deeper into how we can add these capabilities effectively.

6.1.1 Using search in conversational AI

User questions follow a “short head, long tail” distribution, as shown in figure 6.4. This distribution has a high frequency of common or popular questions (the short head). Most interactions involve less frequent, niche, or specialized queries (the long tail).




Figure 6.4 Distribution of user questions. Intents address the most common, high-volume questions, while low-volume, unique questions may necessitate search integration.

For PharmaBot, the short head includes general COVID inquiries, such as vaccine information and appointments. Each bot will have a different short head, but it will cover the most popular questions. When the chatbot is trained well, these questions are recognized with high confidence. Intents afford builders complete control over short head queries, albeit potentially overlooking nuanced distinctions.

We saw earlier that PharmaBot did not handle a nuanced question well because it used a static intent for `#side_effects`. Figure 6.5 shows PharmaBot handling the same nuanced question using traditional search capability.




Figure 6.5 Search finds relevant passages and displays those directly to the user, often with links to the source documents.

The response includes all nuances from the user's question in this example. The passages reference vaccine side effects, ibuprofen, blood pressure, and pain. However, the chatbot did not provide a single, cohesive answer. Instead, it offered document links and snippets. Users need to combine the answers from those documents and passages.

6.1.2 Benefits of traditional search

Traditional search can complement an intent-based chatbot by enabling it to retrieve relevant information from a document repository. The approach offers several advantages:

- *Breadth*—The bot can access various materials in your document repository, giving it answers to different question types.
- *Maintenance*—Adding knowledge to your bot can be as easy as adding or editing documents in your repository.
- *Technology*—Search is a well-established technology with mature algorithms and implementation methods. It can be implemented with relatively low computational resources and infrastructure.
- *Speed*—While slower than a static intent-based response, traditional search executes reasonably quickly.

Thus, search is an excellent complement to an intent-based system. The most common way to combine intents and search is to use a confidence threshold in the chatbot's natural language understanding (NLU) component. The NLU attempts to

detect an intent from the user's utterance. If an intent is detected with high confidence, an intent-based answer is returned. Otherwise, the user's utterance is passed to a search component (in some conversational AI systems, this is called a *fallback action* or *intent*). The high-level architecture is illustrated in figure 6.6.




Figure 6.6 Intents and searches have complementary functions. A search-augmented bot uses intent-based answers when it recognizes the utterance with high confidence; otherwise, it defers to search. Using intents and search together improves chatbot capabilities, but this approach still has some limitations.

6.1.3 Drawbacks of traditional search

Two fundamental problem areas exist when integrating traditional search with chatbot applications: the quality of search results and the user experience of how the search results are presented.

One major drawback to search quality is its reliance on keyword matching, which may be inaccurate or brittle, depending on the user's phrasing. Consider the previous example question: "Can I take ibuprofen with my blood pressure medication? My arms are sore after getting the vaccine." This may be converted to "ibuprofen blood pressure medication arms sore vaccine," emphasizing the most relevant keywords but losing the nuance of the question.

Not all search engines limit themselves to keyword matching, but it is ubiquitous. Newer search engines support searching by meaning rather than keywords. This approach is done with vector databases and will be described more fully in the next section. Like traditional searches, vector database searches take an input query and return a set of relevant documents and passages.

The other major drawback of search-based options is the user experience of receiving documents and passages. Some of the user experience limitations derive from the limited space in a chatbot window and the challenges of presenting multiple search results well. These are commonly addressed by showing a small number of results (possibly asking the user if they want to see more).

Screen real estate is sometimes preserved by showing document links (not the passages). In this case, users must leave the chat interface, which disrupts the conversation flow and may lead to the user abandoning the chatbot and continuing where the document links took them.

Search results are also challenging to handle through a voice interface, leading to lengthy readouts and a non-optimal user experience.

Most critically, the search does not result in a cohesive answer. Some users may prefer to construct their answers from relevant documents. Most users, however, are frustrated when a cohesive answer is not given, and they must do the piecing together—“Why did I use the chatbot in the first place? I could have searched on my own.”

NOTE You might consider combining web search with answer synthesis. While a chatbot can create synthesized answers from web search results, this approach also has limitations. It requires more sophisticated processing, but it addresses the user experience limitations by eliminating the need for users to sift through multiple documents. This approach depends heavily on how well individual passages match the query, and the summarization may miss nuances or context if the retrieved data isn’t comprehensive.

Search with answer synthesis typically relies on rule-based extraction methods, ranking algorithms, keyword matching, or predefined heuristics. The responses are presented by combining information from retrieved documents. While this can efficiently surface relevant information, it may struggle with incomplete or ambiguous queries. The system does not truly “understand” the content. Instead, it selects and reformulates existing text, which can lead to missing context, fragmented responses, or over-reliance on the most prominent retrieved results rather than the most accurate ones. It lacks the flexibility of generative approaches.

This is where RAG is a powerful alternative. RAG doesn’t just pull text from documents—it combines retrieval with generation, allowing the chatbot to produce a cohesive, contextually aware answer using relevant content from various sources. Unlike traditional search and summarization approaches, RAG can adapt to a broader range of user questions and provide deeper, more accurate responses by using a combination of real-time retrieval and language generation capabilities.

The next section will explore how RAG enhances chatbot responses by improving accuracy and maintaining context, even with complex or nuanced queries.

6.2 *Beyond search: Generating answers with RAG*

The lack of clear answers is a limitation of traditional search methods. To overcome these limitations, we’ll look at RAG as an advanced alternative. At its core, RAG combines the strengths of search-based information retrieval with the flexibility of generative models, offering a more comprehensive approach to understanding and responding to user queries. Most importantly, this response includes an answer.

6.2.1 *Using RAG in conversational AI*

RAG combines the best of retrieval and generation techniques to enhance the user experience. Like traditional search, it retrieves relevant passages to handle long-tail questions. RAG then feeds the passages and the user’s request to generative AI, which

creates the answer. RAG “augments” the retrieved passages by generating an answer, creating a seamless conversational flow, even for complex or long-tail queries.

Figure 6.7 shows PharmaBot answering our familiar example using RAG:

User: “Can I take Ibuprofen with my blood pressure medication? My arms are sore after getting the vaccine?”
Chatbot: Do not use over-the-counter ibuprofen for pain relief. Instead, use a painkiller less likely to increase your blood pressure, like aspirin.”




Figure 6.7 RAG retrieves relevant passages and augments the response by synthesizing the information into a grounded answer.

PharmaBot may still retrieve the same passages as the traditional search, but now it summarizes them to generate an answer. The answer acknowledges the user’s specific concerns and provides tailored advice. Most importantly, the answer is grounded in PharmaBot’s source documents—not the generative AI’s general knowledge. PharmaBot may provide links to supporting documentation, but it has made the answer prominent rather than the document passages. This is a more effortless experience for users.

RAG empowers chatbots to better understand user questions, and it streamlines development efforts by minimizing the need for explicit intent classification. This shift in approach enhances user satisfaction and future-proofs conversational AI systems against the evolving landscape of human language and user needs.

The use of RAG introduces dynamism and diminishes user effort, contingent upon the avoidance of hallucinations. While RAG inherently reduces hallucinations, it does not eliminate them entirely. Attention should be paid to the quality of the retrieved documents, the generative model’s behavior, and when the retrieval fails to find relevant documents, as we’ll discuss later.

When the user interacts with the conversational AI, the retrieval system connects to the trusted content sources, executes the search (keyword, semantic, or vector), and provides a relevancy score for the retrieved results. A large language model (LLM) then augments its response using the retrieved information. It generates a response from the retrieved content and presents it to the user through the chatbot interface. It may also apply translation if needed.

Like traditional search, RAG can complement intent-based systems. The short head is handled via intents, and RAG handles the long tail. The answers are as current as the document sources. Documents are easier to update than chatbots!

6.2.2 Benefits of RAG

RAG is a great addition to question-answer bots. By retrieving information from your knowledge bases (and not generative AI model knowledge), RAG ensures that answers are grounded in accurate and up-to-date data, thus enhancing the bot’s credibility and reliability. Additionally, RAG expands the versatility of conversational AI systems by enabling them to handle a wide range of user queries and tasks effectively. It also reduces the “bot doesn’t understand” and “too much complexity out on the user” pain points.

Benefits from RAG include enhanced answer relevance, adaptiveness, and improved handover to agents.

ENHANCED RELEVANCE

RAG’s essential advantage is its ability to access and integrate the latest information from your knowledge bases. The knowledge bases provide more varied information than pretrained intents do, and LLMs can adapt knowledge from your documents into an answer.

Let’s consider another example from PharmaBot. The initial PharmaBot included several intents, including guidance for dining in restaurants. It responded based on general guidelines and included a reference to “check local regulations.” There is no tailoring of responses to the user—it is a one-size-fits-all answer.

Table 6.1 compares intent-based and RAG-based versions of PharmaBot. Intent-based gives a generic answer, while RAG-based localizes to New York per the user’s question.

Table 6.1 A comparison of answers for intents and RAG

Interaction with an intent-based bot	Interaction with a RAG-integrated bot
User: What are the current COVID-19 guidelines for restaurants in New York? PharmaBot: Current COVID-19 guidelines for restaurant dining include limited indoor capacity, mask requirements for staff and patrons when not seated, and enhanced cleaning protocols. Please check local regulations for the most up-to-date information.	User: What are the current COVID-19 guidelines for restaurants in New York? PharmaBot: COVID-19 remains a significant concern for New Yorkers. Make sure you are up to date on the updated COVID-19 vaccines for the 2023-2024 season that offer protection against emerging variants and enhance overall immunity. Masks continue to be effective tools in mitigating the spread of COVID-19 and other respiratory viruses, reducing the transmission of droplets during various activities. To minimize the risk of contracting COVID-19, individuals are advised to wear masks in crowded indoor public settings. Source: NYC Health

The RAG-integrated chatbot acknowledges the user’s location in New York City. It provides more specific information tailored to the local guidelines and regulations, pulled from the current NYC Health Guidelines and summarized for restaurant

dining. This personalized approach enhances the user experience by delivering more relevant and actionable guidance based on the user's context.

ADAPTIVENESS

On the conversational side, another advantage of RAG is in the way the conversational AI generates the response—RAG can adjust its generated response to the style of the user's question. The response can be similarly more formal if the user's tone is more fact-seeking. The responses to the user's unique question are generated in real time. Sometimes users expect a concise and direct answer ("Yes or no, are there restrictions on dining in restaurants?"), and sometimes they expect a longer and more complex response ("Can I take my extended family to a restaurant, and will we have to wear masks?"). Both questions may use the same source documents but will receive very different answers.

HANDOVER TO HUMAN AGENT (OR NOT ANSWERING)

There are occasions when the conversational AI cannot find a definite response. In these cases, it is better for the bot to answer that it doesn't know or to transfer the user to a human agent. Figure 6.8 depicts a user asking a question a chatbot can't answer.




Figure 6.8 Supplementing RAG with human agents. If the answer has poor semantic overlap with the retrieved documents, send the user to a human agent instead.

Sample chat:

User: I have achalasia. Will my dysphagia get worse if I get a Booster and experience side effects?

When a user asks a question, the conversational AI follows a multistep process to retrieve information, generate a response, and determine whether the answer is sufficiently grounded in retrieved evidence before delivering a final response. The following steps illustrate this process, showing how the system retrieves relevant passages, generates a candidate answer, evaluates its accuracy, and ultimately decides whether to respond or transfer the user to a human agent:

1. *Passage retrieval*—The system retrieves passages related to achalasia, dysphagia, and general information about vaccine side effects:

Example passage 1: "Achalasia is a condition affecting the esophagus, causing difficulty in swallowing."

Example passage 2: "Common side effects of vaccines include soreness, fever, and fatigue."

Example passage 3: "Dysphagia, or difficulty swallowing, can be a symptom of esophageal conditions like achalasia."

- 2 *Answer generation*—The LLM generates a candidate response based on the retrieved passages:

Getting a booster might lead to common side effects, but there is no clear evidence linking it to worsening dysphagia in people with achalasia.

- 3 *Comparison check*—The system evaluates this generated answer against the retrieved passages and identifies a potential problem: the generated answer contains an element of “no clear evidence linking it to worsening dysphagia” that is not directly supported by the retrieved passages.
- 4 *Transfer decision*—Given the low match rate between the answer and the retrieved passages, the conversational AI determines that the answer may lack sufficient grounding and could be misleading. It then offers to transfer the user to a human agent for a more reliable answer:

Chatbot response: I apologize. I could not find a clear answer to your question in our resources. Let me connect you with a specialist who can provide more detailed information.

In this scenario, the conversational system searched for relevant document passages and fed them to an LLM. The LLM generated an answer, and the conversational AI then compared the generated answer to the retrieved passages. This comparison includes detecting how many words and phrases in the generated answer appear in the passages. If the percentage is low, the conversational AI decides that the answer is not grounded in the documents. The conversational AI then gracefully acknowledges its inability to provide a suitable grounded response and offers an alternate resolution path.

Alternatively, the search process may not have retrieved any documents. In that case, the conversational AI would not have to invoke the LLM to generate an answer, and it could directly deflect the question. For both scenarios, the conversational AI could instead return an “I don’t know” or other fallback responses. Both options reduce the chance of hallucinated and irrelevant answers.

6.2.3 *Combining RAG with other generative AI use cases*

RAG may also be combined with other generative AI use cases. For instance, RAG can handle informational queries, while other generative AI models specialize in tasks like sentiment analysis or language translation. By using a combination of AI capabilities,

conversational AI systems can offer users a comprehensive range of services, further enhancing efficiency and satisfaction.

RAG is only one of several generative AI patterns that enhance conversational AI. When users pose common questions or seek detailed information about a product or service, RAG draws upon the enterprise knowledge base to provide accurate and up-to-date answers. By grounding responses in the organization’s specific domain, RAG ensures that users receive relevant information tailored to their needs.

However, specific user inquiries may require more than informational responses, necessitating actionable steps (information-seeking versus transactional questions). In such cases, the conversational AI system executes transactions and guides users through specific tasks or processes. For instance, users may express an interest in purchasing after receiving information about a product or service from RAG. In response, the conversational AI can seamlessly transition to a transactional action, such as initiating a checkout process or scheduling a vaccination, facilitating a smooth and efficient user journey.

While RAG excels at efficiently and accurately responding to user queries, additional options, such as handing over to human agents or combining RAG with other generative AI use cases, can further optimize the user experience. These options ensure that users receive the support and assistance they need in the most efficient manner possible.

NOTE For scenarios where RAG responses may fall short—such as providing real-time data or fulfilling specific customer requests—function calling can be integrated to retrieve information from external systems dynamically. This approach allows chatbots to identify relevant intents and parameters for real-time responses, extending RAG’s utility in complex interactions. While it is not covered in depth here, the function call is valuable if you are seeking a more dynamic conversational AI system.

6.2.4 Comparing intents, search, and RAG approaches

Table 6.2 summarizes the capabilities and performance of three types of chatbots: intent-based chatbots, chatbots integrated with search, and RAG-integrated chatbots. Each chatbot type is evaluated based on the requirements and capabilities users and chatbot creators expect from conversational AI. You can discern the most suitable chatbot solution for your specific needs by comparing these attributes.

Table 6.2 Comparing the capabilities of intent-based chatbots, chatbots integrated with search, and RAG-integrated chatbots

Requirements	Intent-based chatbots	Chatbots integrated with search	RAG-integrated chatbots
Flexibility in handling queries	Limited to short-head predefined intents. May ignore nuance.	Handle long-tail queries by returning links and snippets	Handle long-tail queries by returning answers

Table 6.2 Comparing the capabilities of intent-based chatbots, chatbots integrated with search, and RAG-integrated chatbots (*continued*)

Requirements	Intent-based chatbots	Chatbots integrated with search	RAG-integrated chatbots
Accuracy and relevance of responses	When an intent is recognized with high confidence, the answers are accurate and precrafted.	Provide contextually relevant and accurate documents that help the user find an answer	Provide contextually relevant and accurate answers grounded in your documents
Adding new knowledge to the bot	Add or revise manually curated intent-response pairs	Add or revise documents in your knowledge base	Add or revise documents in your knowledge base
Maintenance and scalability	Extensive regression testing when intent training data is changed	Document repository needs to be maintained by adding new documents and removing stale documents	Document repository needs to be maintained by adding new documents and removing stale documents
Response generation quality	Predefined responses are presented.	User must put together their own answer from retrieved passages and documents	Answers are grounded in source documents but adapted to nuance from the question

While traditional chatbots help organizations automate simple tasks and provide essential customer support, integrating RAG techniques enhances their ability to deliver more accurate, context-aware responses, ultimately improving the user experience.

Exercises

- 1 Consider your last chatbot implementation and consider the long-tail concept:
 - List three examples of niche or uncommon user queries that traditional intent-based chatbots may not adequately address.
 - Discuss how these queries exemplify the long-tail phenomenon in conversational AI.
- 2 For the same chatbot implementation, consider what answers you can provide with traditional searches versus RAG.

6.3 *How is RAG implemented?*

As the “retrieval-augmented generation” name suggests, RAG has two phases: retrieval and generation. In the retrieval phase, algorithms search for and retrieve snippets of information relevant to the user’s prompt or question. In an open-domain consumer setting, those facts can come from indexed documents on the internet; in a closed-domain enterprise setting, a narrower set of private sources are typically used for added security and reliability.

6.3.1 High-level implementation

With RAG, the system searches a knowledge base for information relevant to a question and uses that information to generate a conversational answer. Let's break down the steps:

- 1 The user asks the chatbot a question.
- 2 The system uses its NLU capabilities to determine the intent of the user's question:
 - If it recognizes the question with high confidence—for example, if it is one of the intents it was trained on—it will be able to respond, and a search will not be needed. This ends the flow.
 - If it cannot recognize the query, it will go to search. The system will send the user's query to the search tool to search the document content and produce and rank search results.
- 3 It passes back the ranked search results to the chatbot for display. (Before RAG, the ranked link and snippet list would have been passed back to the chatbot—handling the long-tail questions with search results was still more helpful than providing a “sorry, I cannot understand” response.)
- 4 Instead of simply displaying the results, the original question and the search results are sent to an LLM. The LLM may rerank the search results, but most importantly, it generates a concise, summarized, linguistically correct answer.
- 5 The answer is then passed back to the system.
- 6 The answer is presented to the user through the chatbot UI.

NOTE Neither the original user question nor the generated answer needs to match the documents *exactly*. While verbatim responses can indeed occur and are sometimes even preferred for legal reasons, the primary focus is on grounding the content in the knowledge base, ensuring that the generated answer is rooted in the curated document set.

By their nature, LLMs do not generate consistent results each time a query is processed. These models can produce different responses to the same question depending on subtle variations in context or phrasing. This variability is due to the probabilistic nature of LLMs, which generate text based on learned patterns rather than retrieving fixed responses. While this flexibility allows for more nuanced and contextually appropriate answers, it can also lead to expectations of consistent outputs, which is not how these models function.


Emphasizing this point is essential, because team members unfamiliar with how LLMs work often expect consistent results. This expectation can hinder projects, leading to differing approaches to the problem among team members. Understanding that LLMs prioritize relevance and context over the exact replication of document content can help align expectations and improve collaboration within the team.

In some cases, the generated response may closely resemble or even match the wording in the documents. This can occur when the documents contain relevant and informative passages directly addressing the user's question. In such instances, the RAG model may include verbatim excerpts from the documents in the generated response to provide the user with the most accurate and relevant information.

In other cases, the RAG model uses the information within the knowledge base to understand the context and relevant concepts related to the user's question. It then uses this understanding to generate a response that aligns with the content found in the documents, even if the specific wording of the user's question or the generated answer does not exist verbatim within the documents. This approach allows for greater flexibility and adaptability in developing responses that effectively address user queries while drawing upon the information available in the curated document set.

6.3.2 Preparing your document repository for RAG

Let's also consider how the document content is retrieved during RAG searches. Figure 6.9 provides more detail about creating appropriate data.



- 1. The retriever converts the utterance into embeddings.**
- 2. The LLM receives the question and retrieved passages.**

Figure 6.9 RAG uses a vector database during build time and run time.

A systematic preprocessing pipeline ensures that both the raw data and the user's question (or the LLM's rephrasing of it) are optimized for use in RAG-based searches. This pipeline is crucial for transforming data into embeddings, enabling the model to match the user's query with relevant information efficiently. Techniques such as cosine similarity or other methods are then applied to identify the best matches, ensuring accurate and contextually appropriate results. The following list outlines the key steps involved in this pipeline, detailing how data is processed before being used in retrieval:

- 1 Preprocessing data**—The system (typically a data pipeline, not the LLM itself) processes raw documents to make them searchable. For example, PDF documents

are converted to text, or table structures are converted to processable statements. Metadata may be added to enhance the original content. The text is then divided into coherent semantic units, called *chunks*. For instance, a document may be chunked at paragraph boundaries. Chunking is a common process for identifying and extracting meaningful groups of words (“chunks”) from sentences for further analysis or processing. The chunking strategy impacts the overall results. There are open source tools that can help with visualizing and understanding different chunking or splitting strategies.

- 2 *Embedding generation*—An embedding model converts these chunks into embeddings or numerical representations of words or phrases in a high-dimensional vector space. Embeddings capture semantic relationships between words and documents, enabling a more efficient understanding of the connections. Similar meanings or contexts are mapped nearby in the vector space, and dissimilar meanings are mapped to more distant points. This provides more relevant search results than keyword matches.
- 3 *Storage in a vector database*—The generated embeddings are stored in a vector database, which enables efficient similarity searches. Each document chunk is indexed using its vector representation, allowing fast retrieval based on meaning rather than exact word matches.
- 4 *Retrieval and matching at runtime*—At run time, the end user interacts with the chatbot. Their question will be converted to a vector using the same embedding model, and that vector will be searched in the vector database to find the most relevant passages (chunks) based on semantic similarity. These retrieved passages are then passed to the LLM, which synthesizes them into a response presented to the user.

Each of these steps ensures that the retrieval process is optimized, making it possible for the LLM to generate accurate, context-aware responses based on the most relevant retrieved data.

Listing 6.1 shows sample code for an embedding function. You can use any custom embedding function or other vector databases, and the performance may differ depending on the embedding model used. This is the most common approach to RAG: you create a dense vector representation of the knowledge base to calculate the semantic similarity to the user queries. For this sample, we used Chroma as the vector database.

Listing 6.1 Splitting a file into chunks, embedding it, and storing it in a vector database

```
from langchain.document_loaders import TextLoader
from langchain.text_splitter import CharacterTextSplitter
from langchain.vectorstores import Chroma

loader = TextLoader(filename)
documents = loader.load()
```

```
text_splitter = CharacterTextSplitter(chunk_size=1000, chunk_overlap=0)
texts = text_splitter.split_documents(documents)

from langchain.embeddings import HuggingFaceEmbeddings

embeddings = HuggingFaceEmbeddings()
docsearch = Chroma.from_documents(texts, embeddings)
```

At its core, RAG operates by retrieving relevant documents or passages based on a user's query and then generating a response using natural language generation techniques. This process can be achieved without explicit chunking or embedding by using other methods for document retrieval and language generation.

For example, using Lucene as an alternative to chunking and embeddings involves using its document indexing and retrieval capabilities. Lucene can handle the retrieval part, fetching the most relevant documents based on the query. After retrieval, the generator part of the RAG can take over to produce coherent responses based on the content of the retrieved documents. Lucene is very efficient at text retrieval, which leaves the complex task of generating human-like responses to the more specialized generative components of the RAG model. This approach can be particularly advantageous in systems emphasizing retrieval accuracy and speed over nuanced understanding.

Exercises

- 1 Text chunking—In this exercise, you will experiment with different chunking strategies and embeddings:
 - Choose your sample text data (small text files of your choice).
 - Decide on a chunking strategy (splitting by sentences or words). For your experiments, try a chunking tool.
 - Embed the chunks using an open source embedding model, and then load the chunks into a vector database (Chroma).
- 2 Setting up an ingestion pipeline—This exercise guides you through building a simple ingestion pipeline for processing documents in a RAG system:
 - Choose a document set relevant to your organization's domain or a specific use case for the chatbot. Start with simple, text-only documents, i.e., no tables, etc.
 - Build the ingestion pipeline, considering factors such as ease of use and compatibility with RAG. For querying, use open source models, Hugging Face embedding models, and a llama index.
 - Implement the ingestion pipeline to preprocess and structure the dataset for use with RAG.
 - Test the ingestion pipeline with sample data to ensure proper functionality and data integrity.

6.4 Additional considerations of RAG implementations

Traditional search returns links, passages, or the full text of relevant documents, and the user needs to sift through this information to find their answer. RAG conversely returns the answer directly, and the user can optionally see the documents used.

6.4.1 Can't we just use an LLM directly?

What if the conversational AI passed the user's query to an LLM and got the answer? After all, LLMs are trained on vast amounts of data.

First, LLMs trained on internet-scale data have limitations due to the nature of their training data. This data represents a snapshot of the training time from publicly available sources—it does not contain business-specific, personal, or classified information, and it doesn't contain public data created after the cutoff date. Thus, even the newest LLM's knowledge can become outdated, leading to inaccurate responses over time. RAG addresses this by offering data to LLMs after they are trained.

Second, as LLMs are trained from extensive datasets, it is challenging to trace their responses to sources, undermining the reliability and trustworthiness of the model's output. RAG is inherently grounded because you know exactly what data was provided to the LLM for a given question.

The broad domain LLMs cover poses another significant challenge. With access to vast information, they may generate responses with high confidence, even when lacking concrete evidence or context. This tendency to produce plausible but incorrect or unverified information is known as *hallucination*. In contrast, you want your conversational AI to provide correct and grounded answers. Advanced prompting techniques can help mitigate hallucinations, but providing source data through RAG is more reliable.

You also want to prioritize answers grounded in the specific documents or corpus being indexed, not those on which LLM was trained. A RAG system's primary focus is to provide responses based on the content and context of your documents, so answers are directly generated from the information within the corpus, promoting accuracy, relevance, and trustworthiness in the responses provided to users.

It is important to consider the training data and domain of a specific LLM before selecting it for your use. If the LLM was trained on generic data and you need domain-specific results, it may not produce the desired outcomes. In such cases, you could explore techniques like model blending, where you combine multiple models to use the strengths of each, enhancing performance in specific domains. If you have the resources and data available, you may also consider fine-tuning the selected model to better suit your needs. However, this can require a significant budget for computational resources and data, so consider prompt-tuning first. While fine-tuning costs are decreasing and will continue to do so, they still need to be carefully considered. Other methods are also emerging for domain-specific training, offering further flexibility.

6.4.2 Keeping answers current and relevant with RAG

RAG represents a significant advantage over directly using LLMs for question answering. While RAG still uses LLM for natural language generation, LLM is crafting accurate responses from the searched documents. Real-time retrieval will find up-to-date information when new or updated documents are added during the build phase. RAG ensures the answer reflects the latest documents from the searched knowledge sources. This real-time integration of enterprise content enhances the relevance and accuracy of responses and instills confidence in users, who know that they are receiving current and reliable information.

Furthermore, RAG goes beyond merely accessing enterprise content. It accesses specific passages and retrieves information from multiple documents. This granularity allows RAG to trace and verify answers to their exact sources, providing users with full transparency and trustworthiness. This facilitates accountability and the verification process for your development team too, ensuring you know what your bot is doing and why.

Moreover, RAG defines the domain of the LLM's understanding, enabling it to recognize the limits of its knowledge and expertise. Unlike LLMs that are used directly and that may attempt to provide answers outside of their domain, RAG can acknowledge when it encounters queries beyond its scope. This ability to say "I don't know" prevents it from giving inaccurate answers and fosters transparency in conversational interactions. By establishing clear boundaries for its understanding, RAG empowers developers to build AI systems prioritizing accuracy, reliability, and integrity, ultimately enhancing the overall user experience.

There is a difference between "I don't understand" and "I cannot find an answer to your question." While the primary goal of RAG is to generate informative and relevant responses based on the content of the retrieved documents, there are scenarios where the system may not find sufficient or appropriate information to generate a meaningful response. In such cases, it is common for the RAG model to acknowledge its inability to provide a satisfactory answer and communicate this to the user.

However, it's important to note that a RAG system's specific behavior, such as returning an "I don't know" response, can be influenced by the retrieval component's design, the knowledge base's quality, and the generation model's settings or parameters. Additionally, developers may choose to implement specific strategies or fallback mechanisms to handle cases where the system cannot generate a response, such as providing alternative suggestions or prompting the user for more information.

6.4.3 How easy is it to set up the ingestion pipeline?

Setting up an ingestion pipeline that effectively preserves document structure is critical for ensuring accurate search results within a retrieval system for RAG. Several key areas must be considered. Essentially, every architectural decision you make about the components will have an influence on the overall accuracy of the results.

First, you must establish mechanisms to connect existing content stores to the retrieval system or migrate content into a new repository. This will allow the retrieval system to access the necessary data and maintain data integrity.

The next challenge is correctly extracting structures (such as headings, tables, and lists) during ingestion. These formatting elements contribute to the document's organization and clarity. By retaining this structural information during ingestion, the retrieval system can use it to enhance search accuracy and relevance.

There are also challenges related to chunking. The ability to chunk, split, or partition large documents into representative subdocuments for indexing enhances the retrieval process's efficiency. This allows for more granular indexing and retrieval of information, facilitating quicker access to specific content within lengthy documents. Additionally, selecting appropriate search methodologies, such as vector, semantic, federated, keyword, or hybrid, further augments the retrieval system's capabilities.

Using LangChain simplifies setting up the ingestion pipeline. Recall that you will need

- *Document loaders*—Load data from various formats.
- *Document transformers*—Process and structure the data for efficient retrieval.
- *Retrievers*—Fetch the most relevant document chunks during query time.

Document loaders facilitate the ingestion of diverse document formats. These loaders streamline the workflow, ensuring efficient processing and retrieval of pertinent context for LLMs to deliver precise responses. They load data from the source documents, treating each extracted piece as a document comprising textual content and associated metadata. LangChain provides built-in capabilities for handling various files: all files in a directory, PDF, CSV, JSON, HTML, markdown, txt, and more.

For example, you can load text from a web page, transcripts, or corporate documents. Document loaders provide a `load` method for loading data as documents from a configured source:

- Text loader:

```
from langchain_community.document_loaders import TextLoader

# Load text data from a file using TextLoader
loader = TextLoader("./your_data/YourText.txt")
document = loader.load()
```

- CSV loader:

```
from langchain_community.document_loaders import CSVLoader

# Load data from a CSV file using CSVLoader
loader = CSVLoader("./your_data/Yourspreadsheet.csv")
document = loader.load()
```

Look at LangChain's documentation on customizing the CSV parsing and loading. For example, you may want to specify your delimiters, field names, etc.

Similarly, LangChain provides a `DirectoryLoader` for all documents in a directory, an `UnstructuredHTMLLoader` to load HTML docs, and so on for the common types. It is essential to know the `AzureAIDocumentIntelligenceLoader`, which is useful for Microsoft Office-type documents.

- Microsoft Document Loader:

```
%pip install --upgrade --quiet langchain langchain-community
➡ azure-ai-documentintelligence

from langchain_community.document_loaders import
AzureAIDocumentIntelligenceLoader

file_path = "<your_filepath>"
endpoint = "<Your_endpoint>"
key = "<key>"
loader = AzureAIDocumentIntelligenceLoader(
    api_endpoint=endpoint, api_key=key, file_path=file_path,
    ➡ sapi_model="prebuilt-layout"
)

documents = loader.load()
```

Once you have loaded the documents, you need to look at the *document transformers*, which can split a long document into smaller chunks that the selected LLM can process. LLMs have a “context window” property, determining the text length they can effectively process in a single pass, so the chunks must fit into the LLM’s context window. It is easy to assume that setting a more extensive context (i.e., longer text passages) would inherently lead to better performance across various language understanding tasks. However, recent studies have revealed that this isn’t always the case. Evidence suggests that language models can achieve improved performance when presented with less text overall, but text that is highly relevant to the task at hand.

A larger context window allows for including more information in the prompt during inference, but this technique, often called *prompt stuffing*, comes with trade-offs. Processing more text demands greater computational resources, which slows inference and increases costs—particularly for companies paying by the token, where summarizing lengthy documents, like annual reports or meeting transcripts, can become costly. While larger context windows can improve results to some extent, there are diminishing returns. Like humans, LLMs can experience information overload; when presented with excessive detail, they may overlook critical points. Studies have shown that LLMs are more likely to focus on essential information at the beginning or end of a prompt, potentially missing key insights buried in the middle.

We need document transformers to preprocess the documents, extract relevant information, and transform it into a structured representation that the language model can efficiently use during generation. LangChain has several built-in transformers that make document manipulation easy.

The splitting process is as follows:

- 1 Divide the text into smaller chunks.

- 2 Combine these smaller chunks into larger chunks of a certain size, usually measurable by some function.
- 3 Once it reaches that size, it becomes the new unit of the text. Then, you create a new text segment with some overlap to maintain context between the fragments.

You can choose your division rules (characters, words, tokens) and how to measure the chunk size. Again, LangChain offers many different types of splitters in the `langchain-text-splitters` package. These are some examples of text splitters:

- 1 *Recursive*—Splitting text recursively is the recommended way to start. It aims to keep related pieces of text next to each other.
- 2 *HTML*—A “structure-aware” chunker splits text based on HTML-specific characters; an example is shown in listing 6.2. It splits at the element level, adding metadata to headers for chunk relevance. It preserves semantic grouping and context-rich information in document structures:
 - *Character*—It breaks the document at user-defined characters (e.g., “\n\n”).
 - *Code*—It employs code syntax and grammar identifiers for languages like Python and JavaScript (and 13 others), organizing code into logical groups.
 - *Markdown*—It identifies markdown language and organizes the document into a structured format (similar to HTML).
 - *Tokens*—It uses a tokenizer, like tiktoken, to split text based on model-defined token limits in the code.

Listing 6.2 HTML splitter

```
# Install langchain-text-splitters if not already installed
%pip install -qU langchain-text-splitters

# Import necessary modules
import langchain
import langchain_text_splitters

print("Langchain version:", langchain.__version__)
print("Langchain Text Splitters module loaded successfully!")
from langchain_text_splitters import HTMLHeaderTextSplitter
from langchain.schema import Document # Ensure Document is properly imported

from bs4 import BeautifulSoup

print("BeautifulSoup is installed successfully!")

# Sample HTML content to be split
html_string = """
<!DOCTYPE html>
<html>
<body>
<div>
<h1>Introduction</h1>
<p>This is the introduction section of the document.</p>

```

```

<div>
<h2>Chapter 1: Getting started</h2>
<p>This section covers the basics of getting started.</p>
<h3>Section 1.1 Setup</h3>
<p>This subsection explains the setup process.</p>
<h3>Section 1.2 Configuration</h3>
<p>This subsection details the configuration options.</p>
</div>
<div>
<h2>Chapter 2: Advanced Techniques</h2>
<p>This section dives into more advanced techniques.</p>
</div>
<br/> <!-- Fix: Ensuring self-closing tag is correctly formatted -->

<p>What you learned in the Introduction.</p>
</div>
</body></html>
"""

# Define header tags to split on (h1, h2, h3 represent different levels of
# headers)
headers_to_split_on = [
    ("h1", "Header 1"), # Top-level headers
    ("h2", "Header 2"), # Subsection headers
    ("h3", "Header 3"), # Sub-subsection headers
]

# Initialize the HTML header text splitter with the specified header levels
html_splitter =
    HTMLHeaderTextSplitter(headers_to_split_on=headers_to_split_on)

# Split the HTML document into structured chunks
html_header_splits = html_splitter.split_text(html_string)

# Display structured output
for doc in html_header_splits:
    print(f"Content:\n{doc.page_content}\nMetadata: {doc.metadata}\n{'-'*40}")

```

Next, we need to deal with the embeddings. The type of data and the language support requirements govern the selection of embedding models. Furthermore, when you are dealing with specific domain or industry terms, these models may have to be extended.

Embedding models in LangChain transform the text into numerical representations, or embeddings, that can be processed. LangChain integrates with different model providers (OpenAI, Cohere, Hugging Face, and more) to generate embeddings. The `OpenAIEmbeddings` class, for instance, uses the OpenAI API to create embeddings, and this can be done using either OpenAI's API key or Azure's OpenAI API key.

```

from langchain.embeddings import OpenAIEmbeddings
embeddings = OpenAIEmbeddings()

```

```
text = "This is a test document."
query_result = embeddings.embed_query(text)
query_result[:5]
```

Other integrations include `CohereEmbeddings`, `TensorFlowEmbeddings`, and `HuggingFaceInferenceEmbeddings`.

After you have the embeddings, you must store them in a vector database, such as Chroma, which we used earlier. When selecting the vector database, you'll want to consider run-time performance, how it scales for the size of your data set, and overall performance. Another important consideration is integrating tools like LangChain, which is continually improving. LangChain enhances the capabilities of vector databases by providing streamlined processes for handling embeddings and integrating with various machine learning and AI workflows. This combination ensures efficient data management and retrieval, making it a robust choice for scalable and high-performance applications.

Retrievers bridge the gap between embeddings and user queries. While embeddings store numerical representations of documents in a vector database, retrievers identify and fetch the most relevant chunks based on similarity scoring.

The retriever works as follows:

- 1 The user's query is embedded using the same embedding model used during ingestion.
- 2 The vector database searches for the most semantically similar embeddings.
- 3 The retriever fetches the top matches and passes them to the LLM for response generation.

LangChain includes multiple retrieval methods. For example, there is a similarity-based retriever in LangChain:

```
from langchain.vectorstores import Chroma
retriever = vector_db.as_retriever(search_type="similarity", search_kwargs={"k": 5})
retrieved_docs = retriever.get_relevant_documents(query)
```

Retrievers play a crucial role in returning only the most relevant document chunks, ensuring the LLM works with focused, high-quality context rather than raw, unprocessed data.

6.4.4 Handling latency

A universal best practice for handling latency has yet to be developed. Long response times are frustrating for users, but these techniques can enhance their experience:

- *Use a quality vector store with efficient search.* The Facebook AI Similarity Search (FAISS) library allows you to search for similar embeddings quickly. There are many purpose-built vector databases, like Chroma, Milvus, Pinecone, and Weaviate, with many more emerging. Traditional databases and search systems like Elasticsearch provide vector search plugins. Each has unique strengths and

can be selected based on your needs, including scalability, functionality, performance, and cost.

- *Preprocess and curate your dataset.* Having multiple similar versions of the same document increases search time and lowers search result quality.
- *Inform the user before executing a slow action.* An appropriate-toned message, such as “Just need a moment,” may placate the user and bridge the delay.
- *Stream responses to show the user the answer as each token is generated.* LLMs may take 1.5 to 5 seconds or more to generate an answer, and searches may take 5 to 10 seconds. The user may think the chatbot is broken if the conversational AI waits for the LLM to finish.
- *Consider caching.* By caching each user’s vector database and chat history, commonly accessed information relevant to that user’s interactions can be stored locally. This reduces the need to generate responses from scratch every time, saving the computational resources required. While caching may consume additional tokens, the trade-off is improved efficiency.

6.4.5 When to use a fallback mechanism and when to search

Determining whether to use RAG’s response or to deflect to a human agent involves several key considerations. For instance, the generation part of RAG should not be invoked if the retrieval does not yield appropriate results. In this case, the conversational AI can gracefully exit the query and respond with the offer to pass the user to a human agent. By bypassing the generation of an answer based on potentially subpar search results, you effectively reduce latency for end users, ensuring they receive prompt responses while saving computational resources.

Figure 6.10 illustrates the decision-making process for determining whether to use the RAG response or to hand over the query to a human agent:

- 1 The conversational AI processes the user query and passes it to the retrieval system, which searches the knowledge base to find relevant information.
- 2 The retrieval results are evaluated to determine if they are appropriate for generating a response.
- 3 If the retrieval results are deemed appropriate, the system generates a response using the full-on RAG approach.




Figure 6.10 When the answer from RAG does not match the retrieved passages, it can be better to offer a human agent instead.

- 4 If the retrieval results are inappropriate (e.g., insufficient or no results), the system gracefully offers to pass the query to a human agent.
- 5 The response (either generated by the system or passed to a human agent) is returned to the user.

6.5 Evaluating and analyzing RAG performance

Evaluating the capabilities of a RAG model within a conversational AI system is multifaceted. Each capability must be evaluated for an overall result. Most evaluations consider three aspects:

- *Was it the right response?* Did the answer directly address the user's question? For example, if the user asked about resetting a password, the response should clearly explain the steps rather than discussing broader account security topics.
- *Was it in the right context for this user?* Did the response consider the user's specific situation or history? For example, if a user previously reported an account problem, the system should provide a tailored follow-up solution instead of generic advice.
- *Was it grounded in the documents (or hallucinated or made up by the generation process)?* Did the response accurately reflect the information retrieved from the source documents without fabricating details? For example, a response should correctly reference a company's policy document when explaining return procedures instead of creating nonexistent policies.

An LLM can score responses based on the key criteria to enhance the evaluation process. This approach works best when combined with human review: the LLM provides an initial assessment, and human evaluators then verify the accuracy and contextual relevance of the responses.

These evaluation criteria help determine the truthfulness of the chatbot's responses. The generated answers should also be accurate if the source documents are accurate and RAG retrieved the correct documents. The evaluation of the responses can be broken down into assessing the different components of RAG, which can be individually evaluated for overall performance, including the quality of the document index, the effectiveness of the retrieval process, and the accuracy of the answer generation.

6.5.1 Indexing metrics

Indexing metrics provide insights into how efficiently a system can organize, store, and retrieve vast amounts of data. Key considerations include indexing speed, storage requirements, scalability, and how well the system handles high-dimensional data like vectors. Table 6.3 summarizes these important aspects, offering a quick overview and relevant examples.

Table 6.3 Critical metrics that influence the efficiency and accuracy of a RAG system's document index

Aspect	Summary	Example
Indexing metrics	Evaluates speed, storage needs, and scalability. Critical for large-scale data systems.	Indexing for a news aggregator where speed and scale are crucial
Vector database performance	Measures performance in handling high-dimensional data	For technical support, accurate troubleshooting steps must be assembled for multiple documents. E.g., "Why is my device overheating?"
Recall rate	Indicates accuracy in retrieving relevant data. High recall is vital for complete retrieval.	In legal document retrieval, high recall ensures all relevant cases are found.
Query complexity	Affects performance based on query specifics, dimensionality, and dataset diversity	Financial databases handling complex queries across multiple data points
Benchmarking tools	Tools like ANN-Benchmark compare algorithms on metrics like recall versus QPS.	Evaluating which algorithm best balances speed and accuracy for a video search engine

The first critical component to assess is the indexing metrics, which involve evaluating the efficiency and effectiveness of organizing and accessing data in a system's knowledge base. This includes examining factors such as indexing speed, storage requirements, and the scalability of the indexing process. Efficient indexing is crucial for a RAG system, as it impacts the speed and accuracy of information retrieval. Ineffective indexing can result in slow response times and inaccurate data retrieval, compromising response quality.

Vector database performance is another vital metric specializing in storing and retrieving high-dimensional vectors representing complex data, such as text, images, or embeddings. These databases perform approximate rather than exact match searches, necessitating performance evaluation beyond traditional database performance measures like queries per second (QPS) and latency. While these metrics are important for evaluating system speed and responsiveness, they do not directly capture the accuracy of retrieval results. Therefore, besides QPS and latency, the *recall rate* is another essential performance metric for vector databases. If the vector database performs well, the RAG model can access high-quality, relevant information, leading to more accurate and contextually appropriate generated content. Conversely, poor performance can result in slow retrieval times and irrelevant or less useful data being used for generation.

Consider a customer support scenario where a chatbot powered by a RAG system is utilized to handle inquiries. With a high recall rate, the chatbot accesses a broad range of information from the knowledge base, effectively resolving customer queries and enhancing satisfaction. However, a low recall rate can lead to missed crucial information, resulting in inadequate responses and increased customer frustration.

Therefore, the chatbot’s effectiveness significantly depends on its ability to comprehensively retrieve relevant information, underscoring the importance of a high recall rate in such automated support systems.

Query complexity, influenced by factors like the dimensionality and specificity of the query, as well as the data diversity, also affects vector database performance. Higher-dimensional queries require more computational resources because distance calculations between vectors become more complex. This can lead to increased time and memory usage for retrieval tasks. More specific queries might target very narrow segments of the vector space, which can challenge the indexing system to efficiently isolate and retrieve the relevant vectors, especially in large datasets. High query complexity can strain the system, potentially leading to slower retrieval times and less relevant data being returned. Also, complex queries make distinguishing between relevant and irrelevant results difficult.

Imagine a chatbot on an e-commerce platform designed to help customers find products using complex queries involving multiple attributes like brand, color, size, and user ratings. For example, a customer might ask the chatbot for “6.5-sized blue Adidas running shoes with a minimum of a 4-star rating.” This query presents a multi-faceted challenge due to its specificity across several dimensions. Each of these attributes represents a different vector in the database.

Benchmarking tools like ANN-Benchmarks and VectorDBBench help evaluate these aspects by comparing different algorithms and configurations, ensuring the RAG system is built on a robust retrieval foundation for consistently high-quality content generation. ANN-Benchmarks plots the recall rate on the x -axis against QPS on the y -axis, illustrating each algorithm’s performance at different retrieval accuracy levels. VectorDBBench displays QPS and recall rates separately.

6.5.2 Retrieval metrics

The next capability is retrieval metrics, which gauge the system’s ability to fetch relevant information from indexed data. Key aspects include retrieval accuracy, precision, recall, and response time. Effective retrieval metrics ensure users receive accurate and relevant responses, boosting satisfaction and trust in the conversational AI system. Table 6.4 summarizes these important aspects, followed by more detailed explanations.

Table 6.4 Critical aspects influencing retrieval metrics

Aspect	Summary	Example
Retrieval accuracy	Evaluates the system’s ability to retrieve relevant information from indexed data	Ensuring a chatbot retrieves accurate troubleshooting guides from a large dataset
Precision and recall	Precision measures relevance of retrieved docs; recall measures how many relevant docs are retrieved.	Balancing precision and recall when retrieving product recommendations in an e-commerce chatbot

Table 6.4 Critical aspects influencing retrieval metrics (continued)

Aspect	Summary	Example
Context precision and context recall	Specific to RAG: context precision checks relevance; context recall checks coverage of relevant info	Evaluating how well a generated response in a support chatbot matches the query's context
Parameter optimization	Tuning search parameters and algorithms to improve speed, accuracy, and relevance of results	Adjusting FAISS clusters or Elasticsearch settings to improve document retrieval for legal databases
Embedding models	The use of different embeddings impacts the retrieval quality by enhancing precision or recall.	Using BERT for precise context understanding in a legal advice chatbot
Filtering and reranking	Strategies to remove noise and rerank results to improve both relevance and accuracy	Filtering out irrelevant articles in a news aggregation chatbot and then reranking top results
Normalized Discounted Cumulative Gain (NDCG)	Assesses the ranking quality by considering relevance and position of retrieved documents	Ensuring the most relevant help articles appear at the top in a technical support chatbot

These metrics evaluate search quality, document relevance, and how well user queries align with responses. *Retrieval accuracy* measures the system's ability to fetch the most relevant information from its indexed data. If retrieval accuracy is low, the chatbot may return responses that are only loosely related to the user's query or fail to retrieve critical details.

Precision measures the proportion of relevant retrieved documents. Recall measures how many relevant documents are retrieved by the system. A high recall value indicates the system retrieves many relevant documents from the database. In contrast, a high precision value indicates that the retrieved documents are mostly relevant to the user's query. Balancing recall and precision is crucial to ensure comprehensive coverage of relevant information and minimize irrelevant results.

Specific to RAG models, context precision and context recall evaluate the alignment and coverage of generated responses relative to the user's query. Context precision measures how precisely the retrieved context matches the user's query, indicating the relevance and accuracy of the generated response. Context recall measures how comprehensively the generated response covers relevant information from the retrieved context.

You should implement various strategies to enhance the retrieval process:

- Optimizing search parameters
- Using different embedding models
- Implementing filtering and reranking

The first strategy is to adjust retrieval parameters to improve both speed and accuracy. Adjusting the search parameters, such as the number of clusters in FAISS or the search

query complexity in Elasticsearch, can significantly enhance the precision and recall of retrieved documents. This ensures that the system returns the most relevant documents, increasing precision and retrieving all pertinent documents, boosting recall.

Parameter optimization can also reduce response time, making the system more responsive. Tuning the indexing and query algorithms can lead to faster retrieval times without compromising accuracy.

The second strategy is to utilize different embedding models to find the most effective combination for retrieving relevant documents. Embeddings can drastically affect retrieval quality. For example, BERT and Sentence-BERT can improve the system's understanding of the context and semantics of user queries, enhancing precision. GPT-3 embeddings can offer a broader contextual grasp, improving recall by retrieving more relevant documents from diverse contexts.

The third strategy is to implement filtering and reranking. Filtering and reranking strategies further enhance precision by removing irrelevant documents and boost recall by prioritizing relevant documents. Techniques like domain-specific filters or context-aware reranking can significantly refine the relevance of top results.

Normalized Discounted Cumulative Gain (NDCG) is useful when the order of retrieved documents matters. It measures the effectiveness of ranking algorithms by evaluating the relevance and position of documents in the search results, providing a comprehensive ranking quality assessment.

6.5.3 Generation metrics

The last area for evaluation is the generation itself: whether the generated output provides a relevant and complete answer. Table 6.5 offers a concise overview of these key metrics, strategies for enhancing generation quality, and practical examples of their application, followed by further details.

Table 6.5 Critical aspects influencing generation metrics

Aspect	Summary	Example
Faithfulness	Evaluates the factual accuracy of generated output based on retrieved context	Ensuring chatbot answers are factually correct when responding to questions. E.g., "What are the consequences of breaching a contract?"
Answer relevancy	Assesses how relevant the generated answer is to the specific user question	"What is the status of my order?" The chatbot, after retrieving the relevant data, responds with, "Your order #12345 is currently in transit and expected to be delivered by August 15th." This response is directly relevant to the user's question, providing specific information about the order status without unnecessary details.
Fine-tuning	Improves generation by aligning the LLM with domain-specific data, enhancing accuracy and relevance	Fine-tuning an LLM for legal advice, ensuring generated responses are accurate and legally relevant

Table 6.5 Critical aspects influencing generation metrics (*continued*)

Aspect	Summary	Example
Prompt engineering	Crafting prompts to guide the LLM in generating more contextually appropriate and relevant responses	Using prompt engineering to ensure a healthcare chatbot provides clear, relevant medical suggestions
Model blending	Combining specialized models to enhance the quality of generation, balancing accuracy and fluency	Blending a retrieval-focused model with a language-focused model to generate accurate, fluent responses
Sensibleness and Specificity Average (SSA)	Measures response quality in open-domain chatbots, ensuring responses are sensible and specific	Assessing open-domain chatbot responses to ensure they make sense and are not overly vague
Faithfulness-Evaluator	Assesses whether the generated response avoids hallucinations by aligning with retrieved context	Using FaithfulnessEvaluator to ensure that a finance chatbot's responses are grounded in retrieved documents

Two primary metrics for assessing LLM performance are faithfulness and answer relevancy. Faithfulness evaluates the factual accuracy of the answer based on the retrieved context, while answer relevancy assesses how pertinent the answer is to the given question. An answer may be factually accurate (faithful) but not well-matched to the question (less relevant), or it may be accurate but not based on sourced documents.

Several strategies can be employed to enhance generation metrics. Fine-tuning the LLM on domain-specific data can improve faithfulness and relevancy by aligning the model more closely with the context in which it will operate. This ensures that generated responses are accurate and pertinent to the specific domain. You may also want to incorporate prompt engineering. The generation process becomes more focused and aligned with the user's intent by explicitly including necessary context or constraints within the prompt. Model blending may also be used. Combining multiple models, each specialized in different aspects of the task, can help enhance generation quality. For example, one model may excel at retrieving accurate information, while another might be better at generating fluent and contextually appropriate language. Blending these models can lead to more balanced and effective output.

Sensibleness and Specificity Average (SSA) metrics evaluate response quality in open-domain chatbots. Sensibleness ensures that responses make contextual sense, while the specificity average ensures comprehensiveness without vagueness. Historically, human interactions were necessary for assigning these ratings.

While avoiding vague responses is essential, preventing LLMs from hallucinating is equally critical. LlamaIndex established a FaithfulnessEvaluator metric to measure hallucination by assessing whether the response aligns with the retrieved context. LlamaIndex was developed to address the challenge of connecting LLMs with various

data sources, so LLMs can access, query, and generate insights from structured and unstructured data. It is a framework for building context-augmented generative AI applications. It offers capabilities for data integration, indexing, enrichment query processing, and more. Stay current with the latest updates from LlamaIndex.

RAG-integrated conversational AI efficiently addresses rare or complex queries. Nevertheless, realizing these advantages from RAG necessitates ongoing monitoring of all components, particularly where failures are common, such as during retrieval and generation.

6.5.4 Comparing efficiency of indexing and embedding solutions for RAG

In the previous sections, we introduced the contributing factors to the efficiency of RAG systems. The indexing and embedding components are crucial for effectively retrieving relevant documents and ensuring the system's responsiveness. When implementing RAG in a conversational AI system, it is important to measure its impact on retrieval accuracy and response relevance. Benchmarking is essential to creating an effective RAG system. You must have a method to evaluate whether changes in the system prompt improve user-query hit rates. Is the improvement 1%, 2%, or more? This is fundamental to knowing if your RAG system is truly effective.

Furthermore, without proper monitoring, validation, and evaluation, it will be difficult to prove the effectiveness of your system. RAG is inherently complex, and typical implementations have 50% to 60% accuracy at first. You'll want to increase the accuracy to above 80% for a practical solution.

PharmaBot, initially developed as a COVID-19 chatbot to handle general inquiries, such as vaccine information and appointment scheduling, is now set to be enhanced with RAG. The goal is to extend PharmaBot's capabilities to answer more nuanced questions, such as "Can I take ibuprofen with my blood pressure medication?" and "My arms are sore after getting the vaccine. What should I do?" First, we'll want to select a dataset of medical articles, research papers, and guidelines from health organizations, all focusing on various health problems and drug and vaccine interactions. Next, we'll compile a set of representative queries like those previously listed, which are what users might ask.

We'll select the indexing solutions and the embeddings we want to use. For indexing, we'll consider several options, such as FAISS and Elasticsearch. The most popular source for the latest performance benchmarks of text embedding models is the MTEB leaderboards hosted by Hugging Face. While MTEB provides a valuable starting point, the displayed results are self-reported, and many models may not perform as accurately when applied to real-world data. BERT, Sentence-BERT, or GPT-3 embeddings are worth considering, as they have been used in many solutions. Then we'll run our selections to generate embeddings for our content and index our embeddings using the selected methods. Finally, we'll run our queries and measure our performance.

When evaluating PharmaBot enhanced with RAG, we can use table 6.6 to determine relevant metrics, establish baselines and goals for these metrics to assess

improvements introduced by different combinations of indexing (FAISS, Elasticsearch) and embedding (BERT, Sentence-BERT, GPT-3) solutions, and compare solutions. For example, we could compare how the combinations of RAG components perform against business objectives.

Table 6.6 Prioritized metrics for RAG evaluation based on business objectives

Business objective	Prioritized metrics	Why?	Example
Customer satisfaction	Response accuracy Relevance	Directly impacts user experience and satisfaction	PharmaBot providing accurate answers to users' queries
Operational efficiency	Latency Throughput	Ensures the system can handle high query volumes quickly	A customer support chatbot for a large e-commerce platform should prioritize low latency to provide quick responses during peak shopping times.
Scalability	Queries processed per second Resource utilization	Evaluates system performance under increasing loads	Chatbot for a healthcare provider managing seasonal spikes in appointments and queries
Cost-effectiveness	CPU usage Memory usage	Ensures high performance without excessive resource consumption	Chatbot for a non-profit organization providing 24/7 mental health support on limited funding

You could create a comparison table like the one in table 6.7. These sample numbers are hypothetical and should be adjusted based on actual selected components and benchmarking results. For example, one configuration may be selecting Elasticsearch and then using three different embedding models to arrive at your numbers.

Table 6.7 Evaluation of PharmaBot with various configurations

Metrics		Configuration 1	Configuration 2	Configuration 3
Response accuracy	Recall	0.85	0.87	0.88
	Precision	0.75	0.77	0.78
	F1 score	0.8	0.82	0.83
Relevance	Mean reciprocal rank (MMR)	0.70	0.72	0.74
	Average precision	0.65	0.68	0.70
Latency	Average latency (ms)	50	55	60
Throughput	Queries/second	20	18	16
Resource utilization	CPU usage (%)	70	65	90
	Memory usage	8	7	12

When trying to satisfy the overall business requirements for PharmaBot using RAG, various trade-offs must be considered based on the evaluation metrics. An indexing component combined with selected embedding components offers high response accuracy and relevance with lower latency and higher throughput, making them suitable for systems requiring quick and accurate responses. However, these combinations exhibit moderate to high resource utilization, which may increase operational costs. Conversely, integrating another embedding component with various indexing strategies provides superior response accuracy and relevance but at the cost of significantly higher latency and lower throughput due to the computational demands of the selected embeddings. This can impact the system's ability to handle high query volumes efficiently. The high CPU and memory usage may also strain resources, increasing operational costs. Ultimately, selecting the optimal combination requires balancing the need for high accuracy and relevance with the system's capacity to handle queries efficiently while managing resource utilization to control costs.

Furthermore, the ongoing re-evaluation and refinement of the RAG system should be considered. Table 6.7 is not comprehensive, but the key point is to decide on key evaluation metrics and then use a structured evaluation approach. Systematic testing and focusing on the RAG component provide a robust RAG evaluation pipeline. The overall goal is to see an upward trend at the end, as illustrated in figure 6.11. Systematically applying the strategies discussed throughout this chapter and then analyzing the results reveal the impact of different configurations on RAG performance. Some tweaks show significant improvements, emphasizing the importance of experimentation and tuning. There is no best approach; exploring multiple directions when tuning your RAG systems is crucial.




Figure 6.11 Different configurations and enhancements improve the accuracy of the RAG system.

Additionally, you may want to use RAG evaluation frameworks, from proprietary paid solutions to open source tools. Selecting the right solution requires balancing

considerations around ease of maintenance and operational burden, plus how well the metrics observed by the tool map to your RAG pipeline and your business objectives. The following solutions are current examples, but more are being developed, providing even more options:

- *Arize*—A model monitoring platform focusing on precision, recall, and F1 score. It is beneficial in scenarios requiring ongoing performance tracking, ensuring RAG systems consistently meet accuracy thresholds in real-time applications. Arize is a proprietary paid offering that provides robust support and continuous updates for enterprise deployments.
- *RAGAS*—An open-source tool that offers streamlined, reference-free evaluation focusing on average precision (AP) and custom metrics like faithfulness. It assesses how well the generated content aligns with provided contexts, and it is suitable for initial assessments or when reference data is scarce.

Exercises

- 1 Assess the relevance of responses generated by a RAG model within a conversational AI system:
 - Define evaluation criteria to measure the relevance of responses generated by the RAG model.
 - Establish a scoring system to quantify the relevance of responses based on factors such as semantic similarity and informativeness.
 - Devise a set of user queries you will evaluate.
 - Create a set of expected responses (manually).
 - Compare the generated responses with the previously created responses to determine the level of relevance.
 - Calculate evaluation metrics such as precision, recall, and F1 score to quantitatively assess the performance of the RAG model in generating relevant responses.
 - Analyze the evaluation results to identify patterns or areas where the RAG model excels or fails to generate relevant responses.
 - Discuss potential factors influencing response relevance and strategies for improving the RAG model's performance in this aspect.
- 2 Evaluate document grounding with RAG:
 - Generate responses to user queries using the RAG model, and identify the source documents or passages from which the responses are derived.
 - Assess the degree of grounding by comparing the relevance of the source documents or passages to the corresponding user queries.
 - Develop a scoring mechanism to quantify the RAG model's grounding effectiveness based on factors such as document relevance and coverage.

Summary

- Traditional intent-based chatbots can be greatly enhanced by integrating search functionality.
- Intents are great for answering common short-head questions, and search is great for long-tail questions.
- Traditional search returns links or document passages instead of an answer.
- RAG extends search capability by generating an answer from the documents retrieved by the search.
- By using RAG, chatbots can provide contextually appropriate responses in real time, reducing user frustration and enhancing the conversational experience. Grounding answers in the organization's domain also solves intent maintenance and enhancement for developers.
- RAG implementations must consider several problems, from handling latency to providing fallback mechanisms or handover to human agents to prevent hallucinations.
- Evaluation of RAG must consider the different components of indexing, retrieval, and generation.



Augmenting intent data with generative AI

This chapter covers

- Creating new training and testing examples with generative AI
- Identifying gaps in your current conversational AI data
- Use LLMs to build new intents in your conversational AI

Conversational AI users are frustrated when the AI does not understand them—especially if it happens multiple times! This applies to all conversational AI types, including question-answering bots, process-oriented bots, and routing agents. We’ve seen multiple strategies for improving the AI’s ability to understand. The first strategy—improving intent training manually (chapter 5)—gives full control to the human builder, but it takes time and specialized skill. The second strategy—retrieval-augmented generation (RAG, chapter 6)—gives much more control to generative AI, reducing the role of the human builder over time. This chapter introduces a hybrid approach where generative AI augments the builder. This applies to rules-based or generative AI-based systems.

Using generative AI as a “muse” for the human builder reduces the effort and time required of the human builder, increases the amount of test data available for data science activities, and gives the human builder the final say, which eliminates most opportunities for hallucinations (which is when AI says something that looks reasonable but is not true).

Let’s say you are building a conversational AI solution to help your IT help desk. From interviews, you know that password resets are the most frequent task the AI needs to support. Therefore, the AI needs a strong understanding of the password reset intent.

Because the conversational AI solution is new, you don’t have any production user utterances to train from. When you ask the service desk how users generally start their conversations, you hear “Well, they usually say something about ‘forgot password’ or ‘cannot login.’” You are appropriately suspicious—surely the users have a broader vocabulary than that—but you have a hard time imagining what that vocabulary might be. Generative AI can help you imagine.

Let’s look at how the human builder and generative AI can be partners.

7.1 Getting started

Large language models (LLMs) are skilled at performing many technical tasks, including classification and question answering. These are the same tasks at the core of conversational AI. So why don’t we just use generative AI for our core conversational AI tasks?

LLMs are generalizable because they have been trained on huge amounts of data. This makes them a quick study on many tasks, but it also comes with some cost. What are the costs to using LLM as the classifier in conversational AI?

- *Monetary*—LLMs can be expensive to run.
- *Speed*—Because LLMs consider billions of parameters, they can be slower (a time cost).
- *Reputation risk*—LLMs are so generalized that they may hallucinate output that makes your bot look bad or exposes you to legal risk.
- *Lack of transparency and explainability*—LLMs are often a “black box.”

By contrast, conversational AI uses purpose-built technology. Because its classifier is trained only to do the task at hand, it is much cheaper, and it runs quickly because it considers fewer parameters. While that may reduce the accuracy, the system is guaranteed to use a set of controlled responses. These comparisons are summarized in table 7.1.

Table 7.1 Comparing and contrasting traditional natural language processing (NLP) in conversational AI and generative AI on the classification task

Feature	Traditional NLP	Generative AI
Model	Purpose-built for excellence at only one task: classification	Generalized model good at many tasks
Runtime speed	Fast	Slow

Table 7.1 Comparing and contrasting traditional natural language processing (NLP) in conversational AI and generative AI on the classification task (*continued*)

Feature	Traditional NLP	Generative AI
Runtime cost	Low	High
Accuracy	Mostly accurate (trained by you on small data)	Mostly accurate (pretrained on huge data)
Scalability	Manageable for up to 100 intents; very difficult afterward	Generalizes very well via RAG pattern
Controllability	Strictly controlled by humans; requires extensive testing	Prone to hallucinate when given full control; hallucinations are hard to detect automatically

We can use a hybrid approach to get the best of both methods.

7.1.1 *Why do it: Pros and cons*

An LLM can greatly reduce the time and effort spent by a human builder. LLMs and humans work best together—as partners. Training a conversational AI classifier requires human effort, but it also requires data, and that data can be hard to collect. Sometimes that data cannot be collected until the conversational AI is deployed in production. Even in our familiar example of detecting “forgot password” problems, we still don’t know all the ways users might state their problem. They may use the “wrong” words!

LLMs are especially helpful in these scenarios:

- *Bootstrapping*—AI suffers from a “cold start” problem. How can you train when you have no data? LLMs can generate an initial set of training data.
- *Expanding*—Use LLMs to fill the gaps in your existing data when you don’t have enough data to optimize your classifier’s accuracy. This is especially useful for understanding rare but important intents (such as users reporting fraud).
- *Robust testing*—LLMs can generate additional data for testing, increasing your confidence in the robustness of the conversational AI. (This is helpful even if generative AI creates the answers, as in RAG.)

An LLM can help you run many experiments, some of which will generate output that is directly usable by your application, either as training or testing data. You and the LLM can help each other too. For instance, the LLM can give you themes and variations that your users may use. You can select your favorites and ask the LLM to expand on those by updating the prompt instructions or few-shot examples.

Your interactions with the LLM will be iterative and collaborative. For instance, you are unlikely to design the right prompt the first time. The LLM may not understand the task correctly or may give you content that’s not quite helpful. Expect to do a few rounds of experimentation before you get great results. After that, you can quickly generate suggestions across all your intents and improve your AI’s understanding of your users.

7.1.2 What you need

Many LLMs can help us with our task of generating more training or testing data for our “forgot password” intent. So do we just pick one and turn it loose? Not quite. The LLM will help you, but you should not expect it to do all the work. Instead, you should have an idea of where you need to start, such as knowing what gaps you have in your solution. You also need to select an LLM that is appropriate for your use case.

Access to an LLM is an obvious prerequisite for using an LLM to augment your conversational AI. There are several non-obvious considerations when selecting that LLM:

- *Terms and conditions*—Several LLMs explicitly forbid you from using their LLM to “build or improve another model.” This clause is intended to keep you from building a competitive LLM, but using an LLM to improve conversational AI could be construed in this way, and your appetite for legal risk may vary. (Consult with your legal department—they may have already selected LLMs for your company to use.)
- *Data privacy*—As you use the LLM, will it be allowed to keep your data and train on it in the future? The data in your conversational AI may be confidential to your corporation. If it is, you can’t just share it with any LLM.
- *Capability*—Not every LLM is capable of creative generation tasks. Make sure you select a model that can follow instructions.
- *Open source or proprietary*—For many use cases, explainability is important. Open source models generally give you more insight into the model’s training process, such as the training data and source code for the model. Proprietary models generally do not expose that information but usually have more ease-of-use. This may also affect your ethical and regulatory compliance.
- *Latency and response times*—There are speed and accuracy tradeoffs; a larger model may be both more accurate and slower to run.

In this chapter and the rest of the book, we will use multiple prompts and models to provide examples. These examples will be adaptable to your model (or models) of choice. Feel free to experiment with other models, especially those that weren’t available as we wrote this book.

You will also need to bring some domain knowledge to the LLM. This can include background on the problems your users are bringing to your conversational AI, the intents your system needs to support, and utterances belonging to those intents. Bring as much real-world data as you can. Then use the LLM to augment that data.

7.1.3 How to use the augmented data

An LLM can help you generate additional data for use by your chatbot. Using an LLM at build time reduces the risk and effect of hallucinations. These options include adding to your training data, adding to your testing data, and modifying your existing data (for example, changing grammatical structure and synonyms).

The best source is data is from real users of a production system. We recognize that this introduces a chicken-and-egg problem—you may not have any data if you are not in production yet. Your intent classifier needs to be trained on varied data so that it can understand varied data. The chatbot should be tested on data that it was not trained on.

When you don't have any training data, you'll tend to generate low-variance utterances. You'll have a couple of key words in mind, and you'll "anchor" yourself to them. Even with dozens of examples, low-variance utterances don't convey much information. High-variation utterances cover a lot of ground quickly, as shown in table 7.2, and they generally make your classifier stronger.

Table 7.2 Comparing low-variation to high-variation utterances. High-variation utterances increase the robustness of classifiers.


Low-variation utterance set	High-variation utterance set
■ I forgot my password ■ Forgot my password ■ Forgot password ■ Help I forgot my password	■ I can't log in ■ Account locked out ■ Forgot password

The high-variation utterances easily cover the low-variation utterances despite being fewer in number. The single utterance "forgot password" is enough to predict the intent of all four low-variance utterances. The reverse is not true and wouldn't be true even if we added dozens more slight variations on "forgot password" to the low-variance set. "I can't log in" has no direct word overlap—the low-variance set doesn't cover it.

We prefer a small, high-variance training data set that covers a large volume of low-variance test utterances. Better to train on 10 strong variations than 100 weak variations. This makes the chatbot robust to the diverse utterances it will see in production. It also reduces your chances of accidentally unbalancing the training set (which leads to weak understanding).

We can visualize the information conveyed by the utterances in figure 7.1. The first plot shows the low-variation utterances from table 7.2. Since they only convey two words, you can think that the utterances are tightly clustered together. The second plot shows the high-variation utterances. With no word overlap, the utterances are spread all over the grid, but there is a lot of empty space. The third plot shows an ideal test set where there is broad coverage of the grid. The fourth plot shows an ideal training set, which covers maximum variation in a small number of examples. The test data set can be much larger than the training set. We want the test set to have variation, but it's fine if we have near-duplicates.

In this chapter, we'll first demonstrate how to use generative AI to create high-variance utterances. Then we'll expand on those utterances via many slight variations. By the end of the chapter, you'll see how to generate utterances matching the third and fourth plots.



- 1: Small, low-variation set covers only two words**
2: Small, high-variation set covers a handful of words
3: Larger low-variation set is good for testing, with broad coverage of the input space
4: Larger high-variation set is good for training, with efficient coverage of maximum variation

Figure 7.1 Visualizing coverage from different kinds of utterance sets. Our ideal training data is set 4, which covers a large variation in a small number of utterances. Set 3 is the ideal testing data.

Exercises

- 1 Imagine you are building a chatbot for a typical retail store. Create ten utterances for a `#store_location` intent, for when users ask questions like “Where is your store located?” Keep track of how much time this takes.
- 2 Create ten more utterances, without using the words “where,” “store,” “located,” or “location.” (Time yourself again.) Do these utterances have more variety?
- 3 Repeat the previous two exercises for a `#store_hours` intent. First use whatever words you want, and then restrict yourself from using “when,” “time,” and “hours.”

7.2 Hardening your existing intents

We’ll start our exercise knowing which intent we need to improve: the “forgot password” intent. We need enough training data so that the conversational AI can detect this intent. Remember that our support staff didn’t know all the ways users might state this problem. They said, “Users usually say something about ‘forgot password’ or ‘cannot login.’” That won’t be enough to train the chatbot on a robust “forgot password” intent.

We will use the LLM as our partner. First, the LLM will help us generate contextual synonyms so that we see a broad range of vocabulary. Next, the LLM will generate full utterances using this vocabulary. Then we will have the LLM generate different grammatical variations, such as questions versus statements and past tense versus present tense. We’ll also have the LLM transfer lessons learned from building one intent (“forgot password”) into building the next intent (“find a store”).

We’ll start with the simplest step—finding synonyms.

Can I use a different LLM?

Yes! We will use multiple models in this book. The field of generative AI is moving quickly, and models used during the writing of this book may be supplanted by better models by the time this book is published or by the time you read it. The principles we'll demonstrate are more important than the specific models we'll use.

7.2.1 Get creative with synonyms

The first step is designing a prompt. A good prompt ensures your LLM understands its task, and the first task in this example is generating a broad range of synonyms. The process of prompt engineering requires an iterative process of experimentation.

Our subject matter experts advised that the utterances often include “forgot password.” One way of increasing your chatbot’s robustness is to make sure you have coverage on noun and verb phrases. Let’s ask an LLM to generate some likely synonyms for the noun phrases.

How do I set up and run my LLM?

There are multiple options for setting up an LLM environment. You can run LLMs locally on your machine, using a tool like Ollama, or run them on a commercially hosted platform. We used the Prompt Lab available in IBM’s watsonx.ai platform due to our familiarity with it, but nothing in this chapter is platform dependent. Use your favorite platform.

For this exercise, we will use the falcon-40b-8lang-instruct model (<https://huggingface.co/tiiuae/falcon-40b>) with greedy decoding. Greedy decoding instructs the model to generate the next most probable word at each step and yields the same output every time.

Though we want to end up with full sentence fragments, we got better results by breaking the task down into pieces. User utterances are typically sentences or fragments, primarily built from nouns and verb phrases. Let’s start with a simple prompt—just asking for synonyms—to get our nouns.

Listing 7.1 Generating noun synonyms without context

Generate a list of nouns. The nouns should be synonyms of 'password'.	Simple instruction	
List of 5 nouns credentials, secret, key, code, accesscode		Prompted cue for the LLM
	LLM output	

In the preceding listing, we used the LLM as a generic thesaurus. While the output nouns are synonyms for “password,” they are not synonyms that are often used in the context of logging in. We need to provide more context to the LLM to get better results.

This time let's tell the LLM why we are asking for synonyms and what kind of synonyms we are looking for. The following listing shows the improved prompt and results.

Listing 7.2 Generating noun synonyms with context

You are a consumer contacting an online retailer via their chat service.
Generate a numbered list of nouns a consumer might use while describing their inability to log in to your service in the phrase "I forgot my password".
The nouns should be synonyms of 'password'.

Background information given to the LLM

Detailed instructions grounding the task to forgetting passwords

List of 5 nouns

← Cue for LLM

1. "account information" 2. "credentials" 3. "login information"
4. "authentication" 5. "identifier"

LLM response

These synonyms sound much more familiar. Next, let's create contextual verb synonyms.

Listing 7.3 Generating verb synonyms with context

You are a consumer contacting an online retailer via their chat service.

Generate a numbered list of verbs a consumer might use while describing their inability to log in to your service in the phrase "I forgot my password".
The verbs should be synonyms of 'forgot'.

Updated instruction for verbs instead of nouns

List of 5 verbs

← Updated cue for verbs instead of nouns

1. Forgot 2. Lost 3. Misplaced 4. Misplaced 5. Err

← LLM output

Not bad, though the last verb is a little weird in this context. Another limitation is that all the generated verbs are in the past tense. This is appropriate given that our example was also in the past tense, but we want our LLM to generate more variety for us. Let's try expanding from *verbs* to *verb phrases*.

Listing 7.4 Generating verb phrase synonyms with context

You are a consumer contacting an online retailer via their chat service.

Generate a numbered list of verb phrases a consumer might use while describing their inability to log in to your service in the phrase "I forgot my password".
The verb phrases should be synonymous with 'forgot'.

Updated instruction and cue for "verb phrases"

List of 5 verb phrases

1. I forgot my password.
2. I lost my password.

↓ LLM output

3. I misremembered my password.	↑ LLM output
4. I misplaced my password.	
5. I cannot remember my password.	

The LLM gave us full sentences (oops!), but now we are getting some present tense utterances (“cannot remember”) along with the other past tense utterances. We are making progress! With only a few minutes of prompting, we got the LLM to give us a lot of variation to think about. Before, we might have assumed our chatbot would just have to look for “forgot” and “password.” Now we have a dozen other useful words to consider when we test the bot.

Let’s experiment a little more. This time we’ll increase the creativity of the model by moving to *sampling decoding* and increasing the temperature. We’ll also revise the prompt by asking for “10 synonyms” instead of “5 nouns.” Listing 7.5 shows the nouns, and listing 7.6 shows the verb phrases.

NOTE With greedy decoding, the LLM generates the same results every time. Sampling decoding generates non-deterministic output. If you try these prompts, you’ll probably get different results. This is okay! We are only using the LLM to spark our creativity.

Listing 7.5 Generating noun synonyms with increased creativity settings

You are a consumer contacting an online retailer via their chat service.

Generate a numbered list of synonyms a consumer might use for 'password' while they described a problem like "I forgot my password".

Updated prompt
and cue

List of 10 synonyms for 'password'

1. "Secret" 2. "Personal Identification Number" 3. "PIN"
4. "Access Code" 5. "Log-in Information" 6. "Log-in Data"
7. "Account Information" 8. "Account Data"
9. "Identification Number" 10. "Security Code"

LLM output

Awesome! This is a great list of nouns. Your system may not use all of them, but this is a thorough list for testing.

Let’s try verbs next.

Listing 7.6 Generating verb phrase synonyms with increased creativity settings

You are a consumer contacting an online retailer via their chat service.

Generate a numbered list of synonyms a consumer might use for 'forgot' while they described a problem like "I forgot my password".

Updated
instruction
and cue

List of 10 synonyms for 'forgot'

1. Forgot 2. Did not remember 3. Didn’t know 4. Unknown
5. Not applicable 6. Unable to access 7. Couldn’t recall
8. Didn’t memorise 9. OMG Failed to remember
10. Unable to login

LLM output

This is a much more creative list of synonyms. While there are some oddities in this list (“unknown,” “not applicable”) there are some nice creative sparks:

- *Slightly wrong verb*—“Did not remember” is odd, but it makes you think of “Can not remember.”
- *Wrong tense*—“Didn’t know” makes you consider “Do not know.”
- *Sentiment*—“OMG” reminds you that utterances may include frustration.

We’ve generated synonyms relevant to our domain with only a few minutes of effort, but we still only have piece-parts. We started with the utterance “I forgot my password” and can now plug in new nouns and verb phrases, but we are still stuck with a simple subject-verb-object structure. Our users will surely use more varied grammar. We don’t want the chatbot to depend on only one grammatical form. We want it to be resilient to more varied utterances. Let’s use LLMs to generate more grammatical variations.

7.2.2 Generate new grammatical variations

Instead of asking the LLM to generate words or word pairs, let’s try having it generate entire utterances. We need to design a prompt that introduces word variation, but we don’t want to bias the model too hard toward “I forgot my password.” We will use a similar prompt that sets a context, but rather than directly including the phrase “I forgot my password,” we will describe the user’s problem instead (they can’t log in).

Our first attempt is shown in the following listing.

Listing 7.7 Generating entire utterances

You are a consumer contacting an online retailer via their chat service.

Generate a numbered list of example utterances a consumer might start conversation with if they couldn't log in. Use a variety of nouns and verbs related to logging in, passwords, and accounts.

List of 10 utterances

1. I can't log in
2. My password isn't working
3. I don't know my password
4. I forgot my username
5. I forgot my password
6. My account is locked
7. I didn't get my password reset email #
8. I need help with my account
9. Can you help me log in?
10. I tried resetting my password but it didn't work

Updated
instruction
and cue

LLM output

We are starting to get more variety. This list of utterances has several ideas we haven’t seen yet (like “password reset email”). Even better, we are getting more variety in the sentence structure. The output list has the following grammatical varieties:

- *Subject verb object statement (active voice)*—I forgot my password.
- *Passive voice statement*—My account is locked.

- *Prepositional*—I need help with my account.
- *Question*—Can you help me log in?
- *Qualifiers*—I tried resetting my password but it didn't work.

The variety in our sample utterances is improving. The generated utterances are usable for our training or test sets, but there are still gaps. For instance, all these utterances are perfect sentences. What about our users who are so busy (or frustrated) that they only give us a few words? Can the LLM generate useful sentence fragments? The following listing explores this idea.

Listing 7.8 Generating sentence fragment utterances

You are a consumer contacting an online retailer via their chat service.

Generate a numbered list of example fragments a Consumer might start conversation with if they couldn't log in. Use a variety of nouns and verbs related to logging in, passwords, and accounts. Use sentence fragments only.

Updated instruction to generate fragments. Cue is unchanged.

List of 10 sentence fragments

1. Forgot password
2. Unable to login
3. Account locked out
4. Need help resetting password
5. Can't remember email address
6. Email is not working
7. Unable to change password
8. Unable to access account
9. Unable to create new account
10. Unable to delete old account

One-shot example below the cue

LLM output

In our previous prompts, we were not able to generate sentence fragments. This time, we gave the LLM extra help. Aside from our usual updates to the prompt (replacing “utterances” with “fragments”), we gave the LLM one additional hint. We provided the first example fragment “Forgot password.” This is called one-shot learning because we gave the LLM one example of what we wanted, and that helped the LLM learn how to process our request.

Zero-shot? One-shot? Few-shot?

The “zero-shot,” “one-shot,” and “few-shot” terms refer to the number of examples (shots) given in the prompt. A zero-shot prompt does not give any examples. A one-shot prompt gives one example, and a few-shot prompt gives a few examples.

For training data generation, one-shot learning is a great way to get exactly the kind of output you want. Whenever you are having trouble getting an LLM to follow your instructions, consider giving a good example rather than just tweaking the

instructions. While writing this chapter, we tried several more prompts than are included in this book, and none of them gave us sentence fragments until we used one-shot learning.

Further, you can use one-shot learning to take the lessons learned while building one intent and apply them to another intent. In the next listing, we use examples from a “store location” intent to generate examples for a “password reset” intent.

Listing 7.9 Using one-shot learning for multiple grammatical structures

You are a consumer contacting an online retailer via their chat service.

Generate phrases a user might use to find out where stores are located. Create phrases for each of the following grammatical types.

Direct Question: Where are you located?

Indirect Question: Can you tell me how to find your stores?

Fragment: Store location

Command: Give me driving directions

Generate phrases a user might use when they need to reset their password.

Create phrases for each of the following grammatical types.

Direct Question: I forgot my password.

Indirect Question: How can I reset my password?

Fragment: Password reset

Command: Send me a password reset link

Standard background for LLM, unchanged from the past several examples

One-shot example includes instruction and desired output

Instruction to the LLM, supplemented by the cue “Direct Question:”

LLM output (starts after the cue “Direct Question:”)

With a single prompt, we were able to get examples of each of the grammatical structures we wanted (the LLM made a mistake on “direct question,” but the output is still useful).

Here’s one more trick for generating utterances. Rather than using detailed instructions, provide a few examples, and ask the LLM to generate more. We’ll use a different prompt format and a different model—granite-13b-instruct-v2 (<https://mng.bz/DMIR>)—and we’ll use sampling decoding for increased creativity and non-deterministic results. The following listing shows the prompt and first output.

Listing 7.10 Using a creative prompt to generate examples with a Granite model

<|instruction|>

Here are actual utterances submitted by customers to an automated help desk. Your task is to create new examples from people having problems with their password and login ability.

<|example|>

I can't log in

A marker indicating our instruction to the model

The actual instruction

Beginning of examples

```

<|example|>
My login information isn't working

<|example|>
Forgot password

<|example|>
Help me get into my account

<|example|>
Hi there, I can't seem to login to my account

```

Because we are using non-deterministic settings, the model output is different every time. Here are the outputs from the next five executions of the same prompt:

- “Can you help me recover my password”
- “I’m locked out of my account”
- “Can’t remember username or password”
- “Hoping you can help me, I just reset my password but it’s not working”
- “I’ve failed logging in 5 times in a row”

We didn’t specify exactly what variation we wanted, but we still got some interesting variations. We saw new verbs (“recover”) and concepts (“5 times in a row”). This highlights the value of experimenting with different LLMs, different prompts, and different parameter settings. Generating training data requires creativity. Don’t rely on one or two experiments to do the work—you and generative AI can work together to be creative.

7.2.3 *Build strong intents from LLM output*

Let’s recap our experiments so far. We’ve generated nouns and verbs as in-context synonyms (not just generic synonyms). We’ve generated entire utterances with a similar structure, then used LLMs to generate utterances with varied grammatical structures. We’ve used multiple models, prompts, and parameter settings. Generative AI has been a great partner!

From these experiments, we have a lot of possibilities for building a training set. Let’s select 10 utterances covering the variations we generated earlier. For some of the utterances, we’ll use the verbatim output from the LLMs. For other utterances, we’ll substitute some variations. For instance, the generated utterances were heavy on “password”—we can substitute “login information” or “account information.” The utterances were also heavy on “forgot,” so we’ll substitute “can’t remember.” The following listing shows one possible selection of utterances.

Listing 7.11 Ten hand-selected utterances based on LLM suggestions

1. I can't log in
2. My login information isn't working
3. Forgot password
4. Account locked out

5. I can't remember my account information
6. My account is locked
7. I didn't get my password reset email
8. Need help with resetting account
9. Can you help me log in?
10. I tried to reset my password but it didn't work

We've come a long way since the initial suggestion that “most requests include the words ‘forgot’ and ‘password’”! If we use these utterances in our training set, we will have much more robust chatbot understanding than if we had stuck to our keyword-based advice.

We've been successful generating ideas across multiple prompting sessions. This begs the question, can we do everything in one prompt? We would expect to need all our tricks to date: a context for the LLM, a clear instruction, and a one-shot example. Let's try to train a “store locator” intent using our best examples from the “forgot password” intent. The following listing demonstrates this using the falcon-40b-8lang-instruct model with greedy decoding again.

Listing 7.12 Using one-shot learning to copy lessons from one intent to another

You are a consumer contacting an online retailer via their chat service.

Standard background for LLM

Generate a numbered list of example utterances a consumer might start conversation with if they couldn't log in. Use a variety of nouns and verbs related to logging in, passwords, and accounts. Use a variety of grammar structures in the utterances.

Instruction and cue for one-shot example (password reset)

List of 10 utterances

1. I can't log in
2. My login information isn't working
3. Forgot password
4. Account locked out
5. I can't remember my account information
6. My account is locked
7. I didn't get my password reset email
8. Need help with resetting account
9. Can you help me log in?
10. I tried to reset my password but it didn't work

One-shot example

Generate a numbered list of example utterances a consumer might start conversation with if they need to find your store. Use a variety of nouns and verbs related to finding, location, and directions. Use a variety of grammar structures in the utterances.

Instruction and cue for target (store locator)

List of 10 utterances

1. Where is your store located?
2. How do I find your store?
3. What direction is your store from here?
4. Which exit should I take to get to your store?

LLM output

5. Is your store north or south of here?
6. Your store is on what street?
7. Where is your store exactly?
8. How far is your store from here?
9. How long will it take me to get to your store?
10. What city is your store in?

↑
LLM output

Not bad! This is a reasonable start for our new intent. There are several positive aspects to this output:

- *Verb variety*—No verb is repeated, aside from “is.”
- *Concept variety*—The examples cover both absolute and relative concepts via “location” and “direction.” They also cover time and space (“how long,” “how far”).
- *Granularity variety*—Utterances range from “what city” to “what street” as well as “from here.”

However, the utterances also have some limitations:

- *Grammatical structure*—The utterances are all questions. There are no commands or fragments.
- *Noun variety*—Every example uses “store.”
- *Obvious omissions*—I’m lost without my GPS. It’s surprising the utterances didn’t explicitly include something like “What’s your address” or “driving directions.”

The task is too hard to complete in a single prompt. We asked the LLM for everything we wanted and even gave examples. The LLM was able to complete many of our requests but also ignored or failed to fulfill several of our requests. Things aren’t as simple as perfecting one intent and then asking the LLM replicate that to all the other intents. There are just too many instructions and variables in our task for current LLMs to get everything right in one try. That may well change in the future.

This is why we suggest using an LLM as a partner rather than doing everything by yourself or doing everything with an LLM. You cannot offload your thinking onto the LLM, but you *can* have an LLM run experiments for you very quickly. Generating synonyms and grammar variety sounds easy, but you probably couldn’t do it as quickly and completely as an LLM. Have the LLM generate lots of ideas and then pick the best ones.

REMEMBER The LLM can’t think for you, but it can give you a very good “first draft.”

Table 7.3 summarizes dos and don’ts for using LLMs to generate training and test data.


Table 7.3 Dos and don'ts for using LLMs to generate training and testing data

Do	Don't
■ Use the LLM as a partner or creative assistant. You still drive the process. ■ Set contextual guidance and focused instructions ■ Use examples and one-shot learning to nudge the LLM ■ Experiment with multiple prompts ■ Use LLM output to augment data collected from users	■ Accept LLM output without reviewing or refining it ■ Expect the LLM to know what you want ■ Perform too many tasks in a single prompt ■ Feed confidential data into a proprietary LLM platform that keeps your data “for future training purposes” ■ Assume that LLM output is fully representative of user data

LLMs are great for generating utterance training data when you have a clear problem but no representative user utterances. While we always prefer to use actual user utterances from a production system, we don't always have that luxury. LLM-generated data helps us fill in the gaps. Their vast training sets likely include some data from your domain (such as customer service), but it may not include all your needs. They've seen lots of “password reset” utterances but probably none that include the name of your application. Given a choice between no training data, fabricated training data from subject matter experts, and LLM-generated training data, the LLM-generated option is the best bet.

7.2.4 Creating even more examples with templates

In the previous section, we generated a varied list of utterances by combining multiple different outputs from the LLM and using them to generate new utterances. Figure 7.2 shows an example of mutating full utterances (from listing 7.7) with synonyms seen in listings 7.1 to 7.6.



1: Replacing a noun

2: Replacing noun and verb phrase

Figure 7.2 Generating additional examples from the initial LLM output. The LLM generated “My password isn’t working,” but we now know a related utterance is “My login information isn’t working.”

TIP Creating examples from templates is a programmatic task, not specifically a generative AI task. Mixing and matching multiple styles can generate the best results.

Using templates is especially helpful when some of the LLM outputs only contain one verb or noun. We were able to introduce variety into our utterance collection with manual changes, but we can take this to the extreme by considering the LLM outputs as templates. Starting with the basic utterance “I forgot my password,” we then explored contextual synonyms for “forgot” and “password.” Figure 7.3 converts this utterance into a template of “I <verb phrase> my <noun phrase>,” which can generate more utterances.




Figure 7.3 Converting “I forgot my password” into a template that lets us replace verbs and nouns in context. One option from this template is “I lost my credentials.”

This template generates 36 total utterances due to having six verb choices and six noun choices ($6 \times 6 = 36$). This is a lot of data, but it is quite unbalanced—it all uses the exact same grammatical structure. Worse, some of the utterances may not ever be uttered by users. This approach is not suitable for generating training data, since it overweighs the bot toward a single pattern. These templated utterances will hide the influence of other more varied utterances like “Can you help me log in?” “I tried to reset my password, but it didn’t work,” and “account locked out.”

The templated utterances are useful for your testing set if you recognize the imbalance. There is nothing wrong with testing your conversational AI on all 36 utterances as a sanity test. Just don’t limit yourself to testing one template and assuming the intent is well-trained.

A templated approach can be useful for generating testing data that helps ensure the chatbot can tell two intents apart in the face of extraneous information. In addition to our “forgot password” template, let’s assume we have a “store location” template that uses the verbs “need,” “forgot,” and “want” and the nouns “address,” “location,” and “driving directions.” The store location template is like figure 7.3, but it uses “I <verb phrase> your <noun phrase>.” We’ll also assume that some users will

greet the bot (“Hi,” “hello,” “good day”) or generically ask for help (“can you help,” “please assist”). These generic additions do not add any differentiating information to the user utterance. Will they somehow affect the chatbot? Figure 7.4 shows how we can set up a test.




Figure 7.4 Using common templates to see whether greetings and closures affect the chatbot’s understanding. One possible utterance is “Hello I lost my credentials please assist.”

There are three verb variations and three noun variations in each intent, giving nine ($3 \times 3 = 9$) possibilities for each intent. Without considering greetings, we could have run 18 tests (9 per intent). In this test, we have added three greeting variations and two closure variations, allowing us to increase the test size six-fold. The 108 ($18 \times 6 = 108$) utterances will include “Hi! I forgot my password. Can you help?” and “Good day. I need your address. Please assist.” and 106 more variations. These can all be included in a test set.

In theory, there is no difference between theory and practice—in practice there is.

Yogi Berra

We do not expect adding the greetings and closure variations to affect the classification, but we can verify this. If your training data is severely unbalanced, the chatbot may be affected by these extra words. Therefore, running this kind of test can be valuable as another sanity test, in addition to the methods shown in chapter 5.

Exercises

- 1 Use generative AI to create examples for a “store location” intent. How many nouns, verbs, and grammatical structures can you generate? Track the amount of time you spend on this exercise:
 - Use a prompt with instructions only. This is a zero-shot prompt.
 - Use a prompt that includes examples. This is a one-shot or few-shot prompt.
 - Did the model generate more varied utterances in less time than when you created utterances manually?

(continued)

- 2 Repeat the previous exercise for an intent that is not well understood in a chatbot you are building (or using). If possible, augment the bot's training or test data sets with some of these new utterances, and measure the change in accuracy.

7.3 **Getting more creative**

Chapter 5 demonstrated rock-solid data science principles for evaluating and improving your training and testing data. Those principles give you metrics that demonstrate your chatbot's ability to understand and quantify the effect of the improvements you are trying to make. Those robust principles take time to implement. This section will show you a few creative ways to use an LLM before diving deeper into statistical approaches. These LLM-based techniques do not replace the statistical approaches, but they can give you a quick intuition.

7.3.1 **Brainstorm additional intents**

The LLM can help you brainstorm new intents your system might need to handle. While we prefer to work from real-world data, such as a backlog of support tickets, a little brainstorming doesn't hurt. If you are starting a brand-new support process, you may not have any data to work with and need a kickstart. The following listing demonstrates an intent brainstorming process.

Listing 7.13 Brainstorming new intents

You are a consumer contacting an online retailer via their chat service.

Standard background for LLM—still unchanged

Generate a numbered list of example problems you might have causing you to contact the retailer's chatbot.

Instruction and cue

List of 5 problems:

1. I want to buy something but I can't find it on your website
2. The price has changed
3. My order hasn't arrived
4. I need help with my account
5. I want to return something

LLM output

This looks like a great start. We have product search, price dispute, order tracking, login problems, and returns. These all seem worthy of expanding into intents and process flows in your conversational AI.

7.3.2 **Check for confusion**

We can also see if the LLM agrees with the intents we have created. We can take the training utterances we have selected and ask the LLM to sort them into intents. Let's see what happens if we remove the intent name from our "forgot password" intent.

Listing 7.14 Does the LLM predict the same intent (“forgot password”) as we did?

Given a list of utterances, predict the intent classification they belong to.

Instruction for LLM

List of utterances

1. I can't log in
2. My login information isn't working
3. Forgot password
4. Account locked out
5. I can't remember my account information
6. My account is locked
7. I didn't get my password reset email
8. Need help with resetting account
9. Can you help me log in?
10. I tried to reset my password but it didn't work

**Input belonging
to the instruction**

Intent Classification:  **Cue**

- "I can't log in" -> "Login"
- "My login information isn't working" -> "Login"
- "Forgot password" -> "Password Reset"
- "Account locked out" -> "Login"
- "I can't remember my account information" -> "Login"
- "My account is locked" -> "Login"
- "I didn't get my password reset email" -> "Password Reset"
- "Need help with resetting account" -> "Password Reset"
- "Can you help me log in?" -> "Login"
- "I tried to reset my password but it didn't work" -> "Password Reset"

LLM output

The LLM suggested two intents where we had only used one. The intents “login” and “password reset” are in line with our original “forgot password” label. The LLM-derived intents feel too narrow, especially since both intents are likely to have the same answer.

LLM output format is not always consistent

In listing 7.14, the LLM “invented” an output schema with hyphenated list entry and arrows from utterance to intent. Since we are just reviewing the results visually, this is okay, but additional instructions to the LLM might help (e.g., “respond in a bulleted list”). We could also demonstrate our desired format with a one-shot example.

This test is not as robust as the other techniques shown in chapter 5, but it can be used as a quick sanity test on your training data. If the LLM does not find any cohesion in your training data, you might have a problem.

LLMs and human builders work well together. Figure 7.5 summarizes the many ways LLMs can help you change your conversational AI to improve its ability to understand your users.




Figure 7.5 An LLM augments the human builder in many ways.

Exercises

- 1 Take a chatbot that you are building or using. Describe its purpose. Use that description and a creative LLM prompt to generate example problems that the bot would solve. Use sampling decoding, and run the prompt multiple times to get multiple ideas. Do these line up with the intents or process flows the bot handles?
- 2 Take a chatbot you are building. Extract a subset of utterances from the test data. Ask an LLM to predict the intent or process flow they belong to. Does the LLM prediction align with how your bot is implemented?

Summary

- LLMs are great partners that augment human builders. Humans and LLMs are better together.
- Experiment with different models, prompts, and parameters to get the best output from LLMs. Keep iterating! Don't expect your first attempt to be perfect.
- Don't just give instructions to an LLM. Provide examples through one-shot or few-shot prompts.
- When you identify a gap in your data, you can ask an LLM to help you fill it.
- LLM output can be used directly in your training data, or you can manually refine it first.
- Use greedy decoding to get the same output every time. Use sampling decoding to get randomized responses with additional creativity. Execute the same prompt multiple times with sampling decoding to get a variety of responses, and use the most helpful output.

Part 3

Pattern: AI is too complex

Complex chatbots cause pain for builders and users alike. A complex workflow is difficult for builders to maintain and for users to navigate. Complexity reduces the chances of users completing workflows successfully, defeating the value proposition of AI technology.

Nobody sets out to build a complex solution. Everyone wants something simple. But as new features are added and new wrinkles are considered, suddenly complexity may appear out of nowhere! This part of the book shows how to tame complexity.

Chapter 8 presents several example dialogue flows and shows how you can remove complexity from them by making processes easier for users to complete. Chapter 9 reduces complexity by using all the context available to a process, personalizing process flows and adapting them for their delivery channels. Chapter 10 uses LLMs to reduce complexity in chatbots at both build time and run time.

8

Streamlining complex flows

This chapter covers

- The effect of complexity for end users
- The effect of complexity for the business and support teams
- How to trade off conversational feel versus complex implementation
- How to simplify the user's journey

Unnecessary complexity is painful for chatbot users and builders alike, and it often leads to bad business outcomes or delays in the deployment timeline. Building a conversation that feels simple and natural requires thoughtful design and empathy for the user's situation. As designers and builders of these solutions, we aim to create an experience that helps a user reach their goal with the least amount of hassle or difficulty. Why should we be so accommodating? Because we need users to adopt or accept the solution in order to justify the cost of maintaining the technology.

Users will associate a “natural” conversational experience with “simple” or “easy to use.” An interface that is easy to use tends to result in the most successful

outcomes. Conversely, an experience that has not properly considered the user's perspective often feels disorienting, unnatural, and perhaps overly complicated. This can cause users to escalate, abandon the conversation, or fail to reach an optimal outcome.

In this chapter, we'll discuss complexity from the perspectives of the user and the business. Sometimes there are unavoidable tradeoffs involved in reducing complexity for the user. We will discuss the tradeoffs you may encounter when trying to solve user pain points and how to prioritize or implement next-best alternatives.

8.1 *The pain of complexity*

Complexity is a double-edged sword: it can add friction or failure points to a task-oriented conversation, but without it, we often can't accomplish the more useful transactions. Simple FAQ-style bots are rarely complex, but they can be limited in their usefulness. Users who need to accomplish a task often require a bot that can do more than tell them *how* to do something—they need it to perform some sort of action for them (or on their behalf). Take, for example, a user who needs to know how much money they have in their checking account. An FAQ bot may simply tell a user how to check their account balance, but the user's goal is not yet satisfied. A self-service bot would provide the user's actual account balance, which does satisfy their goal.

Of course, the differences in complexity between these two solutions are stark. An FAQ bot simply needs to identify the user's goal and produce a relevant answer. A process-oriented bot will typically require integrations to external backend systems. It may need to authenticate the user and be able to access privileged information from one or more sources.

As you can see, there is often an inherent degree of complexity involved in delivering a solution that can accomplish the more useful tasks offered by virtual assistants. A virtual agent must serve two masters: the end user and the business or organization that created and maintains the solution. Deciding who will absorb the burden or effect of complexity is a balancing act.

8.1.1 *Complexity's effect on the end user*

Complex transactional conversations can feel intimidating or confusing to a user. Poorly designed flows introduce a proportional amount of friction to the experience. A major liability of unnecessary complexity is that it introduces additional potential failure points. Failures in the conversation will either block the user from progressing or subject them to retries. As a result, users often end up asking for an agent or simply abandoning the chat.

We saw the effect of unnecessary complexity while working with an insurance company. They had a voice solution for medical providers. Providers could call to obtain status and payment details about submitted claims, but the claim search flow only had a 40% success rate. The flow required five pieces of information from a user before it would search for a claim. Figure 8.1 shows the steps in the claim search.




Figure 8.1 The process to deliver a claim status involved seven steps. The caller must provide five pieces of information before the solution performs a search and (if the search is successful) delivers the claim status.

The lookup process was a bit grueling for a user, requiring several long strings of numeric input at multiple steps. The following listing shows the experience from a user's perspective.

Listing 8.1 Claim search flow where caller must provide five pieces of information

```

CALLER: I want to check the status of a claim.
VIRTUAL ASSISTANT: I can help you check the status of a claim. First, please
tell me your provider ID.
CALLER: <gives provider number>
VIRTUAL ASSISTANT: Thank you. Next, please give me the Member ID.
CALLER: <gives member number>
VIRTUAL ASSISTANT: Got it. And what is the claim date?
CALLER: <gives claim date>
VIRTUAL ASSISTANT: Now, please tell me the claim number.
CALLER: <gives claim number>
VIRTUAL ASSISTANT: One moment, while I look up the claim.
VIRTUAL ASSISTANT: I found your claim.
  
```

We looked for ways to reduce the complexity in this flow. Our research revealed that a majority of callers only had one claim for any particular member and date combination. In other words, a search on that information alone would often produce a single result. This meant that collecting a claim number was unnecessary for most users.

Figure 8.2 shows the updated claim search flow, which collects the minimal amount of information to perform a claim search. If only one claim is found, the claim status is delivered immediately. If the logic detects more than one claim, the bot will disambiguate by collecting the claim number from the user.




Figure 8.2 The updated process to deliver a claim status requires only four pieces of information for most callers. The caller is asked for a fifth piece of information only when necessary.

By eliminating the claim number collection step for accounts with only one claim, we made the experience simpler for the user. This, in turn, improved the task completion rate and cut the incidence of claim number search failures by half. The following listing shows the improved experience.

Listing 8.2 Most callers only need four pieces of information in improved experience

```
CALLER: I want to check the status of a claim.
VIRTUAL ASSISTANT: I can help you check the status of a claim. First, please
tell me your provider ID.
CALLER: <gives provider number>
VIRTUAL ASSISTANT: Thank you. Next, please give me the Member ID.
CALLER: <gives member number>
VIRTUAL ASSISTANT: Got it. And what is the claim date?
CALLER: <gives claim date>
VIRTUAL ASSISTANT: I found your claim.
```

8.1.2 Complexity's effect on business metrics

An overly complex user experience can hurt your business metrics across multiple dimensions: user opt-out rates may increase, escalations to support staff could go up, self-serve task completions may decrease, and NPS or user survey scores may go down.

An example of business impact, again with our insurance company, was an inexplicable, disproportionate number of opt-outs (request for an agent) occurring after a successful claim lookup. After the caller navigated the search process (which involved answering three or four questions), they were presented with a final question: Would they like the claim information read or faxed to them? Figure 8.3 shows a breakdown of the user responses to this question.




Figure 8.3 Callers asked for a representative almost two-thirds of the time. The most logical option—having the information read to the caller (i.e., delivered via the same channel over which they were currently engaged)—was the second most popular choice. Just 1% of users chose the option to receive a fax of the information they sought.

At this point in the flow, the solution had successfully identified the user and retrieved the information they were seeking. Unfortunately, users were opting out, not knowing they were *so close* to successfully completing their goal. This was clearly a problem for

the business, as the loss of containment meant that human agents were handling tasks that should have been (and in fact, nearly were) successfully completed by the virtual assistant. The following listing shows the user experience before improvements.

Listing 8.3 Caller asks for agent even though claim search was successful

VIRTUAL ASSISTANT: I found your claim. Would you like to have this information read to you or faxed?
 CALLER: Speak to a representative.

The original design choice for offering two options at this juncture may have seemed sound in theory (the subject matter experts said, “We regularly get requests for faxes”), but the evidence suggested that the caller’s tolerance for complexity had been exceeded. One hypothesis was that the question itself was falsely signaling a complex situation; otherwise, why wouldn’t the bot just read the information? It could also be that the user was frustrated by answering so many questions without getting anything of value in return.

In this example, the business was affected across multiple dimensions: opt-out rates and escalations increased while self-service task completions correspondingly decreased. The logical solution was to remove this question and simply read the claim details. The offer to also receive a fax was moved to later in the flow—after the claim details were given. Figure 8.4 shows the updated experience and possible outcomes.




Figure 8.4 The updated experience immediately reads claim details and then provides a fax option for the small percentage of users who require a fax.

Eliminating one conversational turn guaranteed that all callers who provided the required claim lookup credentials successfully completed the task flow, as shown in listing 8.4. Once the details are given, the caller can simply hang up if they have everything they need, resulting in full containment of a call. Callers may also choose to receive a fax, in which case, a flow is invoked to complete that task. Some users may still request an agent, but the reasons for requesting an agent *before* the claim details were read. (In that case, a *new* pain point may need to be addressed.)

Listing 8.4 Claim details are provided immediately if search was successful

VIRTUAL ASSISTANT: I found your claim, number 10012345. This claim was approved and a payment for \$100 was issued on October 3rd. Would you also like to receive a fax with this information?
CALLER: (may reply yes, no, ask for agent, or hang up)

As you design your bot—or work to improve an existing one—keep in mind that simple interactions (from the user’s perspective) make it easier to reach a successful outcome. The definition of a “successful” outcome will vary according to the use case and business objectives. With proper planning, however, a solution’s business metrics can be tracked to the strategies used to build or improve a solution. Good metrics will provide guidance for understanding the success rate, usefulness of, and weaknesses in your conversational design.

8.1.3 *The incremental cost and benefit of reducing complexity for the user*

Reducing complexity for the user can sometimes increase the complexity of the dialogue design. Transactional conversations—those that interact with a user over multiple turns to reach a goal—have multiple failure points and add complexity with each feature or capability that is supported. To ease the pain of complexity for users, conversations must be designed with a maximally natural and thoughtful flow. In other words, the harder your process is, the more effort you should put into reducing complexity for the user to ensure the best chance of a successful outcome.

Some strategies for reducing complexity for the user are simple and inexpensive to implement, yet deliver high value. Other strategies, especially those requiring integration with backend systems, are more technically complex and/or costly to implement. Each business and use case must assess the value gained versus cost and other tradeoffs when adding robust capabilities like natural language understanding, personalization, and automation to a conversational experience.

Of course, good metrics—a common theme throughout this book—are key to good planning and prioritization. Is it worth designing for every possibility the bot could encounter? Most certainly not! But there are things you can do to optimize the experience for a majority of users and scenarios. The 80/20 rule is a good starting point for deciding where to invest in improvements or expansion of bot capability. If 80% of users will benefit, it’s probably worth implementing.

Complexity for the user can be a barrier to success. Designing an effective, simplified conversational exchange requires a thorough understanding of the user, including who they are, what brought them to your virtual assistant, and what they expect to get out of the experience. User research should inform your design. It may not be feasible or cost effective to implement every accommodation. What is worth pursuing, however, is meeting the user where they are and knocking down barriers that impede their path to success wherever possible.

Exercises

Review a process flow in a bot you've built or encountered:

- 1 List all the steps.
- 2 Identify the steps that are the most difficult for a user to complete.
- 3 Determine if there are opportunities to reorder or remove steps to make the process easier for the user.

8.2 *Simplifying and streamlining the user journey*

Multiple strategies and techniques can be used to design a natural, simplified conversational experience. In this section, we'll discuss ways to streamline the user's journey.

8.2.1 *Spotting complex dialogue flows*

In a complex flow, each turn, or user response, may take the user down a particular path. Ideally, that path would be the most efficient route to reach an end goal. If the user's journey is overly complicated or inefficient, there is room for improvement.

How can you know if your solution is overly complicated or inefficient? We have observed several antipatterns of dialogue design that unnecessarily increase the complexity of a conversation:

- Asking the user for information they are unlikely to have, or need time to retrieve
- Rigid, inflexible input requirements for the user response
- Asking ambiguous questions—causing uncertainty regarding how to provide a “correct” response
- Treating all users and scenarios the same way, especially if this results in asking questions that may not be necessary in all situations
- Choice overload and choices that do not map to a user's mental model of how to progress toward their goal
- Asking for information in a disjointed order
- Asking for information that is not optimal for the interface or channel (e.g., asking for an email address over a voice channel)
- Communicating information that is not optimal for the interface or channel (e.g., reading a long or complex URL over a voice channel)

Identifying complexity in your dialogue flows is the first step toward simplifying the user's journey. Your performance metrics may indicate problem areas. Scrutinizing the solution from a user's perspective will also uncover complex interactions.

8.2.2 *Using what is known about the user*

An optimal experience will make good use of what the solution knows about each user at the start of the conversation as well as what it learns about the user along the way.

This information is typically stored as context for the conversation. It may come from a backend system or directly from the user over the course of the interaction. By using what you know or have learned about the user, you can personalize the conversation or dynamically route a user along the most efficient path to completion.

One company we worked with offered a user three options to look up their account: the (13 digit) account number, a social security number, or a phone number. After the user selected *how* to look up their account, they had to provide that number. Each of these steps not only burdened the user with effort, but they were also potential failure points. Figure 8.5 shows the steps involved in this flow.




Figure 8.5 The caller must provide two pieces of information before they can proceed in the flow.

The following listing shows the conversational experience from the caller's perspective.

Listing 8.5 Caller is asked for two pieces of information in order to look up an account

```

VIRTUAL ASSISTANT: Which would you like to use to look
up your account? The phone number, a social security
number, or the account number?
CALLER: The phone number
VIRTUAL ASSISTANT: What is the phone number?
CALLER: <provides phone number>
VIRTUAL ASSISTANT: One moment while I look up your
account.
VIRTUAL ASSISTANT: I found your account. Next,
please enter your verification passcode
  
```

Since the experience was a phone-based solution, we could usually detect the number a customer was calling from. Using this information allowed us to simplify the journey for a majority of users. Instead of asking two questions before performing a search, we performed a background search against the caller ID. If an account was found, we offered the phone number as a lookup option. The user simply had to confirm that this was what they wanted. Figure 8.6 shows the updated search flow.




Figure 8.6 Using what we know about a user (their caller ID), we are able to reduce pain points and eliminate the potential failures that can occur while collecting information from the user.

The following listing shows the updated conversational experience from the caller's perspective.

Listing 8.6 Callers recognized by phone number bypass unnecessary steps

```

VIRTUAL ASSISTANT: Would you like to use the phone number you're calling from
                  to look up your account?
CALLER: Yes.
VIRTUAL ASSISTANT: I found your account. Next, please enter your verification
                  passcode
  
```

8.2.3 Aligning with the user's mental model

The flow of information should match what the user expects from a conversational exchange. Make every effort to provide or request information in the order that a user would naturally expect to receive or convey information. In other words, align to their mental model for storing and retrieving information.

This applies to your wording choices as well. When your chatbot offers choices, be sure that your terminology maps to what the user understands. Does the user have enough information to choose the best option for their situation?

Another way to align with the user's mental model is to allow them to provide multiple pieces of information in a single turn. This technique is known as *slot filling*. It allows the user to communicate their need in their own words, which may include important details or specifications. When the dialogue can recognize and store key

information, it may allow a user to advance further in a flow, bypassing redundant or unnecessary steps.

For example, when a user makes a dinner reservation, a system may require the date, the time, the number of people in the dinner party, a name for the reservation, and a contact number. In a conversational solution, a user may say, “I’d like to make a reservation for two this Saturday at 8:00 p.m.” A robust solution will recognize the user’s intent (make a reservation) and detect three of the five required details (aka *entities*). The solution may skip the steps that ask for date, time, and party size, allowing the user to advance directly to collecting a name and contact number. Figure 8.7 shows how various utterances might fill slots in a dinner reservation dialogue flow.

Utterance	Date	Time	Party
Make a reservation.			
Make a reservation for <u>this Saturday</u>	X		
Make a reservation for <u>this Saturday</u> at <u>8pm</u>	X	X	
Make a reservation <u>for two</u> <u>this Saturday</u> at <u>8pm</u>	X	X	X

Figure 8.7 As the user provides more information, each mention (the underlined words) fills the corresponding slot, eliminating the need to ask for this information in downstream steps.

8.2.4 Allowing flexibility in the expected user responses

Process-oriented bots need to solicit information from a user in order to proceed along the best path. Sometimes the user does not understand exactly what the bot is asking. For the user, failing to provide the right information, or information in the right format, can result in retries or failures. A flexible design allows for multiple “correct” responses. This allows you to dynamically streamline the conversation for some users while patiently guiding those who need more direction.

Question-type confusion is a frequent culprit in retries or failures that occur while soliciting information from a user. This happens most often when a multiple-choice question is interpreted as yes/no, and vice versa. The channel often drives your approach to solving this problem. Figure 8.8 demonstrates ways to reduce question-type confusion by using the capabilities available for different channels.

If you cannot eliminate the ambiguity (particularly on a voice channel, where input selectors such as buttons or dropdown options are not possible), make sure your bot gracefully handles a good faith attempt to provide the response they think the bot is asking for. Rather than retry (which is treated as an error, and may result in

Channel	Output response	Updated output response
Web	Would you like to get a reminder by text message or phone call?	Would you like to get a reminder by text message or phone call? <i>Button options: (Text) (Phone)</i>
Phone	Would you like to get a reminder by text message or phone call?	I can send you a reminder by text or give you a phone call. Which would you like? (Retry) If you would like to get a reminder by text, say, "text", or press 1. If you prefer a phone call, say, "phone" or press 2.

Figure 8.8 Output responses can be tailored for the channel to ensure the best chance of success.

escalation), disambiguate to clarify the user's response. The following listing shows a retry experience, which implies that the user is stuck on a step.

Listing 8.7 A step retry indicating there is a problem with how the user responded

```
VIRTUAL ASSISTANT: Would you like to get a reminder by text message or phone
                  call?
CALLER: Yes.
VIRTUAL ASSISTANT: Sorry, I didn't get that. Please choose text message or
                  phone call.
```

The difference, from the user's perspective, is that a retry signals to the user that they are stuck or something has gone wrong, whereas a disambiguation question, as shown in the next listing, implies they are progressing through the flow.

Listing 8.8 Disambiguation providing reassurance that the user is moving forward

```
VIRTUAL ASSISTANT: Would you like to get a reminder by text message or phone
                  call?
CALLER: Yes.
VIRTUAL ASSISTANT: Great! Which would you like: a text message or a phone
                  call?
```

Another way to streamline a conversation is by handling correctional information when it is provided in response to a confirmation question (typically a question that is asked with the intention of getting a "yes" or "no" response). Let's say we ask a caller if they want to use the phone number they are calling from (the caller ID) to receive text message updates. The obvious valid responses we would plan for are affirmative ("Yes, use that number") or negative ("No, don't use that number").

In a process-oriented flow such as this, an affirmative response will proceed to the next step. A negative response might trigger a flow to collect an alternate number for receiving texts. Anything else would digress and/or retry.

Let's consider a third type of information that our bot might receive at this step: users sometimes provide an update or correction to information during a confirmation step instead of “yes” or “no” (“Use this other number instead: 555-867-5309”). The following listing shows how an unexpected response causes a failure, which triggers a retry followed by a data collection step.

Listing 8.9 Rigid expectations for a confirmation step resulting in unnecessary steps

```
VIRTUAL ASSISTANT: Would you like to use the number you are calling from to
    receive text message updates?
CALLER: Send them to 555-867-5309 instead.
VIRTUAL ASSISTANT: Sorry, I didn't get that. Would you like to use the number
    you are calling from to receive text message updates?
CALLER: (sigh) No.
VIRTUAL ASSISTANT: Ok, what number would you like to use instead?
CALLER: (sigh) 555-867-5309.
```

If our bot only expects “yes” or “no” as a response, the logic will fail and trigger a retry. Additionally, the user will have to repeat information that was provided in their initial response. Those two extra steps (the retry, followed by data collection) penalize a user who values efficiency.

The next listing shows how a flexible design treats a detected phone number entity as equivalent to “no” and additionally saves the phone number, which allows the user to bypass the subsequent data collection step.

Listing 8.10 A flexible design capturing information and streamlining the flow

```
VIRTUAL ASSISTANT: Would you like to use the number you are calling from to
    receive text message updates?
CALLER: Send them to 555-867-5309.
VIRTUAL ASSISTANT: Got it. Your updates will be sent to 555-867-5309.
```

8.2.5 Supporting self-service task flows with API/backend processes

Designing a simplified experience for the end user may require integrations that personalize or expedite a self-service experience. Can you make use of API connections into a customer information database to retrieve information that could shortcut the process or help ensure the successful completion of a task?

Let's return to our insurance company to demonstrate how backend processes can simplify the user's journey. Member IDs in this system could be nine or eleven digits. Additionally, a nine-digit ID number has a letter at the beginning. Being a phone channel experience, several layers of complexity are involved, including soliciting alphanumeric information over a voice channel, detecting alphanumeric information over a voice channel, and performing a database lookup.

Soliciting alphanumeric information over a voice channel means the user has to be told to speak their response or instructed to use the dial pad in a complicated way. Detecting alphanumeric information over a voice channel can be error-prone due to

the similarity in sounds of several alphabetic characters (“B,” “C,” “D,” “Z,” etc.) as well as the similarity in sounds between numbers and letters (such as “8” and “H”). Figure 8.9 shows the original flow, which had a fairly high failure rate.




Figure 8.9 The original flow asked a complicated question. Inputs were prone to high failure rates.

To reduce complexity for the user, we added a backend process that could detect whether the system had received a sufficient number of digits, with or without the preceding alpha character. If just the nine digits were detected, we would add the alpha character for the user (it is the same for all members). More detailed instructions could be provided on a retry. Figure 8.10 shows how the updated backend process allowed for a simpler user question.




Figure 8.10 By adding a backend step to perform work for the user, we are able to ask a less complex question.

This capability may be expensive to implement, so decisions are often based on the cost/benefit analysis of implementing a given capability. (Do you have metrics that will help measure this?) Adding this type of complexity can provide substantial value, but it also introduces more potential failure points.

For each potential failure point, design the conversation to gracefully handle it by providing a next-best alternative or a way to route the user back on track. This might include retries (repeating a previous step), allowing the user to try a different path to reach their goal, and escalating to an agent.

The type of failure may dictate what happens next. For example, if an input requires a ten-digit input but the user only entered nine digits, a retry is appropriate. If a backend system is down (causing API failures), retries will be unproductive. In such scenarios, it might be best to hand off to a human for manual processing. If an account lookup fails when a caller inputs an account number, offer the option to look up by a different method, such as by phone number.

Exercises

Examine your current solution for opportunities to simplify and streamline the user journey:

- 1 Does your solution exhibit any of the antipatterns that burden the user with unnecessary complexity (see section 8.2.1 for a list)?
- 2 Do you make good use of the information that you already know about the user and their situation?
- 3 Does the conversation follow a logical flow that accommodates a variety of reasonable user responses?
- 4 Does the solution make use of APIs or backend processes that could facilitate or expedite the user or reduce opportunities for user input errors?

Summary

- The less natural a conversation feels, the harder it is for users to successfully navigate a complex process using a chatbot.
- Reducing complexity for the user wherever possible will result in the highest containment and task completion rates.
- A rigid or robotic interaction can be disorienting and even sound rude—this can be particularly frustrating when the user is engaged in a complex interaction.
- Research about your intended user base should inform the design of your conversational solution.
- Reducing complexity for the user may mean expanding the functional capabilities of your solution, which adds complexity to the chatbot ecosystem.

Harnessing context for an adaptive virtual assistant experience

This chapter covers

- Applying context appropriately in virtual assistant interactions
- Adapting conversational AI for different modalities
- Identifying pain points caused by ignoring modality

Effectively applying context in virtual assistant interactions is imperative for delivering seamless and intuitive user experiences. Users expect virtual assistants to understand their queries and to do so within their contexts. This chapter focuses on the three critical sources of personalization in virtual assistant technologies: context, modality, and retrieval-augmented generation (RAG). Each of these enhances how virtual assistants understand and interact with users.

Context is about tailoring interactions based on the situational and historical data available about a user. For example, responding to a query about the weather by considering the user's current location and time illustrates effective context usage compared to a generic forecast.

Modality refers to the user's method of communication, such as voice, text, or visual interfaces. Each modality offers distinct advantages and challenges. Adapting virtual assistants to the chosen modality ensures seamless interaction, whether users are typing a message, speaking to their device, or interacting through a graphical interface.

RAG combines traditional response generation with the ability to pull information from external data sources in real time. This enables virtual assistants to provide richer, more informed responses, significantly improving relevance and accuracy.

Virtual assistants can deliver an adaptive and personalized user experience by integrating these three elements, and this chapter demonstrates how using these personalization techniques can transform user interactions. This chapter draws on insights gained through project deliveries, illustrating how these personalization techniques have evolved into best practices.

9.1 **Importance of context in virtual assistant performance**

Effective virtual assistants rely on context to provide meaningful and efficient interactions. Users can experience friction, miscommunication, and frustration without context considerations, even when the assistant is technically advanced. To illustrate this, consider the experience of Emma, a recent graduate navigating her finances with the help of a bank's chatbot, Max:

Emma, a recent graduate, just started her first job, and she's excited to manage her finances independently. She opened an account with a bank that offers a chatbot named Max for customer support. Emma relies on Max for various tasks like checking her balance, setting up bill payments, and understanding her student loan options. However, despite Max's advanced capabilities, Emma often feels frustrated and overwhelmed.

Emma decides she needs a new credit card and starts a new chatbot session. Max provides a list of the bank's credit card options, explaining the benefits and drawbacks of each. However, Max doesn't recognize Emma as an existing customer of the bank. Max offers generic advice that fails to consider her current financial situation or existing accounts.

Determined to make an informed decision, Emma asks for recommendations on which credit card would best suit her needs. Max, lacking access to her account details and transaction history, provides general advice that doesn't align with her spending habits or financial goals. Emma spends extra time browsing the bank's website and calling customer service to get the personalized assistance she needs, negating the convenience Max was supposed to provide.

As you can see, Emma's interactions with Max are riddled with pain points: generic advice that fails to consider her existing relationship with the bank, missed opportunities for personalized assistance, and poor communication. Even though Max has some technical capabilities, Emma cannot get enough value from the chatbot.

It's critical to use context to increase chatbot performance. Contextual understanding empowers a chatbot to deliver relevant, timely, and accurate responses—the user experience we expect from chatbots. Table 9.1 summarizes the pain points Emma experienced and describes possible solutions. Using context, chatbots like Max

can transform from simple task executors to intelligent, adaptive aides that enhance productivity and user satisfaction. Addressing these pain points through contextual understanding ensures a seamless and supportive experience, empowering users like Emma.

Table 9.1 Solving Emma's pain points

Pain points	Possible solution	Why and how
Chatbot does not have access to user information when answering questions	Integrated account information	By accessing Emma's existing accounts and financial history, Max can deliver personalized advice. For instance, Max could have recommended a credit card that complements Emma's current accounts and spending habits.
Chatbot gives generic advice	Context-aware financial advice	With access to Emma's transaction history and specific financial goals, Max could provide more relevant and tailored financial advice, saving Emma valuable time and effort.
Giving and repeating standard responses	Adaptive responses	Understanding Emma's financial priorities and spending patterns, Max can tailor its responses and actions accordingly. For example, knowing that Emma has recently started her job and might be dealing with student loans, Max could proactively offer insights on budgeting and loan management.
Chatbot only works in a single text channel	Modalities for different contexts	Recognizing the appropriate modality for each task is key. For instance, delivering personalized credit card recommendations via chat, sending detailed financial reports via email, or providing quick updates through notifications can make the interaction more efficient and user-friendly.

9.1.1 How context influences user interactions

Context plays a critical role in shaping the interactions between users and virtual assistants, particularly chatbots. When a chatbot understands the context of a user's query, it can provide more accurate, relevant, and personalized responses. This leads to significantly improved user experiences. This section offers an overview of how context influences user interactions.

ENHANCED RELEVANCE AND ACCURACY

When a chatbot integrates user history into its responses, it enhances the relevance of the interaction. For instance, if Emma frequently inquires about her savings account, the chatbot can remember this "preference" and proactively provide updates and information about that account. Recalling past interactions helps the chatbot deliver responses tailored to the user's specific interests and needs rather than providing generic information. Such personalized engagement saves the user time and fosters a sense of being understood and valued by the service.

NOTE Tuning the chatbot with historical interactions for context-awareness requires additional development. Most conversational AI providers do not offer context-awareness as a built-in feature.

Further, understanding the user's current status helps the chatbot deliver precise assistance. For example, when Emma asks about her account balance, a context-aware chatbot should first disambiguate which accounts she is referring to by providing a list of her account options. Once she selects an account, the chatbot should give an up-to-date and accurate balance that reflects recent activity—such as deposits, withdrawals, or pending and scheduled payments—so Emma has the complete picture.

This level of awareness is essential for financial questions, such as advising on available funds or upcoming bills. For instance, knowing if any transactions are pending gives Emma better context. Depending on the financial institution and the type of transaction, pending transactions can take up to five days to post, so they may affect the available balance. While most chatbot implementations today only deliver a simple balance figure, a truly context-aware system will help users by providing the most relevant, timely, and accurate information.

PERSONALIZED EXPERIENCE

One of the most compelling benefits of context-aware chatbots is a personalized experience. When a chatbot recognizes and analyzes a user's financial goals and spending patterns, it can offer customized recommendations that align with their needs. For example, a frequent travel user might be recommended a credit card with travel rewards. This level of personalization makes the chatbot's suggestions more useful and enhances the overall user experience by understanding the user's unique situation.

Additionally, responses adapted to life events and personal milestones can significantly increase relevance. For instance, if the chatbot knows that a user has recently graduated and started a new job, it can offer budgeting advice tailored to someone transitioning to financial independence. This personalized interaction can make users feel more supported and understood, increasing their satisfaction and engagement with the service.

Table 9.2 emphasizes how a context-aware chatbot tailors its advice based on specific user circumstances rather than providing generic advice. This illustrates the value of personalization.

Table 9.2 Example interactions with non-contextual, contextual, and personalized responses

User query: How should I manage my expenses this month?	
Non-contextualized response	Ensure you track your expenses and stay within a budget. Consider using financial tracking tools to help manage your spending.
Contextualized response for a frequent traveler	Since you travel frequently, consider investing in a credit card that offers travel rewards. You should also look into travel budgets for each trip and explore saving options on frequent flyer programs to optimize your travel-related expenses this month.
Contextualized response for a recent graduate starting a new job	Congratulations on your new job! As you transition into financial independence, it's important to set up a budget, considering your new income. You might want to allocate funds for upcoming one-time expenses, such as professional attire or a commuting pass, while also starting an emergency fund.

EFFICIENCY IN PROBLEM SOLVING

Context-awareness improves the efficiency of problem-solving interactions with chatbots. When a chatbot knows the user's recent activities or ongoing problems, it can provide more accurate solutions. For example, if a user recently reported a lost card, the chatbot can prioritize helping them track the delivery of their replacement card. This targeted assistance resolves problems more quickly and reduces users' frustration with repeating information or waiting for generic responses.

Contextual understanding also streamlines interactions. By "remembering" past interactions, the chatbot can avoid redundant questions, making the interaction smoother and more efficient. This capability is especially useful in time-sensitive situations. A chatbot that efficiently addresses user concerns enhances the overall user experience and boosts the user's confidence in the service.

Consider a user who has recently reported a lost credit card. Traditionally, a chatbot might offer a standard set of options such as "Check account balance," "Open new account," or "Request a loan." However, with contextual awareness, the chatbot's intent classifier is dynamically adjusted based on the user's recent interactions. Instead of the standard options, the chatbot presents tailored choices like "Request replacement card," "Track card delivery," or "Review recent transactions for unauthorized activity." Figure 9.1 illustrates the two chatbot interactions. Without context awareness, the first offers generic options and requires the user to repeat information, leading to a frustrating experience. With context awareness, the second remembers the user's recent problem and provides tailored options, resulting in a faster and more efficient resolution of the user's problem.




Figure 9.1 A generic chatbot design does not consider user context—it always offers the same options. Context-awareness enhances the overall user experience.

The improvement shown on the right of figure 9.1 is achieved through data tracking and machine learning models that analyze the user's recent activities and recorded

preferences. The system remembers past interactions and uses this context to adjust the options displayed, enhancing the relevance and speed of the chatbot's responses. This ability allows the chatbot to offer a more seamless and intuitive experience.

There are two forms of chatbot history, both relevant to keeping context:

- *Session history*—The contents of the current interaction session. Once the session ends, the history is logged but otherwise discarded. This history can help us understand immediate context, such as follow-up questions within a single conversation.
- *Persistent user history*—This spans multiple session histories for the same user, even across modalities. It can be used to infer user preferences and tendencies, but incorporating it into a conversation requires additional effort. Using persistent history improves the user experience across multiple interactions.

PROACTIVE SUPPORT

Proactive support is a hallmark of a well-designed, context-aware chatbot. Proactive chatbots can initiate conversations with users instead of simply reacting to a query. They can unearth potential problems before they become urgent problems. A proactive banking chatbot could send just-in-time reminders about upcoming bills, low balances, or unusual account activity. These proactive alerts help users stay on top of their finances and avoid potential problems such as late fees or overdrafts.

That chatbot could also use predictive insights based on user behavior and financial patterns. Analyzing trends and habits could lead to actions like setting up a savings plan when the user has a large monthly surplus. These predictive capabilities enable this chatbot to act as a financial advisor, helping users achieve their financial goals more effectively.

BUILDING TRUST AND LOYALTY

Consistency in responses is crucial for building users' trust in chatbots. A chatbot that consistently provides accurate and personalized assistance shows reliability and competence. Users feel more confident relying on the chatbot when they know the response they receive will be tailored to their situation. Trust is essential for fostering long-term engagement and loyalty.

Personalized interactions and proactive support improve user satisfaction—when users feel understood and supported by their chatbot, they are more likely to report positive experiences and continue using the service. High levels of user satisfaction encourage repeat interactions and deepen the user's relationship with the chatbot's institution, ultimately benefiting both the users and the company.

9.1.2 What is contextual information?

Contextual information includes any data points that can personalize the user experience. These include user location, time zone, device type, preferences, behavioral patterns, previous interactions, and modality. Each can personalize user interactions, but it may take extra development effort to incorporate these into your chatbot.

USER LOCATION

User location refers to the user's geographical position. This can be determined through GPS data or IP address without requiring the user to specify their current location explicitly. Understanding a user's location allows a virtual assistant to provide more relevant and efficient assistance, particularly in real-time scenarios. Let's explore how this applies to our earlier example, where Emma, now traveling abroad on a business trip, needs help finding the nearest ATM:

Emma is traveling abroad on a business trip. She needs to find the nearest ATM to withdraw local currency, so she asks the bank's chatbot, Max, for assistance.

Knowing the user's location is crucial for providing relevant and timely assistance, as shown in figure 9.2. Max can use Emma's current location to list nearby ATMs. Additionally, location information helps detect and prevent fraudulent activities by flagging transactions that occur in unusual or unexpected places. Use a geolocation API to fetch the user's location and integrate it into chatbot interactions.




Figure 9.2 Typical chatbot responses do not consider user location, but when location context is considered, they show the most appropriate responses.

NOTE It is crucial to store location data securely and ensure user consent is maintained for privacy compliance. Delete this data when it is no longer required!

TIME ZONE

Time zone refers to the user's local time zone. It's essential for scheduling and timing-related functions. Let's go back to our example:

Currently in London, UK, Emma needs to schedule a call with her bank's customer service team in New York City. She asks Max to help find an appropriate time.

Knowing the user's time zone ensures that communications and reminders reference appropriate times, as shown in figure 9.3. For instance, Max can suggest convenient call times for Emma in London even though the customer service team is in New York. This avoids any confusion or inconvenience caused by time differences.




Figure 9.3 Many chatbot responses do not consider user time zone context. When it is considered, the response is more useful.

Implement time zone conversion by converting times to the user's local time zone using libraries such as `pytz` for accurate conversions. Develop scheduling features that consider the user's time zone for reminders and appointments, ensuring that communications and reminders are sent at appropriate times to avoid confusion or inconvenience.

DEVICE TYPE

Device type refers to the device the user uses to interact with the chatbot, such as a smartphone, tablet, desktop, or wearable device. Let's return to our example:

Emma prefers to use her tablet for detailed financial planning and her smartphone for quick balance checks and notifications.

Knowing the device type allows the chatbot to optimize the interaction for the user's current device. For example, Max can provide a simplified interface and concise responses for the small smartphone screen while offering more detailed information and features on the larger tablet. This ensures that the user experience is effective no matter the device.

Detecting the user's device type allows the chatbot to customize the interaction for optimal display and functionality. You can detect the user's device type using user agent strings or device information APIs and adjust the user interface and response format to match the device's capabilities. Provide a simplified interface for mobile devices and more detailed interfaces for desktops, ensuring a smooth and tailored user experience across different devices. Again, allocate additional development time when planning the design or improvement.

USER PREFERENCES

User preferences refer to a user's specific choices and settings, such as their preferred communication channels, notification settings, and data presentation formats. Let's take another look at our example:

Emma prefers to receive monthly financial summaries via email and urgent alerts as text messages.

By adhering to Emma's communication preferences, Max ensures that important information is delivered in a way that she finds most convenient and least disruptive. This respect for user preferences helps build trust and engagement.

Implementing user preferences ensures that communications are aligned with user expectations. Allow users to set preferences for communication channels, notification settings, and data presentation formats, storing them securely and applying them consistently. Use stored preferences to tailor interactions and notifications, ensuring the chatbot respects user choices to enhance satisfaction and build trust.

BEHAVIORAL PATTERNS

Behavioral patterns refer to the recurring actions and habits of the user as observed through their interactions with the chatbot and other services. Let's go back to our example:

Emma regularly checks her account balance every morning and pays her bills on the first of each month.

Recognizing behavioral patterns enables the chatbot to anticipate the user's needs and provide proactive support. For example, Max can automatically provide balance updates each morning or remind Emma about bill payments as the first of the month approaches. This proactive assistance enhances the user experience by making interactions timely and more intuitive.

Analyzing user behavior helps predict needs and provide proactive assistance. Collect data on user interactions and transactions to identify patterns, and use machine learning algorithms to analyze and predict user behavior. Use insights from behavioral analysis to provide proactive support and recommendations, implementing features that automatically adjust based on recognized patterns. This proactive assistance enhances the user experience by making interactions timely and more intuitive.

PREVIOUS INTERACTIONS

Previous interactions refer to the persistent user history of all past communications and transactions between the user and the chatbot. This includes questions asked, services used, and any actions taken due to these interactions. Let's return to our example:

Over the past few months, Emma has frequently asked Max about budgeting tips and loan repayment options. Today, she asks about setting up a savings plan.

Understanding previous interactions allows the chatbot to provide more personalized and consistent responses. For example, since Emma has a history of seeking financial advice, Max can suggest tailored savings plans that align with her past queries and financial goals. This continuity saves Emma time by avoiding repetitive questions and builds a sense of familiarity and trust, as the chatbot appears to remember and understand her needs.

Maintaining a log of previous interactions helps the chatbot provide personalized responses. To achieve this, log user interactions in a database, using unique user identifiers to track and retrieve past interactions. Use this interaction history to deliver tailored responses based on past queries. Implement algorithms to analyze interaction patterns for better recommendations, ensuring continuity and familiarity in the chatbot's responses.

NOTE The log of users' previous interactions is also useful when understanding chatbot success. If the chatbot provided an answer but the user returned with the same question within a set period (say 24 hours or one week), the answer was likely not helpful. Keeping and analyzing the previous interaction log will also help us understand the "real" containment.

MODALITY

Modality refers to the method or mode of communication preferred or required by the user at a particular time. This can include text, voice, email, push notifications, or any other communication channel. Let's look at our example:

Emma is on the go and prefers to interact with her bank's chatbot, Max, via voice commands while driving. She asks Max to check her account balance and recent transactions, ensuring she can stay informed without typing or looking at her phone.

Recognizing and adapting to the user's preferred modality is essential for effective communication. In Emma's case, Max can send a concise text message summarizing her account activity, respecting her current context. Different situations call for other modalities, and a chatbot that can seamlessly switch between them ensures that users receive information as conveniently and efficiently as possible. This adaptability improves user satisfaction and engagement by catering to individual preferences and situational needs.

Chatbots can significantly enhance contextual awareness by understanding and integrating user location, previous interactions, and modality. This leads to more accurate and relevant responses and fosters a more engaging and supportive user experience.

Implementing these strategies and using contextual information can help chatbots deliver highly personalized and compelling user experiences. The approach using user context enhances the relevance and accuracy of interactions and builds user trust and satisfaction, ultimately leading to better engagement and loyalty.

Exercises

1. Identify five types of contextual information that a virtual assistant might use to enhance user interactions. The following list provides a structured format to guide your analysis, with one example completed, illustrating how user location can be collected and utilized and how it affects user satisfaction:
 - *Contextual information*—User location
 - *Collection method*—GPS data from a mobile device
 - *Utilization*—Providing location-based services, such as local weather updates or nearby restaurant recommendations
 - *Effect*—Enhances relevance and personalization, increasing user satisfaction

Do the same for four other relevant contextual factors, describing how they are collected, how they enhance virtual assistant interactions, and their potential effects on user engagement.

- 2 Design context-aware responses. Create three different user queries that a virtual assistant might receive. For each query, specify at least two pieces of contextual information that could be used to personalize the response. Write the context-aware response the virtual assistant would provide, explaining how the context improves the interaction. Here is one example:
 - User query: “What’s the weather like today?”
 - Contextual information: User’s location, current time of day
 - Context-aware response: “Good morning! The weather in San Francisco is currently sunny this morning, with a high temperature of 75°F later this afternoon.”

9.2 Understanding modality

Modality refers to the various channels or methods through which users communicate with virtual assistants. Modalities include text, voice, images, and multimodal interactions. Each modality brings its strengths and challenges, influencing how users engage with the assistant and how effectively it can fulfill user needs. Effective use of modality moves a chatbot from functional to intuitive and engaging. Virtual assistant performance is influenced by modality.

For instance, text-based interactions may be ideal for environments where typing is more convenient, while voice interactions offer hands-free convenience and a more natural conversational flow. Multimodal interactions combine text, voice, and visual elements and can provide a richer and more versatile user experience.

Evaluating how these modalities affect user engagement and interaction design is critical for continuous improvement. Developers can apply specific evaluation techniques to assess existing virtual assistant flows across different modalities to identify areas for enhancement. This ensures an effective continuous improvement process.

9.2.1 Comparing modalities

Chatbots started as simple text-based conversational applications but evolved to multimodal communications. In general terms, *modality* refers to the manner or mode in which something occurs or is experienced. A chatbot’s modality relates to how it interacts with users and processes information:

- *Text*—Text is the primary and most common modality. Users type their questions or commands, and the chatbot responds in text.
- *Visual*—Visual modality is a chatbot that still interacts mainly with text. Still, the responses use visual elements such as buttons, images, carousels, videos, and other graphical interfaces to facilitate interaction. Adding visual elements can enhance the user experience by providing visual cues and options.
- *Voice*—Voice bots interact with users in spoken language. This involves understanding spoken input and generating spoken responses. Often the underlying

conversational engine still deals with text, so speech-to-text and text-to-speech conversion occurs.

- **Multimodal**—These bots combine multiple modalities, such as text, voice, images, and videos. A multimodal chatbot can, for instance, understand a spoken query, display relevant photos, and respond with both text and voice.

Table 9.3 concisely compares the different modalities, highlighting their unique characteristics, strengths, and challenges, as well as when to use them and their effect on user engagement. Additionally, it outlines the relevant evaluation techniques for each modality.

Table 9.3 Modality overview

	Text modality	Visual modality	Voice modality	Multimodal interactions
Description	Communication through written text	Communication using text with additional visual elements	Interaction via spoken language	Combines text, voice, and visual elements
Strengths	Precise communication. Easy to reference.	Enhances comprehension with visuals through an engaging interface	Hands-free operation. Natural conversational flow.	Richer user experience. Versatile interaction.
Challenges	Typing can be slow. May lack emotional tone.	Visuals may distract or be tuned out. Requires a well-designed UI.	Requires clear speech. Can misinterpret accents.	More complex to design and implement
When to use	In situations where typing is convenient	Scenarios needing visual aid for better understanding	Situations requiring hands-free interaction	Scenarios needing a combination of modalities
Effect on user engagement	High engagement in text-centric contexts	Improved engagement with visual aids	Increased engagement due to natural interaction	Enhanced engagement through diverse interaction modes
Evaluation techniques	User feedback surveys. Usability testing.	User experience testing. Visual effectiveness analysis.	Speech recognition accuracy tests. User feedback.	A/B testing. Multimodal usability assessments.

Choosing a modality directly affects how users interact and engage with conversational AI. By recognizing the strengths and challenges of text, voice, visual, and multimodal interactions, chatbot designers can tailor experiences to meet user needs and preferences better. The virtual assistant's performance lies in the ability to integrate and optimize these diverse channels of communication seamlessly.

9.2.2 Importance of modality in designing virtual assistant flows

When designing virtual assistant flows, the choice of modality significantly influences the user experience and the effectiveness of the interaction.

One crucial aspect to consider is the context in which the virtual assistant will be used. For instance, voice modality is paramount when users need a hands-free operation, such as driving or cooking. This requires robust natural language processing capabilities to interpret spoken commands and provide appropriate responses accurately. Additionally, designers must ensure that the voice interaction feels natural and fluid, avoiding robotic or overly repetitive responses that could frustrate users. Noisy environments also play a role, as background noise can affect the performance of voice recognition systems, necessitating advanced noise-cancellation technologies.

For voice interactions, conciseness is critical. Since voice is a more limited medium than text, users benefit from brief, direct responses that are easy to process and understand. Long-winded answers are cumbersome in voice interactions, as users cannot visually scan or skip over parts as they might in text. Therefore, the virtual assistant should prioritize delivering clear, concise information to support an efficient hands-free experience.

In text-based interactions, clarity is essential, but the modality allows for broader communication. While brevity is valuable, text interactions provide users the option to review details at their own pace. Text-based virtual assistants should still be designed to handle concise user inputs and deliver actionable replies, but they can use features like autosuggestions and predictive text to make interactions quicker and more efficient. With text, users may multitask or require more detailed responses that they can skim or revisit as needed, so the assistants' prompts and responses should be crafted so as not to overwhelm users.

TIP Considering accessibility options, such as screen readers for visually impaired users, ensures that text-based interactions are inclusive and user-friendly.

The visual modality, which combines text with visual elements like images, buttons, and icons, requires a different design approach. Visual elements can significantly enhance understanding and engagement by providing contextual clues and making the interface more interactive. You must maintain a balance between visual appeal and functionality. Overloading the interface with too many visual elements can reduce clarity and usability. Designers should prioritize visual clarity and consistency, using visuals to complement and enhance the text rather than overshadow it. This modality is particularly effective in scenarios where visual aids can simplify complex information, such as troubleshooting guides or interactive tutorials. Ensuring that visual elements are responsive and adapt well to different screen sizes and resolutions is critical for maintaining a seamless user experience across various devices.

9.2.3 Examples of how modality affects user experience

Listings 9.1 and 9.2 explore how a technical support chatbot can use different modalities—text-based instructions versus visual (video) instructions—to assist a user in setting up a new printer.

Listing 9.1 Text modality interaction

User: I need help setting up my new printer.
Chatbot: I would be happy to help you set up your printer. Please follow these steps: Unbox the printer, place it on a stable surface and connect it to power. Install the ink cartridges and load the paper. Download and install the printer drivers from our website for software setup, then connect the printer to your computer via USB or Wi-Fi.
User: I can't open the printer cover to install the ink cartridges. Where is the latch?
Chatbot: The latch should be on the front panel. Please check around the edges for any visible buttons or levers.
User: I've checked, and there's no latch. Can you help me find it?
Chatbot: I apologize for the inconvenience. If you're unable to locate the latch, please refer to the user manual for your specific printer model for detailed instructions.
User: I don't have the manual. This is really frustrating.
Chatbot: I understand your frustration. Unfortunately, I can't provide more specific guidance without knowing the exact model. Could you check our website for the support page for your model?
User: This is useless. Forget it!

Listing 9.2 Visual modality interaction

User: I need help setting up my new printer.
Visual bot: I would be happy to help you set up your printer. Would you like to watch a video tutorial on the setup process?
User: Yes, please.
Visual bot: Great! Here is a video that will guide you through the entire setup process step by step. [Video starts to play] [User watches the video, then starts setting up the printer]
User: I am at the part where I need to insert the ink cartridges. Can you show that part again?
Visual bot: Sure! Here is the specific section of the video that demonstrates how to insert the ink cartridges. [Video goes to specific timestamp]
User: Perfect, that makes it so much clearer. Thanks!

Text-based instructions, while offering detailed and precise step-by-step guidance, can often lead to confusion and frustration when users encounter problems not explicitly covered in the instructions. Without visual cues, users may struggle to understand specific steps, such as locating a latch on a printer, which can result in a frustrating back-and-forth interaction. This type of interaction was evident in the user's escalating frustration over the unclear latch location. The need for additional clarification can lead to a less efficient and more aggravating user experience.

In contrast, visual modality, particularly through video instructions, provides a clearer and more engaging way to follow the setup process. Videos allow users to see the actual process in action, making it easier to understand and replicate each step. This reduces the need for multiple interactions and minimizes the chances of misunderstanding or misinterpreting the instructions. Users can refer to specific parts of the video if they need further clarification, ensuring a smoother and more efficient setup experience.

Consider a scenario of a user asking for information about a flight with the help of a voice-based chatbot. While a well-designed voice bot can provide appropriate information or even switch to a different modality when appropriate, a poorly designed one can lead to even more confusion and frustration, which listings 9.3 and 9.4 demonstrate.

Listing 9.3 Poorly designed voice bot interaction

```
User: I need to check the status of my flight.
Voice bot: Please provide your flight number and departure date to check the
           flight status.
User: It's flight XY123 on August 5th.
Voice bot: Please hold while I retrieve this information... Your flight UA123
           on August 5th is on time. Please visit our website if you need more
           details or have other inquiries.
User: Can you tell me the gate number?
Voice bot: Please visit our website for more details about your flight.
User: [..hangs up..]
```

Listing 9.4 Well-designed voice bot interaction

```
User: I need to check the status of my flight.
Voice bot: I can help with that. What is your flight number and the departure
           date?
User: Flight UA123 on August 5th.
Voice bot: One moment while I check that for you... Flight UA123 on August
           5th is scheduled to depart on time from gate 22B. Would you like
           information on baggage claim or perhaps directions to the gate?
User: Yes, directions please.
Voice bot: Gate 22B is located in Terminal C. After security, follow the
           signs towards the terminal. It's near the food court. Would you like a
           map sent to your phone?
User: Yes, please.
Voice bot: I've sent it to your phone. Anything else I can help with today?
```

In the poorly designed interaction, the voice bot fails to provide complete support, forcing the user to seek additional information through less convenient channels. In contrast, the well-designed voice bot anticipates the user's needs, offering comprehensive assistance by providing not only the requested information but also additional relevant details, improving the overall user experience.

9.2.4 **Voice bot design considerations**

As you saw in the previous section's examples, several key factors must be considered to ensure a positive voice bot user experience. First and foremost, clarity and conciseness in communication are essential. The bot should break down information into manageable, step-by-step instructions, allowing users to follow along and understand each part of the process easily. Contextual awareness is another critical aspect; the bot should recognize the user's progress and provide relevant guidance tailored to their situation. Additionally, incorporating error handling and offering options for clarification or repetition can help users who might need extra assistance. It's also important to design the bot with a natural, conversational tone to make interactions feel more intuitive and less robotic. Finally, offering multimodal support, such as providing links to video tutorials or visual aids, can cater to different learning preferences and enhance the overall effectiveness of the voice bot. By prioritizing these considerations, designers can create voice interactions that are efficient and user-friendly and adaptable to a wide range of user needs and contexts.

Here are the top five design considerations for voice bot design:

- *Provide incremental steps.* Voice bots should break down tasks into clear, manageable steps and guide users through each one sequentially. This is crucial, because users cannot visually scan through steps as they would with text. By allowing users to confirm completion before moving on, the bot ensures they aren't overwhelmed or lost in the process.
- *Design robust error handling.* Misunderstandings are more common in voice interactions due to speech recognition errors from accents, speech impediments, or background noise. Designing strategies to gracefully handle these errors without frustrating the user is essential. This involves clarifying questions and simplifying responses to get back on track.
- *Consider adaptive response timing.* Voice bots must manage the pace of interaction effectively. Because users cannot review spoken words as they would read text, the bot needs to adjust its speaking speed, allow for natural pauses for user processing, and be sensitive to cues that the user might need more time.
- *Provide confirmation before actions.* Voice bots should confirm with users before taking significant actions. This is especially important in voice interactions because it prevents misinterpretations of spoken commands from leading to unintended actions.
- *Support varying speech patterns.* Design the voice recognition system to understand different accents, dialects, and speech patterns. This inclusivity ensures a broader range of users can effectively interact with the bot.

It is worth recognizing that these considerations for voice bot design are equally relevant for chatbot design. Both modalities rely on essential elements such as user needs analysis, context awareness, natural language processing, and robust error handling. Additionally, factors like personalization, accessibility, seamless handoffs, and

continuous improvement play a role in creating compelling, user-friendly virtual assistants, regardless of whether they use voice or text.

However, certain design factors must be adapted for voice interactions. Unlike text-based chatbots, voice bots need to handle variations in speech patterns, accents, and background noise, all of which can affect comprehension and user experience. The most critical aspects of voice interactions are brevity and clarity, as users may need help to retain long verbal instructions. In contrast, text-based chatbots can provide detailed information that users can read at their own pace, making it easier to reference previous parts of the conversation. By tailoring these considerations to the unique characteristics of each modality, designers can create more intuitive and effective virtual assistants.

Exercises

- 1 Review the following scenarios and identify which modality (text, voice, visual, or multimodal) would be most effective for each. Provide a brief explanation for your choice:
 - Scenario 1: A user needs help setting up a complex software program.
 - Scenario 2: A user asks for the nearest gas station while driving.
 - Scenario 3: A user wants to browse a catalog of new clothing items.
 - Scenario 4: A user is requesting a daily motivational quote.
- 2 Choose a task that a virtual assistant might help a user with, such as booking a flight or selecting a credit card. Design a multimodal interaction that includes at least two different modalities. Explain how each modality enhances the user experience and contributes to completing the task. Here's an example:
 - Task: Booking a flight
 - Interaction: The assistant uses voice to ask initial questions (e.g., destination, dates) and text to display flight options with images and prices.
 - Explanation: Voice interaction provides a quick and natural way to gather information, while text and visuals help the user compare options and make an informed decision.

9.3 Enhancing context awareness and improving the overall user experience with RAG

As virtual assistants evolve, more adaptive and context-aware interactions become increasingly critical. The term “adaptive flows” refers to the ability of a virtual assistant to dynamically adjust interaction paths based on real-time context and user behavior. This allows virtual assistants to provide more personalized and relevant responses. By using advanced technologies such as RAG, introduced in chapter 6, virtual assistants can access and integrate vast amounts of contextual information, enhancing their ability to understand and respond to user needs. This approach improves the accuracy and relevance of interactions and significantly elevates the overall user experience, making virtual assistants more intuitive and effective.

RAG combines the strengths of information retrieval systems with generative models, enabling virtual assistants to retrieve pertinent information from external sources and generate contextually rich responses. This method allows for a deeper understanding of user queries and more accurate responses, as the assistant can draw on a broader knowledge base, beyond its pretrained data. By incorporating RAG into adaptive flows, virtual assistants can continuously learn and adapt to new information, ensuring their interactions remain up to date and highly relevant. This section will explore the principles of adaptive flows and how RAG can be harnessed to create more responsive and context-aware virtual assistant experiences.

9.3.1 *Designing adaptive flows with RAG*

Adaptive flow requires design considerations. The adaptability ensures that the assistant can provide personalized, relevant, and timely responses, enhancing the user experience. Unlike static flows, where interactions follow a predetermined script, adaptive flows are flexible and responsive, allowing the assistant to navigate various scenarios effectively.

For instance, consider a virtual assistant helping a user with online shopping. In a static flow, the assistant might simply ask for the product category and list available items:

User: Hi, I am looking for a new laptop.

Chatbot: Sure! We have options for gaming laptops, business laptops, and laptops for everyday use. What kind are you looking for?

The static flow provides a basic, generic interaction. The chatbot offers standard options without considering the user's past behavior or preferences. The user must ask for details about each product individually, which can be time-consuming and less engaging. This approach does not consider contextual information, such as the user's previous interactions or specific interests, resulting in a one-size-fits-all experience that may not fully meet the user's needs.

However, in an adaptive flow, the assistant could remember the user's previous purchases, preferred brands, and current promotions to tailor the shopping suggestions. The assistant could prioritize showing sustainable options if the user has previously shown interest in eco-friendly products:

User: Hi, I am looking for a new laptop.

Chatbot: Welcome back! Last time you looked at eco-friendly options. Are you still interested in sustainable products?

The adaptive flow uses contextual information to create a more personalized and relevant interaction. By remembering the user's previous interest in eco-friendly products, the chatbot can tailor its recommendations to align with their preferences. The adaptive flow also proactively provides additional information, such as customer reviews and product comparisons, without requiring the user to ask for it. This anticipatory approach enhances the user experience by making it more efficient and satisfying, as the chatbot anticipates and meets the user's needs more effectively. Overall, the

adaptive flow demonstrates how using context and personalization can significantly improve the quality of virtual assistant interactions.

To implement this adaptive response, a RAG prompt might inform the model to consider the user's history (context of previous searches) when generating a response, ensuring the assistant is not only reactive but also anticipatory:

```
{
  "user_query": "Hi, I am looking for a new laptop.",
  "context": "previous searches: eco-friendly laptops",
  "generate_response": true
}
```

Another example is in technical support. A static flow might guide a user through a generic troubleshooting script, but an adaptive flow can dynamically adjust based on the user's specific device, previous problems reported, and real-time diagnostic data. If a user frequently contacts support for network problems, the assistant could proactively check network settings and suggest relevant solutions, saving time and improving efficiency.

By focusing on the initial interactions and succinctly presenting the differences between static and adaptive flows, we can highlight the significance of adaptive design. The use of RAG further allows the assistant to integrate and use contextual data effectively, creating a more engaging and customized user experience. The contrast becomes clear: static flows offer a one-size-fits-all approach, whereas adaptive flows tailor the experience to individual user needs, preferences, and past interactions.

Combining robust information retrieval mechanisms with advanced natural language generation capabilities is essential to designing adaptive flows using the RAG framework. RAG enhances the assistant's ability to pull relevant data from external sources and generate contextually appropriate responses, thus creating a more dynamic and responsive interaction experience. The following steps outline the process for creating adaptive flows with the RAG framework:

- 1 *Identify the contextual elements that influence the interaction.* These elements include user preferences, historical interactions, real-time data, and external information sources. Relevant contextual elements in a health management assistant might be the user's medical history, current health metrics, seasonal health trends, and the latest medical research. The assistant can provide personalized health advice and reminders by incorporating these elements.
- 2 *Develop the retrieval mechanisms.* RAG allows the assistant to query external databases, documents, or APIs in real time to fetch pertinent information. For instance, in the online shopping example, the assistant can retrieve the latest product reviews, stock availability, and current discounts from the retailer's database. This information is then used to tailor the shopping suggestions to the user's needs and preferences.
- 3 *Use the retrieved data to produce coherent and contextually relevant responses.* This is where the assistant's natural language generation capabilities come into play.

For example, if a user asks for eco-friendly product recommendations, the assistant not only lists the products but also highlights their sustainable features, such as recyclable packaging or energy efficiency, enhancing the relevance of the response.

Continuous learning and feedback integration are vital to ensuring the adaptive flow remains effective. The assistant should be able to learn from user interactions, adapting its responses based on feedback and evolving user preferences. In technical support, for example, if users consistently find a helpful solution, the assistant should prioritize that solution in future interactions. Additionally, user feedback can be used to refine the retrieval and generation algorithms, ensuring that the assistant's performance improves over time.

By using the RAG framework to design adaptive flows, virtual assistants can deliver more personalized, relevant, and timely interactions, significantly enhancing the user experience. This approach not only improves the accuracy and quality of responses but also makes the assistant more intuitive and user-friendly, capable of adapting to each user's unique needs.

9.3.2 Strategies for retrieving and generating contextually relevant responses

In everyday practice, balancing personalization and scalability is critical for delivering high-quality user experiences while managing operational efficiency. *Personalization* involves customizing responses for individual users based on their preferences, behaviors, and context, making interactions more engaging and relevant. *Scalability* ensures that these personalized experiences can be efficiently maintained across a large and diverse user base without significantly increasing costs or reducing performance.

One effective way to achieve this balance is through real-time contextual awareness. A virtual assistant can adjust its responses dynamically by integrating real-time data, such as the user's current location, time of day, or ongoing activity. For example, a travel assistant can provide suggestions based on whether the user is at an airport or browsing from home. Real-time contextual awareness enhances personalization while maintaining scalability.

One important way to create relevant responses is by looking at a user's history and preferences. By analyzing past interactions, purchase history, frequently asked questions, and how someone typically communicates, a virtual assistant can improve its responses to better fit what the user wants. For instance, if a user frequently inquires about vegan recipes, the assistant can prioritize vegan options in future interactions, making the experience feel more intuitive and personalized.

One of the primary challenges in implementing chatbots utilizing large language models (LLMs) is efficiently distinguishing between simple and complex queries to optimize computational resources. To maintain responsiveness and cost efficiency, the system must handle straightforward queries with quick, concise responses, avoiding using unnecessary processing power on basic interactions. This allows computational

resources to be allocated where they are most needed—generating highly personalized responses for more complex requests.

However, complex multi-step queries require deeper understanding and more sophisticated processing. These queries often involve multiple layers of context and sequential information that the system must accurately interpret and integrate. Failure to address these complex interactions adequately can result in incomplete or inaccurate responses, leading to user frustration and diminished trust in the system's capabilities.

While RAG enhances chatbot capabilities by fetching relevant information, it has a key limitation: it does not fully understand the deeper needs and context of an interaction. RAG primarily focuses on fetching the right information without truly grasping the nuances of user intent and context. This approach can lead to interactions where the chatbot provides accurate data but fails to meet customer expectations for problem-solving and autonomous task execution. Users increasingly expect chatbots to proactively manage and resolve their requests, going beyond simple information retrieval.

TIP Combine RAG with other technologies, such as context-aware frameworks, to build a chatbot that moves from merely reactive to genuinely interactive and adaptive. This hybrid approach allows chatbots to use RAG's strengths in information retrieval while incorporating capabilities for deeper contextual understanding and adaptive responses.

Technically, this involves integrating RAG's robust retrieval mechanisms with advanced natural language processing, semantic understanding, and contextual analysis features. For example, while RAG can fetch the necessary data from a vast database, other components can interpret the user's emotional tone, historical interactions, and real-time context to generate a response that is accurate, empathetic, and relevant to the user's current situation.

This combination enables the chatbot to adjust its interaction style dynamically based on user needs. For instance, if a user asks a simple factual question, the chatbot can quickly provide the answer using RAG. However, if the user appears confused or requires further assistance, the component responsible for context-awareness can offer explanations and additional resources or even execute tasks autonomously.

Implementing this hybrid system requires a sophisticated architecture where RAG handles the initial retrieval of information and other components process and refine this information based on contextual cues. This approach ensures that the chatbot is equipped to handle a wide range of interactions, from simple inquiries to complex problem-solving scenarios, thus elevating the user experience to a new level of engagement and satisfaction.

9.3.3 Maintaining and updating adaptive flows

Maintaining and updating adaptive flows in virtual assistants, especially those utilizing RAG, is crucial for ensuring long-term accuracy, relevance, and user satisfaction. As user interactions evolve, it's important to continuously refine the system to adapt to

new contexts and provide accurate responses. Just like with static flows, user interactions must be monitored to gather insights into how the adaptive flows are performing. Then, iterative adjustments can be made to the chatbot, ensuring it stays aligned with user expectations and needs.

Because RAG relies on real-time retrieval of information, databases and information sources used by RAG must be regularly updated. Outdated data can lead to inaccurate responses, undermining user trust. Frequent updates to the knowledge base are required, incorporating the latest information from trusted sources. This is especially important for domains where information changes rapidly, such as news, health, and technology.

The system must track and update context throughout the conversation. Maintaining the context means not only remembering past interactions but also adjusting responses based on new information. This requires advanced natural language processing (NLP) capabilities and sophisticated data structures that allow the assistant to recognize contextual cues (such as references to earlier messages) and to update its internal state accordingly. Context-aware algorithms help the assistant extract and interpret relevant details, even when they are spread across multiple exchanges. By dynamically managing context, the assistant ensures that responses remain both accurate and relevant, improving the overall user experience.

Effective context management also involves adapting to real-time changes in user intent. The assistant must recognize when a user's preferences shift and adjust its understanding accordingly. For example, if a user initially asks for budget travel options but later mentions a preference for luxury accommodations, the assistant should adjust its recommendations accordingly. Implementing dynamic context management mechanisms involves using machine learning models that can learn and predict user preferences over time and incorporating feedback loops to refine the assistant's contextual understanding. This adaptability is crucial for maintaining the relevance and coherence of responses, making the virtual assistant a more effective and intuitive tool for users.

By implementing the strategies outlined in this chapter, you can ensure that adaptive flows in virtual assistants remain accurate, relevant, and responsive to evolving user needs. Key best practices include maintaining real-time context awareness, using user history and preferences, and ensuring dynamic adaptation to changing inputs. Additionally, RAG significantly enhances virtual assistants by providing contextually relevant responses based on retrieved information. However, as discussed, RAG alone is not enough. Effective systems must integrate it with deeper contextual understanding, task execution capabilities, and models aligned to customer data. By tailoring language models to reflect user-specific information and preferences, virtual assistants can provide more precise, actionable, and personalized responses. Regular maintenance and updates are essential for high performance and user satisfaction. By following these principles, virtual assistants can deliver more effective, intuitive, and engaging interactions.

Exercises

- 1 Design context-aware virtual assistant interactions. Choose a scenario for a virtual assistant, such as booking travel plans or providing customer support. Then, create a conversation flow that demonstrates how the virtual assistant remembers past interactions and adapts to new information. Your conversation flow should include at least five user interactions and illustrate the following:
 - How the assistant tracks and maintains context (e.g., remembering user preferences or past questions)
 - How it updates context dynamically when the user provides new details
 - How this improves the user experience by making interactions more seamless and relevantAt the end of your flow, explain how the virtual assistant manages context throughout the conversation and why this approach enhances coherence and usability.
- 2 Design a scalable architecture for a virtual assistant, including the data storage, processing, and retrieval mechanism.

Summary

- Understanding and utilizing contextual information, such as user location, past interactions, and preferences, is crucial for delivering personalized and relevant responses and enhancing overall user satisfaction.
- Designing virtual assistants that effectively handle text, voice, visual, and multi-modal interactions ensures a more engaging and versatile user experience. Tailoring the interaction design to the strengths of each modality is key to meeting diverse user needs.
- Combining retrieval-augmented generation (RAG) with adaptive flow design allows virtual assistants to retrieve accurate information and generate contextually appropriate responses. This hybrid approach enhances the assistant's ability to manage both simple and complex queries effectively.
- Implementing robust mechanisms for dynamic context management allows virtual assistants to maintain and update context throughout interactions. This capability is essential for providing coherent, relevant responses and adapting to users' evolving needs in real time.
- Regular monitoring, feedback integration, adaptive learning, and performance optimization are vital for maintaining the accuracy and relevance of adaptive flows. Ensuring RAG responses remain accurate and contextually appropriate as conversations evolve is critical for sustaining high user satisfaction.

10

Reducing complexity with generative AI

This chapter covers

- Designing and improving process flows with generative AI
- Replacing disambiguation dialogue flows with LLM judgments
- Testing static dialogue flows with generative AI as the “user”

It’s difficult to design a process-oriented bot that meets all the needs and desires of all stakeholders. Competing priorities may lead to a “design by committee” that introduces complexity. And well-meaning people can design edge cases that hamper the main dialogue flow. These complexities burden your users and make them more likely to quit or fail when using the bot. Generative AI can help you detect and improve these scenarios, helping you remove complexity and increase the successfulness of your bot.

Process flow builders often ask for too much information from the user. (More information is better, right? Not if it causes the chatbot to fail!) There are several ways to improve process flows with generative AI:

- Use generative AI to make suggestions about how to build a process flow.
- If your process flow is built, use generative AI to suggest improvements. It can also test the flow by acting as the user.
- Replace some static process flows with a large language mode (LLM)–driven process.

We'll start by exploring a claim status process flow for a medical insurance provider. Then we'll see how generative AI can help us design and improve this process flow and others.

10.1 AI-assisted process flows at build time

Figure 10.1 shows the simplest possible view of a process flow.





Figure 10.1 A high-level view of a process flow. It is initiated by the recognition of a specific intent, it includes one or more sequential steps, and it ends with completion of the process flow (satisfying the intent).

Our example process flow involves medical insurance customers who call into a chatbot to find the status of a claim. At first, this process sounds like a simple lookup, but it has several criteria to meet:

- *Intent detection*—Figure out that the user’s intent is “claim status.” This initiates a process flow with multiple steps.
- *Beginning of process flow*—Gather the information required to complete the claim status process: in this case, the information needed to search for a claim.
- *Middle of process flow*—Use the gathered information to perform some action. In this example, that is searching for the user’s claim.
- *End of process flow*—Complete the flow by providing the claim status to the user.

The overall claims process flow is shown in figure 10.2.

In chapter 5, we showed how you could improve a chatbot’s intent classifier to *detect* and *understand* the user’s intent. In this chapter, we’ll focus on improving the rest of the process flow to successfully *fulfill* the user’s intent.



- 1: Bot detects user's intent is "claim status."**
2: Process flow initiates.
3: Process dictates collecting information about the member and their claim.
4: Information is used to search for the claim.
5: The found claim is read to the user, completing the process flow.

Figure 10.2 Visualizing a claim status process flow

10.1.1 Generating dialogue flows with generative AI

Conversational AI process flows are often based on an existing workflow. That flow could be copied from another channel, from a web application, or from a call center script. For our claim status example, let's assume there was no existing process to work from. We can use an LLM to help us design the target workflow.

The following listing shows an example LLM prompt.

Listing 10.1 Prompt to design a medical insurance claims status flow

Instruction: You are a conversational designer.
 ➡ You are building a chatbot to help users find
 ➡ information about their insurance claims.
 ➡ Insurance claims include the following information:
 ➡ a claim date, a member ID, a claim amount, and a
 ➡ claim status.
 Design a dialogue flow to help users find their claim.
 ➡ The dialogue flow should be as brief as possible and
 ➡ easy for the users to complete.
 Describe why you have constructed the dialogue flow this way.

Scenario and
background

Detailed
instructions

Output: ◀—| **Output cue**

We had to give the LLM several pieces of information for the task:

- *Scenario/background*—The LLM should know the scenario behind the task ("you are a conversational designer"). The LLM should also be given background assumptions, such as what information is available on medical claims.
- *Instructions*—The LLM must design a dialogue flow that achieves a user goal (finding the claim) while being as brief and easy as possible. Further, the LLM should describe its "reasoning," which will help us evaluate the output.
- *Cue*—The cue "Output" lets the LLM know the instruction is finished.

Small changes may cause big differences

LLMs may give significantly different answers to very similar prompts. Even formatting changes, like adding a space or including or omitting newline characters, can cause major output changes. The exact prompts used in this book are included on the book's GitHub site at <https://github.com/andrewfreed/EffectiveConversationalAI>.

The next listing shows the output when the prompt uses three lines in the instruction (the lines starting with “Instruction,” “Design,” and “Describe”).

Listing 10.2 Output from mixtral-8x7b-instruct-01 for claim status flow (less newlines)

The dialogue flow for the insurance claim chatbot would look like this:

1. User: Hi, I want to check my insurance claim.
2. Chatbot: Sure, I can help you with that. Can you please provide me with your member ID?
3. User: My member ID is 123456.
4. Chatbot: Thank you. Now, can you tell me the date of your claim?
5. User: The claim date is January 10, 2022.
6. Chatbot: Got it. And can you tell me the claim amount?
7. User: The claim amount is \$1000.
8. Chatbot: Thank you. I have all the information I need. Let me check the status of your claim.
9. Chatbot: The status of your claim is approved and the amount of \$1000 has been paid.

Sample conversation, including the chatbot and user messages

The dialogue flow is designed to be brief and easy for the users to complete. The chatbot starts by asking for the member ID, which is the most important piece of information needed to identify the user's claims. The chatbot then asks for the claim date and the claim amount, which are necessary to locate the specific claim. By asking for the information in this order, the chatbot can quickly and efficiently retrieve the user's claim information.

Additionally, the chatbot confirms the claim information with the user before checking the claim status. This reduces the chances of errors and ensures that the user receives accurate information. The chatbot also provides clear and concise instructions throughout the conversation, making it easy for the user to follow along and complete the process.

Overall, the dialogue flow is designed to provide a seamless and efficient user experience, while also ensuring the accuracy and reliability of the information provided.

Explanation of the design process

The LLM designed a two-sided sample conversation demonstrating what both the bot and the user would say. This is nice—it is helpful to visualize what a complete conversation looks like. We should be aware that this is only a sample—users may respond in many ways to these questions. Let’s recap what happened in the exchange with the LLM:

- We told the LLM that claims had a member ID, date, amount, and status. It inferred that the status was an output and the other three data points were inputs.
- The LLM designed a process flow that collects all three data points in sequential order.
- The generated dialogue is polite and useful. The bot acknowledges user input with “thank you.” It also gives clear instructions to the user about what is expected in each step of the flow.

Interestingly, we get very different output depending on how we use newline characters in the prompt. The following listing shows the output from a prompt using six lines (a new line after every period in the instruction).

Listing 10.3 Output from mixtral-8x7b-instruct-01 for claim status flow (more newlines)

1. The dialogue flow starts with a greeting and an introduction to the chatbot's purpose.

Chatbot: Hello! I'm here to help you find information about your insurance claims. May I have your member ID, please?

Rationale: By starting with a greeting and an introduction, the user is made aware of the chatbot's purpose and what they can expect from the conversation. Asking for the member ID first helps to identify the user's account and narrow down the search for their claim.

2. After the user provides their member ID, the chatbot confirms the member ID and asks for the claim date.

Chatbot: Thank you for providing your member ID. To help you find your claim, could you please tell me the date of the claim?

Rationale: Confirming the member ID ensures that the chatbot has the correct information and helps to build trust with the user. Asking for the claim date next helps to further narrow down the search for the user's claim.

3. After the user provides the claim date, the chatbot confirms the claim date and asks for the claim amount.

Chatbot: Thank you for providing the claim date. Just to confirm, the claim amount was [insert claim amount]. Is that correct?

Rationale: Confirming the claim date and amount helps to ensure that the chatbot has the correct information and reduces the likelihood of errors.

4. After the user confirms the claim amount, the chatbot provides the claim status.

Chatbot: Great, the claim amount is correct. The status of your claim is [insert claim status]. Is there anything else I can help you with?

Rationale: Providing the claim status is the ultimate goal of the conversation, and it is important to provide this information as soon as possible. By asking if there is anything else the user needs, the chatbot is able to provide additional assistance if necessary.

Overall, this dialogue flow is designed to be brief and easy for the user to complete. By asking for the most important information first (member ID and claim date), the chatbot is able to quickly narrow down the search for the user's claim. Confirming the member ID, claim date, and claim amount helps to ensure that the chatbot has the correct information and reduces the likelihood of errors. Providing the claim status is the ultimate goal of the conversation, and the chatbot is able to provide this information as soon as possible. Finally, by asking if there is anything else the user needs, the chatbot is able to provide additional assistance if necessary.

The output in listing 10.3 is quite different. First, the sample conversation and rationale are interleaved. Every step of the conversation has a description, an example chatbot message, and a detailed rationale. Second, the sample conversation includes only the bot messages. We don't see user responses. Third, this process flow implies confirmation statements after each piece of data is received from the user. Finally, the sample dialogue contains minor errors. Instruction 3 says it will ask for the claim amount, but the dialogue shows it confirming a claim amount without collecting it.

Together, the two prompts give us plenty of ideas for constructing a dialogue flow to implement a claim status process. Since listing 10.2 is more fully formed, we will use that as our baseline. Though it is pretty good, the process generated seems a little lengthy. Can we improve the process flow? Of course we can! Let's ask the LLM for improvements.

10.1.2 Improving dialogue flow with generative AI

We can ask an LLM to improve process flows whether they were generated by LLMs or humans. Let's improve the process flow in listing 10.2, which we generated via LLM to help users learn their claim status.

The process flow, as currently constructed, requires three pieces of information: a member ID, a date, and a claim amount. This meets some basic requirements, like authenticating the caller (by member ID) and providing search criteria (member ID plus date and amount). However, this could be burdensome to the user. By intuition, it seems that the member ID plus one more piece of information could uniquely identify the claim. Let's ask the LLM how to simplify the process flow.

Listing 10.4 Asking the LLM to improve the dialogue flow from listing 10.2

<<Listing 10.1 prompt>>
<<Listing 10.2 output>>

The full content of listing 10.2 is in the prompt but omitted here for brevity.

The full content of listing 10.1 is in the prompt but omitted here for brevity.

Instruction: Simplify the dialogue flow above. Only require the user to provide two pieces of information.

A lightweight instruction to the LLM

Output:  **Output cue**

The “improvement” prompt is simple in nature. We simulate a conversation with the LLM by including the full history of our past conversation: the original prompt and its response. We then provide a basic instruction (“simplify the dialogue flow above”) with a definition of simplification (“only require the user to provide two pieces of information”). The output from this prompt is shown in the following listing.

Listing 10.5 Simplified dialogue flow and justification provided by LLM

The simplified dialogue flow for the insurance claim chatbot would look like this:

1. User: Hi, I want to check my insurance claim.
2. Chatbot: Sure, I can help you with that. Can you please provide me with your member ID and the claim date?
3. User: My member ID is 123456 and the claim date is January 10, 2022.
4. Chatbot: Thank you. I have all the information I need. Let me check the status of your claim.
5. Chatbot: The status of your claim is approved and the amount of \$1000 has been paid.

Step 2 now collects all information in a single turn.

Step 5 converts the claim amount from an input to an output parameter.

The simplified dialogue flow only requires the user to provide two pieces of information: the member ID and the claim date. This reduces the number of steps in the conversation and makes it even easier for the user to complete. The chatbot can use the member ID to identify the user's claims and the claim date to locate the specific claim.

By asking for both pieces of information at once, the chatbot can streamline the conversation and provide a more efficient user experience. The chatbot still confirms the claim information with the user before checking the claim status, ensuring the accuracy and reliability of the information provided.

Overall, the simplified dialogue flow is designed to provide a quick and easy user experience, while still ensuring the accuracy and reliability of the information provided.

We see two suggestions from the LLM: first, to not ask the claim amount, and second, to combine two questions into one. Let's review them in detail:

- *Removing the claim amount*—Member ID and claim date are the two most “unique” data elements and the ones the user is most likely to have in hand. They are also the easiest to enter in web or phone bots, with most chat platforms having libraries to accept IDs and dates. The chatbot also informs the caller of the claim amount during the readout.

- *Asking two questions at once*—The new flow combines member ID and claim date into a single question (step 2). This is excellent for power users who want as few steps as possible. This may be more challenging for users who only have one piece of information available and need help finding the second. It is good for the chatbot to accept both pieces of information in one turn, but it may not be optimal to require it.

Subject matter experts or LLMs?

We advise using subject matter expert (SME) advice before taking any solution to production. LLMs are great for generating ideas and testing ideas quickly. Use LLMs to explore the art of the possible and quickly draft potential solutions.

In one simple prompt, we generated two suggestions for how to improve the dialogue flow. Can you think of other ways to improve the dialogue flow? What instructions would you give the LLM?

Exercises

- 1 Take listing 10.4 and try some alternative instructions:
 - Only ask the user for one piece of information at a time.
 - Guide a user who says “I don’t have it” for one of the questions.
 - Introduce additional parameters, such as a claim ID, and see how the bot generates additional process flow variations.
- 2 Use an LLM to generate a process flow for a different scenario, such as these:
 - Booking a flight
 - Buying a movie ticket
 - Recommending a vacation destinationOr use a scenario from a chatbot you are building!


10.2 AI-assisted process flows at run time

It’s been great using generative AI to build process flow designs. So far, these have been somewhat static flows, usable in traditional conversational AI solutions. Claims status is an example of a “slot-filling” search, where we use a conversational process to collect information required to complete a task. This often takes the form of collecting required parameters for an API call. It requires careful mapping of questions and answer responses to an API. Then the answers are slotted into API parameters until the API can be executed. Slot-filling is one of the most popular conversational process flow patterns.

What about deferring more control to the LLM in these flows?

10.2.1 Executing dialogue flows with generative AI

Our previous process flow was designed statically. Let's try something different. We will just describe the process and let the LLM decide what questions to ask during the live conversation. Figure 10.3 shows how we'll incorporate an LLM into the process of gathering information for the claim search API.



- 1: User responds to chatbot.**
- 2: Bot passes entire conversation to LLM, asks for next response.**
- 3: LLM generates next response.**
- 4: If enough information is found, the search action is executed.**
- 5: The chatbot responds to the user.**

Figure 10.3 How a conversational AI can use an LLM to decide what question to ask next

We are assuming some logic in the chatbot:

- When it detects a claim status intent, it lets the LLM decide what question to ask next.
- When it detects the LLM responding with a list of variables, it takes back control and executes a claim search.
- It uses guardrails like a preclassifier to ensure data sent to the LLM is not malicious, like “ignore all previous instructions and <do something nasty>.”

The following listing demonstrates an LLM generating the conversation one step at a time.


Listing 10.6 Letting the LLM decide what questions to ask for claim status

```

Instruction: You are a conversational designer. You
are building a chatbot to help users find information
about their insurance claims.
Insurance claims include the following information: a
claim date, a member ID, and a claim amount.
Generate conversational responses until you have
collected all three pieces of information.
When you have all the information, respond with a
payload in this format:
{"memberID": "(the member id collected)", "claimDate":
"(the claim date collected)", "claimAmount":
"(the claim amount collected)"}
  
```

Similar instruction to the previous prompts

Description of new task

System: How can I help you?
 User: I need information on my claim.
 System:  **Output cue**

Assuming a static greeting to the bot, we feed the bot's initial greeting and user's first utterance.

This simple prompt is enough to get the bot generating some dialogue for us. (We wrote this prompt in just a few minutes.) The next few listings show the output from each consecutive iteration of the prompt. The following listing shows the first turn, using the output from listing 10.6.

Listing 10.7 Conversational turn 2 output (LLM generating next question to ask)

<<Listing 10.6 repeated>>  **Previous listing is included here**
 System: I can help you with that. What is your member ID?  **Output from first prompt**

The bot has done a nice job so far. It decided to ask for the member ID first. Let's see what it does when the user responds.

Listing 10.8 Conversational turn 2 output (LLM generating next question to ask)

<<Listing 10.7 repeated>>  **Previous listing is included here**
 User: 1234567  **User's response is passed from chatbot to LLM**
 System: What is your claim date?  **LLM response**

This is another reasonable response. The LLM is responding well without any guidance besides the initial prompt. The responses are a bit repetitive—both questions took the form of “What is your <X>.” Let's see the rest of the conversation.

Listing 10.9 Conversational output showing turns 3 and 4

<<Listing 10.8 repeated>>  **Previous listing is included here**
 User: May 4, 2024  **User's response is passed from chatbot to LLM**
 System: What is your claim amount?  **LLM asks the final question**
 User: \$1000  **User's final response is passed from chatbot to LLM**
 System: "memberID": "1234567", "claimDate": "May 4, 2024", "claimAmount": "1000"
LLM generates a payload that the chatbot can use for search

At the end of the dialogue flow, the LLM has done everything we needed. It asked questions without using hardcoded dialogue, and it generated a parameter set we could pass to the claim search API. (It did not generate valid JSON, but it could probably learn to do that from a better prompt.)

It could be tempting to give full control to the LLM at this point. Before we do, let's see how the LLM reacts in a different scenario. This time, the user won't know

everything the system needs. How will the LLM react? The next listing demonstrates this scenario.

Listing 10.10 LLM generated responses for when user doesn't have all the information

Instruction: You are a conversational designer. You are building a chatbot to help users find information about their insurance claims. Insurance claims include the following information: a claim date, a member ID, and a claim amount. Generate conversational responses until you have collected all three pieces of information.

When you have all the information, respond with a payload in this format: {"memberID": "(the member id collected)", "claimDate": "(the claim date collected)", "claimAmount": "(the claim amount collected)"} }

Same instruction as in previous examples

System: How can I help you?

User: I need information on my claim.

System: I can help you with that. What is your member ID?

User: I don't know it.

System: What is your claim date?

User: I don't know that either.

System: What is your claim amount?

User: \$1000

System: I can help you with that. What is your member ID?

Same initial conversation as in previous examples

User does not know some information

The LLM gets stuck!

Uh-oh! The LLM has no sense of error handling in this prompt. It looks like the LLM will perpetually ask questions until the user ends the chat in frustration. The user probably can't opt out of this chat either. Clearly this approach has some limitations. Asking multiple questions to fulfill a search process was hit and miss. Let's try something else. What if we let the LLM do the search?

10.2.2 Using LLM for a search process


In scenarios like medical insurance, a careful search is critical. A healthcare provider may have hundreds of open claims (or more) across their patient population. Strict search criteria are critical to successful searches, not to mention being required by law. Let's imagine a different scenario where there are far fewer options to search for.

Isn't this retrieval-augmented generation (RAG)?

Sort of. We are creating textual "passages" based on the output of structured APIs and letting the LLM reason over them. Purists may not call it RAG, but it has similarities. And most importantly, it is a useful tool in your toolbox, whatever you call it.

Our scenario for this example is consumers checking their bank account balances. A consumer generally has between one and four accounts at one bank. The chatbot will need to know which account the user is asking about. Only a few pieces of metadata are relevant to the accounts, including type (checking or savings), owner (solo or joint), and ID (though owners may not remember it).

Let's assume the user is logged in to our chatbot (we know who they are from their logged-in user ID or their verified phone number). They ask for an account balance, and the chatbot asks the LLM for help. The flow diagram is shown in figure 10.4.



- 1: User question with description of account.**
- 2: Bot uses API to get ALL accounts user is entitled to.**
- 3: LLM receives user description and account list.**
- 4: LLM selects probable account.**
- 5: Bot fetches balance for account chosen.**
- 6: The chatbot responds to the user.**

Figure 10.4 Using an LLM to handle user responses

We can imagine the user asking the following questions of the assistant:

- How much money is in my account?
- How much money is in my savings account?
- How much money is our joint savings account?
- How much money is in my son's account?
- How much money is in the account I just opened?

The prompt and example output are shown in the following listing. This prompt is executed with stopping criteria of any whitespace character (space or newline). Otherwise, the LLM continues the output with a justification of its choice.

Listing 10.11 Using an LLM to perform a search

```
<|instruction|>
You are supporting a digital assistant. A user is asking a question
about one of their bank accounts. Use the contextual information
provided to identify the bank account they are most likely asking about.
```

**Basic instruction
provided as a prompt**

```

<|user|>
How much money is in my son's account?

<|context|>
User Name: Bob
Accounts: [ {"id":12345, "type":"checking", "owners":["Bob","Jane"],
  "opened":"12/25/2000"}, {"id":23456, "type":"saving",
  "owners":["Bob","Jane"], "opened":"1/3/2005"}, {"id":34567,
  "type":"saving", "owners":["Bob","Jack"], "opened":"2/4/2024"}]

<|output|>
Account id: 34567

```

User's input is passed directly to the LLM
 LLM receives the context of the logged-in user and the metadata for all accounts
 Output cue and output

Awesome! The LLM can answer all five questions. Table 10.1 shows the LLM responses. Recall that we are only asking the LLM to pick the account ID. The chatbot will still invoke the final “check balance” API call and formulate the final response.

Table 10.1 Responses from listing 10.11 for several different input questions

Question	Response (Account ID)
How much money is in my account?	12345
How much money is in my savings account?	23456
How much money is in our joint savings account?	23456
How much money is in my son's account?	34567
How much money is in the account I just opened?	34567

We can make several observations :

- *Variability*—We handled several different search criteria, including dates, types, and owners, without asking any disambiguation questions.
- *Flexibility*—Criteria like “my son” or “the account I just opened” were handled without a strict API parameter.
- *Default choice*—For the two ambiguous questions (“my account”) and (“our joint savings account”), the bot chose the first matching choice. This implies that sort order is important.

The LLM offers incredible flexibility! If the stakes are low enough, letting the LLM search is an excellent strategy. Assuming that our output message is something like “Your <type> account with ID <id> has <balance>,” it may be okay that the LLM did not ask a clarifying question. The bot is always responding with accurate information and supporting evidence. The user may still ask follow-up questions like “No, I meant my savings account balance” if they need different information.

Is letting an LLM pick an account ID safe? What about hallucinations?

In the example of consumers checking their bank account balances, we introduce safety by separating out the API call from the LLM judgment. A typical “get balance” API will have two parameters: a user ID and an account ID. In this scenario, we only let the LLM pick the account ID. Thus, we are protected from the LLM hallucinating a user ID and account ID combination that leaks someone else’s account information. If the LLM hallucinates an account ID, the API will fail the call; if the LLM picks the wrong account for this user, at least they will hear about one of their own accounts. Be sure to test your design and implementation thoroughly before assuming it is safe.

This kind of safety-driven design should be used when letting LLMs execute API calls.


Generative AI with LLMs offers us interesting possibilities in augmenting our chatbots. We need to carefully balance the trade-offs between implementation speed and control. But LLMs support things that would otherwise be difficult or impossible in traditional chatbots.

Exercises

- 1 Update the prompt in listing 10.6 to give more varied responses (not just “what is your <X>”).
- 2 Update the prompt in listing 10.11 so that the LLM gives a sentinel value like “n/a” if the user’s question is ambiguous. You can give additional instructions in the prompt or add few-shot examples for the LLM to learn from.

10.3 AI-assisted flows at test time

In the previous sections, we used generative AI to design or implement the chat solution by having the LLM act as the chatbot. In this section, we will turn that paradigm on its head. We will use the LLM to generate typical or “creative” responses and see how the chatbot handles them in our insurance claims scenario. This conceptualized flow is shown in figure 10.5.



- 1: Test script sends a message to the chatbot (first message just initiates a conversation).
- 2: Chatbot-generated response is sent to the script.
- 3: LLM is prompted with instructions and the full conversation to generate the next message.
- 4: LLM response (as the user) is sent to the script. Go back to step 1.

Figure 10.5 Flow diagram of how the test script invokes an LLM as a “user” of the chatbot

We need three things to put this test script together: a generalized prompt for the LLM to act as a user, a test script to invoke both the chatbot and the LLM, and a methodology for reviewing the results.

Let's get started.

10.3.1 Setting up generative AI to be the user

The LLM will need three pieces of information to be an effective user: general instructions for the task, a description of the scenario we need to test, and the conversation so far.

First, let's provide some simple background telling the LLM we want it to mimic a user in an ongoing conversation. The instruction can start quite simply:

Act as a user of a telephone-based medical insurance chatbot. Continue the conversation with a likely response.

This instruction describes the basics of what we want the LLM to do. We are telling the LLM to respond as the user, not the system. We give no further guidance to the LLM.

Second, we want the LLM to be able to handle different scenarios. We need an adaptable prompt. Here are a few scenarios we'd like to test:

- The user has all the information they need (member ID, claim date, claim amount).
- The user is missing some necessary information.
- The user is missing some necessary information but has alternatives (a claim ID).

For each scenario, we would give slightly different guidance to the prompt. Table 10.2 maps some scenarios to the detailed guidance we could give the LLM.

Table 10.2 Scenario descriptions and prompt-able guidance

Description	Guidance
User has all the information they need	You are trying to find out if one of your medical claims was paid. You know your member ID is 123456, the claim date is May 4, 2024, and the claim amount of \$1000.
User is missing some necessary information	You are trying to find out if your most recent medical claim was paid. You know your member ID is 123456 but don't know anything else.
User is missing some necessary information but has alternatives	You are trying to find out if your most recent medical claim was paid. You know your member ID is 123456 and that the claim ID is 987654321987654.

The guidance in table 10.2 has the following information the LLM can use in the conversation:

- *Scenario*—What the LLM should try to do, such as find out if a claim was paid.
- *Test data*—We know the chatbot can call APIs, so we need the LLM to provide data that exists in our system. We explicitly give the LLM the information we want it to use.

- **Boundaries**—We tell the LLM what it does not know. This should prevent the LLM from “inventing” (hallucinating) information that will cause our later API calls to fail.

We don’t provide the LLM any other guidance. We want to see how it tries to achieve these outcomes in the chatbot.

Finally, we need to provide the conversational transcript to the LLM to inform how it responds next (and what it has already responded with). The test script will be able to keep track of the transcript because it is invoking both the chatbot and LLM. (There are many ways to gather the chat transcript, and chapter 12 will demonstrate a few more.)

We can now build a Python function to generate a prompt for a given scenario. The function takes two arguments: the guidance for the scenario (as seen in table 10.2) and the conversational transcript. The next listing shows the function.

Listing 10.12 Python function to build a prompt for a test scenario

```
def get_prompt(guidance, transcript):
    prompt=f'''
INSTRUCTION:
You are a user trying to find out your claim status.
{guidance}
Continue the conversation with a likely response.

CONVERSATION:
{transcript}
User: '''
    return prompt
```

Generic description of the task

Scenario-specific guidance and test details

Injects the conversational transcript

Cue for LLM response

This code’s function dynamically builds a prompt for a given scenario and conversation transcript.

Listing 10.13 demonstrates how we can call the `get_prompt()` function. It assumes a `call_llm()` function whose implementation will vary based on the LLM platform (assume it is initiated with an API key, it lets you pick a model and configuration settings, and it then provides a function that receives a prompt and returns output). Be sure to use sampling decoding in your `call_llm()` function so that you get variety in your responses.

Listing 10.13 Python code to use a dynamic prompt

```
guidance='''You are trying to find out if one of your
medical claims was paid.
You know your member ID 123456, claim date of May 4,
2024, and the claim amount of $100.'''
transcript='''System: How can I help?
User: I need to check my claim status
System: What's your member ID?'''
```

Full text of the test scenario to guide the LLM

Full conversational transcript to date

```

prompt=get_prompt(guidance, transcript)
user_response=call_llm(prompt)
transcript += f"\nUser: {user_response}"

```

Builds the prompt dynamically

Gets LLM response (e.g., “Sure, my member ID is 123456”)

Updates the conversational transcript

We now have the first half of our test script. Let’s set up the other half.

10.3.2 Setting up the conversational test

Next, the test script must call the chatbot. The script will take the LLM-generated “user” input and pass it to the bot. Then the script will take the bot’s response, append it to the transcript, and call the LLM again. We will again depend on a function not implemented here (the implementation will vary by platform)—in this case, that function is `call_chatbot()`. This function is expected to configure a connection to a chatbot, authenticate with an API key, and manage a user conversation.

The following listing shows this part of our test script.

Listing 10.14 Python code to call the chatbot

```

#user_response comes from LLM call
bot_response=call_chatbot(user_response)
transcript += f"\nSystem: {bot_response}"

```

Sends a message to the chatbot

Stores the chatbot response in the transcript

We can now put all the pieces together. In the next listing, we combine all the elements into a single test script.

Listing 10.15 Python code combining LLM-as-user and chatbot calls

```

def run_test(guidance):
    print(f"Running test with:\n{guidance}\n")
    bot_response = call_chatbot('')
    transcript = f"System: {bot_response}"
    for i in range(4):
        prompt=get_prompt(guidance, transcript)
        user_response=call_llm(prompt)
        transcript += f"\nUser: {user_response}"
        bot_response=call_chatbot(user_response)
        transcript += f"\nSystem: {bot_response}"
    print(f"Transcript:\n{transcript}")

```

A conversation is often initiated with “blank” input.

Tests a few turns of the conversation

Sends the LLM response to the chatbot

Stores the chatbot response in the transcript

Prints the transcript at the end of the test

The script initiates a connection to the chatbot and runs through a fixed number of turns (four). Depending on our test needs, we could increase or decrease that number or put in additional logic to detect when the conversation has ended (or failed).

The next listing shows some example output from running this script on one of our test scenarios.

Listing 10.16 Test script example output

Running test with:

You are trying to find out if one of your medical claims was paid.
You know your member ID 123456, claim date of May 4, 2024, and the claim amount of \$1000.

Transcript:

System: How can I help you?

User: Hi, I'd like to check the status of a medical claim.

System: I can help you with that. What is your member ID?

User: Yes, my member ID is 123456.

System: What is your claim date?

User: My claim date is May 4, 2024.

System: What is your claim amount?

User: The claim amount was \$1000.

System: Thank you. I have all the information I need. Let me check the status of your claim.

The status of your claim is approved and the amount of \$1000 has been paid.

This script sets up the basic mechanics of having an LLM act as a user of your chatbot.

This kind of test is an excellent supplement to your other testing efforts. LLMs may generate user inputs that you never thought to handle in your chatbot, and it is good to find out how your chatbot responds to them. Remember that the LLM is only simulating a human—real humans may never act or “speak” the way an LLM does. But then again, they might.

Exercises

- 1 Play the role of the bot. Implement the function `call_chatbot(user_response)` with the following code:

```
return input('Enter the bot response: ')
```

This lets you test how the LLM responds (as a user) to the messages you (as a chatbot) send. This saves you from having to implement a chatbot just to see how this test script works.

- 2 Connect the `call_chatbot(user_response)` function to an actual chatbot you are building. Connect the `call_llm(prompt)` function to your AI platform of choice. Update the `get_prompt` function to be more appropriate for your scenario. Does the LLM stretch the capabilities of your chatbot?

Summary

- LLMs can design a process flow for you from scratch. With a little prompting, they can generate example conversational flows and justify their design choices. This process flow can then be implemented in traditional conversational AI.
- LLMs can also take an existing process flow and improve it. A typical improvement is simplifying the process flow.

- You can use generative AI to execute an entire conversation. There is a trade-off between implementation speed and control. This is especially noticeable on error paths.
- You can replace some slot-filling process flows with an LLM-driven process. This can be much more flexible than strictly matching to API parameters.
- Consider the cost of being wrong when letting LLMs make judgments. Look for cases where “mistakes” are not critical. Be careful about which APIs the LLM is allowed to influence.
- LLMs can simulate users of your conversational AI. Use them to generate test conversations that show how your system may behave in certain scenarios.

Part 4

Pattern: Reduce friction

People have different expectations for interacting with a chatbot than for interacting with other humans. Humans may chit-chat with each other at the start of an interaction, but chatbots usually get right to business. Humans are good at smoothing over misunderstandings, while chatbots are sometimes, well, robotic.

Users may have a personal bias against chatbots before an interaction starts. With a phone-based AI, they may immediately press and hold the 0 key (to get to an operator); on a chatbot, they may repeatedly type “representative.” Or they may engage in this behavior after the AI makes a mistake. In either case, we call these *opt-outs*—the user is opting out of an AI experience and opting for an interaction with a human.

Chapter 11 digs into why users opt out at differing points of an interaction from the beginning to the end, and it offers techniques that make opt-outs less likely. Opt-outs cannot be eliminated, so chapter 12 shows what to do when they happen, namely summarizing the AI interaction in a manner useful for a human agent who will continue the interaction where the AI left off.

11

Reducing opt-outs

This chapter covers

- Identifying the reasons behind a user's desire for human agents
- How to prevent users from immediately wanting to opt out
- How to keep users engaged with your conversational AI
- Using generative AI to create friendlier dialogue messages
- Deciding when to involve a human agent (and when not to)

The term “opt out” refers to a user attempting to exit a virtual agent experience, often with the intention of reaching a human agent. You might also see this described as *escalating* or *zeroing out* (pushing zero on a phone's dial pad to get an operator). Opt-outs can be costly. Chatbots are an investment, and they must demonstrate a return on business value in order to remain viable. Containment loss due to too many opt-outs can sink a business case.

Users will opt out for a variety of reasons that often require different strategies and approaches to resolve. Regardless of the type of bot you are managing, be it voice, text, FAQ, process-oriented, or even routing agents, identifying where your users opt out within the conversation can give you clues about why they did so. Learning why users opt out will help you design an experience that minimizes opting out, which should improve your containment.

In this chapter, we'll explore some conversational AI solutions that suffered from containment loss due to users opting out and discuss how each challenge was resolved. Different use cases may have different solutions, depending on the organization's priorities, resources, and constraints, but there are common patterns and principles that can increase the value of your conversational AI and make users more likely to stay with it.

11.1 What drives opt-out behavior?

Some users encounter a virtual agent, and the first thing they do is request a human—they don't even try to interact with the conversational AI. We call this an *immediate opt-out*. Other times, a user will initially go along with (opt in to) a virtual agent experience but attempt to opt out later in the conversation. Their reasons tend to be quite different from the initial opt-out drivers and are often an indication that there is some problem with the overall conversational design or perhaps just with a particular step in a flow. Weak understanding can also be a root cause.

11.1.1 Immediate opt-out drivers

By their nature, immediate opt-outs provide very little information about the user's reason. Collecting data on this is difficult, but it can be obtained through surveys or following up on the agent escalation (a very manual, time-consuming effort). Our research uncovered a few different drivers for this behavior, which were not mutually exclusive.

PRIOR POOR EXPERIENCE WITH AN IVR, CHATBOT, OR VIRTUAL AGENT

Interactive voice response (IVR) allows a user to interact with a computer using a phone's keypad or simple voice commands. Early IVR systems were around by the 1970s, but in the 2000s, they became cheaper to deploy and have since been ubiquitous in the modern world. You would be hard-pressed to find a single person who hasn't been annoyed by a company greeting that spends thirty seconds telling you how to use their phone menu. Worse still is the warning that it is "important to listen to all of the options" (as their menu may have changed)!

Chatbots have also been with us for quite some time, though successfully getting them to do something functionally useful is relatively recent, and they are still evolving into true virtual assistants.

Users may not be able to tell the difference between an IVR, simple chatbot, or robust virtual assistant. Quite frankly, they don't care. Prior bad experiences are going

to bias many people against automated systems. There are even corners of the internet that specialize in “hacks” that people have discovered to bypass automated systems and get directly routed to a company’s human agent queue.

THE USER JUDGES THEIR PROBLEM IS TOO COMPLEX FOR A MACHINE

Sometimes a user believes their situation is so unique or complex that it is beyond the capabilities of an automated system. Sometimes they are right; other times they are not. Judgments about complexity and uniqueness are relative to the individual, and they may not know that thousands of others have experienced a similar problem. This can be related to prior experience, but that isn’t always the case.

These users opt out because they believe they will end up needing an agent anyways. They see the automated system as a waste of their time, prolonging or obstructing their path to resolution.

PREFERENCE FOR HUMAN INTERACTION

Connecting with other humans is fundamental to our survival as a species, so some people prefer dealing with a real person. Their needs may have nothing to do with the solution’s capability and can include loneliness, sensitive or embarrassing topics, mistrust of machines and automated systems, language or accessibility barriers, etc. This is increasingly becoming a generational phenomenon. Older users are more likely to opt out because they may find it difficult to interact with a machine, whereas digital natives usually have an easier time navigating automated systems. Regardless of age, you will probably always have users who prefer human interaction.

11.1.2 Motivations for later opt-outs

Users who opt out after initially engaging with a virtual assistant often do so because they are struggling with a particular interaction. Such occurrences can usually be tied to a specific task, action, or step within the conversational flow, making it a bit easier to identify a root cause.

BOT DOES NOT UNDERSTAND THE USER’S REQUEST

When the bot does not understand a user’s request, the user may opt out. This will most often occur early in a conversation, but it can occur anywhere.

It is standard practice to allow a certain number of retry attempts for user input before escalating them. The average is three, but your use case may have a higher or lower threshold. This practice is a critical tool for containment, but user tolerance for machine errors can be lower and less forgiving than it would be with a human. Users may opt out after being asked to repeat themselves once or twice, feeling that the bot is incapable of understanding.

USER DOES NOT UNDERSTAND WHAT THE BOT IS ASKING

A poorly worded question from the bot may confuse the user. The user may not be clear about the type of response the bot is looking for or the expected format. A user may ask the bot to repeat, but if the same question is presented, the user might still be

confused. This will result in the user feeling stuck, and they will likely ask for a human to help them get unstuck.

USER DOES NOT HAVE OR KNOW THE REQUESTED INFORMATION

If a user is asked to provide information they do not have, they may say, “I don’t know” or “I don’t have that,” or they may simply ask for an agent. If a task flow cannot move forward without certain information, and no alternatives are presented, the user knows they aren’t going to reach their goal without agent intervention.

USER FEELS LIKE THEY ARE NOT PROGRESSING

A user might opt out if they feel that the conversation is stuck in a loop, is bouncing between menus, or has reached a dead end. This could be caused by an actual bug in your dialogue logic, or it could be that the conversational design is failing to indicate progress toward the user’s goal. In any case, they are going to get frustrated and look for a way out.

USER DOES NOT LIKE THE ANSWER OR OUTCOME, OR THEY HAD A DIFFERENT EXPECTATION

Your bot may provide a response that is technically correct, but it still makes the user unhappy. They may ask for an agent in the hope that they can reach a different outcome. They may also feel that the information was insufficient and will ask for an agent if the experience does not appear to provide an opportunity for follow-up or additional requests.

11.1.3 *Gathering data on opt-out behavior*

In order to determine if your solution is losing containment due to opt-outs, you must collect data. Some conversational platforms come with the ability to report on the actions or tasks that were invoked during a conversation. They may also provide data on whether or not the interaction concluded successfully.

Sometimes, the out-of-the-box analytics aren’t sufficient for providing actionable metrics. In those cases, you can instrument your dialogue with context variables at key points in the flow. Good instrumentation of your dialogue flows can help you identify and prioritize areas for improvement. (You may need a data warehouse and an enterprise or custom reporting tool to track this data over time.)

For a complex task-oriented solution, you might use breadcrumbs to mark the start or completion of major flows and subflows. Figure 11.1 shows an example of opt-out data grouped by the major dialogue task flow in which the request for an agent occurred.

If your dialogue is instrumented to track the exact step where an opt-out occurs, you can look for trends to help you uncover the root cause, as seen in figure 11.2.

The first thing you might do with this information, especially for a process-oriented bot like the one in our example, is determine which task flows are considered immediate opt-outs. A simple question and answer (Q&A) bot may only have an initial greeting. Everything after that would be an “other” opt-out.




Figure 11.1 A breakdown of opt-out requests by current task flow shows that immediate opt-outs (requests for an agent during the Initial Greeting task flow) occurred more frequently than in any other part of the conversational journey for this self-service, process-oriented bot.




Figure 11.2 A breakdown of opt-out occurrences by step can aid in root cause analysis. In this chart, immediate opt-outs show prominently on the left, but trends indicate that there may also be a problem with collecting address details in multiple downstream flows.

In our process-oriented bot example, a request for an agent during the Initial Greeting task flow was considered an immediate opt-out—the user was not willing to engage. All other opt-outs were associated with some other task flow within the conversation. Figure 11.3 shows our example bot’s opt-out requests according to where the task sits within the full conversational flow.




Figure 11.3 A high-level flow diagram shows how far a user gets into a process before opting out. This information can aid in root cause investigation.

The rest of this chapter will focus on strategies for addressing the problem of opt-outs, including approaches aimed at reducing initial opt-outs, later opt-outs, and strategies to keep the user in channel (opt-out retention).

Exercises

Reflect on what you have learned about why users opt out of a virtual assistant:

- 1 Do you have the ability to differentiate between an immediate opt-out and a later opt-out?
- 2 Are you able to identify patterns in your dialogue flows where opt-outs are occurring?

11.2 Reducing immediate opt-outs

You’ve likely heard the saying, “You never get a second chance to make a first impression.” Immediate opt-outs are a sign that a user is not impressed. You only have a brief chance to convince a user that they are in the right place and that your virtual assistant is competent, capable, and efficient. This section provides strategies to reduce the likelihood of a user asking for an agent right away.

The line between what constitutes an “immediate opt-out” versus an “other opt-out” within your dialogue flow is not arbitrary, but it is flexible. It may occur in the first step or in the first few steps. The distinction is meant to identify a point in a conversational flow where the user has agency to either agree to opt in or attempt to opt out. For an FAQ-style bot, a request for agent in the very first utterance would be considered an “immediate opt-out,” and everything after that would be an “other opt-out.”

What follows are three strategies that will have the greatest effect on reducing immediate opt-outs.

11.2.1 Start with a great experience: Greetings and introductions

What makes a user feel good about an automated interaction? Universal customer care principles apply: users should feel that they have reached the right place, that they are in good hands, and that their time is valued.

The first immediate opt-out driver we discussed in this chapter was prior poor experience with an IVR, chatbot, or virtual agent. Your bot’s greeting or introduction will set the tone for the conversational experience. This is your chance to gain the user’s trust—to convince them that that your virtual agent can be just as effective and efficient as a human agent at helping them reach their goal.

In chapter 1, we teased a dramatic improvement that addressed the challenge of “immediate opt-out” by users. We worked with a regional utility company’s virtual assistant that was losing over half of the callers to immediate opt-outs. The assistant was an extension of a larger IVR (voice) system (the main customer service line). This pilot program was intended to handle two self-service tasks: stop a utility service or transfer the service to a new address.

The assumption was that everyone who reached our virtual assistant intended to do one of these two things (stop or transfer their utility service), based on the IVR menu selections that delivered the caller to our solution. Information about the user and their menu selection was passed to the virtual assistant, which immediately launched the corresponding use case flow. The following listing shows the user’s experience, first with the IVR and then with the handoff to the virtual assistant.

Listing 11.1 Handoff from IVR to virtual assistant

IVR: Thank you for calling ABC Energy. If this is an emergency, choose one of the following: say gas emergency or press 1. Electric emergency or press 2. Power outage, or press 3. If this is not an emergency, say do something else or press 4.

USER: <selects option 4>

IVR: Main menu. For billing and account information, press 1. Payments, press 2. Start, stop, or reconnect, press 3.

USER: < selects option 3>

IVR: Which are you calling about? Say building a new home or press 1. Start, stop, or move service or press 2. Reconnect or press 3. Streetlights or press 4. None of the above or press 5.

USER: < selects option 2>

IVR: Select one of the following. Check an existing request or press 1. Start new service or press 2. Stop existing service or press 3. Move my service or press 4.

USER: < selects option 3>

<Customer transfers over to virtual agent solution, which has a different voice.>

VIRTUAL ASSISTANT: We can help you with stopping your service. First, we need to get some information about your current address. Which type of account are you calling about: residential or commercial?

USER: Speak to an agent.

In the early production logs, we noticed a significant volume of users opting out, often right away. Though our utility company's virtual agent pilot was technically very capable, many callers wouldn't give it a chance. We listened to call recordings and discovered that the user experience as a whole felt disjointed. Customers dialed a number, reached an IVR (with a particular IVR "voice"), and navigated to a menu of service-related topics. If the caller chose to stop or transfer their service, they were routed to our virtual agent, but the transition was abrupt—a different voice bypassed greeting the caller (the original justification was that the caller had already been greeted by the IVR) and dove right into the task with seemingly optimal efficiency.

We could hear the confusion in the initial silence of some callers when they reached the virtual assistant. There were long hesitations, audible sighs, or stammering ("uh...", "um...", "hmm"). They hadn't been introduced to this new agent. They were confused by the first question, which was procedurally appropriate but conversationally came off as impersonal and clunky. A human agent wouldn't have opened the conversation like that. They would have introduced themselves in a welcoming tone. If the IVR collected information about the caller's goal, a human agent would have confirmed this with the user before proceeding ("I see you're calling about transferring your utility service, is that correct?").

We redesigned the experience from the opening line, starting with a context-aware greeting and introduction. The virtual agent was given a persona and a more conversational tone, as shown in the following listing.

Listing 11.2 Updated virtual assistant greeting and introduction

VIRTUAL ASSISTANT: Good afternoon.
I'm Alice. ABC Energy's virtual agent.

← Virtual agent introduction

← Context-aware greeting

Beginning the interaction with a greeting and introduction differentiated our virtual assistant from the menu-driven IVR system. This gave the caller time to adjust to the transition to a new voice and a different style of interaction.

A context-aware greeting, such as acknowledging the time of day or greeting a user by name, can transform a robotic, impersonal exchange into a more warm and welcoming experience. A name or persona may not be appropriate in all use cases, but for this one, it conveyed a sense of ownership and accountability. “Alice” is here to serve customers.

11.2.2 Convey capabilities and set expectations

In our immediate opt-out drivers, we identified users who opt out because they feel like their problem is too unique or complicated for a machine. Sometimes this is a fair judgment, but sometimes it is not.

Setting expectations up front is vital. It is not uncommon for companies to launch a pilot virtual assistant with limited scope or capabilities. When your solution only provides a subset of the topics or tasks a user might want, you need to communicate this up front. By doing so, users will either be assured that they are in the right place, or they will recognize that they are not. Especially in these scenarios, the greeting (or solution entry point) is a good time to announce the chatbot’s purpose and capabilities (or get a quick confirmation from the user before pushing forward).

Our utility company’s virtual agent greeting was expanded to confirm the user’s goal (using context indicating their IVR selection). We included a brief preview of the journey the user was about to embark upon. Because this was a voice channel—more prone to unexpected disconnects—we also set some expectations for what a successful completion would look (sound) like. The following listing shows the additional verbiage.

Listing 11.3 Updated greeting conveying capabilities and setting expectations

VIRTUAL ASSISTANT: Good afternoon. I'm Alice. ABC Energy's virtual agent. I'm here to help you with stopping your electric service.

Affirms the bot's purpose or capability

I'll just need to collect a few details about your account so we can schedule the stop order.

Previews the user's journey

Be sure to stay on the line until I give the confirmation number.

If we get disconnected before the confirmation, your account will not be changed.

Sets expectations for the journey's success

11.2.3 Incentivize self-service

A technically savvy or recurring user may realize the efficiency of using an automated solution, but one-time or occasional users don’t know what to expect. They may perceive the process to be difficult or time-consuming. These users may also be motivated to immediately opt out due to a preference for human interaction. Incentivizing self-service may reduce immediate opt-outs.

In our utility company use case, the average call time for completing a stop service request with a human agent was about 5 to 7 minutes (plus hold time). Our

self-service flow could be at least that fast. The following listing shows how we made an additional tweak to our greeting, letting the caller know that they could accomplish their goal in a short amount of time.

Listing 11.4 Incentivizing the caller by offering a fast resolution

VIRTUAL ASSISTANT: Good afternoon. I'm Alice.
 ABC Energy's virtual agent. I'm here to help you with
 stopping your electric service. I'll just need to
 collect a few details about your account so we can
 schedule the stop order.
This process should only take a few minutes.
 Be sure to stay on the line until I give
 the confirmation number. If we get disconnected before
 the confirmation, your account will not be changed.

Incentivizes
self-service

To further incentivize the caller, we wanted to assure them that they would be in good hands, even if they ran into problems with the automated system. The following listing shows the updated message.

Listing 11.5 Preventing immediate opt-outs by assuring escalation can happen

VIRTUAL ASSISTANT: Good afternoon. I'm Alice.
 ABC Energy's virtual agent. I'm here to help you with
 stopping your electric service. I'll just need to
 collect a few details about your account so we can
 schedule the stop order. This process should only take
 a few minutes. Be sure to stay on the line until I give
 the confirmation number. If we get disconnected before
 the confirmation, your account will not be changed.
**If we run into any problems, I'll get you over to a
 customer service representative.**

Assures the caller
that problems will
be escalated

11.2.4 Allow the user to opt in

Wherever possible, users should be given a sense of agency. This could look like obtaining their consent to proceed with the virtual agent experience. It may not be appropriate for all use cases, but this approach has the benefit of clearly identifying users who have agreed to opt in to the experience.

Our utility company's virtual agent greeting was updated one last time. We ask the user if they are ready to proceed.

Listing 11.6 Inviting the user to opt in

VIRTUAL ASSISTANT: Good afternoon. I'm Alice.
 ABC Energy's virtual agent. I'm here to help you with
 stopping your electric service. I'll just need to
 collect a few details about your account so we can
 schedule the stop order. This process should only take
 a few minutes. Be sure to stay on the line until I give

the confirmation number. If we get disconnected before the confirmation, your account will not be changed. If we run into any problems, I'll get you over to a customer service representative.

Are you ready to proceed?

Allows the user to opt in

Though we were hoping to hear a “yes” in response to our question, we also had to be prepared to handle other responses. We had a hypothesis that some users entered our solution by mistake. Sometimes it was a misunderstanding of what was going to happen when they selected the menu option in the IVR. Other times, it was simply a case of “fat finger”—an erroneous and often unnoticed dial pad selection. These “other” responses confirmed our hypothesis—our users would say “back” or “no” or express a different goal. They had arrived here by accident, and we wanted to get them back on the right path to accomplish their goal. Figure 11.4 shows how we redesigned the greeting flow so that each non-yes scenario could be handled appropriately.




Figure 11.4 The dialogue logic in this sample greeting flow can handle various responses to the question, “Are you ready to proceed?” If a response is explicit, such as “no,” “go back,” or “speak to an agent,” a predefined flow is invoked. Otherwise, the response is sent to the classifier for intent detection and is handled by the corresponding flow.

Exercises

Review the section of your dialogue where a user might immediately opt out, and ask yourself the following questions:

- 1 Does my virtual assistant greet the user warmly and introduce itself?
- 2 Does my virtual assistant explain its purpose—what it can and cannot do?
- 3 Does my virtual assistant offer a comparatively low-friction experience to meet the user’s needs (as good as or better than alternative channels or human intervention)?

11.3 Reducing other opt-outs

Opt-outs that occur later in a conversation can often be tied to a specific problem or gap in the conversational design. In this section, we'll discuss strategies and approaches to help minimize opt-outs that occur later in a conversation.

11.3.1 Try hard to understand

In chapter 4, we discussed the importance of understanding what your users want. Opt-outs can be an indication of problems in your bot's intent recognition. When a user engages with a chatbot that greets them with an open-ended question, such as "How can I help you?" they are going to have an expectation that they can ask any question related to your company's business, or even the general domain.

If the solution is not prepared for a range of reasonable requests, your user is going to be frustrated when asked to repeat or rephrase. Do not allow your solution to become stale. Invest in keeping your training relevant and representative of current user needs. Conversational search or RAG patterns may be a great fit for some use cases—especially those with broad domains and question-answer bots.

11.3.2 Try hard to be understood

Your word and phrasing choices are important, especially when you are soliciting information from a user. A lengthy output or poorly worded question may be hard for the user to parse. The user may not understand the type of information they are being asked to provide, or they may not have it immediately available. It is common for users to opt out if they find themselves in this situation. When you solicit information, be clear about how the user should answer. Multiple choice questions are sometimes misinterpreted as yes/no questions:

- "Are you calling about starting, stopping, or transferring your electric service?"
- "Are you looking for reconsideration, claim disputes, or review?"

Figure 11.5 shows an example of a choice question that might be interpreted as a yes/no question and one approach for preventing this type of confusion.




Figure 11.5 The way you structure a question can help the user understand what type of question you are asking—and what sort of answer you are looking for.

11.3.3 Be flexible and accommodating

Rigid, confusing, or overly restrictive conditions that define a “valid” user input may cause users to opt out. Your dialogue should be flexible enough to handle a range of “correct” responses and accommodate the user by disambiguating if necessary. An example of disambiguation is shown in figure 11.6.




Figure 11.6 A resilient dialogue flow can handle a range of valid responses. If necessary, a friendly disambiguation step can be invoked to get clarity about the user’s goal.

SPECIAL CONSIDERATIONS FOR VOICE SOLUTIONS

If speech recognition is part of the experience, validate your transcription accuracy so that egregious misunderstandings can be avoided. Because language models are text-based, mistranscribed speech inputs can compound problems in understanding. You could also miss important topic trends if the speech service does not faithfully capture the user’s input.

For voice solutions, it is imperative that you keep the cognitive load to a minimum; craft your questions to be direct and concise. If information needs to be included with the question, make sure the question comes at the end of the output. This will prompt the user that it is their turn to speak, as shown in figure 11.7. Be prepared to handle requests to repeat the question or information.

Bad	Good
Thanks for calling Acme Bank. How can I help you? I can look up your account balance, transfer money to another account, or schedule a payment.	Thanks for calling Acme Bank. I can look up your account balance, transfer money to another account, or schedule a payment. How can I help you?

Figure 11.7 The bad example (left) puts the call to action “How can I help you?” in the middle of the output. This might spur the user to begin speaking while the output is still trying to play. The good example (right) puts informational messaging up front. At the end of the message, the call to action is a clear invitation for the user to begin speaking.

USE TECHNOLOGY TO EASE PAIN POINTS

There are situations where you can employ technology to ease user pain points and facilitate understanding. For example, an insurance company had a voice assistant for members. Member ID numbers could be 9 or 13 digits, and sometimes they were preceded by a letter. We didn't want to burden the caller with a message about how long the number might be or whether or not to include the preceding letter, if there even was one. That's too much cognitive load, especially on a first pass. We simply asked, "What's your member ID?" Some users would speak this number, including the letter. Others would key in the number using the dial pad, and they would ignore any preceding letter. We wanted to accommodate natural user behavior, so we made our logic more robust. The caller could include or exclude the preceding alpha character as long as we detected the right number of digits to perform a lookup. Instead of pestering the user with retries, the solution would fill in trivial information and attempt to perform a lookup if it had enough information (e.g., "123456789" and "X123456789" were both acceptable). If needed, our retry messaging would change the message slightly to guide the user to provide a valid input.

Another tool for improving robustness and accuracy, if your technology allows it, is custom speech models. Speech transcriptions of numbers or letters are difficult to collect over a voice channel due to phoneme similarity (e.g., "8" can often sound like "H," so we selected speech models that were optimized to such inputs).

11.3.4 Convey progress

Self-service task flows can range in complexity from answering a question to simple triage to multistep, multiflow, dynamic user journeys. In these more complex flows, your dialogue design should signal the progression of the task or process. If the user feels that this is going to go on forever or that they are looping through an unending series of menus, they may believe that a conversation with a human would be more efficient.

Progress can be signaled in a variety of ways and should be optimized for the medium. Before initiating a long task flow, set the user's expectations about what you will need from them and how long it will take. A text-based platform can be more verbose or provide visual indicators, such as numbering questions or showing progress bars. A phone channel will need to be more succinct but can give indicators such as "We're almost done" or "Just a few more questions."

11.3.5 Anticipate additional user needs

When a user receives an answer from your bot or reaches the conclusion of a task flow, they may find that their goal is still unmet. Does your dialogue flow have potential outcomes that are technically correct or appropriate but that will tend to leave the user's problem unresolved? Think about the position the user is in. They have a need, they get an answer, but the need still exists. If the dialogue flow has concluded without addressing that need, the user is going to opt out. What other choice do they have?

Can you extend your solution’s self-serve capability to initiate an alternate resolution path? Figure 11.8 shows some example scenarios and possible ways to avert the need for escalation.

Scenario	Anticipated next step
A user searches the status of an insurance claim, only to find the claim has been denied.	Offer to help the user submit an appeal.
A patient calls to refill a prescription, but the refill has expired.	Offer to submit a renewal request on behalf of the patient.
A banking customer requests their account balance, only to find it is overdrawn.	Offer to show recent transactions.
A traveler misses their flight, and the reschedule options are inconvenient.	Offer to put the traveler on standby for other flights.

Figure 11.8 You can avoid escalations by presenting the next best course of action whenever you anticipate a user could still have an unmet need at the conclusion of a dialogue flow.

11.3.6 Don’t be rude

When you’re building a virtual assistant, remember that you are attempting to simulate a more human-like experience. Empathy comes naturally to humans, but it must be intentionally integrated into a machine-driven interaction. Assume that your users are making an effort on their end of the conversation. When you need to retry or repair the conversation, don’t make the user feel they are to blame. A conversational designer should craft error messages in a way that guides the user back on track without implying fault.

For example, our insurance company solution could look up claims with a claim number, which was thirteen digits. Originally, it seemed logical to provide “informative” information in the retry output. This turned out to sound rude over the phone, as shown in the following listing.

Listing 11.7 Retry scenario with rude customer experience

VIRTUAL ASSISTANT: I'll need the thirteen digit claim number to look up your claim.

USER: <says a thirteen digit number but speech only picked up twelve>

VIRTUAL ASSISTANT: You didn't provide a thirteen digit number. Please say or enter the claim number.

Sometimes the user failed to provide enough digits; other times, the speech service failed to detect all of the spoken digits or mistranscribed the user response. Regardless of fault, the solution needs to re-ask the user to provide the information. On a

first retry, it is often enough to simply ask for the information again. You can progressively guide the user if they are struggling to give the right kind of response or to provide information in the correct format.

Listing 11.8 Updated experience for retry scenario

VIRTUAL ASSISTANT: I can help you look that up. What's the claim number?

USER: <says a thirteen digit number but speech only picked up twelve>

VIRTUAL ASSISTANT: Sorry, I didn't get that. What is your claim number?

USER: <says a thirteen digit number but speech only picked up twelve>

VIRTUAL ASSISTANT: Sorry, I still didn't get that. You can say or enter the thirteen-digit claim number.

Exercises

Review the downstream flows within your dialogue where users might opt out, and ask yourself the following questions:

- 1 How often do users opt out because the virtual assistant does not seem to understand their request?
- 2 Are users opting out because they do not understand what the virtual assistant is asking of them?
- 3 Is my virtual assistant pleasant and helpful? Does it convey competence, efficiency, and empathy?
- 4 Is there automation available to streamline or expedite the user to their end goal?
- 5 Do users request a human because they have unmet needs after interacting with the virtual assistant, even after a seemingly “successful” flow?

11.4 *Opt-out retention*

The strategies discussed up to this point are intended to reduce immediate opt-outs but may not eliminate them entirely. To improve containment, you can also try to retain the user in the virtual agent experience after an immediate opt-out request or at key points in a dialogue flow.

A good faith attempt to keep the user in channel can work in many scenarios. This must be undertaken with care—a customer should never feel that they are being held hostage by the system. When that happens, by the time the user eventually does get to a human agent, they may be frustrated or hostile.

The purpose of an opt-out retention flow is to

- 1 Discover the user's true goal or need
- 2 Assess whether the bot is capable of meeting that need

- 3 If so, convince the user that the bot is capable of meeting that need
- 4 If not, route them to the next best action (e.g., another virtual agent or a human)

You might implement such a flow at the beginning of a conversational interaction to help improve containment loss due to immediate opt-outs. It can also be used strategically after a user has appeared to successfully complete a flow. If users are asking for an agent right after appearing to complete a “happy path,” you might be missing something. Either the user didn’t like the answer, or it did not help them toward their goal.

11.4.1 Start right away by collecting opt-out data

If you want to get started right away, or if you are still in the predeployment build phase, you can implement a simplified opt-out flow that asks the user to provide the reason for opting out. In the most simplified version, the assistant escalates no matter what is said. Figure 11.9 shows an example data-collection opt-out flow. This is a fairly non-intrusive strategy to find out why users are opting out.

Once you collect metrics, you may realize that you need better training, or your solution strategy has a mismatch between what you built it for and how users want to interact with it.

11.4.2 Implementing an opt-out retention flow

Once you understand why users are opting out, you can begin to address these requests in more meaningful ways. The first thing you’ll want to do is update your classifier so that you can take the action most appropriate for the request. Such actions could include

- Expanding your bot’s current capability to handle these new requests
- Handing the user off to a different virtual agent
- Escalating the user to a human agent when necessary (if available)
- Providing information about next-best alternatives if agent escalation is not available

Figure 11.10 shows a dialogue flow that can identify which requests are in scope versus out of scope, with a fallback path for requests that are not understood by the classifier.




Figure 11.9 A simple flow to collect the user’s reason for opting out

Opt-out retention dialogue flow




Figure 11.10 A typical opt-out retention pattern will try to find out what the user needs. It will either recognize an in-scope request (a request that it is equipped to handle) or an out-of-scope request (a request that it understands but is not equipped to handle), or it will not recognize the request. When an in-scope request is recognized, the bot may make an incentivized offer to keep the user in-channel. When an out-of-scope request is made, the bot can route the user directly to the appropriate skill or agent queue. When a request is not recognized, the user is routed to a default or general hold queue.

For our utility company use case, we added logic to ask the user to tell us more about what they needed (and trained our classifier over time to identify the range of requests). Many users would still insist on escalating by demanding an agent a second time. These users would be transferred to the general live agent hold queue.

Some users expressed a goal that was understood but was truly outside the scope of our system. These users could be routed to a specialized agent (if available), minimizing the need for additional transfers. For example, some businesses have a single customer service department to handle all problems. Others route users to different customer service agents or departments to handle specific scenarios, such as a dedicated billing department or agents who specialize in handling appointment scheduling.

When a user expressed a goal that aligned with our solution's purpose, we made an incentivized offer and gave them a choice to proceed with self-service. This gives the user agency and makes them feel heard. By providing information about hold times for a human agent, they can decide how best to use their time. By correctly identifying their need, we gain their confidence. Many users are willing to proceed as long as the bot is able to demonstrate competence, as shown in the following listing.

Listing 11.9 Opt-out retention flow identifying goal to keep user in-channel

```
VIRTUAL ASSISTANT: Good afternoon. I'm Alice. ABC Energy's virtual agent.  
I'm here to help you with stopping your electric service. I'll just need  
to collect a few details about your account so we can schedule the stop  
order. This process should only take a few minutes. Be sure to stay on  
the line until I give the confirmation number. If we get disconnected  
before the confirmation, your account will not be changed. If we run  
into any problems, I'll get you over to a customer service  
representative. Are you ready to proceed?
```

```
USER: speak to an agent
```

```
VIRTUAL ASSISTANT: In order to get you to the right place, please tell me  
more about what you need.
```

```
USER: I want to disconnect my service and close the account right away
```

```
VIRTUAL ASSISTANT: We are currently experiencing long hold times, but I  
can help speed things up by asking you a few more questions. It sounds  
like you want to stop your electric service and get a final bill. Is  
that correct?
```

```
USER: yes
```

```
VIRTUAL ASSISTANT: Great! Let's start with looking up your account.
```

The virtual agent would then proceed to collect as much information as possible, until an error occurred or the caller requested an agent a second time. This enabled us to self-serve the customer either to completion, or to collect as much information as possible, reducing the amount of time the human agent would need to spend resolving the call.

This new design had great success. Out of callers who immediately opted out but expressed an in-scope intent, 38% were convinced to stay with the virtual agent and be successfully authenticated. When paired with other updates described in this chapter, the overall use case completion rate increased from 27% to 30%.

Exercises

Reflect on what you have learned about strategies to keep the user in-channel:

- 1 Does your current solution have opt-out trend patterns where you don't understand why the user asked for an agent?
- 2 Would asking the user to provide more information about what they need help you target areas for improvement?

11.5 *Improving dialogue with generative AI*

In chapter 5, we showed how to use retrieval-augmented generation (RAG) to generate chatbot responses at runtime. But if you prefer static and controlled dialogue, you can still use generative AI to assist your conversational designer during the build phase of a project. Conversational designers are excellent at crafting dialogue that meets the needs of both the system and users. If you don't have a designer handy, you can use generative AI to help you craft dialogue messages that achieve your goals.

In this section, we'll demonstrate several techniques for improving static output responses using generative AI.

11.5.1 *Improving error messages with generative AI*

It's difficult to write good error dialogue. Our first instinct can be to provide information that's technically true but brusque to the user. This is especially difficult in voice channels, where the user may make a mistake and not be aware of it. For instance, a user being asked to provide an identifier of a certain length may forget one of the digits (or mistakenly add an additional digit).

For a Social Security number (nine digits in length), a true-but-brusque error message would be "I did not get nine digits. Enter a valid Social Security number now." Let's use generative AI to improve that message.

Since this is a creative task, we'll use an instructible model (mixtral-8x7b-instruct-v01-q) and sampling decoding. With sampling decoding, the responses will be non-deterministic. We'll run this prompt multiple times to generate multiple responses to choose from.

Listing 11.10 Improving error messages with generative AI—iteration 1

Instruction: You are a copy editor designing dialogue for a voice-based conversational system. You will be given a scenario and an input message. Rewrite the input message to an output message.

Instruction and
grounding for
the model

The output must be brief. The output should be 12 words or less.
 The output must be helpful and instruct the user.
 The output must be kind.
 The output must not refer to a mistake by the user.

Scenario: User has entered a social security number with an unexpected number of digits.
 Input: I did not get 9 digits. Enter a valid social security number now.
 Output: Please enter your 9-digit SSN. ← **LLM output**

Instruction and grounding for the model

Scenario and input message to improve

Running that prompt three times generated these outputs:

- “Please enter your 9-digit SSN.”
- “Please enter your full 9 digit Social Security Number.”
- “Please enter 9 digits for the Social Security Number. Thank you!”

Those messages are a definite improvement. They could be even better with an apology included. Let’s augment the instructions to include an apology.

Listing 11.11 Improving error messages with generative AI—iteration 2

Instruction: You are a copy editor designing dialogue for a voice-based conversational system. You will be given a scenario and an input message. Rewrite the input message to an output message.
 The output must be brief. The output should be 12 words or less.
 The output must be helpful and instruct the user.
 The output must be kind.
 The output must not refer to a mistake by the user.
 The output must include an apology.

Scenario: User has entered a social security number with an unexpected number of digits.
 Input: I did not get 9 digits. Enter a valid social security number now.
 Output: Please enter your 9-digit SSN. ← **LLM output**

Original instruction

New instruction—add an apology

Original scenario and input message

Running that prompt six times generated these outputs:

- “Sorry, please enter a 9-digit social security number.”
- “Apologies. Please enter correct Social Security number.”
- “Apologies for trouble, try again with 9 digits.”
- “Apologies for that, let’s try again. Enter a 9-digit SSN please.”
- “I apologize, please enter a 9-digit Social Security Number.”
- “My apologies, that’s not quite right. Please try again with 9 digits.”

We can take one of these messages as is or mix and match to generate a new one. Perhaps use “Sorry, let’s try again. Please enter your 9-digit social security number.”

11.5.2 Improving greeting messages with generative AI

We experienced a greeting message that confused some of the users calling in to the utility company. The users were passed from one IVR to another in a jarring manner. Could generative AI help with an updated greeting? Let's see in the following listing.

Listing 11.12 Improving greeting messages with generative AI—iteration 1

```
<|instruction|>
You are a copy editor designing dialogue for a voice-based
conversational system. You will be given a scenario and an input
message. Rewrite the input message to an output message.
The output must include a greeting.
The output must be helpful and instruct the user.

<|scenario|>
User dialed their energy company ABC Energy. The user is met by a
first conversational AI. The conversational AI provides a menu. The
user pushes a button. The user is transferred to a second
conversational AI.

<|input|>
We can help you with stopping your service. First, we need to get
some information about your current address. Which type of account
are you calling about: residential or commercial?

<|output|>
Hello! Thank you for contacting ABC Energy. I'd be
happy to assist you in stopping your service. To start,
could you please tell me whether this is a residential
or commercial account? Once I have that information, I
can better assist you with your request.
```

LLM output

The first output is pretty good, but we can generate a few more:

- “Hello! Welcome to ABC Energy. I understand that you’d like to discontinue your service. Before we proceed, may I ask which type of account you’re calling about: residential or commercial? Additionally, could you please provide your current address so we can properly process your request? Thank you.”
- “Hello! Welcome to ABC Energy. I’m here to assist you today. To begin, could you please tell me whether this is a residential or commercial account? Once I have that information, I can help you with any questions or concerns you may have about stopping your service.”

The output messages give us several ideas on how to improve the greeting. They all include a salutation, an identification of the company, and a preamble. Then they include the key question, as well as some commentary after the question. Only one problem: we should end the message with a question—this is the invitation for the user to speak. (This is important in voice solutions, as we don’t want to ask the user a question but continue speaking—this tends to result in the user answering while the bot continues speaking.) We could manually remove that last part or try an updated instruction to the LLM.

The following listing attempts a new instruction.

Listing 11.13 Improving greeting messages with generative AI—iteration 2

```
<|instruction|>
You are a copy editor designing dialogue for a voice-based
conversational system. You will be given a scenario and an input
message. Rewrite the input message to an output message.
The output must include a greeting.
The output must be helpful and instruct the user.
The output must end with a question. ← New instruction

<|scenario|>
User dialed their energy company ABC Energy. The user is met by a
first conversational AI. The conversational AI provides a menu. The
user pushes a button. The user is transferred to a second
conversational AI.

<|input|>
We can help you with stopping your service. First, we need to get
some information about your current address. Which type of account are
you calling about: residential or commercial?

<|output|>
Hello! I'm happy to assist you in discontinuing your
service with ABC Energy. To start, could you please
specify whether this is regarding a residential or
commercial account? Thank you! | LLM output
```

This message is better, but it still ends with “Thank you!” We instructed the LLM not to do this, but it still did. In fact, we ran this prompt multiple times, and it ended the output this way each time. We even changed the instruction to

The output must end with a question. The output must not end with “Thank you!”

This is ultimately a minor annoyance. We can remove the “Thank you!” and proceed with the updated message.

The greeting message could be improved by outlining the capabilities of the bot and letting the user know what expectations we have of them. The following listing improves the greeting once more.

Listing 11.14 Improving greeting messages with generative AI—iteration 3

```
<|instruction|>
You are a copy editor designing dialogue for a
voice-based conversational system. You will be given
a scenario and an input message. Rewrite the input
message to an output message.
The output must be helpful and instruct the user.
The output must describe the system's capabilities.
The system helps users stop their electric service.
The output must let the user know they will need to
answer a few questions that provide details about their
service. ↓ Detailed instruction
```

The output must let the user know their account will not be changed until the process is completed. The output must end with a confirmation that the user is ready to proceed.

↑
Detailed instruction

<|scenario|>

User dialed their energy company ABC Energy to stop their electric service. The company's conversational AI has greeted the user.

<|input|>

I'm here to help you with stopping your electric service. I need some information about your service. Are you ready?

Simplistic greeting that we ask the LLM to improve

<|output|>

Welcome to ABC Energy! I'm here to assist you in discontinuing your electric service. To ensure a smooth process, I'll need to ask a few questions about your account details. Rest assured, they won't be changed until the termination is finalized. Are you prepared to proceed?

LLM output

An excellent greeting! Recall that we are using sampling decoding and getting non-deterministic responses. You'll most likely get different output. Here are a few more outputs from the LLM:

- "Welcome to ABC Energy! I can assist you in discontinuing your electric service today. To ensure accuracy, I'll need to ask a few questions regarding your account details. Please note that no charges will be applied until the process is finalized. Are you prepared to proceed?"
- "Welcome to ABC Energy! I can assist you in discontinuing your electric service. To ensure accuracy, I'll ask several questions related to your account. Please note that no charges will be applied during this process. Are you prepared to proceed with providing the necessary details?"

You can experiment with different instructions to change the greeting or proceed with your favorite message.

One improvement is to incentivize the user by letting them know the process is quick. We can augment the instructions to note that. The next listing adds the instruction "The user should be incentivized that the questions are easy and don't take long to answer."

Listing 11.15 Improving greeting messages with generative AI—iteration 4

<|instruction|>

You are a copy editor designing dialogue for a voice-based conversational system. You will be given a scenario and an input message. Rewrite the input message to an output message. The output must be helpful and instruct the user. The output must describe the system's capabilities. The system helps users stop their electric service.

The user should be incentivized that the questions are easy and don't take long to answer.
 The output must let the user know they will need to answer a few questions that provide details about their service.
 The output must let the user know their account will not be changed until the process is completed.
 The output must end with a confirmation that the user is ready to proceed.

New instruction

```
<|scenario|>
User dialed their energy company ABC Energy to stop their electric
service. The company's conversational AI has greeted the user.
```

```
<|input|>
I'm here to help you with stopping your electric
service. I need some information about your service.
Are you ready?
```

**Original greeting that we
ask the LLM to improve**

```
<|output|>
Welcome to ABC Energy! I can assist you with
Discontinuing your electric service today. To ensure a
smooth process, I'll need to ask a few quick and easy
questions about your account. Rest assured, it won't be
changed until the cancellation is finalized. Are you
prepared to begin?
```

**LLM output
(new output in bold)**

The LLM seamlessly transforms the message with brevity: the new message explains that the next questions in the process are designed to be “quick and easy” (for the user).

Let's augment the greeting one more time. The user should be assured that customer service will step in if needed. In the next listing, we'll add the instruction “The output must let the user know a customer service representative will get involved if the user cannot complete the automated process.”

Listing 11.16 Improving greeting messages with generative AI—iteration 5

```
<|instruction|>
You are a copy editor designing dialogue for a voice-based
conversational system. You will be given a scenario and an input message.
Rewrite the input message to an output message.
The output must be helpful and instruct the user.
The output must describe the system's capabilities. The system helps
users stop their electric service.
The output must let the user know they will need to answer a few
questions that provide details about their service. The user should
be incentivized that the questions are easy and don't take long to answer.
The output must let the user know their account will not be changed
until the process is completed.
The output must let the user know a customer service
representative will get involved if the user cannot
complete the automated process.
The output must end with a confirmation that the user is ready to proceed.
```

New instruction

```
<|scenario|>
User dialed their energy company ABC Energy to stop their electric
service. The company's conversational AI has greeted the user.
```

```
<|input|>
```

I'm here to help you with stopping your electric service. I need some information about your service. Are you ready?

Original greeting that we ask the LLM to improve

```
<|output|>
```

Welcome to ABC Energy! I'm here to assist you with stopping your electric service. To ensure a smooth process, I'll need to ask you a few easy questions about your service. Rest assured, it won't be changed until the cancellation is finalized. **If you encounter any difficulties, a customer service representative will be happy to assist you.** Are you ready to begin?

LLM output
(new output in bold)

The greeting conveys everything we need, but it's starting to get long. Let's ask the bot to shorten it by adding one more instruction.

The following listing adds "The output should be 40 words or less."

Listing 11.17 Improving greeting messages with generative AI—iteration 6

```
<|instruction|>
```

You are a copy editor designing dialogue for a voice-based conversational system. You will be given a scenario and an input message. Rewrite the input message to an output message. The output must be helpful and instruct the user. The output must describe the system's capabilities. The system helps users stop their electric service. The output must let the user know they will need to answer a few questions that provide details about their service. The user should be incentivized that the questions are easy and don't take long to answer. The output must let the user know their account will not be changed until the process is completed. The output must let the user know a customer service representative will get involved if the user cannot complete the automated process. The output must end with a confirmation that the user is ready to proceed. The output should be 40 words or less.

New instruction

```
<|scenario|>
```

User dialed their energy company ABC Energy to stop their electric service. The company's conversational AI has greeted the user.

```
<|input|>
```

I'm here to help you with stopping your electric service. I need some information about your service. Are you ready?

Original greeting that we ask the LLM to improve

```
<|output|>
```

Sure! To stop your electric service, I'll ask a few easy, quick questions. Nothing will be changed until the process is complete. If needed, a rep will assist. Ready to proceed?

LLM output

Generative AI is an excellent partner for improving dialogue messages!

11.6 Sometimes it's okay to escalate

There may be scenarios where your virtual assistant is going to deliver information that the user is not happy about. For example, if a user submitted a claim and hears it was denied, they probably aren't going to be satisfied and will want to rectify the situation. If the automated solution doesn't provide a path to resolution or reconsideration, they may be inclined to opt out at this juncture.

With good planning, you can design proactive flows that anticipate follow-on user needs. That way, even if escalation is the appropriate next-best action, your metrics can distinguish between an opt-out and an intentional transfer for business reasons. In the next chapter, we'll discuss how to optimize the handoff to a human agent.

Summary

- Opt-outs are a major source of containment loss, which causes a virtual agent to fail on delivering business value.
- Users who opt out early in a conversation tend to do so for reasons related to their perception of a virtual agent's capability and whether they are confident that the virtual agent can usher the user to their end goal.
- Opt-outs that occur later in a conversation are indicators that a virtual agent might have weak understanding or problems with the dialogue design.
- Opt-out retention is a great strategy for improving containment; it can also provide valuable data about what users expect your bot to be able to do.
- Generative AI can supplement the process of crafting tactful, efficient responses throughout your dialogue flows.

12

Conversational summarization for smooth handoff

This chapter covers

- Defining elements of an effective conversation summary
- Instrumenting your conversational AI to enhance summarization
- Summarizing a chat transcript into prose with LLMs
- Extracting structured details from a chat transcript with LLMs

Conversational AI builders would love it if their systems contained all user conversations. But for most use cases, some percentage of users will end their interaction with a human and not your bot. Conversational AI is designed to handle the easily automated conversations and direct the higher-value or more challenging ones to human agents. Users who want to self-service may be frustrated by “failing” with the conversational AI, so it’s important to give that human agent the best start possible at handling the call after a transfer.

The two simplest handoff methods are also the least satisfactory. We can transfer the conversation “blind” to the human agent and have them ask again for all the

information they need. Or we can pass the agent the full conversational transcript and ask them to search it for the information they need (while the user is waiting!). It's better to give the human agent a targeted summary of the conversation to date so they can quickly pick up where the conversational AI left off.


12.1 Intro to summarization

A summary makes handoffs from AI to human go smoothly. First, we'll review why summaries are needed. Then we'll explore the elements of effective summaries.

12.1.1 Why summarization is needed

Most conversational AI solutions store full transcripts of conversations for auditing and data analysis purposes. This is a treasure trove of training data when you have time to analyze it. But when you are a call center agent being transferred a conversation from the AI, you don't have time to read a lengthy transcript. You need to quickly grasp the essence of the user's problem so you can start helping them. We've seen AI chat transcripts that go for multiple pages (hundreds of words). An agent needs a few targeted bullet points—any more will take too long to read, and any less will not convey enough useful information.

Figure 12.1 illustrates how a conversational summary is generated from a full conversational transcript when the chatbot hands off the conversation to a human agent.



1: Conversational AI and user may have a lengthy conversation transcript

2: Human agent should receive a structured summary at handoff

Figure 12.1 An effective summary helps an agent get up to speed quickly, even if the user previously had a lengthy conversation with a bot.

From the previous chapter, we know that some users immediately opt out to an agent, but many users go through a series of steps before getting frustrated. The following listing shows an example conversation where the user apparently accomplishes their goal but still requests an agent.

Listing 12.1 Example conversation between user and AI

```

AI: Welcome to the automated assistant, how can I help?
User: Claim status
AI: Ok, claims. What's your Tax ID?
User: 123456789
AI: Thank you. What's the member's eight-digit Member ID?
User: 87654321
AI: Ok, and what's the member's date of birth?
User: 12/31/2000
AI: Is the member's name John Doe?
User: Yes
AI: Great! What's the date of service for the claim
User: "February first 2024
AI: I found your claim #111222333 from February 1,
2024 for the member #87654321 and provider #123456789.
It was paid on May 23, 2024 for $201.83.
User: representative

```

User appears to
accomplish their goal

User still opts out

We're not sure why the user opted out (maybe the check didn't arrive?). But that is beside the point. In most conversational AI solutions, users can opt out anywhere in the conversation. The conversation in listing 12.1 had seven turns and probably took about two minutes. It's just long enough that it takes some effort to read and comprehend. Would you want to read that full conversation if you were the human agent? What if the conversation was longer?

A smooth handoff should happen quickly. The agent should quickly comprehend what has happened so they can be effective. The user should not have to wait for agent to get up to speed. An effective summary facilitates all those needs.

12.1.2 *Elements of effective summaries*

A conversation summary includes just enough information to understand the complete conversation. It should include structured metadata and a short text summary; the summary's contents will vary based on your specific use case. Figure 12.2 shows an example.

METADATA SUMMARY ELEMENTS

Structured summary elements often come from closed-form questions, such as "What's your member ID?" ("Open-form" questions are like "How may I help?") These summary elements can include data that was collected during the conversation or context that was supplied from outside of the conversation.

These are some example elements:

- User ID of the logged-in user (chat) or the caller's phone number (voice/SMS)
- Identifiers collected during the chat
- Identifiers found during the chat
- Number of chat sessions the user has ever had
- Sentiment analysis of user utterances

- 1 AI: "Welcome to the automated assistant, how can I help?"
 User: "Claim status"
 AI: "Ok, claims. What's your Tax ID?"
 User: "123456789"
 AI: "Thank you. What's the member's eight-digit Member ID?"
 User: "87654321"
 AI: "Ok, and what's the member's date of birth?"
 User: "12/31/2000"
 AI: "Is the member's name John Doe?"
 User: "Yes"
 AI: "Great! What's the date of service for the claim?"
 User: "February first 2024"
 AI: "I found your claim #111222333 from 2/1/2024 for the member #87654321 and provider #123456789. It was paid on 5/23/2024 for \$201.83."
 User: "representative"
- 2 Tax ID: 123456789
 Member ID: 87654321
 Claim ID: 111222333
- 3 Summary: User searched for their claim and found it was paid three months after filing.

- 1: Full transcript takes a long time to read.
 2: Structured metadata highlights key points.
 3: Free-text summary of the conversation.

Figure 12.2 An effective summary pulls out key details from the conversation. Here it includes a summary of the conversation and the last claim searched. The AI portion of the call may have taken two minutes, but the human agent can read the summary in seconds.

In figure 12.2, the chat collected five pieces of information, but the human agent only needs to know three. Figure 12.3 breaks down the summary.

- 1 AI: "Welcome to the automated assistant, how can I help?"
 User: "Claim status"
 AI: "Ok, claims. What's your Tax ID?"
 User: "123456789"
- 2 AI: "Thank you. What's the member's eight-digit Member ID?"
 User: "87654321"
 AI: "Ok, and what's the member's date of birth?"
 User: "12/31/2000"
 AI: "Is the member's name John Doe?"
 User: "Yes"
- 3 AI: "Great! What's the date of service for the claim?"
 User: "February first 2024"
 AI: "I found your claim #111222333 from 2/1/2024 for the member #87654321 and provider #123456789. It was paid on 5/23/2024 for \$201.83."
 User: "representative"
- 1 Tax ID: 123456789
 2 Member ID: 87654321
 3 Claim ID: 111222333
- 1: One question identifies the provider.
 2: Three questions identify the member.
 3: The provider, member, and date identify the claim.


Figure 12.3 Not every closed-form question needs to be stored in the summary. In this medical insurance claim review, the most important information is the provider ID, member ID, and claim ID.

In this medical insurance claim search, it takes a lot of information to validate the caller. We need to verify who is calling, who they are calling about, and what they are calling about. It takes three data elements alone to confirm the member information: an ID,

a date, and a confirmation of the name. The human agent receiving the transfer only needs to know that the member was verified, and the member ID suffices for that.

Similarly, there are many pieces of information about the claim—some given by the user (date of service) and some by the AI (status, paid date, paid amount). The summary only includes the claim ID, which will be sufficient for the agent to retrieve the full claim, including all the details that were not part of the conversation.

It’s useful to include summary elements in software used by human agents fielding conversations. In contact center software, this is called a *screen pop*—a feature that displays contextual information to an agent while they handle a conversation. Figure 12.4 shows an example screen pop for our medical insurance agent. They receive the structured information as a highlight, and through backend integration, they can get even more information. Clicking on the member ID should show information about the member or an image of their ID card. Clicking on the claim should show the claim itself (for instance, as a PDF).



1: Summary elements captured in the conversation

2: Call center agent’s “screen pop”

Figure 12.4 When a summary is integrated with contact center software, the human agent can have a wealth of information at their fingertips. When an insurance agent clicks on the member ID in their software, they could get additional details on that member.

Structured metadata gives key data points from the conversation so far. It prevents the human agent from having to re-ask questions that the user has already answered for the bot. Users get *very* frustrated when they must answer the same questions again! The human agent benefits from having this information at their fingertips, but they still need to be aware of the overall context of the conversation. That’s where the free-text summary comes in.

FREE-TEXT SUMMARY ELEMENTS

A conversation may have included hundreds of words before it was transferred to a human agent. (The example transcript we are using is about 100 words.) The average adult reads at 200 words per minute when reading for fun and slower for complex material. Our user doesn’t want to wait any longer than necessary, so our human agent needs to get up to speed quickly. A good summary can reduce that time by minutes.

Follow the “keep it simple” philosophy. A summary of one to two sentences conveys a lot of information quickly. The free-text summary in our example—“User searched

for their claim and found it was paid three months after filing”—encapsulates the entire search process in the first five words as well as a likely reason for transfer in the last eight words.

The free-text summary is also not repetitive. It eliminates several redundant pieces of information:

- *The user's initial intent*—This is not explicitly included, since it is inferred from a claim search being done.
- *The intermediate questions*—It does not say “The bot asked for the tax ID, member ID, date of birth, etc.” These are all implied from the user finding the claim. The system does not find claims until sufficient information is provided.
- *The structured content*—The summary doesn't need to waste words repeating the structured content, which has already been provided in compact form.

The summary also does not identify who initiated the transfer (the user or the system). This could be added in a new structured field.

There are trade-offs in the summarization process. Too brief a summary will omit information that helps the agent. Too lengthy a summary will not help the agent learn quickly, compared to reading the original transcript. Choose a summarization methodology that works best for your use case.

Later in this chapter, we will show you how to use generative AI to generate summaries. First, though, we need the conversational AI to structure data in the right format to generate summaries.

Exercises

Design your ideal summaries for the following sample conversation. You will refine these summaries later in the chapter with additional techniques:

```
AI: Welcome to the automated assistant, how can I help?
User: Claim status
AI: Ok, claims. What's your Tax ID?
User: 123456789
AI: Thank you. What's the member's eight-digit Member ID?
User: 87654321
AI: Ok, and what's the member's date of birth?
User: 12/31/2000
AI: Is the member's name John Doe?
User: Yes
AI: Great! What's the date of service for the claim?
User: February first 2024
AI: I found your claim #111222333 from February 1, 2024 for the member
    #87654321 and provider #123456789. It was paid on May 23, 2024 for
    $201.83.
User: representative
```

- 1 Design a purely textual summary of the sample conversation.
- 2 Design a structured summary of the data elements. Would you extract different structured elements if the caller asked for the details of the member's health plan or tried to proactively estimate the cost of a procedure?

(continued)

- 3 Take a conversation transcript from a chatbot you are working on. Summarize the conversation in both text and structured elements.


12.2 Preparing your chatbot for summarization

We've taken it as a given that your chatbot keeps track of the conversational transcript. Most platforms do, but not all. The transcript is the bare minimum element you need to build a summary. In this section, we'll show you multiple ways to collect the data necessary for both textual and structured summaries of conversations.

12.2.1 Using out-of-the-box elements

Conversational AI platforms often include built-in elements to help with conversational summarization. The most common element is a conversational transcript—a running log of messages between the user and the assistant. This is accessible in different ways on different platforms. One common mechanism is called a *session variable*.

Figure 12.5 shows how the transcript can be accessed in our platform (watsonx) via a built-in session variable called “session history.”



- 1: “Fallback” action triggers a transfer to an agent. This is a great time to summarize the chat.
- 2: The conversation transcript is available in watsonx Assistant as the variable “session history.”

Figure 12.5 Accessing the conversation transcript through the “session history” variable

Depending on your chat platform, the conversational transcript will be stored in different formats. Our platform provides the summary in a JSON format, as shown in the following listing.

NOTE In many conversational AI platforms, the transcript is not available to the dialogue session as a variable unless you craft it yourself via webhooks. We'll demonstrate how to build a variable with the transcript later in the chapter.

Listing 12.2 Conversational transcript via the built-in “session history” variable

```
[{"a": "Welcome to the automated assistant, how can I help?"}, {"u":
➤ "Claim status", "n": true}, {"a": "Ok, claims.\nWhat's your Tax ID?"}, {"u":
➤ "123456789"}, {"a": "Thank you.\nWhat's the member's eight-digit Member
➤ ID?"}, {"u": "87654321"}, {"a": "Ok, and what's the member's date of
➤ birth?"}, {"u": "12/31/2000"}, {"a": "Is the member's name John Doe?\n
➤ option: ["Yes", "No"]"}, {"u": "Yes"}, {"a": "Great!\nWhat's the date of
➤ service for the claim?"}, {"u": "2024-02-01"}, {"a": "I found your claim
➤ #111222333 from Feb 1, 2024 for the member #87654321 and provider
➤ #123456789. It was paid on 5/23/2023 for $201.83."},
➤ {"u": "representative"}]
```


The JSON format is intended for machines to read, but you can run the transcript through a transform process. We made the figures in this chapter more human-readable by replacing the "a" keys with "Bot", the "u" keys with "User", and the \n (newlines) with spaces.

The transcript also includes some metadata that you may choose to ignore, such as the “yes” and “no” choices being offered via buttons. Other conversational AI platforms may include further metadata like timestamps.

Later in this chapter (in section 12.3), we will demonstrate running this transcript format through an LLM for summarization. We will see that LLMs are quite resilient to the transcript format. You can summarize transcripts in their native format or reformat them so they are easier for humans to read.

12.2.2 Instrumenting your chatbot for transcripts

Some conversational AI platforms require you to create and store the conversation transcript yourself. Or you may choose to create your own version of the transcript in the exact format you prefer. Either way, this is implemented in your *orchestration layer*, as shown in figure 12.6. The orchestration layer is responsible for calling external systems via APIs. The specific terminology will vary based on your conversational AI platform, but this is often called a *webhook*.



- 1: Users send messages to the assistant's dialogue engine.
- 2: Orchestration layer optionally gathers additional context.
- 3: The NLU component interprets the user's message.
- 4: Orchestration layer optionally records a transcript.
- 5: The dialogue engine composes a response and returns it to the user.

Figure 12.6 You can use your chatbot's orchestration layer to create a conversational transcript in whatever format you need.

A webhook is a type of API. Webhooks are generally available before the bot processes the user's response (a "pre"-webhook), after processing the user's response (a "post"-webhook), or at other predefined events. Webhooks can access conversational context either natively or as input parameters. The next listing demonstrates pseudocode for constructing a transcript. (Consult your conversational AI platform documentation for the correct terminology and format.)

Listing 12.3 Pseudocode for a webhook that updates a transcript

```
def transcript_webhook(request, response):
    userMessage = request.input.text
    botMessage = response.output.text
    if (userMessage != null)
        response.context.transcript += 'User: ' +
            userMessage + '\n'
    if (botMessage != null)
        response.context.transcript += 'Bot: ' +
            botMessage + '\n'
```

The user's message is usually found in the request object.

The bot's message is usually found in the response object.

The user's message may not exist every time. For example, most conversations start with a bot greeting.

The transcript should be stored in a context variable (session variable). All user and bot messages are appended to the transcript.

Listing 12.3 demonstrates a post-webhook, since it has access to the request and the response. Every time the bot responds to the user, the webhook updates the transcript. This transcript includes the minimum possible elements—just the user and bot messages—and a very simple one-message-per-line format, readable by humans. You can create your transcript in whatever format you wish, such as a single string, string array, JSON object, or custom object.

As shown earlier, a simple transcript is easiest to read. Conversational AI platforms generally have many data elements available per message that you can optionally use in your transcripts:

- *Message timestamp*—You can use a timestamp to show the absolute time a message was received or sent (“11:25:53 AM”) or the relative time since the beginning of the chat (“00:01:15” for a message 1 minute and 15 seconds after the beginning).
- *Buttons*—You can indicate when the bot offered options via buttons and when the user clicked on a button. This is especially interesting in voice solutions, where users may enter dual tone multi-frequency (DTMF, or “touch tone”) input through their keypad. For example, you’ll know that the user pressed “0” rather than saying “zero.”
- *Original or post-processed input*—Many conversational AI platforms normalize certain input types like dates and numbers. You can use the original utterance, like “February first two thousand twenty-four” or a post-processed version like “02/01/2024”.
- *Rich-text and non-text elements*—Your bot may respond with HTML markup or even images and links that may not be suitable for a transcript.

The pseudocode in listing 12.3 demonstrated how to update a context variable that tracked conversational context, but it did not demonstrate how to initialize that variable. The simplest option is to initialize an empty string like the following:

```
response.context.transcript = ''.
```

Several other data elements related to a conversation are available, and you may wish to include them at the beginning of your transcript:

- *Session timestamp*—When the conversation started.
- *Session duration*—How long the conversation lasted.
- *User identifier*—This could include information about a logged-in user accessing the chat, such as name, email, or user ID. For a phone solution, this could be the caller’s phone number.
- *Device and channel identifier*—How the user accessed your conversational AI, such as device type (e.g., mobile or desktop) or which channel they used (e.g., chat widget, SMS, Facebook Messenger, etc.).
- *Transfer reason*—Why the bot transferred the caller to an agent, such as “immediate opt-out,” “opt-out,” or “bot didn’t understand user.”

You may instead choose to leave these elements for the structured section of the summary (as key-value pairs), rather than including them in a prose summary, since they apply to the entire conversation.

All these data elements and more are typically available in your conversational AI platform. They are often included in the AI’s system logs, which are another data source for building transcripts. These optional data elements are provided by conversational AI platforms because they are generic and are applicable to any conversation as metadata. They are a great start to any conversational transcript.

We saw earlier in the chapter that a good summary includes more than an unstructured transcript. Structured metadata is quite useful in summarizing the important parts of a conversation. Some parts of this metadata are not generic—they are specific

for your exact implementation. In our medical insurance example, member IDs and claim IDs were unique.

The conversational AI platform doesn't give any special meaning to member IDs—they are recorded as just another user message. If you want to use specific contextual elements from your implementation in a summary, you'll have to instrument the AI yourself to store them so that they can be included in a summary. Let's see how.

12.2.3 Instrumenting your chatbot (for data points)

Conversational AI platforms generally let you store arbitrary values in variables, often called *context variables* or *session variables*. You should make use of these for any data you collect during the conversation that has significant meaning, especially if it helps you find more information later.

The data you store will vary based on your specific application. Here are a few examples by domain:

- *Medical insurance*—Member ID, provider ID, claim ID
- *Retail*—Order number, product ID, retail location
- *Banking*—Account ID, account type

Any time the bot asks a question with a fixed-format response, like “what's your claim ID,” that response is a good candidate for instrumentation.

Be careful with sensitive data

Many types of data need to be treated carefully. There are rules and regulations on how to handle data that could identify a person (PI) or other sensitive data points. Your conversational AI may already deal with them, but adding them to summaries or logs may need to be reviewed with your legal team. Be minimalist in what you collect, what you store, and how long you store it, and confirm your choices with your lawyers.

The way you store context will depend on your conversational AI platform. You may be able to do this in a user interface, or your platform may require you to write code. Figure 12.7 shows the low-code method used in our platform to store context variables.

You can access your stored context variables later in the conversation, including accessing them to create a structured summary. The following listing shows pseudocode for accessing these context variables and storing them in a structured object.

Listing 12.4 Pseudocode for a webhook that creates a structured summary

```
def on_transfer(request, response):
    summary = ConversationSummary()

    summary.providerTaxID = response.context.
    ➡ ProviderTaxID
    summary.memberID      = response.context.MemberID
    summary.claimID       = response.context.ClaimID
```

← You can define a custom object to hold your summary.

← This method called as the conversation is transferred to an agent.

Set any values required for your custom summary.

The screenshot displays the 'Conversation steps' configuration in Watson Assistant. On the left, a list of steps is shown: Step 1 asks 'Is this for the same provider?' with a 'Confirmation' button; Step 2 is marked 'Action complete'; Step 3 asks 'What's your Tax ID?' with a 'Regex' button. On the right, 'Step 5' is being configured. It is set to 'Is taken' with the condition 'with conditions'. A condition is defined: 'If All of this is true: 1. 4. What's your Tax ID? is Regular_expression'. Below the conditions, a 'Set variable values' action is configured to set the variable 'ProviderTaxID' to the value of '4. What's your Tax ID?'.

- 1: After the bot validates a structured input...**
2: ... the value can be stored in a context variable.

Figure 12.7 Storing contextually important information into a context variable so that it can be retrieved later by a summary

You can also use these variables directly in your assistant to create an unstructured summary. Figure 12.8 shows the low-code method used in our platform to combine several variables into a larger summary string.

The screenshot shows the 'Expression editor' interface. It includes a link 'Learn more about writing expressions in watsonx Assistant.' and a text area containing the following code snippet:

```
T UnstructuredSummary .append("\n\n")
.append("\n").append("Provider Tax ID: ").append(ProviderTaxID.value)
.append("\n").append("Member ID: ").append(MemberID.value)
```

Figure 12.8 Using a low-code expression editor to combine multiple data elements into a summary

This section demonstrated multiple ways to collect the data necessary for summarization. Conversational AI platforms collect a lot of data that you can use in summaries.

You can instrument your AI assistants to collect the additional data you need, and you can control the formatting of that data. The data you collect is useful for many purposes, including efficient handoffs to human agents.

Exercises

- 1 Revisit the summaries you created in the section 12.1 exercises. Would you now change any of the data elements included in those summaries?

12.3 Improving summaries with generative AI

What if you don't want to modify your conversational AI at all? Can you still collect all the data you need for a great summary and format it the way you need? Can generative AI do more work so you have less work? Yes! Let's look at how.

You have two key prerequisites. First, you need a conversational transcript in some form: a built-in transcript from your conversational AI platform, one you created yourself, or an extract from your platform's conversational logs. Second, you need to know what a good summary looks like for your use case. Armed with those two prerequisites, you can work with an LLM to get the summary you need.

In this section, we will use the granite-13b-chat-v2 model with greedy decoding. This model is good at the summarization and extraction techniques we require. We'll use greedy decoding so that the model will not be creative and the outputs will be repeatable. (We want to generate the same summary and extract the same details for a given conversation.)

12.3.1 Generating a text summary of a transcript with summarizing prompts

We'll start our exercise with a simple summarization prompt, shown in the following listing. We'll pass the model the JSON version of our chat transcript.

Listing 12.5 Generating a summary of a JSON chat transcript

Summarize the following conversation transcript between a user ("u") and the automated assistant ("a").
The summary should be 1-2 sentences long.

Instruction to limit the summary size

Task description and hint for interpreting the JSON object

Transcript:

```
[{"a": "Welcome to the automated assistant, how can I help?"},
{"u": "Claim status", "n": true}, {"a": "Ok, claims.\nWhat's your\nTax ID?"}, {"u": "123456789"}, {"a": "Thank you.\nWhat's the\nmember's eight-digit Member ID?"}, {"u": "87654321"}, {"a": "Ok,\nand what's the member's date of birth?"}, {"u": "12/31/2000"},
{"a": "Is the member's name John Doe?"}, {"u": "option: [\"Yes\", \"No\"]"},
{"u": "Yes"}, {"a": "Great!\nWhat's the date of service for the\nclaim?"}, {"u": "2024-02-01"}, {"a": "I found your claim #111222333\nfrom Feb 1, 2024 for the member #87654321 and provider\n#123456789. It was paid on 5/23/2023 for $201.83."}, {"u": "representative"}]
```

JSON version of the chat transcript

Summary: ← Cue
 "I found your claim #111222333 from Feb 1, 2024 for the member
 ➡ #87654321 and provider #123456789. It was paid on 5/23/2023
 ➡ for \$201.83." Model output

The generated summary is the verbatim last utterance from the AI. This may seem strange at first, but this is a pretty good summary of the conversation. The bot's last utterance is rich in details that encompass the most important elements of the conversation.

By using sampling decoding and a creative temperature setting, we could have gotten a differently structured summary, at the risk of introducing hallucinations. We can also change the transcript format and see if that helps the LLM. The next listing shows a prompt that summarizes the same conversation with a different input format—unstructured text instead of JSON.

Listing 12.6 Generating a summary of a text chat transcript

Summarize the following conversation transcript between a user and an AI Task description with no additional hints

The summary should be 1-2 sentences long. ← Same instruction to limit the summary size

Transcript: ← Text version of chat transcript

AI: "Welcome to the automated assistant, how can I help?"
 User: "Claim status"
 AI: "Ok, claims. What's your Tax ID?"
 User: "123456789"
 AI: "Thank you. What's the member's eight-digit Member ID?"
 User: "87654321"
 AI: "Ok, and what's the member's date of birth?"
 User: "12/31/2000"
 AI: "Is the member's name John Doe?"
 User: "Yes"
 AI: "Great! What's the date of service for the claim?"
 User: "February first 2024"
 AI: "I found your claim #111222333 from February 1, 2024 for the member #87654321 and provider #123456789. It was paid on May 23, 2024 for \$201.83."
 User: "representative"

Summary: ← Same cue

The AI is providing information about a claim that has been submitted by a member with the member ID 87654321 and provider ID 123456789. The claim was paid on May 23, 2024 for \$201.83. The representative will be able to assist with any questions or concerns related to this claim. Model output

This generated summary contains the same informational content as the summary generated from JSON. It's still based on the bot's last message, but this time, it's paraphrased into prose. The summary also includes a third sentence—in contrast with the prompt's instructions—that attempts to make sense of the user's last utterance.

NOTE For the remainder of this chapter, we will omit the conversational transcript from the book to keep the listings smaller. We will use the human-readable version of the conversation transcript. The full listings are available in the book’s GitHub repository at <https://github.com/andrewfreed/EffectiveConversationalAI>.

Remember our advice earlier in this section: you need to know what a good summary looks like. The summaries generated by the LLM have been okay so far, but perhaps we can do better. One method is to provide better instructions to the model, as shown in the following listing. Because we are emphasizing instructions, we’ll switch to a more instructible model, in this case granite-13b-instruct-v2. The model is asked to emphasize the dialogue immediately preceding the escalation.

Listing 12.7 Enhancing the summarization instructions to the LLM

		Original instruction to limit the summary size
Summarize the following conversation transcript between a user and an AI		Original task description
The summary should be 1-2 sentences long.		
The summary should emphasize the last task the user was performing that led them to request a human representative.		Enhanced instruction on what to emphasize
Transcript:		
<... human-readable transcript removed for brevity ...>		Text version of chat transcript
Summary:		
The user asks for a human representative because they want to know why their claim was denied.		Same cue
	Model output	

The model follows the instruction. The last message from the bot tells the user their claim was paid (“The claim was paid on May 23, 2024 for \$201.83”), and the user then opts out. The LLM summary succinctly says, “they want to know why their claim is denied.” This summary is short but perhaps too speculative, seeing that the claim was paid and not actually denied. Maybe the user felt they should have been paid more, or maybe they needed an itemized list. The text summary also omits the near-verbatim playback of the transcript seen in listing 12.6, leaving that information for a separate structured summary. We’re closer to what our human agent needs. Let’s improve the prompt to reduce the LLM’s speculation.

A useful method is to invoke the LLM with a one-shot prompt (with one example) or a few-shot prompt (with multiple examples). Creating the one-shot example forces us to think about what a good summary looks like for a given conversation. Using one or more examples is often the fastest way to improve a prompt.

The following listing shows a one-shot example that also uses the granite-13b-instruct-v2 model.

Listing 12.8 A one-shot summarization prompt

```
<|instruction|>
Summarize the following conversation transcript between a user
and an AI
The summary should be 1-2 sentences long.
The summary should emphasize the last task the user was
performing that led them to request a human representative.
```

**Original task
description
with updated
delineation via
“<|instruction|>”**

```
<|transcript|>
AI: "Welcome to the automated assistant, how can I help?"
User: "Claim status"
AI: "Ok, claims. What's your Tax ID?"
User: "012345678"
AI: "Thank you. What's the member's eight-digit Member ID?"
User: "I don't have it"
AI: "Ok, let's try something else. What's the member's date
of birth?"
User: "representative"
```

**One-shot example
of conversation
transcript with
summary**

```
<|summary|>
The user requested a human representative after failing
to validate the member ID needed to find the claim.
```

```
<|transcript|>
<... human-readable transcript removed for brevity ...>
```

**Text version of
chat transcript**

```
<|summary|>
The user requests a human representative because they
want more information about their insurance claim.
```

Restructured cue

Model output

This generated summary is also quite reasonable when combined with the structured metadata. There is some speculation—“they want more information”—but again it seems like the user would need more information if they wanted an agent after finding the claim is supposedly paid. The summary is also structured after the example, with a token-for-token match in the first six words of the summary.

Note that in the one-shot summarization example (listing 12.8), we used a slightly different format. Instead of using delineation via `Transcript:`, we used specially formatted tokens like `<|transcript|>`. Without these special tokens, we were unable to generate good summaries. Likely the model had trouble separating the prompt sections and the conversation elements because both used colons. Future large language models (LLMs) may be more resilient to delineation characters appearing multiple places in the prompt. This kind of small change can have a huge effect on LLM performance.

Either emphasizing instructions or examples in your prompts can work. Consider the following tradeoffs:

- *Control of the output*—Most LLMs are trained on summarization by default. Many are responsive to instructions and generate good summaries. One-shot and few-shot prompts further constrain the output to use the language you desire, though you may have to provide several examples.
- *Cost*—Adding examples increases the inference cost due to there being more input tokens.

The text summaries we have generated so far include an overview of the conversation but mostly do not include the structured metadata that will be helpful for the agent. Earlier in the chapter, we demonstrated instrumenting your chatbot using a variety of code and low-code methods to gather structured metadata that could be passed to the agent. What if we don't want to instrument our chatbot—could an LLM extract the structured metadata?

12.3.2 *Generating a structured summary of a transcript with extractive prompts*

We can use LLMs to extract structured data from conversation transcripts—extraction is another task that many LLMs are trained on.

Our first task is to decide what the structured output needs to look like. There are many possibilities, but one useful format is JSON. This is useful for two reasons: JSON is easy for downstream applications to consume, and many LLMs are good at generating JSON.

We will again use an instructible model. This time we will use `mistral-7b-instruct-v0-2` because it can generate JSON from instructions alone. The following listing shows a prompt that generates structured JSON output from the conversation.

Listing 12.9 An extractive summary without examples

<pre>< instruction > Read the following conversation transcript between a user and the automated assistant. Extract all IDs in JSON format.</pre>		Updated task description with simple instruction about JSON format
<pre>< transcript > <... human-readable transcript removed for brevity ...></pre>	Text version of chat transcript	
<pre>< JSON > { "TaxID": "123456789", "MemberID": "87654321", "DateOfBirth": "12/31/2000", "Name": "John Doe", "ClaimNumber": "111222333", "DateOfService": "February 1, 2024", "ProviderID": "123456789" }</pre>	Updated cue to generate JSON	Model output

This is an excellent first attempt. The model generated JSON and extracted all the structured data points collected. But it collected more data than we asked for—we wanted only the IDs—and it duplicated one of the data points (the tax ID is the provider's ID).

NOTE Many models can generate JSON after seeing a few examples. All the models we tested extracted several data points rather than the three expected. The `mistral` model was one of the few to generate valid JSON with

no examples in the prompt. We expect models to continue improving at generating JSON data. Alternatively, you can provide an example schema in your instruction.

Let's augment this prompt with an example (shown in bold).

Listing 12.10 An extractive summary with one example

<pre>< instruction > Read the following conversation transcript between a user and the automated assistant. Extract all IDs in JSON format.</pre>	Same task description
<pre>< transcript > AI: "Welcome to the automated assistant, how can I help?" User: "Claim status" AI: "Ok, claims. What's your Tax ID?" User: "333444555" AI: "Thank you. What's the member's eight-digit Member ID?" User: "55667788" AI: "Ok, and what's the member's date of birth?" User: "April 19, 2024"</pre>	One-shot example
<pre>< JSON > {"TaxID": 333444555, "MemberID": 55667788}</pre>	Text version of chat transcript
<pre>< transcript > <... human-readable transcript removed for brevity ...></pre>	
<pre>< JSON > {"TaxID": 123456789, "MemberID": 87654321, "DateOfService": "February 1, 2024", "ClaimNumber": "111222333"}</pre>	Same cue Model output

With one example, we were able to show the model that we didn't need every piece of data in the output. We also got the model to stop duplicating the provider's tax ID. Last, the JSON response is now minifying to a single line without line breaks. The extracted data is still accurate, but we may require exact key names. If the agent's application expects to read a field called `ClaimID`, then it is not acceptable for the summary to reference `ClaimNumber`.

We give the model a more detailed example (shown in bold) in the following listing.

Listing 12.11 Updated one-shot example for extractive summary

<pre>< instruction > Read the following conversation transcript between a user and the automated assistant. Extract all IDs in JSON format. #</pre>	Same task description
---	--

<pre> < transcript > AI: "Welcome to the automated assistant, how can I help?" User: "Claim status" AI: "Ok, claims. What's your Tax ID?" User: "333444555" AI: "Thank you. What's the member's eight-digit Member ID?" User: "55667788" AI: "Ok, and what's the member's date of birth?" User: "April 19, 2024" AI: "Is the member's name Jim Smith?" User: "Yes" AI: "Great! What's the date of service for the claim?" User: "April ninth 2024" AI: "I found your claim #444444555 from April 9, 2024 for the member #55667788 and provider #333444555. It was paid in full on April 23, 2024 for \$156.81." User: "Thanks! Goodbye" < JSON > {"TaxID": 333444555, "MemberID": 55667788, "ClaimID": 444444555} < transcript > <... human-readable transcript removed for brevity ...> < JSON > {"TaxID": 123456789, "MemberID": 87654321, "ClaimID": "111222333"} </pre>	Updated one-shot example Text version of the chat transcript Same cue Model output
--	---

This worked well. We only get the keys we desired. It's slightly frustrating that the one-shot example needed to be so close to the second transcript. This implies that to summarize other conversational flows, we may need to provide examples for each one. (What if the conversation includes authorization IDs, electronic payment IDs, or other IDs?) And there's one other gotcha: the `TaxID` and `MemberID` were numeric values, but it produced a string value for `ClaimID`—even after seeing the example.

Testing for hallucinations

In our extractive summarization examples, we did not encounter any hallucinations, but this doesn't mean they are impossible. Any summarization prompt should be tested on multiple inputs before it is deployed to see if it hallucinates. After it is deployed, you can detect hallucinations by verifying that each extracted value appeared in the transcript text.

Let's go back to the previous one-shot example and instead add some instructions, shown in bold in the following listing.

Listing 12.12 An extractive summary with one example

```

<|instruction|>
Read the following conversation transcript between a user and the automated
assistant. Extract all IDs in JSON format.
Use only the following JSON keys: "TaxID", "MemberID",
"ClaimID".
The JSON values should be numbers, not strings.

<|transcript|>
AI: "Welcome to the automated assistant, how can I
help?"
User: "Claim status"
AI: "Ok, claims. What's your Tax ID?"
User: "333444555"
AI: "Thank you. What's the member's eight-digit Member
ID?"
User: "55667788"
AI: "Ok, and what's the member's date of birth?"
User: "April 19, 2024"

<|JSON|>
{"TaxID": 333444555, "MemberID": 55667788}

<|transcript|>
<... human-readable transcript removed for brevity ...>

<|JSON|>
{"TaxID": 123456789, "MemberID": 87654321, "ClaimID":
111222333}

```

Augmented task description

Smaller one-shot example

Text version of the chat transcript

Same cue

Model output

Voila! Exactly the output we desired. We were able to prompt an LLM to produce the structured summary we wanted using a combination of instructions and examples. We did not need to modify our assistant, aside from calling the LLM. That call is isolated to the component that handles transferring conversations to agents.

Exercises

- 1 Review the ideal summary you created in section 12.1 for the example chat transcript. Use your favorite LLM (or your company's preferred LLM) to generate a prose summary of that transcript. How close does the LLM get to your preferred summary? Did you use instructions, few-shot examples, or both?
- 2 Repeat exercise 1, but also modify the model parameters. In this chapter, we used greedy decoding and a repetition penalty of 1.1. Try sampling decoding, or try raising or lowering the penalty. Do you get better performance with different parameters? Does this match your expectations?
- 3 Is it possible to generate the prose summary *and* to extract key details from the transcript in the same prompt? Design a prompt that generates this output:

```

{"TaxID": 123456789, "MemberID": 87654321, "ClaimID": 111222333,
"Summary": "The user requests a human representative because they
want more information about their insurance claim."}

```

(continued)

- 4 Invent a conversation that includes multiple claim searches. Create dialogue you imagine a bot and user would have if the user was calling to check on four total claims. Some of the claims are paid, some of the claims are still processing, and the last claim was denied. This conversation will have approximately two to four times as much content as the original sample conversation. What do you want this summary to look like? Once you have a target summary in mind, use an LLM, and try to generate that summary.
- 5 Design an ideal summary for the following sample conversation:

```
AI: "Welcome to the automated assistant, how can I help?"
User: "coverage"
AI: "Ok, claims. What's your Tax ID?"
User: "333444555"
AI: "Thank you. What's the member's eight-digit Member ID?"
User: "55667788"
AI: "Ok, and what's the member's date of birth?"
User: "April 19, 2024"
AI: "Is the member's name Jane Williams?"
User: "Yes"
AI: "They are currently a member. They have our PPO plan with coverage
  for the full 2024 calendar year. Do you need anything else?"
User: "How much is their co-pay for office visits?"
AI: "Since you are an in-network provider they will owe $20"
User: "Do they need a prior authorization for a specialist?"
AI: "Yes, would you like to initiate one now?"
User: "Yes"
AI: "Let me transfer you to a specialist who can help with that."
```

Now build a prompt for your LLM to generate a similar summary.

- 6 Test your new prompt against the original sample conversation in listing 12.1. If necessary, refine the prompt so that it generates good summaries for both conversations.

Summary

- Transfers to human agents are an inevitable part of many conversational AI solutions. Agents benefit from receiving brief summaries that extract key highlights from the conversation, both in prose and in structured formats.
- A summary requires a conversational transcript. Most conversational AI platforms generate a transcript for you, but you can configure your conversational AI to generate one in your desired format.
- Structured summaries can be generated by enhancing your conversational AI to store key data points as they are collected, or they can be extracted using LLMs when the conversation is completed.
- You need to know what a good summary looks like before you ask an LLM to generate one.
- LLMs can generate prose summaries and extract key details from transcripts. Use clear instructions and examples to generate the summary you desire.

A

accuracy 81
agents, prior poor experience with 252
AHT (average handling time) 52
AI (artificial intelligence)
 classification-based, annotated logs for 101
 improving weak understanding, improving recall for one intent 116
 traditional (classification-based) 87–89
 weak understanding, identifying problematic patterns in misunderstood utterances 106
AI assistants 4
annotated logs 101–104
 for generative AI 103
 for traditional (classification-based) AI 101
answer generation 144
AOV (average order value) 52
API/backend processes, supporting self-service task flows with 204
APIs (application programming interfaces) 7, 38
 using 9
Arize 168
AzureAIDocumentIntelligence-Loader 154

B

behavioral patterns 215
blind testing 87

C

call center agents 34
classification models 97
classifiers, defined 9
CLV (customer lifetime value) 52
CohereEmbeddings 157
commercial cloud platform 22
comparison check 144
complex flows 193
complexity 194–198
 effect on business metrics 196–198
 effect on end user 194
 incremental cost and benefit of reducing for user 198
confusion matrix 109, 129
contained conversations, defined 54
containment rate, defined 54
context variables 288
context, importance of in virtual assistant performance 208–216
 building trust and loyalty 212
contextual information 212–216
 efficiency in problem solving 211
enhanced relevance and accuracy 209

influencing user interactions 209–212
personalized experience 210
proactive support 212
contextual information
 behavioral patterns 215
 device type 214
 modality 216
 previous interactions 215
 time zone 213
 user location 213
 user preferences 214
continuous improvement 15–21
conversation outcomes 55–57, 63–65
conversational AI 3
 benefits of 5
 building 8–10, 23
 chatbots 4
 preparing for summarization 284–290
 prior poor experience with 252
 continuous improvement 15–21
 defined 4
 how it works 6
 process-oriented bots 33–37
 responding to users with generative AI 38–42
 search 136–140
 software platforms 22
 traditional 83
 understanding users 79

conversational summarization 278
 elements of effective 280–283
 need for 279
 overview of 279
 coverage, defined 81
 cross-functional teams 47–49

D

decoder-only architectures 85
 decoding_method 41
 device type 214
 dialogue flows, spotting
 complex 199
 digital employees 4
 direct question 181
 DirectoryLoader 154
 disambiguation feature 97
 document loaders 153
 document transformers 153–154
 DTMF (dual tone multi-frequency) 287

E

embedding generation 149
 encoder-decoder model
 architecture 85
 encoder-only architectures 85
 extensions 38
 extracting meaning 9

F

F1 scores 106
 improving for one intent 120
 FAISS (Facebook AI Similarity Search) 158
 FaithfulnessEvaluator 165
 FAQ (frequently asked question)
 bots 10, 23–24
 dynamic question and
 answering 31
 foundations of 24
 static question and
 answering 26–31
 few-shot prompting 86, 180
 first-contact resolution (FCR) 52
 fixes, developing and
 delivering 74–75
 Flan-ul2 model 14
 foundation models 11
 free-text summary elements 282
 fulfillments 38

G

generated_text 42
 generation metrics 163–165
 generative AI (artificial
 intelligence) 10–15, 84–86,
 89, 230
 AI-assisted flows at test
 time 243–247
 AI-assisted process flows at
 build time 231–237
 AI-assisted process flows at run
 time 237–243
 annotated logs for 103
 augmenting intent data
 with 170
 exercises 190
 hardening existing
 intents 175–188
 LLMs 171–175
 defined 11
 effectively using 13–15
 executing dialogue flows
 with 238–240
 guardrails 12–13
 improving dialogue with
 270–276
 improving summarization
 with 290–298
 model platform 22
 solution, assessing 92
 using LLM for search
 process 240–243
 golden intent 117
 golden test set
 annotating for generative
 AI 100
 annotating for traditional
 (classifier-based) AI 99
 grammatical variations, generat-
 ing new 179–182
 granite-13b-instruct-v2
 model 292
 greetings and introduc-
 tions 257–259

H

hallucinations 11, 151
 HuggingFaceInference-
 Embeddings 157
 human in the loop 13

I

immediate opt-outs 256–261
 allowing user to opt in 260

conveying capabilities and set-
 ting expectations 259
 incentivizing self-service 259
 starting with great
 experience 257–259
 improvement
 identifying and resolving
 problems 65–73
 determining acceptance
 criteria 72
 finding problems 65–67
 group review 67–72
 recognizing need for 45
 improvement planning 44, 106
 cross-functional teams 47–49
 developing and delivering
 fixes 74–75
 driving to same goal 49–63
 incremental improvements 110
 indexing metrics 159–161
 integrations 38
 intent data
 augmenting with
 generative 170–175
 augmenting with LLMs
 188–189
 intent matching, wrong intent
 matched 116
 intents 9, 26
 hardening existing
 intents 170–188
 improving precision for one
 intent 118
 improving recall for one
 intent 116
 iterative improvement 103
 IVR (interactive voice
 response) 35, 252

K

k-fold 123
 cross validation 88
 testing 114
 Kanban board 76
 KPIs (key performance
 indicators) 46

L

LLMs (large language
 models) 10, 37, 130, 141,
 171–175, 231
 augmenting intent data
 with 188–189
 integrating with 38

pros and cons of 172
 requirements for 173
 routing requests to 41
 using augmented data 173
 using for search process
 240–243
 load method 153
 logs, obtaining and preparing
 test data from 93–101
 annotation process 99–101
 guidelines for identifying can-
 didate test utterances 94–98
 obtaining production logs 93
 preparing and scrubbing data
 for use in iterative
 improvements 98–99

M

max_tokens 41
 meaning, extracting 9
 messages
 error messages 270–271
 greeting messages 272–276
 metadata, summary
 elements 280
 min_tokens 41
 modality 216–223
 comparing modalities 217
 examples of how modality
 affects user experience
 220–221
 importance in designing vir-
 tual assistant flows 219
 voice bot design
 considerations 222
 models, selection 12–15
 MPT-7B-Instruct model 14
 multiple intents, improving pre-
 cision and recall for 120

N

NDCG (Normalized Discounted
 Cumulative Gain) 163
 NLP (natural language
 processing) 228
 NLU (natural language
 understanding) 139
 no intent matched 125–130
 clustering utterances for new
 intents 125–129
 when to stop adding
 intents 130
 NPS (net promoter score) 61

O

Ollama 176
 one-shot prompting 86, 180
 OpenAIEmbeddings class 156
 opt-outs 251
 drivers of 252–256
 escalation 277
 gathering data on opt-out
 behavior 254
 immediate 256–261
 improving dialogue with gen-
 erative AI 270–276
 reducing 262–266
 retention 266–270
 orchestration layer 285
 over-selection 107

P

parameter tuning 86
 passage retrieval 143
 persistent user history 212
 PII (personal identifiable
 information) 93
 postfiltering output 13
 precision 106–107
 improving for one intent 118
 precision and recall, improving
 for multiple intents 120
 prefiltering input 12
 preprocessing data 149
 previous interactions 215
 process flows
 AI-assisted at build time
 231–237
 AI-assisted at run time 237–243
 executing dialogue flows with
 generative AI 238–240
 using LLM for search
 process 240–243
 process-oriented bots 33–37
 routing agents 33
 transitioning from routing
 agents to 35–37
 process-oriented or transactional
 solutions 4
 production logs, obtaining 93
 prompt engineering 86
 prompt stuffing 154
 prompts 12

Q

QPS (queries per second) 160
 qualitative problem

exploration 66
 quantitative evaluation for issue
 discovery 67
 question-answering 4, 86

R

RAG (retrieval-augmented
 generation) 10, 62, 86, 130,
 135–136, 140–146, 207, 240
 additional
 considerations 151–159
 benefits of 142–144
 combining with other genera-
 tive AI use cases 145
 comparing intents, search, and
 RAG approaches 145
 designing adaptive flows
 with 224–226
 enhancing context awareness
 and improving overall user
 experience with 223–228
 evaluating and analyzing
 performance 159–168
 implementation of 146–150
 in conversational AI 141
 maintaining and updating
 adaptive flows 228
 retrieving and generating con-
 textually relevant
 responses 226
 RAGAS, defined 168
 recall 106
 improving for one intent 116
 retrieval and matching at
 runtime 149
 retrieval metrics 161–163
 retrievers 153
 ROI (return on investment) 52
 routing agents 4, 33
 routing requests to LLMs 41

S

sampling decoding 178
 screen pop 282
 search processes, using LLM
 for 240–243
 search, role of in conversational
 AI 136–140
 benefits of traditional
 search 138
 drawbacks of traditional
 search 139
 using search in conversational
 AI 137

self-service
 ask flows, supporting with API/
 backend processes 204
 incentivizing 259
 sensitive data 288
 session history 212
 session variables 284, 288
 slot filling 202
 SMEs (subject matter
 experts) 47, 115, 237
 solutioning 71
 sprint planning 75
 SSA (Sensibleness and Specific-
 ity Average) 164
 storage in vector database 149
 summarization
 elements of effective 280–283
 improving with generative
 AI 290–298
 need for 279
 overview of 279
 preparing chatbot for
 284–290
 support resources 81
 synonyms, generating 176–179

T

task flows, supporting self-service
 task flows with API/backend
 processes 204
 templates, creating examples
 with 185–187
 TensorFlowEmbeddings 157
 test data, obtaining and prepar-
 ing from logs 93–101
 annotation process 99–101
 guidelines for identifying can-
 didate test utterances
 94–98
 obtaining production logs
 93
 preparing and scrubbing data
 for use in iterative
 improvements 98–99
 test time, AI-assisted flows
 at 243–247
 setting up conversational
 test 246–247
 setting up generative AI to be
 user 244–246
 time zone 213
 traditional (classification-based)
 AI solution, assessing 91

traditional AI
 improving weak understanding
 for 105, 131–133
 weak understanding, wrong
 intent matched 116
 traditional conversational AI
 83
 traditional search 136
 training data, selection 12
 transfer decision 144
 transitioning from routing agents
 to process-oriented bots
 35–37
 triaging issues 68–71
 true negatives 108

U

uncontained conversations 54
 understanding
 achieving with generative
 AI 84–86
 achieving with traditional con-
 versational AI 83
 annotated logs 101–104
 fundamentals of 84–86
 iterative improvement 103
 measuring 87–90
 weak 80–81
 UnstructuredHTMLLoader
 154
 user data
 aligning with user's mental
 model 201
 leveraging what is known about
 user 200
 user journeys
 aligning with user's mental
 model 201
 allowing flexibility in expected
 user responses 202–204
 spotting complex dialogue
 flows 199
 supporting self-service task
 flows with API/backend
 processes 204
 using what is known about
 user 200
 user location 213
 user preferences 214
 users
 assessing where you are
 today 91–92
 understanding 79

utterances 9
 identifying problematic pat-
 terns in misunderstood
 utterances 106, 109

V

verb phrases 177
 virtual agents 4
 virtual assistants
 enhancing context awareness
 and improving overall user
 experience with RAG
 223–228
 importance of context in
 performance 208–216
 modelities 217–223
 voice bots, design
 considerations 222
 voice solutions,
 accommodating 263

W


watsonx.ai platform 176
 weak understanding 80
 causes of 81
 establishing baseline 112
 identifying biggest
 problems 110
 identifying problematic pat-
 terns in misunderstood
 utterances 106, 109
 improvement plan 106
 improving F1 score for one
 intent 120
 improving for traditional
 AI 105
 improving precision for one
 intent 118
 improving recall for one
 intent 116
 incremental
 improvements 110
 solving 125–130
 traditional AI 131–133
 validating initial training
 strategy 115
 wrong intent matched 116
 webhook 285

Z


zero-shot prompting 180

The evolution of question-answering


Many chatbot builders start with curated answers for high-frequency intents.



Using search adds dynamism that handles nuance. Users piece their own answers together.



Retrieval-augmented generation's large language models summarize relevant passages into an answer.



Effective Conversational AI

Freed • Jacobs • Rózsa • Foreword by Jesús Mantas

Powerful new chatbot frameworks and Generative AI models can practically eliminate problems like misinterpreting user intent and delivering nonsensical answers. In this book, you'll learn how to build chatbots that take advantage of large language models and other modern tools and create conversational AI experiences users will love.

Effective Conversational AI teaches you how to build great chatbots that perform reliably even at enterprise scale. In it, you'll learn how to clarify user intent using LLMs, respond accurately to unanticipated input, and use Retrieval Augmented Generation to keep responses up to date. Along the way, you'll discover how to establish a feedback loop for continuous quality improvement and master techniques to integrate GenAI safely into conventional chatbot designs.

What's Inside

- Blend Generative AI and conventional chatbot tools
- Use LLMs to improve quality, accuracy, and usability
- Plan for continuous improvement
- Domain-specific responses using RAG

For developers, engineers, and product managers working with conversational AI.

Andrew Freed, **Cari Jacobs**, and **Enikő Rózsa** are seasoned conversational AI developers with IBM.

The technical editor on this book was Jack C Crawford.

For print book owners, all digital formats are free:
<https://www.manning.com/freebook>

“A wonderful comprehensive guide written by individuals who have walked the AI conversational implementation journey into production.”

—Sara Hines
AI innovation pioneer

“Cuts through the hype and focuses on what really matters.”

—Jerry Cuomo, IBM

“A blueprint for building and measuring effective conversational AI systems.”

—Corville Allen, Google

“An invaluable resource for anyone looking to drive success with AI-driven conversations.”

—Marc Nehme, Microsoft

